

LATTICE-BASED DATA STRUCTURES FOR DETERMINISTIC PARALLEL AND DISTRIBUTED PROGRAMMING

Lindsey Kuper

Submitted to the faculty of the University Graduate School
in partial fulfillment of the requirements
for the degree
Doctor of Philosophy
in the School of Informatics and Computing
Indiana University
September 2015

Accepted by the Graduate Faculty, Indiana University, in partial fulfillment of the requirements for the degree of Doctor of Philosophy.

Ryan R. Newton, Ph.D.

Lawrence S. Moss, Ph.D.

Amr Sabry, Ph.D.

Chung-chieh Shan, Ph.D.

09/08/2014

Copyright 2015
Lindsey Kuper
All rights reserved

Acknowledgements

Not long ago, someone asked me if I knew I wanted to work on LVars at the time I applied to grad school, and I had to laugh. Not only did I not know then that I wanted to work on LVars, I didn't even know what subfield of computer science I wanted to study. I had an idea of what kinds of problems I thought were interesting, but I didn't really grasp that there was actually a field of study devoted to those problems and that "programming languages" was what it was called. As it happened, one of the Ph.D. programs to which I applied turned out to be a very good place to study PL, and that program was also the only one that accepted me. So the first acknowledgement I want to make here is to the Indiana University computer science program for taking a chance on me and making this dissertation possible.

I took Dan Friedman's programming languages course in my first semester at IU in fall 2008, and the following spring, Dan invited me to be a teaching assistant for the undergraduate edition of the course, known as C311, together with his student and collaborator Will Byrd. This led to my first PL research experience, which consisted of playing with a fantastic little language called miniKanren¹ during the summer of 2009 with Dan, Will, and a group of relational programming aficionados fueled by intense curiosity and Mother Bear's pizza. At the time, I was utterly unprepared to do research and was in way over my head, but we had a lot of fun, and I went on to help out with C311 and another course, H211, for three more semesters under Dan and Will's expert guidance.

Amal Ahmed joined the IU faculty in 2009 and was a tremendous influence on me as I began truly learning how to read research papers and, eventually, write them. Although Amal and I never worked on determinism as such, it was nevertheless from Amal that I learned how to design and prove properties of calculi much like the λ_{LVar} and λ_{LVish} calculi in this dissertation. Amal also had a huge influence on the culture of the PL group at IU. Weekly "PL Wonks" meetings are a given these days, but it wasn't so very long ago that the PL group didn't have regular meetings or even much of a group identity. It was

¹<http://www.minikanren.org> — not that any such website existed at the time!

Amal who rounded us all up and told us that from then on, we were all going to give talks every semester. (Happily, this process has now become entirely student-driven and self-sustaining.) It was also at Amal's urging that I attended my first PL conference, ICFP 2010 in Baltimore, and got to know many of the people who have since become my mentors, colleagues, collaborators, and dear friends in the PL community, including Neel Krishnaswami, Chung-chieh Shan, Tim Chevalier, Rob Simmons, Jason Reed, Ron Garcia, Stevie Strickland, Dan Licata, and Chris Martens.

When Amal left IU for Northeastern in 2011, I had to make the difficult decision between staying at IU or leaving, and although I chose to stay and work with Ryan Newton on what would eventually become this dissertation, I think that Amal's influence on the flavor of the work I ended up doing is evident. (In fact, it was Amal who originally pointed out the similarity between the Independence Lemma and the frame rule that I discuss in Section 2.5.5.)

Ryan's first act as a new faculty member at Indiana in 2011 was to arrange for our building to get a fancy espresso machine (and to kick-start a community of espresso enthusiasts who lovingly maintain it). Clearly, we were going to get along well! That fall I took Ryan's seminar course on high-performance DSLs, which was how I started learning about deterministic parallel programming models. Soon, we began discussing the idea of generalizing single-assignment languages, and Ryan suggested that I help him write a grant proposal to continue working on the idea. We submitted our proposal in December 2011, and to my surprise, it was funded on our first try. I gratefully acknowledge everyone at the National Science Foundation who was involved in the decision to fund grant CCF-1218375, since without that early vote of confidence, I'm not sure I would have been able to keep my spirits up during the year that followed, in which it took us four attempts to get our first LVars paper published. Ryan's non-stop energy and enthusiasm were crucial ingredients, too. (And, although I didn't feel thankful at the time, I'm now also thankful to the anonymous reviewers of POPL 2013, ESOP 2013, and ICFP 2013, whose constructive criticism helped us turn LVars from a half-baked idea into a convincing research contribution.)

In addition to Ryan, I was lucky to have a wonderful dissertation committee consisting of Amr Sabry, Larry Moss, and Chung-chieh Shan. Amr, Larry, and Ken all played important roles in the early development of LVars and helped resolve subtle issues with the semantics of λ_{LVar} and its determinism proof. What appears now in Chapter 2 is relatively simple — but it is only simple because of considerable effort spent making it so! I’m also grateful to all my committee members for showing up to my thesis proposal at eight-thirty in the morning in the middle of an Indiana December snowstorm. Additionally, I want to offer special thanks to Sam Tobin-Hochstadt, Neel Krishnaswami, and Aaron Turon — all of whom gave me so much guidance, encouragement, and support throughout this project that I think of them as *de facto* committee members, too.

By fall 2012, working on LVars had gotten me interested in the idea of developing a separation logic for deterministic parallelism. Aaron and Neel liked the idea, too, and expressed interest in working on it with me. When we all met up in Saarbrücken in January 2013, though, we realized that we should first focus on making the LVars programming model more expressive, and we put the separation logic idea on hold while we went on an exciting side trip into language design. That “side trip” eventually took us all the way to event handlers, quiescence, freezing, and quasi-determinism, the topics of Chapter 3 and of our POPL 2014 paper. It also led to the LVish Haskell library, the topic of Chapter 4, which was largely implemented by Aaron and Ryan. Neel and I collaborated on the original proof of quasi-determinism for λ_{LVish} , and Neel also provided lots of essential guidance as I worked on the revised versions of the proofs that appear in this dissertation. I also want to acknowledge Derek Dreyer for helping facilitate our collaboration, as well as for giving me an excuse to give a *Big Lebowsky*-themed talk at MPI-SWS.

One of the most exciting and rewarding parts of my dissertation work has been pursuing the connection between LVars and distributed data consistency, the topic of Chapter 5. It can be a challenge for people coming from different subfields to find common ground, and I’m indebted to the many distributed systems researchers and practitioners who have met me much more than halfway, particularly the BOOM

group at Berkeley and the Riak folks at Basho. Speaking at RICON 2013 about LVars was one of the highlights of my Ph.D. experience; thanks to everyone who made it so.

Katerina Barone-Adesi, José Valim, and Zach Allaun — not coincidentally, all members of the Recurse Center community, which is the best programming community in the world — gave feedback on drafts of this dissertation and helped me improve the presentation. Thanks to all of them.

Jason Reed contributed the wonderful spot illustrations that appear throughout. My hope is that they'll make the pages a bit more inviting and offer some comic relief to the frustrated student. They've certainly done that for me.

Finally, this dissertation is dedicated to my amazing husband, Alex Rudnick, without whose love, support, advice, and encouragement I would never have *started* grad school, let alone finished.

For Alex, who said,
“You’re gonna destroy grad school.”

Lindsey Kuper

LATTICE-BASED DATA STRUCTURES FOR DETERMINISTIC PARALLEL AND DISTRIBUTED
PROGRAMMING

Deterministic-by-construction parallel programming models guarantee that programs have the same observable behavior on every run, promising freedom from bugs caused by schedule nondeterminism. To make that guarantee, though, they must sharply restrict sharing of state between parallel tasks, usually either by disallowing sharing entirely or by restricting it to one type of data structure, such as single-assignment locations.

I show that *lattice-based* data structures, or *LVars*, are the foundation for a guaranteed-deterministic parallel programming model that allows a more general form of sharing. LVars allow multiple assignments that are inflationary with respect to a given lattice. They ensure determinism by allowing only inflationary writes and “threshold” reads that block until a lower bound is reached. After presenting the basic LVars model, I extend it to support *event handlers*, which enable an event-driven programming style, and non-blocking “freezing” reads, resulting in a *quasi-deterministic* model in which programs behave deterministically modulo exceptions.

I demonstrate the viability of the LVars model with *LVish*, a Haskell library that provides a collection of lattice-based data structures, a work-stealing scheduler, and a monad in which LVar computations run. LVish leverages Haskell’s type system to index such computations with *effect levels* to ensure that only certain LVar effects can occur, hence statically enforcing determinism or quasi-determinism. I present two case studies of parallelizing existing programs using LVish: a k -CFA control flow analysis, and a bioinformatics application for comparing phylogenetic trees.

Finally, I show how LVar-style threshold reads apply to the setting of *convergent replicated data types* (CvRDTs), which specify the behavior of eventually consistent replicated objects in a distributed system. I extend the CvRDT model to support deterministic, strongly consistent *threshold queries*. The technique

generalizes to any lattice, and hence any CvRDT, and allows deterministic observations to be made of replicated objects before the replicas' states converge.

Contents

Chapter 1. Introduction	1
1.1. The deterministic-by-construction parallel programming landscape	2
1.2. Monotonic data structures as a basis for deterministic parallelism	4
1.3. Quasi-deterministic and event-driven programming with LVars	5
1.4. The LVish library	6
1.5. Deterministic threshold queries of distributed data structures	8
1.6. Thesis statement, and organization of the rest of this dissertation	8
1.7. Previously published work	10
Chapter 2. LVars: lattice-based data structures for deterministic parallelism	11
2.1. Motivating example: a parallel, pipelined graph computation	14
2.2. LVars by example	17
2.3. Lattices, stores, and determinism	21
2.4. λ_{LVar} : syntax and semantics	27
2.5. Proof of determinism for λ_{LVar}	31
2.6. Generalizing the put and get operations	42
Chapter 3. Quasi-deterministic and event-driven programming with LVars	50
3.1. LVish, informally	52
3.2. LVish, formally	57
3.3. Proof of quasi-determinism for λ_{LVish}	68
Chapter 4. The LVish library	80
4.1. The big picture	81
4.2. The LVish library interface for application writers	82

4.3. Par-monad transformers and disjoint parallel update	93
4.4. Case study: parallelizing k -CFA with LVish	97
4.5. Case study: parallelizing PhyBin with LVish	104
Chapter 5. Deterministic threshold queries of distributed data structures	108
5.1. Background: CvRDTs and eventual consistency	111
5.2. Adding threshold queries to CvRDTs	113
5.3. Determinism of threshold queries	118
5.4. Discussion: reasoning about per-query consistency choices	119
Chapter 6. Related work	121
6.1. Deterministic Parallel Java (DPJ)	122
6.2. FlowPools	122
6.3. Bloom and Bloom ^L	123
6.4. Concurrent Revisions	124
6.5. Frame properties and separation logics	124
Chapter 7. Summary and future work	126
7.1. Another look at the deterministic parallel landscape	126
7.2. Distributed programming and the future of LVars and LVish	128
Bibliography	130
Appendix A. Proofs	133
A.1. Proof of Lemma 2.1	133
A.2. Proof of Lemma 2.2	136
A.3. Proof of Lemma 2.3	138
A.4. Proof of Lemma 2.4	140
A.5. Proof of Lemma 2.5	141
A.6. Proof of Lemma 2.6	146
A.7. Proof of Lemma 2.8	149
A.8. Proof of Lemma 2.9	160

A.9. Proof of Lemma 2.10	161
A.10. Proof of Lemma 3.2	162
A.11. Proof of Lemma 3.3	171
A.12. Proof of Lemma 3.4	176
A.13. Proof of Lemma 3.5	180
A.14. Proof of Lemma 3.6	183
A.15. Proof of Lemma 3.7	186
A.16. Proof of Lemma 3.8	191
A.17. Proof of Lemma 3.10	196
A.18. Proof of Lemma 3.11	224
A.19. Proof of Lemma 3.12	226
A.20. Proof of Theorem 5.2	229
Appendix B. A PLT Redex Model of λ_{LVish}	233
Curriculum Vitae	

CHAPTER 1

Introduction

Parallel programming—that is, writing programs that can take advantage of parallel hardware to go faster—is notoriously difficult. A fundamental reason for this difficulty is that programs can yield inconsistent results, or even crash, due to unpredictable interactions between parallel tasks.

Deterministic-by-construction parallel programming models, though, offer the promise of freedom from subtle, hard-to-reproduce nondeterministic bugs in parallel code. Although there are many ways to construct individual deterministic programs and verify their determinism, deterministic-by-construction programming models provide a *language-level* guarantee of determinism that holds for all programs written using the model.

A deterministic program is one that has the same *observable behavior* every time it is run. How do we define what is observable about a program’s behavior? Certainly, we do *not* wish to preserve behaviors such as running time across multiple runs—ideally, a deterministic parallel program will run faster when more parallel resources are available. Moreover, we do not want to count scheduling behavior as observable—in fact, we want to specifically *allow* tasks to be scheduled dynamically and unpredictably, without allowing such *schedule nondeterminism* to affect the observable behavior of a program. Therefore, in this dissertation I will define the observable behavior of a program to be *the value to which the program evaluates*.

This definition of observable behavior ignores side effects other than *state*. But even with such a limited notion of what is observable, schedule nondeterminism can affect the outcome of a program. For instance, if a computation writes 3 to a shared location while another computation writes 4, then a subsequent third computation that reads and returns the location’s contents will nondeterministically return 3 or 4, depending on the order in which the first two computations ran. Therefore, if a parallel

programming model is to guarantee determinism by construction, it must necessarily limit sharing of mutable state between parallel tasks in some way.

1.1. The deterministic-by-construction parallel programming landscape

There is long-standing work on deterministic-by-construction parallel programming models that limit sharing of state between tasks. The possibilities include:

- *No-shared-state parallelism.* One classic approach to guaranteeing determinism in a parallel programming model is to allow *no* shared mutable state between tasks, forcing tasks to produce values independently. An example of no-shared-state parallelism is pure functional programming with function-level task parallelism, or *futures*—for instance, in Haskell programs that use the `par` and `pseq` combinators [35]. The key characteristic of this style of programming is lack of side effects: because programs do not have side effects, expressions can evaluate simultaneously without affecting the eventual value of the program. Also belonging in this category are parallel programming models based on *pure data parallelism*, such as Data Parallel Haskell [44, 14] or the River Trail API for JavaScript [25], each of which extend existing languages with *parallel array* data types and (observably) pure operations on them.
- *Data-flow parallelism.* In *Kahn process networks* (KPNs) [29], as well as in the more restricted *synchronous data flow* systems [32], a network of independent “computing stations” communicate with each other through first-in, first-out (FIFO) queues, or *channels*. In this model, each computing station is a task, and channels are the only means of sharing state between tasks. Furthermore, reading data from a channel is a *blocking* operation: once an attempt to read has started, a computing station cannot do anything else until the data to be read is available. Each station computes a sequential, monotonic function from the *history* of its input channels (*i.e.*, the input it has received so far) to the history of its output channels (the output it has produced so far). KPNs are the basis for deterministic stream-processing languages such as StreamIt [24].
- *Single-assignment parallelism.* In parallel *single-assignment* languages, “full/empty” bits are associated with memory locations so that they may be written to at most once. Single-assignment

1. INTRODUCTION

locations with blocking read semantics are known as *IVars* [3] and are a well-established mechanism for enforcing determinism in parallel settings: they have appeared in Concurrent ML as *SyncVars* [46]; in the Intel Concurrent Collections (abbreviated “CnC”) system [11]; and have even been implemented in hardware in Cray MTA machines [5]. Although most of these uses of *IVars* incorporate them into already-nondeterministic programming environments, the *monad-par* Haskell library [36] uses *IVars* in a deterministic-by-construction setting, allowing user-created threads to communicate through *IVars* without requiring the *IO* monad. Rather, operations that read and write *IVars* must run inside a *Par* monad, thus encapsulating them inside otherwise pure programs, and hence a program in which the only effects are *Par* effects is guaranteed to be deterministic.

- *Imperative disjoint parallelism*. Finally, yet another approach to guaranteeing determinism is to ensure that the state accessed by concurrent threads is *disjoint*. Sophisticated permissions systems and type systems can make it possible for imperative programs to mutate state in parallel, while guaranteeing that the same state is not accessed simultaneously by multiple threads. I will refer to this style of programming as *imperative disjoint parallelism*, with Deterministic Parallel Java (DPJ) [8, 7] as a prominent example.

The four parallel programming models listed above—no-shared-state parallelism, data-flow parallelism, single-assignment parallelism, and imperative disjoint parallelism—all seem to embody rather different mechanisms for exposing parallelism and for ensuring determinism.



If we view these different programming models as a toolkit of unrelated choices, though, it is not clear how to proceed when we want to implement an application with multiple parallelizable components that are best suited to different programming models. For example, suppose we have an application in which we want to exploit data-flow pipeline parallelism via FIFO queues, but we also want to mutate disjoint slices of arrays. It is not obvious how to compose two programming models that each only allow communication through a single type of shared data structure—and if we do manage to compose them, it is not obvious whether the determinism guarantee of the individual

models is preserved by their composition. Therefore, we seek a general, broadly-applicable model for deterministic parallel programming that is not tied to a particular data structure.

1.2. Monotonic data structures as a basis for deterministic parallelism

In KPNs and other data-flow models, communication takes place over blocking FIFO queues with ever-increasing *channel histories*, while in LVar-based programming models such as CnC and monad-par, a shared data store of blocking single-assignment memory locations grows monotonically. Hence *monotonic data structures*—data structures to which information can only be added and never removed, and for which the timing of updates is not observable—emerge as a common theme of guaranteed-deterministic programming models.

In this dissertation, I show that *lattice-based* data structures, or *LVars*, offer a general approach to deterministic parallel programming that takes monotonicity as a starting point. The states an LVar can take on are elements of a given *lattice*. This lattice determines the semantics of the put and get operations that comprise the interface to LVars (which I will explain in detail in Chapter 2):

- The put operation can only make the state of an LVar “grow” with respect to the lattice, because it updates the LVar to the *least upper bound* of the current state and the new state. For example, on LVars of collection type, such as sets, the put operation typically inserts an element.
- The get operation allows only limited observations of the contents of an LVar. It requires the user to specify a *threshold set* of minimum values that can be read from the LVar, where every two elements in the threshold set must have the lattice’s greatest element \top as their least upper bound. A call to get blocks until the LVar in question reaches a (unique) value in the threshold set, then unblocks and returns that value.

Together, least-upper-bound writes via put and threshold reads via get yield a programming model that is deterministic by construction. That is, a program in which put and get operations on LVars are the only side effects will have the same observable result every time it is run, regardless of parallel

execution and schedule nondeterminism. As we will see in Chapter 2, no-shared-state parallelism, data-flow parallelism and single-assignment parallelism are all subsumed by the LVars programming model, and as we will see in Section 4.3, imperative disjoint parallel updates are compatible with LVars as well.

Furthermore, as I show in Section 2.6, we can generalize the behavior of the `put` and `get` operations while retaining determinism: we can generalize from the least-upper-bound `put` operation to a set of arbitrary *update operations* that are not necessarily idempotent (but are still inflationary and commutative), and we can generalize the `get` operation to allow a more general form of threshold reads. Generalizing from `put` to arbitrary inflationary and commutative updates turns out to be a particularly useful extension to the LVars model; I formally extend the model to support these update operations in Chapter 3, and in Chapter 4 I discuss how arbitrary update operations are useful in practice.

1.3. Quasi-deterministic and event-driven programming with LVars

The LVars model described above guarantees determinism and supports an unlimited variety of shared data structures: anything viewable as a lattice. However, it is not as general-purpose as one might hope. Consider, for instance, an algorithm for unordered graph traversal. A typical implementation involves a monotonically growing set of “seen nodes”; neighbors of seen nodes are fed back into the set until it reaches a fixed point. Such fixpoint computations are ubiquitous, and would seem to be a perfect match for the LVars model due to their use of monotonicity. But they are not expressible using the threshold `get` and least-upper-bound `put` operations, nor even with the more general alternatives to `get` and `put` mentioned above.

The problem is that these computations rely on *negative* information about a monotonic data structure, *i.e.*, on the *absence* of certain writes to the data structure. In a graph traversal, for example, neighboring nodes should only be explored if the current node is *not yet* in the set; a fixpoint is reached only if no new neighbors are found; and, of course, at the end of the computation it must be possible to learn exactly which nodes were reachable (which entails learning that certain nodes were not). In Chapter 3, I describe two extensions to the basic LVars model that make such computations possible:

1. INTRODUCTION

- First, I add the ability to attach *event handlers* to an LVar. When an event handler has been registered with an LVar, it causes a callback function to run, asynchronously, whenever events arrive (in the form of monotonic updates to the LVar). Crucially, it is possible to check for *quiescence* of a group of handlers, discovering that no callbacks are currently enabled—a transient, negative property. Since quiescence means that there are no further changes to respond to, it can be used to tell that a fixpoint has been reached.
- Second, I extend the model with a primitive operation `freeze` for *freezing* an LVar, which allows its contents to be read immediately and exactly, rather than the blocking threshold read that `get` allows. The `freeze` primitive imposes the following trade-off: once an LVar has been frozen, any further writes that would change its value instead raise an exception; on the other hand, it becomes possible to discover the exact value of the LVar, learning both positive and negative information about it, without blocking. Therefore, LVar programs that use `freeze` are *not* guaranteed to be deterministic, because they could nondeterministically raise an exception depending on how writes and `freeze` operations are scheduled. However, such programs satisfy *quasi-determinism*: all executions that produce a final value (instead of raising an exception) produce the *same* final value.

Writes to an LVar could cause more events to occur after a group of handlers associated with that LVar has quiesced, and those events could trigger more invocations of callback functions. However, since the contents of the LVar can only be read through `get` or `freeze` operations anyway, early quiescence poses no risk to determinism or quasi-determinism, respectively. In fact, freezing and quiescence work particularly well together because freezing provides a mechanism by which the programmer can safely “place a bet” that all writes to an LVar have completed. Hence freezing and handlers make it possible to implement fixpoint computations like the graph traversal described above. Moreover, if we can ensure that a freeze does indeed happen after all writes to the LVar in question have completed, then we can ensure that the entire computation is deterministic, and it is possible to enforce this “freeze-last” idiom at the implementation level, as I discuss below (and, in more detail, in Section 4.2.5).

1.4. The LVish library

To demonstrate the practicality of the LVars programming model, in Chapter 4 I will describe *LVish*, a Haskell library¹ for deterministic and quasi-deterministic programming with LVars.

LVish provides a `Par` monad for encapsulating parallel computations.² A `Par` computation can create lightweight, library-level threads that are dynamically scheduled by a custom work-stealing scheduler. LVar operations run inside the `Par` monad, which is indexed by an *effect level*, allowing fine-grained specification of the effects that a given computation is allowed to perform. For instance, since `freeze` introduces quasi-determinism, a computation indexed with a deterministic effect level is not allowed to use `freeze`. Thus, the *type* of an LVish computation reflects its determinism or quasi-determinism guarantee. Furthermore, if a `freeze` is guaranteed to be the *last* effect that occurs in a computation, then it is impossible for that `freeze` to race with a write, ruling out the possibility of a runtime write-after-freeze exception. LVish exposes a `runParThenFreeze` operation that captures this “freeze-last” idiom and has a deterministic effect level.

LVish also provides a variety of lattice-based data structures (e.g., sets, maps, arrays) that support concurrent insertion, but not deletion, during `Par` computations. Users may also implement their own lattice-based data structures, and LVish provides tools to facilitate the definition of user-defined LVars. I will describe the proof obligations for LVar implementors and give examples of applications that use user-defined LVars as well as those that the library provides.

Finally, Chapter 4 illustrates LVish through two case studies, drawn from my collaborators’ and my experience using the LVish library, both of which make use of handlers and freezing:

- First, I describe using LVish to parallelize a control flow analysis (*k*-CFA) algorithm. The goal of *k*-CFA is to compute the flow of values to expressions in a program. The *k*-CFA algorithm proceeds in two phases: first, it explores a graph of *abstract states* of the program; then, it summarizes the results of the first phase. Using LVish, these two phases can be pipelined; moreover, the original

¹Available at <http://hackage.haskell.org/package/lvish>.

²The `Par` monad exposed by LVish generalizes the original `Par` monad exposed by the *monad-par* library (<http://hackage.haskell.org/package/monad-par>, described by Marlow *et al.* [36]), which allows determinism-preserving communication between threads, but only through IVars, rather than LVars.

graph exploration phase can be internally parallelized. I contrast our LVish implementation with the original sequential implementation from which it was ported and give performance results.

- Second, I describe using LVish to parallelize *PhyBin* [40], a bioinformatics application for comparing sets of phylogenetic trees that relies heavily on a parallel tree-edit distance algorithm [52]. The *PhyBin* application crucially relies on the aforementioned ability to perform arbitrary inflationary and commutative (but not idempotent) updates on LVars (in contrast to the idempotent put operation). We show that the performance of the parallelized *PhyBin* application compares favorably to existing widely-used software packages for analyzing collections of phylogenetic trees.

1.5. Deterministic threshold queries of distributed data structures

The LVars model is closely related to the concept of *conflict-free replicated data types* (CRDTs) [49] for enforcing *eventual consistency* [56] of replicated objects in a distributed system. In particular, *state-based* or *convergent* replicated data types, abbreviated as *CvRDTs* [49, 48], leverage the mathematical properties of lattices to guarantee that all replicas of an object (for instance, in a distributed database) eventually agree.

Although *CvRDTs* are provably eventually consistent, queries of *CvRDTs* (unlike threshold reads of LVars) nevertheless allow inconsistent intermediate states of replicas to be observed. That is, if two replicas of a *CvRDT* object are updated independently, reads of those replicas may disagree until a (least-upper-bound) *merge* operation takes place.

Taking inspiration from LVar-style threshold reads, in Chapter 5 I show how to extend *CvRDTs* to support deterministic, *strongly consistent* queries using a mechanism called *threshold queries* (or, seen from another angle, I show how to port threshold reads from a shared-memory setting to a distributed one). The threshold query technique generalizes to any lattice, and hence any *CvRDT*, and allows deterministic observations to be made of replicated objects before the replicas' states have converged. This work has practical relevance since, while real distributed database applications call for a combination of eventually consistent and strongly consistent queries, *CvRDTs* only support the former. Threshold queries

extend the CvRDT model to support both kinds of queries within a single, lattice-based reasoning framework. Furthermore, since threshold queries behave deterministically regardless of whether all replicas agree, they suggest a way to save on synchronization costs: existing operations that require all replicas to agree could be done with threshold queries instead, and retain behavior that is *observably* strongly consistent while avoiding unnecessary synchronization.



1.6. Thesis statement, and organization of the rest of this dissertation

With the above background, I can state my thesis:

Lattice-based data structures are a general and practical unifying abstraction for deterministic and quasi-deterministic parallel and distributed programming.

The rest of this dissertation supports my thesis as follows:

- *Lattice-based data structures*: In Chapter 2, I formally define LVars and use them to define λ_{LVar} , a call-by-value parallel calculus with a store of LVars that support least-upper-bound put and threshold get operations. In Chapter 3, I extend λ_{LVar} to add support for arbitrary update operations, event handlers, and the `freeze` operation, calling the resulting language λ_{LVish} . Appendix B contains runnable versions of λ_{LVar} and λ_{LVish} implemented using the PLT Redex semantics engineering system [19] for interactive experimentation.
- *general*: In Chapter 2, I show how previously existing deterministic parallel programming models (single-assignment languages, Kahn process networks) are subsumed by the lattice-generic LVars model. Additionally, I show how to generalize the put and get operations on LVars while preserving their determinism.
- *deterministic*: In Chapter 2, I show that the basic LVars model guarantees determinism by giving a proof of determinism for the λ_{LVar} language with put and get.
- *quasi-deterministic*: In Chapter 3, I define quasi-determinism and give a proof of quasi-determinism for λ_{LVish} , which adds arbitrary update operations, the `freeze` operation, and event handlers to the λ_{LVar} language of Chapter 2.

1. INTRODUCTION

- *practical and parallel*: In Chapter 4, I describe the interface and implementation of the LVish Haskell library, which is based on the LVars programming model, and demonstrate how it is used for practical programming with the two case studies described above, including performance results on parallel hardware.
- *distributed programming*: In Chapter 5, I show how LVar-style threshold reads apply to the setting of distributed, replicated data structures. In particular, I extend convergent replicated data types (CvRDTs) to support strongly consistent threshold queries, which take advantage of the existing lattice structure of CvRDTs and allow deterministic observations to be made of their contents without requiring all replicas to agree.

1.7. Previously published work

The material in this dissertation is based on research done jointly with several collaborators, some of which appears in the following previously published papers:

- Lindsey Kuper and Ryan R. Newton. 2013. LVars: lattice-based data structures for deterministic parallelism. In *Proceedings of the 2nd ACM SIGPLAN Workshop on Functional High-Performance Computing* (FHPC '13).
- Lindsey Kuper, Aaron Turon, Neelakantan R. Krishnaswami, and Ryan R. Newton. 2014. Freeze after writing: quasi-deterministic parallel programming with LVars. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (POPL '14).
- Lindsey Kuper and Ryan R. Newton. 2014. Joining forces: toward a unified account of LVars and convergent replicated data types. In the *5th Workshop on Determinism and Correctness in Parallel Programming* (WoDet '14).
- Lindsey Kuper, Aaron Todd, Sam Tobin-Hochstadt, and Ryan R. Newton. 2014. Taming the parallel effect zoo: extensible deterministic parallelism with LVish. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation* (PLDI '14).

CHAPTER 2

LVars: lattice-based data structures for deterministic parallelism


Programs written using a *deterministic-by-construction* model of parallel computation are guaranteed to always produce the same observable results, offering programmers freedom from subtle, hard-to-reproduce nondeterministic bugs. While a number of popular languages and language extensions (*e.g.*, Cilk [22]) *encourage* deterministic parallel programming, few of them guarantee determinism at the language level—that is, for *all* programs that can be written using the model.

Of the options available for parallel programming with a language-level determinism guarantee, perhaps the most mature and broadly available choice is pure functional programming with function-level task parallelism, or *futures*. For example, Haskell programs using futures by means of the `par` and `pseq` combinators can provide real speedups on practical programs while guaranteeing determinism [35].¹ Yet pure programming with futures is not ideal for all problems. Consider a *producer/consumer* computation in which producers and consumers can be scheduled onto separate processors, each able to keep their working sets in cache. Such a scenario enables *pipeline parallelism* and is common, for instance, in stream processing. But a clear separation of producers and consumers is difficult with futures, because whenever a consumer forces a future, if the future is not yet available, the consumer immediately switches roles to begin computing the value (as explored by Marlow *et al.* [36]).

Since pure programming with futures is a poor fit for producer/consumer computations, one might then turn to *stateful* deterministic parallel models. Shared state between computations allows the possibility for race conditions that introduce nondeterminism, so any parallel programming model that hopes to guarantee determinism must do something to tame sharing—that is, to restrict access to mutable state shared among concurrent computations. Systems such as Deterministic Parallel Java [8, 7], for instance,

¹When programming with `par` and `pseq`, a language-level determinism guarantee obtains if user programs are written in the *Safe Haskell* [53] subset of Haskell (which is implemented in GHC Haskell by means of the `Safe Haskell` language pragma), and if they do not use the `IO` monad.

accomplish this by ensuring that the state accessed by concurrent threads is *disjoint*. Alternatively, a programming model might allow *data* to be shared, but limit the *operations* that can be performed on it to only those operations that commute with one another and thus can tolerate nondeterministic thread interleavings. In such a setting, although the order in which side-effecting operations occur can differ on multiple runs, a program will always produce the same observable result.²



In Kahn process networks (KPNs) [29] and other *data-flow parallel* models—which are the basis for deterministic stream-processing languages such as StreamIt [24]—communication among processes takes place over blocking FIFO queues with ever-increasing *channel histories*. Meanwhile, in *single-assignment* [55] or *IVar-based* [3] programming models, such as the Intel Concurrent Collec-

tions system (CnC) [11] and the *monad-par* Haskell library [36], a shared data store of blocking single-assignment memory locations grows monotonically. Hence both programming models rely on *monotonic data structures*: data structures to which information can only be added and never removed, and for which the timing of updates is not observable.

Because state modifications that only add information and never destroy it can be structured to commute with one another and thereby avoid race conditions, it stands to reason that diverse deterministic parallel programming models would leverage the principle of monotonicity. Yet systems like StreamIt, CnC, and monad-par emerge independently, without recognition of their common basis. Moreover, since each one of these programming models is based on a single type of shared data structure, they lack *generality*: IVars or FIFO streams alone cannot support all producer/consumer applications, as I discuss in Section 2.1.

Instead of limiting ourselves to a single type of shared data structure, though, we can take the more general notion of monotonic data structures as the basis for a new deterministic parallel programming model. In this chapter, I show how to generalize IVars to *LVars*, thus named because the states an LVar

²There are many ways to define what is observable about a program. As noted in Chapter 1, I define the observable behavior of a program to be the value to which it evaluates.

can take on are elements of a given *lattice*.³ This lattice determines the semantics of the put and get operations that comprise the interface to LVars (which I will explain in detail in the sections that follow): the put operation takes the least upper bound of the current state and the new state with respect to the lattice, and the get operation performs a *threshold read* that blocks until a lower bound in the lattice is reached.

Section 2.2 introduces the concept of LVars through a series of small code examples. Then, in Sections 2.3 and 2.4 I formally define λ_{LVar} , a deterministic parallel calculus with shared state, based on the call-by-value λ -calculus. The λ_{LVar} language is general enough to subsume existing deterministic parallel languages because it is parameterized by the choice of lattice. For example, a lattice of channel histories with a prefix ordering allows LVars to represent FIFO channels that implement a Kahn process network, whereas instantiating λ_{LVar} with a lattice with one “empty” state and multiple “full” states (where $\forall i. \text{empty} < \text{full}_i$) results in a parallel single-assignment language. Different instantiations of the lattice result in a family of deterministic parallel languages.

Because lattices are composable, any number of diverse monotonic data structures can be used together safely. Moreover, as long as a data structure presents the LVar interface, it is fine to wrap an existing, optimized concurrent data structure implementation; we need not rewrite the world’s data structures to leverage the λ_{LVar} determinism result.

The main technical result of this chapter is a proof of determinism for λ_{LVar} (Section 2.5). The key lemma, Independence (Section 2.5.5), gives a kind of *frame property* that captures the commutative effects of LVar computations. Such a property would *not* hold in a typical language with shared mutable state, but holds in the setting of λ_{LVar} because of the semantics of put and get.

Finally, in Section 2.6, I consider some alternative semantics for the put and get operations that generalize their behavior while retaining the determinism of the original semantics: I generalize the put operation from least-upper-bound writes to inflationary, commutative writes, and I generalize the get operation to allow a more general form of threshold reads.

³This “lattice” need only be a *bounded join-semilattice* augmented with a greatest element \top , in which every two elements have a least upper bound but not necessarily a greatest lower bound; see Section 2.3.1. For brevity, I use the term “lattice” in place of “bounded join-semilattice with a designated greatest element” throughout this dissertation.

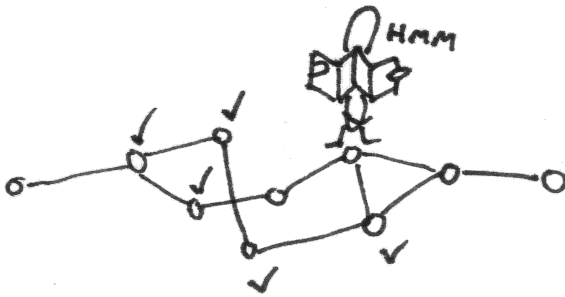
2.1. Motivating example: a parallel, pipelined graph computation

What applications motivate going beyond IIVars and FIFO streams? Consider applications in which independent subcomputations contribute results to shared mutable data structures. Hindley-Milner type inference is one example: in a parallel type-inference algorithm, each type variable monotonically acquires information through unification. Likewise, in control-flow analysis, the *set* of locations to which a variable refers monotonically *shrinks*. In logic programming, a parallel implementation of conjunction might asynchronously add information to a logic variable from different threads.

To illustrate the issues that arise in computations of this nature, consider a specific problem, drawn from the domain of *graph algorithms*, where issues of ordering create a tension between parallelism and determinism:

In a directed graph, find the connected component containing a vertex v , and compute a (possibly expensive) function f over all vertices in that component, making the set of results available asynchronously to other computations.

For example, in a directed graph representing user profiles on a social network and the connections between them, where v represents a particular user’s profile, we might wish to find all (or the first k degrees of) profiles connected to v , then map a function f over each profile in that set in parallel.



This is a challenge problem for deterministic-by-construction parallel programming models. Existing parallel solutions (such as, for instance, the parallel breadth-first graph traversal implementation from the Parallel Boost Graph Library [1]) often traverse the connected component in a non-deterministic order—even though the outcome of

the program, that is, the final connected component, is deterministic. Neither IIVars nor streams provide a way to orchestrate this traversal. For example, IIVars cannot accumulate sets of visited nodes, nor can they be used as “mark bits” on visited nodes, since they can only be written once and not tested

```

-- Traverse each level of the graph `g` in parallel, maintaining at
-- each recursive step a set of nodes that have been seen and a set of
-- nodes left to process. `nbrs g n` is the neighbor nodes of node
-- `n` in graph `g`.
bf_traverse :: Graph -> Set NodeLabel -> Set NodeLabel -> Set NodeLabel
bf_traverse g seen nu =
  if nu == empty
  then seen
  else let seen' = union seen nu
        allNbr = parFold union (parMap (nbrs g) nu)
        nu'     = difference allNbr seen'
        in bf_traverse g seen' nu'

-- Find the connected component containing the vertex `profile0`, and
-- map `analyze` over all nodes in it:
connected_component = bf_traverse profiles empty (singleton profile0)
result = parMap analyze connected_component

```

Listing 2.1. A purely functional Haskell program that finds the connected component of the profiles graph that is reachable from the node `profile0`, then maps the `analyze` function over the nodes found. Although component discovery proceeds in parallel, we cannot call `analyze` on any of the nodes in the connected component until they have all been discovered.

for emptiness. Streams, on the other hand, impose an excessively strict ordering for computing the unordered *set* of vertex labels in a connected component. Before turning to an LVar-based approach, though, let us consider whether a purely functional (and therefore deterministic by construction) program can meet the specification.

Listing 2.1 gives a Haskell implementation of a *level-synchronized* breadth-first graph traversal that finds the connected component reachable from a starting vertex. Nodes at distance one from the starting vertex are discovered—and set-unioned into the connected component—before nodes of distance two are considered. Level-synchronization necessarily sacrifices some parallelism: parallel tasks cannot

continue discovering nodes in the component before synchronizing with all other tasks at a given distance from the start. Nevertheless, level-synchronization is a popular strategy for implementing parallel breadth-first graph traversal. (In fact, the aforementioned implementation from the Parallel Boost Graph Library [1] uses level-synchronization.)

Unfortunately, the code given in Listing 2.1 does not quite implement the problem specification given above. Even though connected-component discovery is parallel, members of the output set do not become available to analyze until component discovery is *finished*, limiting parallelism: the first nodes in the connected component to be discovered must go un-analyzed until all of the nodes in the connected component have been traversed. We could manually push the `analyze` invocation inside the `bf_traverse` function, allowing the `analyze` computation to start sooner, but doing so just passes the problem to the downstream consumer, unless we are able to perform a heroic whole-program fusion.

If `bf_traverse` returned a list, lazy evaluation could make it possible to *stream* results to consumers incrementally. But since it instead returns a *set*, such pipelining is not generally possible: consuming the results early would create a proof obligation that the determinism of the consumer does not depend on the order in which results emerge from the producer.⁴

A compromise would be for `bf_traverse` to return a list of “level-sets”: distance one nodes, distance two nodes, and so on. Thus level-one results could be consumed before level-two results are ready. Still, at the granularity of the individual level-sets, the problem would remain: within each level-set, one would not be able to launch all instances of `analyze` and asynchronously use those results that finished first. Moreover, we would still have to contend with the problem of separating the scheduling of producers and consumers that pure programming with futures presents [36].

⁴As intuition for this idea, consider that purely functional set data structures, such as Haskell’s `Data.Set`, are typically represented with balanced trees. Unlike with lists, the structure of the tree is not known until all elements are present.

2.1.2. An LVar-based solution. Consider a version of `bf_traverse` written using a programming model with limited effects that allows *any* data structure to be shared among tasks, including sets and graphs, so long as that data structure grows monotonically. Consumers of the data structure may execute as soon as data is available, enabling pipelining, but may only observe irrevocable, monotonic properties of it. This is possible with a programming model based on LVars. In the rest of this chapter, I will formally introduce LVars and the λ_{LVar} language and give a proof of determinism for λ_{LVar} . Then, in Chapter 3, I will extend the basic LVars model with additional features that make it easier to implement parallel graph traversal algorithms and other fixpoint computations, and I will return to `bf_traverse` and show how to implement a version of it using LVars that makes pipelining possible and truly satisfies the above specification.

2.2. LVars by example

IVars [3, 41, 36] are a well-known mechanism for deterministic parallel programming.⁵ An IVar is a *single-assignment* variable [55] with a blocking read semantics: an attempt to read an empty IVar will block until the IVar has been filled with a value. The “I” in IVar stands for “immutable”: an IVar can only be written to once, and thereafter its contents cannot change. LVars are a generalization of IVars: an LVar allows multiple writes, but only so long as those writes are monotonically increasing with respect to a given lattice of states.

Consider a program in which two parallel computations write to an LVar *lv*, with one thread writing the value 2 and the other writing 3:

(Example 2.1)

```

let par _ = put lv 3
      _ = put lv 2
in get lv

```

Here, `put` and `get` write and read LVars, respectively, and the expression

```
let par  $x_1 = e_1$ ;  $x_2 = e_2$ ; ... in  $e_3$ 
```

⁵IVars are so named because they are a special case of *I-structures* [3]—namely, those with only one cell.

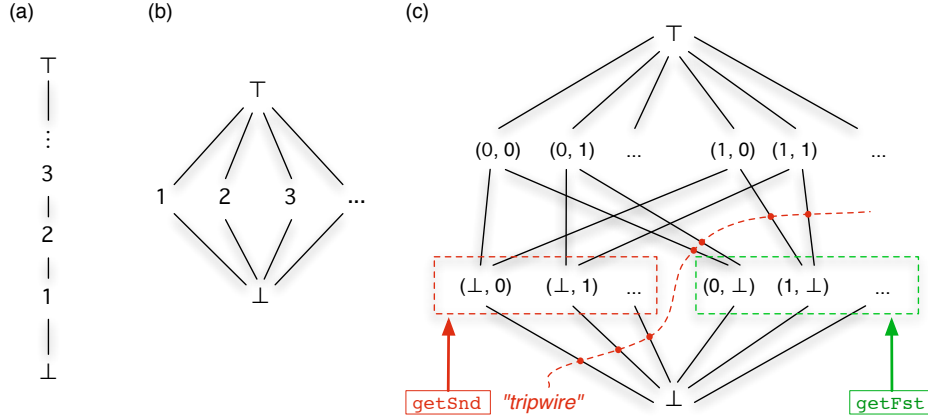


Figure 2.1. Example LVar lattices: (a) positive integers ordered by \leq ; (b) IVar containing a positive integer; (c) pair of natural-number-valued IVars, annotated with example threshold sets that would correspond to a blocking read of the first or second element of the pair. Any state transition crossing the “tripwire” for `getSnd` causes it to unblock and return a result.

has *fork-join* semantics: it launches concurrent subcomputations e_1, e_2, \dots whose executions arbitrarily interleave, but must all complete before the expression e_3 runs.

The `put` operation is defined in terms of the specified lattice of LVar states; it updates the LVar to the *least upper bound* (lub) of its current state and the new state being written. If lv ’s lattice is the \leq ordering on positive integers, as shown in Figure 2.1(a), then lv ’s state will always be $\max(3, 2) = 3$ by the time `get lv` runs, since the lub of two positive integers n_1 and n_2 is $\max(n_1, n_2)$. Therefore Example 2.1 will deterministically evaluate to 3, regardless of the order in which the two `put` operations occur.

On the other hand, if lv ’s lattice is that shown in Figure 2.1(b), in which the lub of any two distinct positive integers is \top , then Example 2.1 will deterministically raise an exception, indicating that conflicting writes to lv have occurred. This exception is analogous to the “multiple put” error raised upon multiple writes to an IVar. Unlike with a traditional IVar, though, multiple writes of the *same* value (say, `put lv 3` and `put lv 3`) will *not* raise an exception, because the lub of any positive integer and itself is that integer—corresponding to the fact that multiple writes of the same value do not allow any nondeterminism to be observed.

2.2.1. Threshold reads. However, merely ensuring that writes to an LVar are monotonically increasing is not enough to guarantee that programs behave deterministically. Consider again the lattice of Figure 2.1(a) for lv , but suppose we change Example 2.1 to allow a direct, non-blocking read of the LVar to be interleaved with the two puts:

(Example 2.2)

```

let par _ = put lv 3
      _ = put lv 2
      x = read lv
in x

```

Since the two puts and the read can be scheduled in any order, Example 2.2 is nondeterministic: x might be either 2 or 3, depending on the order in which the LVar effects occur. Therefore, to maintain determinism, LVars allow only blocking, restricted get operations. Rather than observing the exact value of the LVar, the get operation can only observe that the LVar has reached one of a specified set of *lower bound* states. This set of lower bounds, which we provide as an extra argument to get, is called a *threshold set* because the values in it form a “threshold” that the state of the LVar must cross before the call to get is allowed to unblock and return. When the threshold has been reached, get unblocks and returns not the exact contents of the LVar, but instead, the (unique) element of the threshold set that has been reached or surpassed.

We can then change Example 2.2 to behave deterministically using get with a threshold set:

(Example 2.3)

```

let par _ = put lv 3
      _ = put lv 2
      x = get lv {3}
in x

```

Here we chose the singleton set $\{3\}$ as the threshold set. Since lv ’s value can only increase with time, we know that once it is at least 3, it will remain at or above 3 forever; therefore the program will deterministically evaluate to 3. Had we chosen $\{2\}$ as the threshold set, the program would deterministically evaluate to 2; had we chosen $\{4\}$, it would deterministically block forever.

As long as we only access LVars with `put` and `get`, we can arbitrarily share them between threads without introducing nondeterminism. That is, the `put` and `get` operations in a given program can happen in any order, without changing the value to which the program evaluates.

2.2.2. Incompatibility of threshold sets. While Example 2.3 is deterministic, the style of programming it illustrates is only useful for synchronization, not for communicating data: we must specify in advance the single answer we expect to be returned from the call to `get`. In general, though, threshold sets do not have to be singleton sets. For example, consider an LVar lv whose states form a lattice of *pairs* of natural-number-valued IVars; that is, lv is a pair (m, n) , where m and n both start as \perp and may each be updated once with a non- \perp value, which must be some natural number. This lattice is shown in Figure 2.1(c).

We can then define `getFst` and `getSnd` operations for reading from the first and second entries of lv :

$$\begin{aligned} \text{getFst } p &\triangleq \text{get } p \{ (m, \perp) \mid m \in \mathbb{N} \} \\ \text{getSnd } p &\triangleq \text{get } p \{ (\perp, n) \mid n \in \mathbb{N} \} \end{aligned}$$

This allows us to write programs like the following:

(Example 2.4)

```

let par _ = put lv (⊥, 4)
      _ = put lv (3, ⊥)
      x = getSnd lv
in x

```

In the call `getSnd lv`, the threshold set is $\{(\perp, 0), (\perp, 1), \dots\}$, an infinite set. There is no risk of non-determinism because the elements of the threshold set are *pairwise incompatible* with respect to lv 's lattice: informally, since the second entry of lv can only be written once, no more than one state from the set $\{(\perp, 0), (\perp, 1), \dots\}$ can ever be reached. (I formalize this incompatibility requirement in Section 2.3.3.)

In the case of Example 2.4, `getSnd lv` may unblock and return $(\perp, 4)$ any time after the second entry of `lv` has been written, regardless of whether the first entry has been written yet. It is therefore possible to use LVars to safely read parts of an incomplete data structure—say, an object that is in the process of being initialized by a constructor.

2.2.3. The model versus reality. The use of explicit threshold sets in the LVars model should be understood as a mathematical modeling technique, *not* an implementation approach or practical API. The core of the LVish library (which I will discuss in Chapter 4) provides unsafe operations to the authors of LVar data structure libraries, who can then export operations like `getFst` and `getSnd` as a safe interface for application writers, implicitly baking in the particular threshold sets that make sense for a given data structure without ever explicitly constructing them.

To put it another way, operations on a data structure exposed as an LVar must have the *semantic effect* of a lub for writes or a threshold for reads, but none of this need be visible to clients (or even written explicitly in the code). Any data structure API that provides such a semantics is guaranteed to provide deterministic concurrent communication.

2.3. Lattices, stores, and determinism

As a minimal substrate for LVars, I introduce λ_{LVar} , a parallel call-by-value λ -calculus extended with a *store* and with communication primitives `put` and `get` that operate on data in the store. The class of programs that I am interested in modeling with λ_{LVar} are those with explicit effectful operations on shared data structures, in which parallel subcomputations may communicate with each other via the `put` and `get` operations.

In λ_{LVar} , stores contain LVars. Whereas IVars are single-assignment variables—either empty or filled with an immutable value—an LVar may have an arbitrary number of states forming a set D , which is partially ordered by a relation \sqsubseteq . An LVar can take on any sequence of states from D , so long as that sequence respects the partial order—that is, so long as updates to the LVar (made via the `put` operation)

are *inflationary* with respect to \sqsubseteq . Moreover, the `get` operation allows only limited observations of the LVar’s state. In this section, I discuss how lattices and stores work in λ_{LVar} and explain how the semantics of `put` and `get` together enforce determinism in λ_{LVar} programs.

2.3.1. Lattices. The definition of λ_{LVar} is parameterized by D : to write concrete λ_{LVar} programs, we must specify the set of LVar states that we are interested in working with, and an ordering on those states. Therefore λ_{LVar} is actually a *family* of languages, rather than a single language.

Formally, D is a *bounded join-semilattice* augmented with a greatest element \top . That is, D has the following structure:

- D has a least element \perp , representing the initial “empty” state of an LVar.
- D has a greatest element \top , representing the “error” state that results from conflicting updates to an LVar.
- D comes equipped with a partial order \sqsubseteq , where $\perp \sqsubseteq d \sqsubseteq \top$ for all $d \in D$.
- Every pair of elements in D has a lub, written \sqcup . Intuitively, the existence of a lub for every two elements in D means that it is possible for two subcomputations to independently update an LVar, and then deterministically merge the results by taking the lub of the resulting two states.

We can specify all these components as a 4-tuple $(D, \sqsubseteq, \perp, \top)$ where D is a set, \sqsubseteq is a partial order on the elements of D , \perp is the least element of D according to \sqsubseteq , and \top is the greatest element. However, I use D as a shorthand for the entire 4-tuple $(D, \sqsubseteq, \perp, \top)$ when its meaning is clear from the context.

Virtually any data structure to which information is added gradually can be represented as a lattice, including pairs, arrays, trees, maps, and infinite streams. In the case of maps or sets, \sqcup could be defined as union; for pointer-based data structures like tries, \sqcup could allow for unification of partially-initialized structures.

The simplest example of a useful D is one that represents the states that a single-assignment variable (that is, an IVar) can take on. The states of a natural-number-valued IVar, for instance, are the elements of the lattice in Figure 2.1(b), that is,

$$D = (\{\top, \perp\} \cup \mathbb{N}, \sqsubseteq, \perp, \top),$$

where the partial order \sqsubseteq is defined by setting $\perp \sqsubseteq d \sqsubseteq \top$ and $d \sqsubseteq d$ for all $d \in D$. This is a lattice of height three and infinite width, where the naturals are arranged horizontally. After the initial write of some $n \in \mathbb{N}$, any further conflicting writes would push the state of the IVar to \top (an error). For instance, if one thread writes 2 and another writes 1 to an IVar (in arbitrary order), the second of the two writes would result in an error because $2 \sqcup 1 = \top$.

In the lattice of Figure 2.1(a), on the other hand, the \top state is unreachable, because the lub of any two writes is just the maximum of the two. If one thread writes 2 and another writes 1, the resulting state will be 2, since $2 \sqcup 1 = 2$. Here, the unreachability of \top models the fact that no conflicting updates can occur to the IVar.

2.3.2. Stores. During the evaluation of a λ_{LVar} program, a *store* S keeps track of the states of LVars. Each LVar is represented by a *binding* that maps a location l , drawn from a countable set Loc , to a *state*, which is some element $d \in D$. Although each LVar in a program has its own state, the states of all the LVars are drawn from the same lattice D .⁶



Definition 2.1 (store, λ_{LVar}). A *store* is either a finite partial mapping $S : Loc \xrightarrow{\text{fin}} (D - \{\top\})$, or the distinguished element \top_S .

⁶In practice, different LVars in a program might correspond to different lattices (and, in the LVish Haskell library that I will present in Chapter 4, they do). Multiple lattices can in principle be encoded using a sum construction, so this modeling choice is just to keep the presentation simple.

I use the notation $S[l \mapsto d]$ to denote extending S with a binding from l to d . If $l \in \text{dom}(S)$, then $S[l \mapsto d]$ denotes an update to the existing binding for l , rather than an extension. Another way to denote a store is by explicitly writing out all its bindings, using the notation $[l_1 \mapsto d_1, \dots, l_n \mapsto d_n]$. The set of states that a store can take on forms a lattice, just as D does, with the empty store \perp_S as its least element and \top_S as its greatest element. It is straightforward to lift the \sqsubseteq and \sqcup operations defined on elements of D to the level of stores:

Definition 2.2 (store ordering, λ_{LVar}). A store S is *less than or equal to* a store S' (written $S \sqsubseteq_S S'$) iff:

- $S' = \top_S$, or
- $\text{dom}(S) \subseteq \text{dom}(S')$ and for all $l \in \text{dom}(S)$, $S(l) \sqsubseteq S'(l)$.

Definition 2.3 (lub of stores, λ_{LVar}). The lub of stores S_1 and S_2 (written $S_1 \sqcup_S S_2$) is defined as follows:

- $S_1 \sqcup_S S_2 = \top_S$ iff there exists some $l \in \text{dom}(S_1) \cap \text{dom}(S_2)$ such that $S_1(l) \sqcup S_2(l) = \top$.
- Otherwise, $S_1 \sqcup_S S_2$ is the store S such that:
 - $\text{dom}(S) = \text{dom}(S_1) \cup \text{dom}(S_2)$, and
 - For all $l \in \text{dom}(S)$:

$$S(l) = \begin{cases} S_1(l) \sqcup S_2(l) & \text{if } l \in \text{dom}(S_1) \cap \text{dom}(S_2) \\ S_1(l) & \text{if } l \notin \text{dom}(S_2) \\ S_2(l) & \text{if } l \notin \text{dom}(S_1). \end{cases}$$

Equivalence of stores is also straightforward. Two stores are equal if they are both \top_S or if they both have the same set of bindings:

Definition 2.4 (equality of stores, λ_{LVar}). Two stores S and S' are *equal* iff:

- (1) $S = \top_S$ and $S' = \top_S$, or
- (2) $\text{dom}(S) = \text{dom}(S')$ and for all $l \in \text{dom}(S)$, $S(l) = S'(l)$.

By Definition 2.3, if $d_1 \sqcup d_2 = \top$, then $[l \mapsto d_1] \sqcup_S [l \mapsto d_2] = \top_S$. Notice that a store containing a binding $l \mapsto \top$ can never arise during the execution of a λ_{LVar} program, because (as I will show in Section 2.4) an attempted write that would take the state of some location l to \top would raise an error before the write can occur.

2.3.3. Communication primitives. The `new`, `put`, and `get` operations create, write to, and read from LVars, respectively. The interface is similar to that presented by mutable references:

- `new` extends the store with a binding for a new LVar whose initial state is \perp , and returns the location l of that LVar (*i.e.*, a pointer to the LVar).
- `put` takes a pointer to an LVar and a new state and updates the LVar’s state to the lub of the current state and the new state, potentially pushing the state of the LVar upward in the lattice. Any update that would take the state of an LVar to \top results in an error.
- `get` performs a blocking “threshold” read that allows limited observations of the state of an LVar. It takes a pointer to an LVar and a *threshold set* T , which is a non-empty subset of D that is *pairwise incompatible*, meaning that the lub of any two distinct elements in T is \top . If the LVar’s state d_1 in the lattice is *at or above* some $d_2 \in T$, the `get` operation unblocks and returns d_2 . Note that d_2 is a unique element of T , for if there is another $d'_2 \neq d_2$ in the threshold set such that $d'_2 \sqsubseteq d_1$, it would follow that $d_2 \sqcup d'_2 \sqsubseteq d_1$, and so $d_2 \sqcup d'_2$ cannot be \top , which contradicts the requirement that T be pairwise incompatible.

The intuition behind `get` is that it specifies a subset of the lattice that is “horizontal”: no two elements in the threshold set can be above or below one another. Intuitively, each element in the threshold set is an “alarm” that detects the activation of itself or any state above it. One way of visualizing the threshold set for a `get` operation is as a subset of edges in the lattice that, if crossed, set off the corresponding alarm. Together these edges form a “tripwire”. Figure 2.1(c) shows what the “tripwire” looks like for an example `get` operation. The threshold set $\{(\perp, 0), (\perp, 1), \dots\}$ (or a subset thereof) would pass the

incompatibility test, as would the threshold set $\{(0, \perp), (1, \perp), \dots\}$ (or a subset thereof), but a combination of the two would not pass.

The requirement that the elements of a threshold set be pairwise incompatible limits the expressivity of threshold sets. In fact, it is a stronger requirement than we need to ensure determinism. Later on, in Section 2.6, I will explain how to generalize the definition of threshold sets to allow more programs to be expressed. For now, I will proceed with the simpler definition above.

2.3.4. Monotonic store growth and determinism. In LVar-based languages, a store can only change in one of two ways: a new, empty location (pointing to \perp) is created, or a previously \perp binding is permanently updated to a meaningful value. It is therefore straightforward in such languages to define an ordering on stores and establish determinism based on the fact that stores grow monotonically with respect to the ordering. For instance, *Featherweight CnC* [11], a single-assignment imperative calculus that models the Intel Concurrent Collections (CnC) system, defines ordering on stores as follows:⁷

Definition 2.5 (store ordering, Featherweight CnC). A store S is *less than or equal to* a store S' (written $S \sqsubseteq_S S'$) iff $\text{dom}(S) \subseteq \text{dom}(S')$ and for all $l \in \text{dom}(S)$, $S(l) = S'(l)$.

Our Definition 2.2 is reminiscent of Definition 2.5, but Definition 2.5 requires that $S(l)$ and $S'(l)$ be *equal*, instead of our weaker requirement that $S(l)$ be *less than or equal to* $S'(l)$ according to the given lattice \sqsubseteq . In λ_{LVar} , stores may grow by updating existing bindings via repeated puts, so Definition 2.5 would be too strong; for instance, if $\perp \sqsubset d_1 \sqsubseteq d_2$ for distinct $d_1, d_2 \in D$, the relationship $[l \mapsto d_1] \sqsubseteq_S [l \mapsto d_2]$ holds under Definition 2.2, but would not hold under Definition 2.5. That is, in λ_{LVar} an LVar could take on the state d_1 , and then later the state d_2 , which would not be possible in Featherweight CnC.

⁷A minor difference between λ_{LVar} and Featherweight CnC is that, in Featherweight CnC, no store location is explicitly bound to \perp . Instead, if $l \notin \text{dom}(S)$, then l is defined to point to \perp .

Given a lattice $(D, \sqsubseteq, \perp, \top)$ with elements $d \in D$:

configurations $\sigma ::= \langle S; e \rangle \mid \mathbf{error}$
 expressions $e ::= x \mid v \mid ee \mid \mathbf{get} \, ee \mid \mathbf{put} \, ee \mid \mathbf{new}$
 values $v ::= () \mid d \mid l \mid T \mid \lambda x. e$
 threshold sets $T ::= \{d_1, d_2, \dots\}$
 stores $S ::= [l_1 \mapsto d_1, \dots, l_n \mapsto d_n] \mid \top_S$
 evaluation contexts $E ::= [] \mid Ee \mid eE \mid \mathbf{get} \, Ee \mid \mathbf{get} \, eE \mid \mathbf{put} \, Ee \mid \mathbf{put} \, eE$

Figure 2.2. Syntax for λ_{LVar} .

I establish in Section 2.5 that λ_{LVar} remains deterministic despite the relatively weak \sqsubseteq_S relation given in Definition 2.2. The key to maintaining determinism is the blocking semantics of the `get` operation and the fact that it allows only *limited* observations of the state of an LVar.

2.4. λ_{LVar} : syntax and semantics

The syntax of λ_{LVar} appears in Figure 2.2, and Figures 2.3 and 2.4 together give the operational semantics. Both the syntax and semantics are parameterized by the lattice $(D, \sqsubseteq, \perp, \top)$.

A *configuration* $\langle S; e \rangle$ comprises a store and an expression. The *error configuration*, written **error**, is a unique element added to the set of configurations, but $\langle \top_S; e \rangle$ is equal to **error** for all expressions e . The metavariable σ ranges over configurations.

Stores are as described in Section 2.3.2, and expressions may be variables, values, function applications, `get` expressions, `put` expressions, or `new`. The value forms include the unit value $()$, elements d of the specified lattice, store locations l , threshold sets T , or λ -expressions $\lambda x. e$. A threshold set is a set $\{d_1, d_2, \dots\}$ of one or more elements of the specified lattice.

Given a lattice $(D, \sqsubseteq, \perp, \top)$ with elements $d \in D$:

$$\text{incomp}(T) \triangleq \forall d_1, d_2 \in T. (d_1 \neq d_2 \Rightarrow d_1 \sqcup d_2 = \top)$$

$$\boxed{\sigma \longmapsto \sigma'}$$

E-Beta

$$\frac{}{\langle S; (\lambda x. e) v \rangle \longmapsto \langle S; e[x := v] \rangle}$$

E-New

$$\frac{}{\langle S; \text{new} \rangle \longmapsto \langle S[l \mapsto \perp]; l \rangle} \quad (l \notin \text{dom}(S))$$

E-Put

$$\frac{S(l) = d_1 \quad d_1 \sqcup d_2 \neq \top}{\langle S; \text{put } l \ d_2 \rangle \longmapsto \langle S[l \mapsto d_1 \sqcup d_2]; () \rangle}$$

E-Put-Err

$$\frac{S(l) = d_1 \quad d_1 \sqcup d_2 = \top}{\langle S; \text{put } l \ d_2 \rangle \longmapsto \text{error}}$$

E-Get

$$\frac{S(l) = d_1 \quad \text{incomp}(T) \quad d_2 \in T \quad d_2 \sqsubseteq d_1}{\langle S; \text{get } l \ T \rangle \longmapsto \langle S; d_2 \rangle}$$

Figure 2.3. Reduction semantics for λ_{LVar} .

$$\boxed{\sigma \mapsto \sigma'}$$

E-Eval-Ctxt

$$\frac{\langle S; e \rangle \longmapsto \langle S'; e' \rangle}{\langle S; E[e] \rangle \mapsto \langle S'; E[e'] \rangle}$$

Figure 2.4. Context semantics for λ_{LVar} .

The operational semantics is split into two parts: a *reduction semantics*, shown in Figure 2.3, and a *context semantics*, shown in Figure 2.4.

The reduction relation \longmapsto is defined on configurations. There are five rules in the reduction semantics: the E-Beta rule is standard β -reduction, and the rules E-New, E-Put/E-Put-Err, and E-Get respectively

express the semantics of the `new`, `put`, and `get` operations described in Section 2.3.3. The E-New rule creates a new binding in the store and returns a pointer to it; the side condition $l \notin \text{dom}(S)$ ensures that l is a fresh location. The E-Put rule updates the store and returns $()$, the unit value. The E-Put-Err rule applies when a put to a location would take its state to \top ; in that case, the semantics steps to **error**. The incompatibility of the threshold set argument to `get` is enforced in the E-Get rule by the $\text{incomp}(T)$ premise, which requires that the lub of any two distinct elements in T must be \top .⁸

The context relation \mapsto is also defined on configurations. It has only one rule, E-Eval-Ctxt, which is a standard context rule, allowing reductions to apply within a context. The choice of context determines where evaluation can occur; in λ_{LVar} , the order of evaluation is nondeterministic (that is, a given expression can generally reduce in more than one way), and so it is generally *not* the case that an expression has a unique decomposition into redex and context.⁹ For example, in an application $e_1 e_2$, either e_1 or e_2 might reduce first. The nondeterminism in choice of evaluation context reflects the nondeterminism of scheduling between concurrent threads, and in λ_{LVar} , the arguments to `get`, `put`, and application expressions are *implicitly* evaluated concurrently.

2.4.1. Fork-join parallelism. λ_{LVar} has a call-by-value semantics: arguments must be fully evaluated before function application (β -reduction, via the E-Beta rule) can occur. We can exploit this property to define the syntactic sugar `let par` for *parallel composition* that we first saw earlier in Example 2.1. With `let par`, we can evaluate two subexpressions e_1 and e_2 in parallel before evaluating a third subexpression e_3 :

$$\text{let par } x_1 = e_1; x_2 = e_2 \text{ in } e_3 \triangleq ((\lambda x_1. (\lambda x_2. e_3)) e_1) e_2$$

⁸Although $\text{incomp}(T)$ is given as a premise of the E-Get reduction rule (suggesting that it is checked at runtime), as I noted earlier in Section 2.2.3, in a real implementation of LVars threshold sets need not have any runtime representation, nor do they need to be written explicitly in the code. Rather, it is the data structure author's responsibility to ensure that any operations provided for reading from LVars have threshold semantics.

⁹In fact, my motivation for splitting the operational semantics into a reduction semantics and a context semantics is to isolate the nondeterminism of the context semantics, which simplifies the determinism proof of Section 2.5.




Figure 2.5. A series-parallel graph induced by parallel λ -calculus evaluation (a); a non-series-parallel graph induced by put/get operations (b).

Although e_1 and e_2 can be evaluated in parallel, e_3 cannot be evaluated until both e_1 and e_2 are values, because the call-by-value semantics does not allow β -reduction until the operand is fully evaluated, and because it further disallows reduction under λ -terms (sometimes called “full β -reduction”). In the terminology of parallel programming, a `let par` expression executes both a *fork* and a *join*. Indeed, it is common for fork and join to be combined in a single language construct, for example, in languages with parallel tuple expressions such as Manticore [21].

Since `let par` expresses *fork-join* parallelism, the evaluation of a program comprising nested `let par` expressions would induce a runtime dependence graph like that pictured in Figure 2.5(a). The λ_{LVar} language (minus `put` and `get`) can support any *series-parallel* dependence graph. Adding communication through `put` and `get` introduces “lateral” edges between branches of a parallel computation, as in Figure 2.5(b). This adds the ability to construct arbitrary non-series-parallel dependency graphs, just as with *first-class futures* [50].

Because we do not reduce under λ -terms, we can sequentially compose e_1 before e_2 by writing `let _ = e_1 in e_2` , which desugars to $(\lambda_. e_2) e_1$. Sequential composition is useful for situations in which expressions must run in a particular order, e.g., if we want to first allocate a new LVar with `new` and then write to it using `put`.

2.4.2. Errors and observable determinism. Is the `get` operation deterministic? Consider two lattice elements d_1 and d_2 that have no ordering and have \top as their lub, and suppose that puts of d_1 and d_2 and a `get` with $\{d_1, d_2\}$ as its threshold set all race for access to an LVar lv :

(Example 2.5)

$$\begin{aligned} & \text{let } \text{par } _ = \text{put } lv \ d_1 \\ & \quad _ = \text{put } lv \ d_2 \\ & \quad x = \text{get } lv \ \{d_1, d_2\} \\ & \text{in } x \end{aligned}$$

Eventually, Example 2.5 is guaranteed to raise **error** by way of the E-Put-Err rule, because $d_1 \sqcup d_2 = \top$. Before that happens, though, `get` $lv \ \{d_1, d_2\}$ could return either d_1 or d_2 . Therefore, `get` *can* behave nondeterministically—but this behavior is not observable in the final outcome of the program, since one of the two puts will raise **error** before the x in the body of the `let par` can be evaluated, and under our definition of observable determinism, only the final outcome of a program counts.

2.5. Proof of determinism for λ_{LVar}

The main technical result of this chapter is a proof of determinism for the λ_{LVar} language. The determinism theorem says that if two executions starting from a given configuration σ terminate in configurations σ' and σ'' , then σ' and σ'' are the same configuration, up to a permutation on locations. (I discuss permutations in more detail below, in Section 2.5.1.)

In order to prove determinism for λ_{LVar} , I first prove several supporting lemmas. Lemma 2.1 (Permutability) deals with location names, and Lemma 2.3 (Locality) establishes a useful property for dealing with expressions that decompose into redex and context in multiple ways. After that point, the structure of the proof is similar to that of the proof of determinism for Featherweight CnC given by Budimlić *et al.* [11]. I reuse the naming conventions of Budimlić *et al.* for Lemmas 2.4 (Monotonicity), 2.5 (Independence), 2.6 (Clash), 2.7 (Error Preservation), and 2.8 (Strong Local Confluence). However, the statements and proofs of those properties differ considerably in the setting of λ_{LVar} , due to the generality of LVars and other differences between the λ_{LVar} language and Featherweight CnC.

On the other hand, Lemmas 2.9 (Strong One-Sided Confluence), 2.10 (Strong Confluence), and 2.11 (Confluence) are nearly identical to the corresponding lemmas in the Featherweight CnC determinism proof. This is the case because, once Lemmas 2.4 through 2.8 are established, the remainder of the determinism proof does not need to deal specifically with the semantics of LVars, lattices, or the store, and instead deals only with execution steps at a high level.

2.5.1. Permutations and permutability. The E-New rule allocates a fresh location $l \in Loc$ in the store, with the only requirement on l being that it is not (yet) in the domain of the store. Therefore, multiple runs of the same program may differ in what locations they allocate, and therefore the reduction semantics is nondeterministic with respect to which locations are allocated. Since this is not a kind of nondeterminism that we care about, we work modulo an arbitrary *permutation* on locations.

Recall from Section 2.3.2 that we have a countable set of locations Loc . Then, a permutation is defined as follows:

Definition 2.6 (permutation, λ_{LVar}). A *permutation* is a function $\pi : Loc \rightarrow Loc$ such that:

- (1) π is invertible, that is, there is an inverse function $\pi^{-1} : Loc \rightarrow Loc$ with the property that $\pi(l) = l'$ iff $\pi^{-1}(l') = l$; and
- (2) π is the identity on all but finitely many elements of Loc .

Condition (1) in Definition 2.6 ensures that we only consider location renamings that we can “undo”, and condition (2) ensures that we only consider renamings of a finite number of locations. Equivalently, we can say that π is a bijection from Loc to Loc such that it is the identity on all but finitely many elements.

Definitions 2.7, 2.8, and 2.9 lift Definition 2.6 to expressions, stores, and configurations, respectively.

There is nothing surprising about these definitions: to apply a permutation π to an expression, we just apply π to any locations that occur in the expression. We can also lift π to evaluation contexts, structurally: $\pi([]) = []$, $\pi(E e) = \pi(E) \pi(e)$, and so on. To lift π to stores, we apply π to all locations in the domain of the store. (We do not have to do any renaming in the codomain of the store, since

location names cannot occur in elements of the lattice D and hence cannot occur in the contents of other store locations.) Since π is a bijection, it follows that if some location l is not in the domain of some store S , then $\pi(l) \notin \text{dom}((\pi(S)))$, a fact that will be useful to us shortly.

Definition 2.7 (permutation of an expression, λ_{LVar}). A *permutation* of an expression e is a function π defined as follows:

$$\begin{aligned}
\pi(x) &\triangleq x \\
\pi(()) &\triangleq () \\
\pi(d) &\triangleq d \\
\pi(l) &\triangleq \pi(l) \\
\pi(T) &\triangleq T \\
\pi(\lambda x. e) &\triangleq \lambda x. \pi(e) \\
\pi(e_1 e_2) &\triangleq \pi(e_1) \pi(e_2) \\
\pi(\text{get } e_1 e_2) &\triangleq \text{get } \pi(e_1) \pi(e_2) \\
\pi(\text{put } e_1 e_2) &\triangleq \text{put } \pi(e_1) \pi(e_2) \\
\pi(\text{new}) &\triangleq \text{new}
\end{aligned}$$

Definition 2.8 (permutation of a store, λ_{LVar}). A *permutation* of a store S is a function π defined as follows:

$$\begin{aligned}
\pi(\top_S) &\triangleq \top_S \\
\pi([l_1 \mapsto d_1, \dots, l_n \mapsto d_n]) &\triangleq [\pi(l_1) \mapsto d_1, \dots, \pi(l_n) \mapsto d_n]
\end{aligned}$$

Definition 2.9 (permutation of a configuration, λ_{LVar}). A *permutation* of a configuration $\langle S; e \rangle$ is a function π defined as follows: if $\langle S; e \rangle = \mathbf{error}$, then $\pi(\langle S; e \rangle) = \mathbf{error}$; otherwise, $\pi(\langle S; e \rangle) = \langle \pi(S); \pi(e) \rangle$.

With these definitions in place, I can prove Lemma 2.1, which says that the names of locations in a configuration do not affect whether or not that configuration can take a step: a configuration σ can step to σ' exactly when $\pi(\sigma)$ can step to $\pi(\sigma')$.

Lemma 2.1 (Permutability, λ_{LVar}). *For any finite permutation π ,*

- (1) $\sigma \longrightarrow \sigma'$ *if and only if* $\pi(\sigma) \longrightarrow \pi(\sigma')$.
- (2) $\sigma \longmapsto \sigma'$ *if and only if* $\pi(\sigma) \longmapsto \pi(\sigma')$.

Proof. See Section A.1. The forward direction of part 1 is by cases on the rule in the reduction semantics by which σ steps to σ' ; the only interesting case is the E-New case, in which we make use of the fact that if $l \notin \text{dom}(S)$, then $\pi(l) \notin \text{dom}(\pi(S))$. The reverse direction of part 1 relies on the fact that if π is a permutation, then π^{-1} is also a permutation. Part 2 of the proof builds on part 1. \square

Because the names of locations in a configuration do not affect whether it can step, we can rename locations as needed, which will be important later on when proving the *confluence* lemmas of Section 2.5.8.¹⁰

2.5.2. Internal determinism. My goal is to show that λ_{LVar} is deterministic according to the definition of *observable determinism* that I gave in Chapter 1—that is, that a λ_{LVar} program always evaluates to the same value. In the context of λ_{LVar} , a “program” can be understood as a configuration, and a “value” can be understood as a configuration that cannot step, either because the expression in that configuration is actually a λ_{LVar} value, or because it is a “stuck” configuration that cannot step because no rule of the operational semantics applies. In λ_{LVar} , the latter situation could occur if, for instance, a configuration contains a blocking get expression and there are no other expressions left to evaluate that might cause it to unblock.

This definition of observable determinism does *not* require that a configuration takes the same sequence of steps on the way to reaching its value at the end of every run. Borrowing terminology from Blelloch *et al.* [6], I will use the term *internally deterministic* to describe a program that does, in fact, take the same

¹⁰For another example of using permutations in the metatheory of a language to account for an allocator’s nondeterministic choice of locations in an otherwise deterministic setting, see Krishnaswami [31].

sequence of steps on every run.¹¹ Although λ_{LVar} is not internally deterministic, all of its internal non-determinism is due to the E-Eval-Ctxt rule! This is the case because the E-Eval-Ctxt rule is the only rule in the operational semantics by which a particular configuration can step in multiple ways. The multiple ways in which a configuration can step via E-Eval-Ctxt correspond to the ways in which the expression in that configuration can be decomposed into a redex and an evaluation context. In fact, it is exactly this property that makes it possible for multiple subexpressions of a λ_{LVar} expression (a `let par` expression, for instance) to be evaluated in parallel.

But, leaving aside evaluation contexts for the moment—we will return to them in the following section—let us focus on the rules of the reduction semantics in Figure 2.3. Here we can see that if a given configuration can step by the reduction semantics, then there is only one rule by which it can step, and only one configuration to which it can step. The only exception is the E-New rule, which nondeterministically allocates locations and returns pointers to them—but we can account for this by saying that the reduction semantics is internally deterministic up to a permutation on locations. Lemma 2.2 formalizes this claim, which we will later use in the proof of Strong Local Confluence (Lemma 2.8).

Lemma 2.2 (Internal Determinism, λ_{LVar}). *If $\sigma \longrightarrow \sigma'$ and $\sigma \longrightarrow \sigma''$, then there is a permutation π such that $\sigma' = \pi(\sigma'')$.*

Proof. Straightforward by cases on the rule of the reduction semantics by which σ steps to σ' ; the only interesting case is for the E-New rule. See Section A.2. \square

2.5.3. Locality. In order to prove determinism for λ_{LVar} , we will have to consider situations in which we have an expression that decomposes into redex and context in multiple ways. Suppose that we have an

¹¹I am using “internally deterministic” in a more specific way than Blelloch *et al.*: they define an internally deterministic program to be one for which the *trace* of the program is the same on every run, where a trace is a directed acyclic graph of the operations executed by the program and the control dependencies among them. This definition, in turn, depends on the definition of “operation”, which might be defined in a fine-grained way or a coarse-grained, abstract way, depending on which aspects of program execution one wants the notion of internal determinism to capture. The important point is that internal determinism is a stronger property than observable determinism.

expression e such that $e = E_1[e_1] = E_2[e_2]$. The configuration $\langle S; e \rangle$ can then step in two different ways by the E-Eval-Ctxt rule: $\langle S; E_1[e_1] \rangle \mapsto \langle S_1; E_1[e'_1] \rangle$, and $\langle S; E_2[e_2] \rangle \mapsto \langle S_2; E_2[e'_2] \rangle$.

The interesting case is the one where E_1 and E_2 are different. The key observation we can make here is that the \mapsto relation acts “locally”. That is, when e_1 steps to e'_1 within its context, the expression e_2 will be left alone, because it belongs to the context. Likewise, when e_2 steps to e'_2 within its context, the expression e_1 will be left alone. Lemma 2.3 formalizes this claim.

Lemma 2.3 (Locality, λ_{LVar}). *If $\langle S; E_1[e_1] \rangle \mapsto \langle S_1; E_1[e'_1] \rangle$ and $\langle S; E_2[e_2] \rangle \mapsto \langle S_2; E_2[e'_2] \rangle$ and $E_1[e_1] = E_2[e_2]$, then:*

If $E_1 \neq E_2$, then there exist evaluation contexts E'_1 and E'_2 such that:

- $E'_1[e_1] = E_2[e'_2]$, and
- $E'_2[e_2] = E_1[e'_1]$, and
- $E'_1[e'_1] = E'_2[e'_2]$.

Proof. Let $e = E_1[e_1] = E_2[e_2]$. The proof is by induction on the structure of the expression e . See Section A.3. □

2.5.4. Monotonicity. The Monotonicity lemma says that, as evaluation proceeds according to the \longrightarrow relation, the store can only grow with respect to the \sqsubseteq_S ordering.

Lemma 2.4 (Monotonicity, λ_{LVar}). *If $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$, then $S \sqsubseteq_S S'$.*

Proof. Straightforward by cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$. The interesting cases are for the E-New and E-Put rules. See Section A.4. □

2.5.5. Independence. Figure 2.6 shows a *frame rule*, due to O’Hearn *et al.* [43], which captures the idea of *local reasoning* about programs that alter state. In it, C is a program, and $\{p\} C \{q\}$ is a *Hoare triple* (in the style of Hoare logic [27]) specifying the behavior of C : it says that if the assertion p is true before

Frame rule:

$$\frac{\{p\} C \{q\}}{\{p * r\} C \{q * r\}}$$

Lemma 2.5 (Independence), simplified:

$$\frac{\langle S; e \rangle \longleftrightarrow \langle S'; e' \rangle}{\langle S \sqcup_S S''; e \rangle \longleftrightarrow \langle S' \sqcup_S S''; e' \rangle}$$

Figure 2.6. Comparison of O’Hearn *et al.*’s frame rule [43] and a simplified version of the Independence lemma. The separating conjunction connective $*$ in the frame rule requires that its arguments be disjoint; the Independence lemma uses the \sqcup_S operation in place of $*$.

C runs, then the assertion q will be true afterwards. For example, p and q might respectively describe the state of the heap before and after a heap location is updated by C .

Given a program C with precondition p and postcondition q , the frame rule tells us that running C starting from a state satisfying the precondition $p * r$ will result in a state satisfying the postcondition $q * r$. These two assertions use the *separating conjunction* connective $*$, which combines two assertions that can be satisfied in a *non-overlapping* manner. For instance, the assertion $p * r$ is satisfied by a heap if the heap can be split into two non-overlapping parts satisfying p and r , respectively.

Therefore, if C can run safely starting from a state satisfying p and end in a state satisfying q , then it does not do any harm to also have the disjoint property r be true when C runs: the truth of r will not interfere with the safe execution of C . Furthermore, if r is true to begin with, running C will not interfere with the truth of r . The frame rule gets its name from the fact that r is a “frame” around C : everything that is not explicitly changed by C is part of the frame and is inviolate.¹² O’Hearn *et al.* refer

¹²The “frame” terminology was originally introduced in 1969 by McCarthy and Hayes [37], who observed that specifying only what is changed by an action does not generally allow an intelligent agent to conclude that nothing else is changed; they called this dilemma the *frame problem*.

to the resources (such as heap locations) actually used by C as the “footprint” of C ; r is an assertion about resources outside of that footprint.

The Independence lemma establishes a similar “frame property” for λ_{LVar} that captures the idea that independent effects commute with each other. Consider an expression e that runs starting in store S and steps to e' , updating the store to S' . The Independence lemma provides a double-edged guarantee about what will happen if we evaluate e starting from a larger store $S \sqcup_S S''$: we know both that e will update the store to $S' \sqcup_S S''$, and that e will step to e' as it did before. Here, $S \sqcup_S S''$ is the lub of the original S and some other store S'' that is “framed on” to S ; intuitively, S'' is the store resulting from some other independently-running computation.¹³

Lemma 2.5 requires as a precondition that the store S'' must be *non-conflicting* with the original transition from $\langle S; e \rangle$ to $\langle S'; e' \rangle$, meaning that locations in S'' cannot share names with locations newly allocated during the transition; this rules out location name conflicts caused by allocation.

Definition 2.10 (non-conflicting store). A store S'' is *non-conflicting* with the transition $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ iff $(\text{dom}(S') - \text{dom}(S)) \cap \text{dom}(S'') = \emptyset$.

Lemma 2.5 (Independence). If $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ (where $\langle S'; e' \rangle \neq \mathbf{error}$), then for all S'' such that S'' is non-conflicting with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ and $S' \sqcup_S S'' \neq \top_S$:

$$\langle S \sqcup_S S''; e \rangle \longrightarrow \langle S' \sqcup_S S''; e' \rangle.$$

Proof. By cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$. The interesting cases are for the E-New and E-Put rules. Since $\langle S'; e' \rangle \neq \mathbf{error}$, we do not need to consider the E-Put-Err rule. See Section A.5. □

2.5.6. Clash. The Clash lemma, Lemma 2.6, is similar to the Independence lemma, but handles the case where $S' \sqcup_S S'' = \top_S$. It establishes that, in that case, $\langle S \sqcup_S S''; e \rangle$ steps to **error**.

¹³See Section 6.5 for a more detailed discussion of frame properties and where they manifest in the LVars model.

Lemma 2.6 (Clash). *If $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ (where $\langle S'; e' \rangle \neq \mathbf{error}$), then for all S'' such that S'' is non-conflicting with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ and $S' \sqcup_S S'' = \top_S$:*

$\langle S \sqcup_S S''; e \rangle \longrightarrow^i \mathbf{error}$, where $i \leq 1$.

Proof. By cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$. As with Lemma 2.5, the interesting cases are for the E-New and E-Put rules, and since $\langle S'; e' \rangle \neq \mathbf{error}$, we do not need to consider the E-Put-Err rule. See Section A.6. \square

2.5.7. Error preservation. Lemma 2.7, Error Preservation, says that if a configuration $\langle S; e \rangle$ steps to **error**, then evaluating e in the context of some larger store will also result in **error**.

Lemma 2.7 (Error Preservation, λ_{LVar}). *If $\langle S; e \rangle \longrightarrow \mathbf{error}$ and $S \sqsubseteq_S S'$, then $\langle S'; e \rangle \longrightarrow \mathbf{error}$.*

Proof. Suppose $\langle S; e \rangle \longrightarrow \mathbf{error}$ and $S \sqsubseteq_S S'$. We are required to show that $\langle S'; e \rangle \longrightarrow \mathbf{error}$.

By inspection of the operational semantics, the only rule by which $\langle S; e \rangle$ can step to **error** is E-Put-Err. Hence $e = \text{put } l \ d_2$. From the premises of E-Put-Err, we have that $S(l) = d_1$. Since $S \sqsubseteq_S S'$, it must be the case that $S'(l) = d'_1$, where $d_1 \sqsubseteq d'_1$. Since $d_1 \sqcup d_2 = \top$, we have that $d'_1 \sqcup d_2 = \top$. Hence, by E-Put-Err, $\langle S'; \text{put } l \ d_2 \rangle \longrightarrow \mathbf{error}$, as we were required to show. \square

2.5.8. Confluence. Lemma 2.8, the Strong Local Confluence lemma, says that if a configuration σ can step to configurations σ_a and σ_b , then there exists a configuration σ_c that σ_a and σ_b can each reach in at most one step, modulo a permutation on the locations in σ_b . Lemmas 2.9 and 2.10 then generalize that result to arbitrary numbers of steps.

The structure of this part of the proof differs from the Budimlić *et al.* determinism proof for Featherweight CnC in two ways. First, Budimlić *et al.* prove a *diamond* property, in which σ_a and σ_b each step to σ_c in *exactly* one step. They then get a property like Lemma 2.8 as an immediate consequence of the diamond property, by choosing $i = j = 1$. But a true diamond property with exactly one step “on each

side of the diamond” is stronger than we need here, and, in fact, does not hold for λ_{LVar} ; so, instead, I prove the weaker “at most one step” property directly.

Second, Budimlić *et al.* do not have to deal with permutations in their proof, because the Featherweight CnC language does no allocation; there is no counterpart to λ_{LVar} ’s `new` expression in Featherweight CnC. Instead, Featherweight CnC models the store as a pre-existing array of locations, where every location has a default initial value of \perp . Because there is no way (and no need) to allocate new locations in Featherweight CnC, it is never the case that two subexpressions independently happen to allocate locations with the same name—which is exactly the situation that requires us to be able to rename locations in λ_{LVar} . In fact, that situation is what makes the entire notion of permutations described in Section 2.5.1 a necessary part of the metatheory of λ_{LVar} .

Taking the approach of Featherweight CnC, and therefore avoiding allocation entirely, would simplify both the λ_{LVar} language and its determinism proof. On the other hand, when programming with the LVish Haskell library of Chapter 4, one does have to explicitly create and allocate new LVars by calling the equivalent of `new`, and so by modeling the `new` operation in λ_{LVar} , we keep the semantics a bit more faithful to the implementation.

Lemma 2.8 (Strong Local Confluence). *If $\sigma \mapsto \sigma_a$ and $\sigma \mapsto \sigma_b$, then there exist σ_c, i, j, π such that $\sigma_a \mapsto^i \sigma_c$ and $\pi(\sigma_b) \mapsto^j \sigma_c$ and $i \leq 1$ and $j \leq 1$.*

Proof. Since the original configuration σ can step in two different ways, its expression decomposes into redex and context in two different ways: $\sigma = \langle S; E_a[e_{a_1}] \rangle = \langle S; E_b[e_{b_1}] \rangle$, where $E_a[e_{a_1}] = E_b[e_{b_1}]$, but E_a and E_b may differ and e_{a_1} and e_{b_1} may differ. In the special case where $E_a = E_b$, the result follows by Internal Determinism (Lemma 2.2).

If $E_a \neq E_b$, we can apply the Locality lemma (Lemma 2.3); at a high level, it shows that e_{a_1} and e_{b_1} can be evaluated independently within their contexts. The proof is then by a double case analysis on the rules of the reduction semantics by which $\langle S; e_{a_1} \rangle$ steps and by which $\langle S; e_{b_1} \rangle$ steps. In order to

combine the results of the two independent steps, the proof makes use of the Independence lemma (Lemma 2.5). The most interesting case is that in which both steps are by the E-New rule and they allocate locations with the same name. In that case, we can use the Permutability lemma (Lemma 2.1) to rename locations so as not to conflict. See Section A.7. \square

Lemma 2.9 (Strong One-Sided Confluence). *If $\sigma \mapsto \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq m$, then there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq 1$.*

Proof. By induction on m ; see Section A.8. \square

Lemma 2.10 (Strong Confluence). *If $\sigma \mapsto^n \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq n$ and $1 \leq m$, then there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq n$.*

Proof. By induction on n ; see Section A.9. \square

Lemma 2.11 (Confluence). *If $\sigma \mapsto^* \sigma'$ and $\sigma \mapsto^* \sigma''$, then there exist σ_c and π such that $\sigma' \mapsto^* \sigma_c$ and $\pi(\sigma'') \mapsto^* \sigma_c$.*

Proof. Strong Confluence (Lemma 2.10) implies Confluence. \square

2.5.9. Determinism. Finally, the determinism theorem, Theorem 2.1, is a direct result of Lemma 2.11:

Theorem 2.1 (Determinism). *If $\sigma \mapsto^* \sigma'$ and $\sigma \mapsto^* \sigma''$, and neither σ' nor σ'' can take a step, then there exists π such that $\sigma' = \pi(\sigma'')$.*

Proof. We have from Confluence (Lemma 2.11) that there exists σ_c and π such that $\sigma' \mapsto^* \sigma_c$ and $\pi(\sigma'') \mapsto^* \sigma_c$. Since σ' cannot step, we must have $\sigma' = \sigma_c$.

By Permutability (Lemma 2.1), σ'' can step iff $\pi(\sigma'')$ can step, so since σ'' cannot step, $\pi(\sigma'')$ cannot step either.

Hence we must have $\pi(\sigma'') = \sigma_c$. Since $\sigma' = \sigma_c$ and $\pi(\sigma'') = \sigma_c$, $\sigma' = \pi(\sigma'')$. \square

2.5.10. Discussion: termination. I have followed Budimlić *et al.* [11] in treating *determinism* separately from the issue of *termination*. Yet one might legitimately be concerned that in λ_{LVar} , a configuration could have both an infinite reduction path and one that terminates with a value. Theorem 2.1 says that if two runs of a given λ_{LVar} program reach configurations where no more reductions are possible, then they have reached the same configuration. Hence Theorem 2.1 handles the case of *deadlocks* already: a λ_{LVar} program can deadlock (e.g., with a blocked `get`), but it will do so deterministically.

However, Theorem 2.1 has nothing to say about *livelocks*, in which a program reduces infinitely. It would be desirable to have a *consistent termination* property which would guarantee that if one run of a given λ_{LVar} program terminates with a non-**error** result, then every run will. I conjecture (but do not prove) that such a consistent termination property holds for λ_{LVar} . Such a property could be paired with Theorem 2.1 to guarantee that if one run of a given λ_{LVar} program terminates in a non-**error** configuration σ , then every run of that program terminates in σ . (The “non-**error** configuration” condition is necessary because it is possible to construct a λ_{LVar} program that can terminate in **error** on some runs and diverge on others. By contrast, the existing determinism theorem does not have to treat **error** specially.)

2.6. Generalizing the put and get operations

The determinism result for λ_{LVar} shows that adding LVars (with their accompanying `new/put/get` operations) to an existing deterministic parallel language (the λ -calculus) preserves determinism. But it is not the case that the `put` and `get` operations are *the most general* determinism-preserving operations on LVars. In this section, I consider some alternative semantics for `put` and `get` that generalize their behavior while retaining the determinism of the model.

2.6.1. Generalizing from least-upper-bound writes to inflationary, commutative writes. In the LVars model as presented in this chapter so far, the only way for the state of an LVar to evolve over time is through a series of `put` operations. Unfortunately, this way of updating an LVar provides no efficient way to model, for instance, an atomically incremented counter that occupies one memory location.

Consider an LVar based on the lattice of Figure 2.1(a). Under the semantics of `put`, if two independent writes each take the LVar’s contents from, say, 1 to 2, then after both writes, its contents will be 2, because `put` takes the maximum of the previous value and the current value. Although this semantics is deterministic, it is not the desired semantics for every application. Instead, we might want each write to *increment* the contents of the LVar by one, resulting in 3.

To support this alternative semantics in the LVars model, we generalize the model as follows. For an LVar with lattice $(D, \sqsubseteq, \perp, \top)$, we can define a family of *update operations* $u_i : D \rightarrow D$, which must meet the following two conditions:

- $\forall d, i. d \sqsubseteq u_i(d)$
- $\forall d, i, j. u_i(u_j(d)) = u_j(u_i(d))$

The first of these conditions says that each update operation is inflationary with respect to \sqsubseteq . The second condition says that update operations commute with each other. These two conditions correspond to the two informal criteria that we set forth for monotonic data structures at the beginning of this chapter: the requirement that updates be inflationary corresponds to the fact that monotonic data structures can only “grow”, and the requirement that updates be commutative corresponds to the fact that the timing of updates must not be observable.¹⁴

In fact, the `put` operation meets the above two conditions, and therefore can be viewed as a special case of an update operation that, in addition to being inflationary and commutative, also happens to compute a lub. However, when generalizing LVars to support update operations, we must keep in mind that `put` operations do not necessarily mix with arbitrary update operations on the same LVar. For example, consider a family of update operations $\{u_{(+1)}, u_{(+2)}, \dots\}$ for atomically incrementing a counter represented by a natural number LVar, with a lattice ordered by the usual \leq on natural numbers. The $u_{(+1)}$ operation increments the counter by one, $u_{(+2)}$ increments it by two, and so on. It is easy to see that

¹⁴Of course, commutativity of updates alone is not enough to assure that the timing of updates is not observable; for that we also need threshold reads.

these operations commute. However, a `put` of 4 and a $u_{(+1)}$ do not commute: if we start with an initial state of 0 and the `put` occurs first, then the state of the LVar changes to 4 since $\max(0, 4) = 4$, and the subsequent $u_{(+1)}$ updates it to 5. But if the $u_{(+1)}$ happens first, then the final state of the LVar will be $\max(1, 4) = 4$. Furthermore, multiple distinct families of update operations only commute among themselves and cannot be combined.

In practice, the author of a particular LVar data structure must choose which update operations that data structure should provide, and it is the data structure author’s responsibility to ensure that they commute. For example, the LVish Haskell library of Chapter 4 provides a set data structure, `Data.LVar.Set`, that supports only `put`, whereas the counter data structure `Data.LVar.Counter` supports only increments; an attempt to call `put` on a `Counter` would be ruled out by the type system. However, *composing* LVars that support different families of update operations is fine. For example, an LVar could represent a monotonically growing collection (which supports `put`) of counter LVars, where each counter is itself monotonically increasing and supports only increment. Indeed, the PhyBin case study that I describe in Section 4.5 uses just such a collection of counters.

In Chapter 3, in addition to extending λ_{LVar} to support the new features of *freezing* and *event handlers*, I generalize the `put` operation to allow arbitrary update operations. More precisely, I replace the `put` operation with a family of operations `puti`, with each corresponding to an update operation u_i . The resulting generalized language definition is therefore parameterized not only by a given lattice, but also by a given family of update operations. Furthermore, as we will see in Section 3.3.5, we will need to generalize the Independence lemma (Lemma 2.5) in order to accommodate this change to the language.

2.6.2. A more general formulation of threshold sets. Certain deterministic computations are difficult to express using the definition of threshold sets presented in Section 2.3.3. For instance, consider an LVar that stores the result of a parallel logical “and” operation on two Boolean inputs. I will call this data structure an `AndLV`, and its two inputs the *left* and *right* inputs, respectively.

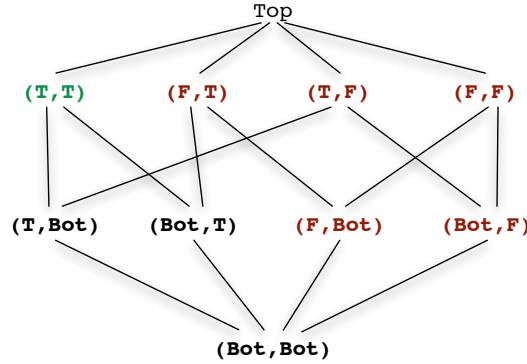


Figure 2.7. The lattice of states that an AndLV can take on. The five red states in the lattice correspond to a false result, and the one green state corresponds to a true one.

We can represent the states an AndLV can take on as pairs (x, y) , where each of x and y are T, F, or Bot. The (Bot, Bot) state is the state in which no input has yet been received, and so it is the least element in the lattice of states that our AndLV can take on, shown in Figure 2.7. An additional state, Top, is the greatest element of the lattice; it represents the situation in which an error has occurred—if, for instance, one of the inputs writes T and then later changes its mind to F.

The lattice induces a lub operation on pairs of states; for instance, the lub of (T, Bot) and (Bot, F) is (T, F) , and the lub of (T, Bot) and (F, Bot) is Top since the overlapping T and F values conflict. The put operation updates the AndLV’s state to the lub of the incoming state and the current state.

We are interested in learning whether the result of our parallel “and” computation is “true” or “false”. Let us consider what observations it is possible to make of an AndLV under our existing definition of threshold reads. The states (T, T) , (T, F) , (F, T) , and (F, F) are all pairwise incompatible with one another, and so $\{(\text{T}, \text{T}), (\text{T}, \text{F}), (\text{F}, \text{T}), (\text{F}, \text{F})\}$ —that is, the set of states in which both the left and right inputs have arrived—is a legal threshold set argument to get. The trouble with this threshold read is that it does not allow us to get *early answers* from the computation. It would be preferable to have a get operation that would “short circuit” and unblock immediately if a single input of, say, (F, Bot) or (Bot, F) was written, since no later write could change the fact that the result of the whole computation

would be “false”.¹⁵ Unfortunately, we cannot include (F, Bot) or (Bot, F) in our threshold set, because the resulting threshold set would no longer be pairwise incompatible, and therefore would compromise determinism.

In order to get short-circuiting behavior from an AndLV without compromising determinism, we need to make a slight generalization to how threshold sets and threshold reads work. In the new formulation, we divide up threshold sets into subsets that we call *activation sets*, each consisting of *activation states*. In the case of the observation we want to make of our AndLV, one of those activation sets is the set of states that the data structure might be in when a state containing at least one F value has been written—that is, the set $\{(F, \text{Bot}), (\text{Bot}, F), (F, T), (T, F), (F, F)\}$. When we reach a point in the lattice that is at or above any of those states, we know that the result will be “false”. The other activation set is the singleton set $\{(T, T)\}$, since we have to wait until we reach the state (T, T) to know that the result is “true”; a state like (T, Bot) does not appear in any of our activation sets.

We can now redefine “threshold set” to mean *a set of activation sets*. Under this definition, the entire threshold set that we would use to observe the contents of our AndLV is:

$$\{\{(F, \text{Bot}), (\text{Bot}, F), (F, T), (T, F), (F, F)\}, \{(T, T)\}\}$$

We redefine the semantics of *get* as follows: if an LVar’s state reaches (or surpasses) any state or states in a particular activation set in the threshold set, *get* returns *that entire activation set*, regardless of which of its activation states was reached. If no state in any activation set in the threshold set has yet been reached, the *get* operation will block. In the case of our AndLV, as soon as either input writes a state containing an F, our *get* will unblock and return the first activation set, that is, $\{(F, \text{Bot}), (\text{Bot}, F), (F, T), (T, F), (F, F)\}$. Hence AndLV has the expected “short-circuit” behavior and

¹⁵Actually, this is not quite true: a write of (F, Bot) followed by a write of (T, Bot) would lead to a result of *Top*, and to the program stepping to the **error** state, which is certainly different from a result of “false”. But, even if a write of (T, Bot) is due to come along sooner or later to take the state of the AndLV to *Top* and thus raise **error**, it should still be fine for the *get* operation to allow “short-circuit” unblocking, because the result of the *get* operation does not count as observable under our definition of observable determinism (as discussed in Section 2.4.2).

does not have to wait for a second input if the first input contains an F. If, on the other hand, the inputs are (T, Bot) and (Bot, T), the `get` will unblock and return $\{(T, T)\}$.

In practice, the value returned from the `get` could be more meaningful to the client—for instance, a Haskell implementation could return `False` instead of returning the actual activation set that corresponds to “false”. However, the translation from $\{(F, \text{Bot}), (\text{Bot}, F), (F, T), (T, F), (F, F)\}$ to `False` could just as easily take place on the client side. In either case, the activation set returned from the threshold read is the same regardless of *which* of its activation states caused the read to unblock, and it is impossible for the client to tell whether the actual state of the lattice is, say, (T, F), (F, F), or some other state containing F.

As part of this activation-set-based formulation of threshold sets, we need to adjust our criterion for pairwise incompatibility of threshold sets. Recall that the purpose of the pairwise incompatibility requirement (see Section 2.3.3) was to ensure that a threshold read would return a unique result. We need to generalize this requirement, since although more than one element *in the same activation set* might be reached or surpassed by a given write to an LVar, it is still the case that writes should only unblock a *unique* activation set in the threshold set. The pairwise incompatibility requirement then becomes that elements in an activation set must be *pairwise incompatible* with elements in every other activation set. That is, for all distinct activation sets Q and R in a given threshold set:

$$\forall q \in Q. \forall r \in R. q \sqcup r = \top$$

In our AndLV example, there are two distinct activation sets, so if we let $Q = \{(T, T)\}$ and $R = \{(F, \text{Bot}), (\text{Bot}, F), (F, T), (T, F), (F, F)\}$, the lub of (T, T) and r must be `Top`, where r is any element of R . We can easily verify that this is the case.

To illustrate why we need pairwise incompatibility to be defined this way, consider the following (illegal) “threshold set” that does not meet the pairwise incompatibility criterion:

$$\{\{(F, \text{Bot}), (\text{Bot}, F)\}, \{(T, \text{Bot}), (\text{Bot}, T)\}\}$$

A `get` corresponding to this so-called threshold set will unblock and return $\{(F, \text{Bot}), (\text{Bot}, F)\}$ as soon as a state containing an F is reached, and $\{(T, \text{Bot}), (\text{Bot}, T)\}$ as soon as a state containing a T is reached. If, for instance, the left input writes (F, Bot) and the right input writes (Bot, T) , and these writes occur in arbitrary order, the threshold read will return a nondeterministic result, depending on the order of the two writes. But if `get` uses the properly pairwise-incompatible threshold set that has Q and R as its two activation sets, it will block until the write of (F, Bot) arrives, and then will deterministically return Q , the “false” activation set, regardless of whether the write of (Bot, T) has arrived yet. Hence “short-circuit” evaluation is possible.

Finally, we can mechanically translate the old way of specifying threshold sets into activation-set-based threshold sets and retain the old semantics (and therefore the new way of specifying threshold sets generalizes the old way). In the translation, every member of the old threshold set simply becomes a singleton activation set. For example, if we wanted a *non*-short-circuiting threshold read of our `AndLV` under the activation-set-based semantics, our threshold set would simply be

$$\{\{(T, T)\}, \{(T, F)\}, \{(F, T)\}, \{(F, F)\}\},$$

which is a legal threshold set under the activation-set-based semantics, but has the same behavior as the old, non-short-circuiting version.

I use the activation-set-based formulation of threshold sets in Chapter 5, where I bring threshold reads to the setting of replicated, distributed data structures. I prove that activation-set-based threshold queries of distributed data structures behave deterministically (according to a definition of determinism that is particular to the distributed setting; see Section 5.3 for the details). That said, there is nothing about activation-set-based threshold sets that makes them particularly suited to the distributed setting; either the original formulation of threshold sets or the even more general *threshold functions* I discuss in the following section would have worked as well.

2.6.3. Generalizing from threshold sets to threshold functions. The previous section’s generalization to activation-set-based threshold sets prompts us to ask: are further generalizations possible while retaining determinism? The answer is yes: both the original way of specifying threshold sets and the more general, activation-set-based formulation of them can be described by *threshold functions*. A threshold function is a partial function that takes a lattice element as its argument and is undefined for all inputs that are not at or above a given element in the lattice (which I will call its *threshold point*), and *constant* for all inputs that *are* at or above its threshold point. (Note that “not at or above” is more general than “below”: a threshold function is undefined for inputs that are neither above nor below its threshold point.)

Threshold functions capture the semantics of both the original style of threshold sets and the activation-set-based style:

- In the original style of threshold sets, every element d of a threshold set can be described by a threshold function that has d as its threshold point and returns d for all inputs at or above that point.
- In the activation-set-based style of threshold sets, every element d of an activation set Q can be described by a threshold function that has d as its threshold point and returns Q for all inputs at or above that point.

In both cases, inputs for which the threshold functions are undefined correspond to situations in which the threshold read blocks.

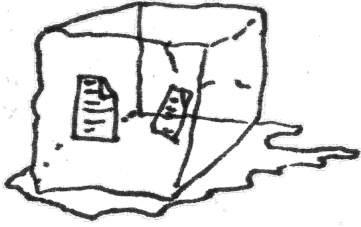
Seen from this point of view, it becomes clear that the key insight in generalizing from the original style of threshold sets to the activation-set-based style of threshold sets is that, for inputs for which a threshold function is defined, its return value need not be its threshold point. The activation set Q is a particularly useful return value, but *any* constant return value will suffice.

CHAPTER 3

Quasi-deterministic and event-driven programming with LVars

The LVars programming model presented in Chapter 2 is based on the idea of *monotonic data structures*, in which information can only be added, never removed, and the timing with which information is added (and hence the *order* in which it is added) is not observable. A paradigmatic example is a set that supports insertion but not removal, but there are many others. In the LVars model, all shared data structures (called LVars) are monotonic, and the states that an LVar can take on form a *lattice*. Writes to an LVar must correspond to a lub operation in the lattice, which means that they monotonically increase the information in the LVar, and that they commute with one another. But commuting writes are not enough to guarantee determinism: if a read can observe whether or not a concurrent write has happened, then it can observe differences in scheduling. So, in the LVars model, the answer to the question “has a write occurred?” (*i.e.*, is the LVar above a certain lattice value?) is always yes; the reading thread will block until the LVar’s contents reach a desired threshold. In a monotonic data structure, the absence of information is transient—another thread could add that information at any time—but the presence of information is forever.

We want to use LVars to implement fixpoint computations like the parallel graph traversal of Section 2.1. But we cannot do so using only least-upper-bound writes and threshold reads, because in order to determine when the set of traversed nodes in the graph has reached a fixpoint, we need to be able to see the exact contents of that set, and it is impossible to learn the exact contents of the set using only threshold reads.



In this chapter, I describe two extensions to the basic LVar model of Chapter 2 that give us a new way to approach problems like the parallel graph traversal problem. First, I add the ability to attach *event handlers* to an LVar that allow callback functions to run in response to updates to the LVar. We say

that a group of event handlers is *quiescent* when no callbacks are currently enabled to run. Second, I add a new primitive operation, *freeze*, that returns the exact contents of an LVar without blocking. Using *freeze* to read an LVar comes with the following tradeoff: once an LVar has been read, it is *frozen*, and any further writes that would change its value instead throw an exception.

The threshold reads that we have seen so far encourage a synchronous, *pull* model of programming in which threads ask specific questions of an LVar, potentially blocking until the answer is “yes”. The addition of handlers, quiescence, and freezing, by contrast, enables an asynchronous, *push* model of programming. We will refer to this extended programming model as the *LVish* programming model. Because quiescence makes it possible to tell when the fixpoint of a computation has been reached, the LVish model is particularly well suited to problems like the graph traversal problem that we saw in Section 2.1.

Unfortunately, freezing does not commute with writes that change an LVar.¹ If a freeze is interleaved before such a write, the write will raise an exception; if it is interleaved afterwards, the program will proceed normally. It would appear that the price of negative information is the loss of determinism!

Fortunately, the loss is not total. Although LVar programs with freezing are not guaranteed to be deterministic, they do satisfy a related property that I call *quasi-determinism*: all executions that produce a final value produce the *same* final value. To put it another way, a quasi-deterministic program can be trusted to never change its answer due to nondeterminism; at worst, it might raise an exception on

¹The same is true for quiescence detection, as we will see in Section 3.1.2.

some runs. This exception can in principle pinpoint the exact pair of freeze and write operations that are racing, greatly easing debugging.

In general, the ability to make exact observations of the contents of data structures is in tension with the goal of guaranteed determinism. Since pushing towards full-featured, general monotonic data structures leads to flirtation with nondeterminism, perhaps the best way of ultimately getting deterministic outcomes is to traipse a short distance into nondeterministic territory, and make our way back. The identification of quasi-deterministic programs as a useful intermediate class of programs is a contribution of this dissertation. That said, in many cases the freeze construct is only used as the very final step of a computation: after a global barrier, freezing is used to extract an answer. In this common case, determinism is guaranteed, since no writes can subsequently occur.

I will refer to the LVars model, extended with handlers, quiescence, and freezing, as the *LVish model*. The rest of this chapter introduces the LVish programming model, first informally through a series of examples, and then formally, by extending the λ_{LVar} calculus of Chapter 2 to add support for handlers, quiescence, freezing, and the arbitrary update operations described in Section 2.6.1, resulting in a calculus I call λ_{LVish} . I will also return to our parallel graph traversal problem and show an solution implemented using the LVish Haskell library.

Finally, the main technical result of this chapter is a proof of quasi-determinism for λ_{LVish} . The key to the proof is a generalized version of the Independence lemma of Chapter 2 that accounts for both freezing and the arbitrary update operations that λ_{LVish} allows.

3.1. LVish, informally

While LVars offer a deterministic programming model that allows communication through a wide variety of data structures, they are not powerful enough to express common algorithmic patterns, like fixpoint computations, that require both positive and negative queries. In this section, I explain our extensions to the LVars model at a high level; Section 3.2 then formalizes them.

3.1.1. Asynchrony through event handlers. Our first extension to LVars is the ability to do asynchronous, event-driven programming through event handlers. An *event* for an LVar can be represented by a lattice element; the event *occurs* when the LVar’s current value reaches a point at or above that lattice element. An *event handler* ties together an LVar with a callback function that is asynchronously invoked whenever some events of interest occur.

To illustrate how event handlers work, consider again the lattice of Figure 2.1(a) from Chapter 2. Suppose that *lv* is an LVar whose states correspond to this lattice. The expression

(Example 3.6) $\text{addHandler } lv \{1, 3, 5, \dots\} (\lambda x. \text{put } lv \ x + 1)$

registers a handler for *lv* that executes the callback function $\lambda x. \text{put } lv \ x + 1$ for each odd number that *lv* is at or above. When Example 3.6 is finished evaluating, *lv* will contain the smallest even number that is at or above what its original value was. For instance, if *lv* originally contains 4, the callback function will be invoked twice, once with 1 as its argument and once with 3. These calls will respectively write $1 + 1 = 2$ and $3 + 1 = 4$ into *lv*; since both writes are ≤ 4 , *lv* will remain 4. On the other hand, if *lv* originally contains 5, then the callback will run three times, with 1, 3, and 5 as its respective arguments, and with the latter of these calls writing $5 + 1 = 6$ into *lv*, leaving *lv* as 6.

In general, the second argument to `addHandler`, which I call an *event set*, is an arbitrary subset Q of the LVar’s lattice, specifying which events should be handled.² Event handlers in the LVish model are somewhat unusual in that they invoke their callback for *all* events in their event set Q that have taken place (*i.e.*, all values in Q less than or equal to the current LVar value), even if those events occurred prior to the handler being registered. To see why this semantics is necessary, consider the following, more subtle example (written in a hypothetical language with a semantics similar to that of λ_{LVar} , but with the addition of `addHandler`):

²Like threshold sets, these event sets are a mathematical modeling tool only; they have no explicit existence in the LVish library implementation.

(Example 3.7)

```

let par _ = put lv 0
      _ = put lv 1
      _ = addHandler lv {0, 1} ( $\lambda x.$  if  $x = 0$  then put lv 2)
in get lv {2}

```

Can Example 3.7 ever block? If a callback only executed for events that arrived after its handler was registered, or only for the largest event in its event set that had occurred, then the example would be nondeterministic: it would block, or not, depending on how the handler registration was interleaved with the puts. By instead executing a handler’s callback once for *each and every* element in its event set below or at the LVar’s value, we guarantee quasi-determinism—and, for Example 3.7, guarantee the result of 2.

The power of event handlers is most evident for lattices that model collections, such as sets. For example, if we are working with lattices of sets of natural numbers, ordered by subset inclusion, then we can write the following function:

$$\text{forEach} = \lambda lv. \lambda f. \text{addHandler } lv \{ \{0\}, \{1\}, \{2\}, \dots \} f$$

Unlike the usual `forEach` function found in functional programming languages, this function sets up a *permanent*, asynchronous flow of data from *lv* into the callback *f*. Functions like `forEach` can be used to set up complex, cyclic data-flow networks, as we will see in Chapter 4.

In writing `forEach`, we consider only the singleton sets to be events of interest, which means that if the value of *lv* is some set like $\{2, 3, 5\}$ then *f* will be executed once for each singleton subset ($\{2\}$, $\{3\}$, $\{5\}$)—that is, once for each element. In Chapter 4, we will see that this kind of event set can be specified in a lattice-generic way, and that it corresponds closely to our implementation strategy.

3.1.2. Quiescence through handler pools. Because event handlers are asynchronous, we need a separate mechanism to determine when they have reached a *quiescent* state, *i.e.*, when all callbacks for the events that have occurred have finished running. Detecting quiescence is crucial for implementing

fixpoint computations. To build flexible data-flow networks, it is also helpful to be able to detect quiescence of multiple handlers simultaneously. Thus, our design includes *handler pools*, which are groups of event handlers whose collective quiescence can be tested.

The simplest way to program with handler pools is to use a pattern like the following:

```
let  $h$  = newPool
  in addHandlerInPool  $h$   $lv$   $Q$   $f$ ;
  quiesce  $h$ 
```

where lv is an LVar, Q is an event set, and f is a callback. Handler pools are created with the `newPool` function, and handlers are registered with `addHandlerInPool`, a variant of `addHandler` that takes a



handler pool as an additional argument. Finally, `quiesce` takes a handler pool as its argument and blocks until all of the handlers in the pool have reached a quiescent state.

Whether or not a handler is quiescent is a non-monotonic property: we can move in and out of quiescence as more writes to an LVar occur, and even if all states at or below the current state have been handled, there is no way to know that more writes will not arrive to move the LVar’s state upwards in the lattice and trigger more callbacks. Early quiescence poses no risk to quasi-determinism, however, because `quiesce` does not yield any information about *which* events have been handled—any such questions must be asked through LVar functions like `get`. In practice, `quiesce` is almost always used together with `freezing`, which I explain next.

3.1.3. Freezing and the “freeze-after” pattern. Our final addition to the LVar model is the ability to *freeze* an LVar, which forbids further changes to it, but in return allows its exact value to be read. We expose freezing through the function `freeze`, which takes an LVar as its sole argument and returns the exact value of the LVar as its result. Any writes that would change the value of a frozen LVar instead raise an exception, and it is the potential for races between such writes and `freeze` that makes the LVish model quasi-deterministic, rather than fully deterministic.

```

traverse g startNode = do
  seen <- newEmptySet
  h <- newHandler seen
  (\node -> do
    mapM (\v -> insert v seen) (neighbors g node)
    return ())
  insert startNode seen -- Kick things off
  quiesce h
  freeze seen

```

Listing 3.1. A deterministic parallel graph traversal that uses `runParThenFreeze`.

Putting all the above pieces together, we arrive at a particularly common pattern of programming in the LVish model:

```

freezeAfter =  $\lambda lv. \lambda Q. \lambda f.$  let h = newPool
                        in addHandlerInPool h lv Q f;
                        quiesce h;
                        freeze lv

```

In this pattern, an event handler is registered for an LVar, subsequently quiesced, and then the LVar is frozen and its exact value is returned.

3.1.4. A parallel graph traversal using handlers, quiescence, and freezing. We can use the new features in LVish to write a parallel graph traversal in the simple fashion shown in Listing 3.1. This code, written using the LVish Haskell library, discovers (in parallel) the set of nodes in a graph `g` reachable from a given node `startNode`, and is guaranteed to produce a deterministic result. It works by first creating a new LVar, `seen`, to represent the set of seen nodes, then adding `startNode` to the set. `newHandler` is a helper function similar to `addHandlerInPool`. It takes the LVar `seen` as its first argument, and its second argument is the callback to be run whenever an event occurs (that is, whenever a new element is added to the set of seen nodes): for each new element that is seen, we look up its neighbors in `g` and then insert each of those elements into the set of seen nodes as well. The computation continues until there are no more events to handle and `quiesce h` returns. We will return to this example in Section 4.2, which discusses the LVish library API in more detail.

3.2. LVish, formally

In this section, I present λ_{LVish} , a core calculus for the LVish programming model. It extends the λ_{LVar} language of Chapter 2. Rather than modeling the full ensemble of event handlers, handler pools, quiescence, and freezing as separate primitives in λ_{LVish} , though, I instead formalize the “freeze-after” pattern—which combined them—directly as a primitive. This simplifies the calculus while still capturing the essence of the programming model. I also generalize the `put` operation to allow the arbitrary *update operations* of Section 2.6.1, which are inflationary and commutative but do not necessarily compute a lub.

3.2.1. Freezing. To model freezing, we need to generalize the notion of the state of an LVar to include information about whether it is “frozen” or not. Thus, in λ_{LVish} an LVar’s *state* is a pair (d, frz) , where d is an element of the set D and frz is a “status bit” of either true or false. A state where frz is false is “unfrozen”, and one where frz is true is “frozen”.

I define an ordering \sqsubseteq_p on LVar states (d, frz) in terms of the given ordering \sqsubseteq on elements of D . Every element of D is “freezable” except \top . Informally:

- Two unfrozen states are ordered according to the given \sqsubseteq ; that is, $(d, \text{false}) \sqsubseteq_p (d', \text{false})$ exactly when $d \sqsubseteq d'$.
- Two frozen states do not have an order, unless they are equal: $(d, \text{true}) \sqsubseteq_p (d', \text{true})$ exactly when $d = d'$.
- An unfrozen state (d, false) is less than or equal to a frozen state (d', true) exactly when $d \sqsubseteq d'$.
- The only situation in which a frozen state is less than an unfrozen state is if the unfrozen state is \top ; that is, $(d, \text{true}) \sqsubseteq_p (d', \text{false})$ exactly when $d' = \top$.

Adding status bits to each element (except \top) of the lattice $(D, \sqsubseteq, \perp, \top)$ results in a new lattice $(D_p, \sqsubseteq_p, \perp_p, \top_p)$. (The p stands for pair, since elements of this new lattice are pairs (d, frz) .) I write \sqcup_p for the lub operation that \sqsubseteq_p induces. Definitions 3.1 and 3.2 and Lemmas 3.1 and 3.2 formalize this notion.

Definition 3.1 (lattice with status bits). Suppose $(D, \sqsubseteq, \perp, \top)$ is a lattice. We define an operation $\text{Freeze}(D, \sqsubseteq, \perp, \top) \triangleq (D_p, \sqsubseteq_p, \perp_p, \top_p)$ as follows:

(1) D_p is a set defined as follows:

$$D_p \triangleq \{(d, \text{frz}) \mid d \in (D - \{\top\}) \wedge \text{frz} \in \{\text{true}, \text{false}\}\} \cup \{(\top, \text{false})\}$$

(2) $\sqsubseteq_p \in \mathcal{P}(D_p \times D_p)$ is a binary relation defined as follows:

$$\begin{aligned} (d, \text{false}) \sqsubseteq_p (d', \text{false}) &\iff d \sqsubseteq d' \\ (d, \text{true}) \sqsubseteq_p (d', \text{true}) &\iff d = d' \\ (d, \text{false}) \sqsubseteq_p (d', \text{true}) &\iff d \sqsubseteq d' \\ (d, \text{true}) \sqsubseteq_p (d', \text{false}) &\iff d' = \top \end{aligned}$$

(3) $\perp_p \triangleq (\perp, \text{false})$.

(4) $\top_p \triangleq (\top, \text{false})$.

Lemma 3.1 (Partition of D_p). *If $(D, \sqsubseteq, \perp, \top)$ is a lattice and $(D_p, \sqsubseteq_p, \perp_p, \top_p) = \text{Freeze}(D, \sqsubseteq, \perp, \top)$, and $X = D - \{\top\}$, then every member of D_p is either*

- (d, false) , with $d \in D$, or
- (x, true) , with $x \in X$.

Proof. Immediate from Definition 3.1. □

Definition 3.2 (lub of states, λ_{LVish}). We define a binary operator $\sqcup_p \in D_p \times D_p \rightarrow D_p$ as follows:

$$\begin{aligned}
 (d_1, \text{false}) \sqcup_p (d_2, \text{false}) &\triangleq (d_1 \sqcup d_2, \text{false}) \\
 (d_1, \text{true}) \sqcup_p (d_2, \text{true}) &\triangleq \begin{cases} (d_1, \text{true}) & \text{if } d_1 = d_2 \\ (\top, \text{false}) & \text{otherwise} \end{cases} \\
 (d_1, \text{false}) \sqcup_p (d_2, \text{true}) &\triangleq \begin{cases} (d_2, \text{true}) & \text{if } d_1 \sqsubseteq d_2 \\ (\top, \text{false}) & \text{otherwise} \end{cases} \\
 (d_1, \text{true}) \sqcup_p (d_2, \text{false}) &\triangleq \begin{cases} (d_1, \text{true}) & \text{if } d_2 \sqsubseteq d_1 \\ (\top, \text{false}) & \text{otherwise} \end{cases}
 \end{aligned}$$

Lemma 3.2 says that if (D, \leq, \perp, \top) is a lattice, then $(D_p, \sqsubseteq_p, \perp_p, \top_p)$ is as well:

Lemma 3.2 (Lattice structure). *If $(D, \sqsubseteq, \perp, \top)$ is a lattice and $(D_p, \sqsubseteq_p, \perp_p, \top_p) = \text{Freeze}(D, \sqsubseteq, \perp, \top)$, then:*

- (1) \sqsubseteq_p is a partial order over D_p .
- (2) Every nonempty finite subset of D_p has a lub.
- (3) \perp_p is the least element of D_p .
- (4) \top_p is the greatest element of D_p .

Therefore $(D_p, \sqsubseteq_p, \perp_p, \top_p)$ is a lattice.

Proof. See Section A.10. □

3.2.2. Update operations. λ_{LVish} generalizes the put operation of λ_{LVar} to a family of operations put_i , in order to allow the generalized update operations of Section 2.6.1 that are commutative and inflationary, but do not necessarily compute a lub. To make this possible we parameterize λ_{LVish} not only by the lattice $(D, \sqsubseteq, \perp, \top)$, but also by a set U of *update operations*, as discussed previously in Section 2.6.1:

Definition 3.3 (set of update operations). Given a lattice $(D, \sqsubseteq, \perp, \top)$ with elements $d \in D$, a *set of update operations* U is a set of functions $u_i : D \rightarrow D$ meeting the following conditions:

- $\forall d, i. d \sqsubseteq u_i(d)$, and
- $\forall d, i, j. u_i(u_j(d)) = u_j(u_i(d))$.

The first of the conditions in Definition 3.3 says that each update operation is inflationary with respect to \sqsubseteq , and the second condition says that update operations commute with each other. Every set of update operations always implicitly contains the identity function.

If we want to recover the original semantics of `put`, we can do so by instantiating U such that there is one u_i for each element d_i of the lattice D , and defining $u_i(d)$ to be $d \sqcup d_i$. On the other hand, if D is a lattice of natural numbers and we want increment-only counters, we can instantiate U to be a singleton set $\{u\}$ where $u(d) = d + 1$. (As described in Section 2.6.1, we could also have a set of update operations $\{u_{(+1)}, u_{(+2)}, \dots\}$, where $u_{(+1)}(d)$ increments d 's contents by one, $u_{(+2)}(d)$ increments by two, and so on.) Update operations are therefore general enough to express lub writes as well as non-idempotent increments. (When a write is specifically a lub write, I will continue to use the notation `put`, without the subscript.)

In λ_{LVish} , the `put` operation took two arguments, a location l and a lattice element d . The `puti` operations take a location l as their only argument, and `puti l` performs the update operation $u_i(l)$ on the contents of l .

More specifically, since l points to a state (d, frz) instead of an element d , `puti l` must perform u_{p_i} , a lifted version of u_i that applies to states. Given U , we define the set U_p of lifted operations as follows:

Definition 3.4 (set of state update operations). Given a set U of update operations u_i , the corresponding *set of state update operations* U_p is a set of functions $u_{p_i} : D_p \rightarrow D_p$ defined as follows:

$$\begin{aligned} u_{p_i}((d, \text{false})) &\triangleq (u_i(d), \text{false}) \\ u_{p_i}((d, \text{true})) &\triangleq \begin{cases} (d, \text{true}) & \text{if } u_i(d) = d \\ (\top, \text{false}) & \text{otherwise} \end{cases} \end{aligned}$$

Because every set U of update operations implicitly contains the identity function, the same is true for the set U_p of state update operations. Furthermore, it is easy to show that state update operations commute, just as update operations do; that is, $\forall d, i, j. u_{p_i}(u_{p_j}(p)) = u_{p_j}(u_{p_i}(p))$.

3.2.3. Stores. During the evaluation of λ_{LVish} programs, a *store* S keeps track of the states of LVars. Each LVar is represented by a binding from a location l , drawn from a set Loc , to its state, which is some pair (d, frz) from the set D_p . The way that stores are handled in λ_{LVish} is very similar to how they are handled in λ_{LVar} , except that store bindings now point to states (d, frz) , that is, elements of D_p , instead of merely to d , that is, elements of D .

Definition 3.5 (store, λ_{LVish}). A *store* is either a finite partial mapping $S : Loc \xrightarrow{\text{fin}} (D_p - \{\top_p\})$, or the distinguished element \top_S .

I use the notation $S[l \mapsto (d, frz)]$ to denote extending S with a binding from l to (d, frz) . If $l \in \text{dom}(S)$, then $S[l \mapsto (d, frz)]$ denotes an update to the existing binding for l , rather than an extension. Another way to denote a store is by explicitly writing out all its bindings, using the notation $[l_1 \mapsto (d_1, frz_1), l_2 \mapsto (d_2, frz_2), \dots]$.

We can lift the \sqsubseteq_p and \sqcup_p operations defined on elements of D_p to the level of stores:

Definition 3.6 (store ordering, λ_{LVish}). A store S is *less than or equal to* a store S' (written $S \sqsubseteq_S S'$) iff:

- $S' = \top_S$, or
- $\text{dom}(S) \subseteq \text{dom}(S')$ and for all $l \in \text{dom}(S)$, $S(l) \sqsubseteq_p S'(l)$.

Definition 3.7 (lub of stores, λ_{LVish}). The lub of two stores S_1 and S_2 (written $S_1 \sqcup_S S_2$) is defined as follows:

- $S_1 \sqcup_S S_2 = \top_S$ if $S_1 = \top_S$ or $S_2 = \top_S$.
- $S_1 \sqcup_S S_2 = \top_S$ if there exists some $l \in \text{dom}(S_1) \cap \text{dom}(S_2)$ such that $S_1(l) \sqcup_p S_2(l) = \top_p$.
- Otherwise, $S_1 \sqcup_S S_2$ is the store S such that:

– $\text{dom}(S) = \text{dom}(S_1) \cup \text{dom}(S_2)$, and

– For all $l \in \text{dom}(S)$:

$$S(l) = \begin{cases} S_1(l) \sqcup_p S_2(l) & \text{if } l \in \text{dom}(S_1) \cap \text{dom}(S_2) \\ S_1(l) & \text{if } l \notin \text{dom}(S_2) \\ S_2(l) & \text{if } l \notin \text{dom}(S_1) \end{cases}$$

If, for example,

$$(d_1, \text{frz}_1) \sqcup_p (d_2, \text{frz}_2) = \top_p,$$

then

$$[l \mapsto (d_1, \text{frz}_1)] \sqcup_S [l \mapsto (d_2, \text{frz}_2)] = \top_S.$$

Just as a store containing a binding $l \mapsto \top$ can never arise during the execution of a λ_{LVar} program, a store containing a binding $l \mapsto (\top, \text{frz})$ can never arise during the execution of a λ_{LVish} program. An attempted write that would take the value of l to (\top, false) —that is, \top_p —will raise an error, and there is no (\top, true) element of D_p .

3.2.4. λ_{LVish} : syntax and semantics. The syntax of λ_{LVish} appears in Figure 3.1, and Figures 3.2 and 3.3 together give the operational semantics. As with λ_{LVar} in Chapter 2, both the syntax and semantics are parameterized by the lattice $(D, \sqsubseteq, \perp, \top)$, and the operational semantics is split into two parts, a *reduction semantics*, shown in Figure 2.3, and a *context semantics*, shown in Figure 2.4. The reduction semantics is also parameterized by the set U of update operations.

The λ_{LVish} grammar has most of the expression forms of λ_{LVar} : variables, values, application expressions, get expressions, and new. Instead of put expressions, it has put_i expressions, which are the interface to the specified set of update operations. λ_{LVish} also adds two new language forms, the **freeze** expression and the **freeze – after – with** expression, which I discuss in more detail below.

Values in λ_{LVish} include all those from λ_{LVar} —the unit value $()$, lattice elements d , locations l , threshold sets P , and λ expressions—as well as states p , which are pairs (d, frz) , and event sets Q . Instead of

Given a lattice $(D, \sqsubseteq, \perp, \top)$ with elements $d \in D$:

configurations	σ	$::=$	$\langle S; e \rangle \mid \mathbf{error}$
expressions	e	$::=$	$x \mid v \mid ee \mid \mathbf{get}_i ee \mid \mathbf{put}_i e \mid \mathbf{new} \mid \mathbf{freeze} e$ $\mid \mathbf{freeze} e \text{ after } e \text{ with } e$ $\mid \mathbf{freeze} l \text{ after } Q \text{ with } \lambda x. e, \{e, \dots\}, H$
values	v	$::=$	$() \mid d \mid p \mid l \mid P \mid Q \mid \lambda x. e$
threshold sets	P	$::=$	$\{p_1, p_2, \dots\}$
event sets	Q	$::=$	$\{d_1, d_2, \dots\}$
“handled” sets	H	$::=$	$\{d_1, \dots, d_n\}$
stores	S	$::=$	$[l_1 \mapsto p_1, \dots, l_n \mapsto p_n] \mid \top_S$
states	p	$::=$	(d, \mathbf{frz})
status bits	\mathbf{frz}	$::=$	$\mathbf{true} \mid \mathbf{false}$
evaluation contexts	E	$::=$	$[] \mid Ee \mid eE \mid \mathbf{get} Ee \mid \mathbf{get} eE \mid \mathbf{put}_i E$ $\mid \mathbf{freeze} E \mid \mathbf{freeze} E \text{ after } e \text{ with } e$ $\mid \mathbf{freeze} e \text{ after } E \text{ with } e \mid \mathbf{freeze} e \text{ after } e \text{ with } E$ $\mid \mathbf{freeze} v \text{ after } v \text{ with } v, \{e, \dots, E, e, \dots\}, H$

Figure 3.1. Syntax for λ_{LVish} .

T , I now use the metavariable P for threshold sets, in keeping with the fact that in λ_{LVish} , members of threshold sets are states p .

As with λ_{LVar} , the λ_{LVish} context relation \mapsto has only one rule, E-Eval-Ctxt, which allows us to apply reductions within a context. The rule itself is identical to the corresponding rule in λ_{LVar} , although the set of evaluation contexts that the metavariable E ranges over is different.

Given a lattice $(D, \sqsubseteq, \perp, \top)$ with elements $d \in D$, and a set of U of update operations $u_i : D \rightarrow D$:

$$\text{incomp}(P) \triangleq \forall p_1, p_2 \in P. (p_1 \neq p_2 \Rightarrow p_1 \sqcup_p p_2 = \top_p) \quad \boxed{\sigma \hookrightarrow \sigma'}$$

E-Beta		E-New		
$\frac{}{\langle S; (\lambda x. e) v \rangle \hookrightarrow \langle S; e[x := v] \rangle}$		$\frac{}{\langle S; \mathbf{new} \rangle \hookrightarrow \langle S[l \mapsto (\perp, \mathbf{false})]; l \rangle} \quad (l \notin \mathit{dom}(S))$		
E-Put	$S(l) = p_1 \quad u_{p_i}(p_1) \neq \top_p$	E-Put-Err	E-Get	
	$\langle S; \mathbf{put}_i l \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$	$S(l) = p_1 \quad u_{p_i}(p_1) = \top_p$		
$\langle S; \mathbf{put}_i l \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$		$S(l) = p_1 \quad \mathit{incomp}(P) \quad p_2 \in P \quad p_2 \sqsubseteq_p p_1$		
E-Freeze-Init		$\langle S; \mathbf{get} l P \rangle \hookrightarrow \langle S; p_2 \rangle$		
$\langle S; \mathbf{freeze} l \text{ after } Q \text{ with } \lambda x. e \rangle \hookrightarrow \langle S; \mathbf{freeze} l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle$				
E-Spawn-Handler		$S(l) = (d_1, \mathit{frz}_1) \quad d_2 \sqsubseteq d_1 \quad d_2 \notin H \quad d_2 \in Q$		
$\langle S; \mathbf{freeze} l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \hookrightarrow \langle S; \mathbf{freeze} l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle$				
E-Freeze-Final	$S(l) = (d_1, \mathit{frz}_1) \quad \forall d_2. (d_2 \sqsubseteq d_1 \wedge d_2 \in Q \Rightarrow d_2 \in H)$	E-Freeze-Simple		$S(l) = (d_1, \mathit{frz}_1)$
$\langle S; \mathbf{freeze} l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \hookrightarrow \langle S[l \mapsto (d_1, \mathbf{true})]; d_1 \rangle$		$\langle S; \mathbf{freeze} l \rangle \hookrightarrow \langle S[l \mapsto (d_1, \mathbf{true})]; d_1 \rangle$		

Figure 3.2. Reduction semantics for λ_{LVish} .

$$\sigma \mapsto \sigma'$$

$$\frac{\text{E-Eval-Ctxt} \quad \langle S; e \rangle \hookrightarrow \langle S'; e' \rangle}{\langle S; E[e] \rangle \mapsto \langle S'; E[e'] \rangle}$$

Figure 3.3. Context semantics for λ_{LVish} .

3.2.5. Semantics of new, put_i, and get. Because of the addition of status bits to the semantics, the E-New and E-Get rules have changed slightly from their counterparts in λ_{LVar} :

- **new** (implemented by the E-New rule) extends the store with a binding for a new LVar whose initial state is (\perp, false) , and returns the location l of that LVar (*i.e.*, a pointer to the LVar).
- **get** (implemented by the E-Get rule) performs a blocking threshold read. It takes a pointer to an LVar and a threshold set P , which is a non-empty set of LVar states that must be pairwise incompatible, expressed by the premise $\text{incomp}(P)$. A threshold set P is pairwise incompatible iff the lub of any two distinct elements in P is \top_p . If the LVar's state p_1 in the lattice is at or above some $p_2 \in P$, the get operation unblocks and returns p_2 .

λ_{LVish} replaces the λ_{LVar} put operation with the put_i operation, which is actually a set of operations that are the interface to the provided update operations u_i . For each update operation u_i , put_i (implemented by the E-Put rule) takes a pointer to an LVar and updates the LVar's state to the result of calling u_{p_i} on the LVar's current state, potentially pushing the state of the LVar upward in the lattice. The E-Put-Err rule applies when a put_i operation would take the state of an LVar to \top_p ; in that case, the semantics steps to **error**.

3.2.6. Freezing and the freeze — after — with primitive. The E-Freeze-Init, E-Spawn-Handler, E-Freeze-Final, and E-Freeze-Simple rules are all new additions to λ_{LVish} . The E-Freeze-Simple rule gives the semantics for the freeze expression, which takes an LVar as argument and immediately freezes and returns its contents.

More interesting is the `freeze` — `after` — `with` primitive, which models the “freeze-after” pattern I described in Section 3.1.3. The expression

$$\text{freeze } lv \text{ after } Q \text{ with } f$$

has the following semantics:

- It attaches the callback f to the LVar lv . The callback will be executed, once, for each element of the event set Q that the LVar’s state reaches or surpasses. The callback is a function that takes a lattice element as its argument. Its return value is ignored, so it runs solely for effect. For instance, a callback might itself do a `puti` to the LVar to which it is attached, triggering yet more callbacks.
- If execution reaches a point where there are no more elements of Q left to handle and no callbacks still running, then we have reached a quiescent state, the LVar lv is frozen, and its *exact* state is returned (rather than an underapproximation of the state, as with `get`).

To keep track of the running callbacks, λ_{LVish} includes an auxiliary form,

$$\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H$$

where:

- The value l is the LVar being handled/frozen;
- The set Q (a subset of the lattice D) is the event set;
- The value $\lambda x. e_0$ is the callback function;
- The set of expressions $\{e, \dots\}$ is the set of running callbacks; and
- The set H (a subset of the lattice D) represents those values in Q for which callbacks have already been launched; we call H the “handled” set.

Due to λ_{LVish} ’s use of evaluation contexts, any running callback can execute at any time, as if each is running in its own thread. The rule E-Spawn-Handler launches a new callback thread any time the

LVar's current value is at or above some element in Q that has not already been handled. This step can be taken nondeterministically at any time after the relevant put_i has been performed.



The rule E-Freeze-Final detects quiescence by checking that two properties hold. First, every event of interest (lattice element in Q) that has occurred (is bounded by the current LVar state) must be handled (be in H). Second, all existing callback threads must have terminated with a value. In other words, every enabled callback has completed. When such a quiescent state is detected,

E-Freeze-Final freezes the LVar's state. Like E-Spawn-Handler, the rule can fire at any time, nondeterministically, that the handler appears quiescent—a transient property! But after being frozen, any further put_i updates that would have enabled additional callbacks will instead fault, causing the program to step to **error**.

Therefore, freezing is a way of “betting” that once a collection of callbacks have completed, no further updates that change the LVar's value will occur. For a given run of a program, either all updates to an LVar arrive before it has been frozen, in which case the value returned by freeze — after — with is the lub of those values, or some update arrives after the LVar has been frozen, in which case the program will fault. And thus we have arrived at *quasi-determinism*: a program will always either evaluate to the same answer or it will fault.

To ensure that we will win our bet, we need to guarantee that quiescence is a *permanent* state, rather than a transient one—that is, we need to perform all updates either prior to freeze — after — with, or by the callback function within it (as will be the case for fixpoint computations). In practice, freezing is usually the very last step of an algorithm, permitting its result to be extracted. As we will see in Section 4.2.5, our LVish library provides a special `runParThenFreeze` function that does so, and thereby guarantees full determinism.

3.3. Proof of quasi-determinism for λ_{LVish}

In this section, I give a proof of quasi-determinism for λ_{LVish} that formalizes the claim made earlier in this chapter: that, for a given program, although some executions may raise exceptions, all executions that produce a final result will produce the same final result.

The quasi-determinism theorem I show says that if two executions starting from a configuration σ terminate in configurations σ' and σ'' , then either σ' and σ'' are the same configuration (up to a permutation on locations), or one of them is **error**. As with the determinism proof for λ_{LVar} in Section 2.5, quasi-determinism follows from a series of supporting lemmas. The basic structure of the proof follows that of the λ_{LVar} determinism proof closely. However, instead of the Independence property that I showed for λ_{LVar} (Lemma 2.5), here I prove a more general property, Generalized Independence (Lemma 3.7), that accounts for the presence of both freezing and arbitrary update operations in λ_{LVish} . Also, in the setting of λ_{LVish} , the Strong Local Confluence property (Lemma 2.8) becomes Strong Local *Quasi*-Confluence (Lemma 3.10), which allows the possibility of an **error** result, and the quasi-confluence lemmas that follow—Strong One-sided Quasi-Confluence (Lemma 3.11), Strong Quasi-Confluence (Lemma 2.10), and Quasi-Confluence (Lemma 2.11)—all follow this pattern as well.

3.3.1. Permutations and permutability. As with λ_{LVar} , the λ_{LVish} language is nondeterministic with respect to the names of locations it allocates. We therefore prove quasi-determinism up to a permutation on locations. We can reuse the definition of a permutation verbatim from Section 2.5.1:

Definition 3.8 (permutation, λ_{LVish}). A *permutation* is a function $\pi : \text{Loc} \rightarrow \text{Loc}$ such that:

- (1) it is invertible, that is, there is an inverse function $\pi^{-1} : \text{Loc} \rightarrow \text{Loc}$ with the property that $\pi(l) = l'$ iff $\pi^{-1}(l') = l$; and
- (2) it is the identity on all but finitely many elements of Loc .

We can lift π to apply expressions, stores, and configurations. Because expressions and stores are defined slightly differently in λ_{LVish} than they are in λ_{LVar} , we must update our definitions of permutation of a store and permutation of an expression:

Definition 3.9 (permutation of an expression, λ_{LVish}). A *permutation* of an expression e is a function π defined as follows:

$$\begin{aligned}
\pi(x) &\triangleq x \\
\pi(()) &\triangleq () \\
\pi(d) &\triangleq d \\
\pi(p) &\triangleq p \\
\pi(l) &\triangleq \pi(l) \\
\pi(P) &\triangleq P \\
\pi(Q) &\triangleq Q \\
\pi(\lambda x. e) &\triangleq \lambda x. \pi(e) \\
\pi(e_1 e_2) &\triangleq \pi(e_1) \pi(e_2) \\
\pi(\text{get } e_1 e_2) &\triangleq \text{get } \pi(e_1) \pi(e_2) \\
\pi(\text{put}_i e) &\triangleq \text{put}_i \pi(e) \\
\pi(\text{new}) &\triangleq \text{new} \\
\pi(\text{freeze } e) &\triangleq \text{freeze } \pi(e) \\
\pi(\text{freeze } e_1 \text{ after } e_2 \text{ with } e_3) &\triangleq \text{freeze } \pi(e_1) \text{ after } \pi(e_2) \text{ with } \pi(e_3) \\
\pi(\text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{e, \dots\}, H) &\triangleq \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e), \{\pi(e), \dots\}, H
\end{aligned}$$

Definition 3.10 (permutation of a store, λ_{LVish}). A *permutation* of a store S is a function π defined as follows:

$$\begin{aligned}
\pi(\top_S) &\triangleq \top_S \\
\pi([l_1 \mapsto p_1, \dots, l_n \mapsto p_n]) &\triangleq [\pi(l_1) \mapsto p_1, \dots, \pi(l_n) \mapsto p_n]
\end{aligned}$$

And the definition of permutation of a configuration is as it was before:

Definition 3.11 (permutation of a configuration, λ_{LVish}). A *permutation* of a configuration $\langle S; e \rangle$ is a function π defined as follows: if $\langle S; e \rangle = \mathbf{error}$, then $\pi(\langle S; e \rangle) \triangleq \mathbf{error}$; otherwise, $\pi(\langle S; e \rangle) \triangleq \langle \pi(S); \pi(e) \rangle$.

We can then prove a Permutability lemma for λ_{LVish} , which says that a configuration σ can step to σ' exactly when $\pi(\sigma)$ can step to $\pi(\sigma')$.

Lemma 3.3 (Permutability, λ_{LVish}). *For any finite permutation π ,*

- (1) $\sigma \longrightarrow \sigma'$ if and only if $\pi(\sigma) \longrightarrow \pi(\sigma')$.
- (2) $\sigma \longmapsto \sigma'$ if and only if $\pi(\sigma) \longmapsto \pi(\sigma')$.

Proof. Similar to the proof of Lemma 2.1 (Permutability for λ_{LVar}); see Section A.11. □

3.3.2. Internal Determinism. In Chapter 2, we saw that the reduction semantics for λ_{LVar} is internally deterministic: that is, if a configuration can step by the reduction semantics, there is only one rule by which it can step, and only one configuration to which it can step, modulo location names. For λ_{LVish} , we can also show an internal determinism property, but with a slight additional wrinkle.

In λ_{LVish} , the E-Spawn-Handler rule picks out an eligible element d_2 from the set Q (that is, an event) and launches a new callback thread to handle that event. But, since there could be more than one eligible element in Q that E-Spawn-Handler could choose, the choice of event is a source of nondeterminism. Therefore, we show that λ_{LVish} is internally deterministic modulo *choice of events*, as well as modulo location names. This property will be useful to us later on in the proof of Strong Local Quasi-Confluence (Lemma 3.10).

Lemma 3.4 (Internal Determinism, λ_{LVish}). *If $\sigma \longrightarrow \sigma'$ and $\sigma \longrightarrow \sigma''$, then there is a permutation π such that $\sigma' = \pi(\sigma'')$, modulo choice of events.*

Proof. Straightforward by cases on the rule of the reduction semantics by which σ steps to σ' ; the only interesting case is for the E-New rule. See Section A.12. □

3.3.3. Locality. Just as with the determinism proof for λ_{LVar} , proving quasi-determinism for λ_{LVish} will require us to handle expressions that can decompose into redex and context in multiple ways. An expression e such that $e = E_1[e_1] = E_2[e_2]$ can step in two different ways by the E-Eval-Ctxt rule: $\langle S; E_1[e_1] \rangle \mapsto \langle S_1; E_1[e'_1] \rangle$, and $\langle S; E_2[e_2] \rangle \mapsto \langle S_2; E_2[e'_2] \rangle$.

The Locality lemma says that the \mapsto relation acts “locally” in each of these steps. The statement of the Locality lemma is the same as that of Lemma 2.3 (Locality for λ_{LVar}), but the proof must account for the set of possible evaluation contexts in λ_{LVish} being different (and larger) than the set of evaluation contexts in λ_{LVar} .

Lemma 3.5 (Locality, λ_{LVish}). *If $\langle S; E_1[e_1] \rangle \mapsto \langle S_1; E_1[e'_1] \rangle$ and $\langle S; E_2[e_2] \rangle \mapsto \langle S_2; E_2[e'_2] \rangle$ and $E_1[e_1] = E_2[e_2]$, then:*

If $E_1 \neq E_2$, then there exist evaluation contexts E'_1 and E'_2 such that:

- $E'_1[e_1] = E_2[e'_2]$, and
- $E'_2[e_2] = E_1[e'_1]$, and
- $E'_1[e'_1] = E'_2[e'_2]$.

Proof. Let $e = E_1[e_1] = E_2[e_2]$. The proof is by induction on the structure of the expression e . See Section A.13. □

3.3.4. Monotonicity. The Monotonicity lemma says that, as evaluation proceeds according to the \hookrightarrow relation, the store can only grow with respect to the \sqsubseteq_S ordering.

Lemma 3.6 (Monotonicity, λ_{LVish}). *If $\langle S; e \rangle \hookrightarrow \langle S'; e' \rangle$, then $S \sqsubseteq_S S'$.*

Proof. Straightforward by cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$. The interesting cases are for the E-New and E-Put rules. See Section A.14. □

3.3.5. Generalized Independence. Recall from Chapter 2 that in order to prove determinism for λ_{LVar} , we needed to establish a “frame property” that captures the idea that independent effects commute with each other. For λ_{LVar} , the Independence lemma (Lemma 2.5) established that property. It shows that, if a configuration $\langle S; e \rangle$ can step to $\langle S'; e' \rangle$, then it is possible to “frame on” an additional store S'' without interfering with the ability to take a step—that is, $\langle S \sqcup_S S''; e \rangle \longrightarrow \langle S' \sqcup_S S''; e' \rangle$, subject to certain restrictions on S'' .

For λ_{LVish} , we need to establish a similar frame property. However, since we have generalized from put to put_i , we also need to generalize our frame property. In fact, the original Lemma 2.5 does not hold for λ_{LVish} . As an example, consider an LVar whose states form a lattice $\perp < 0 < 1 < \top$. Consider the transition

$$\langle [l \mapsto (0, \text{false})]; \text{put}_i l \rangle \longrightarrow \langle [l \mapsto (1, \text{false})]; () \rangle,$$

where the update operation u_i happens to increment its argument by one. Now suppose that we wish to “frame” the store $[l \mapsto (1, \text{false})]$ onto this transition using Lemma 2.5; that is, we wish to show that

$$\langle [l \mapsto (0, \text{false})] \sqcup_S [l \mapsto (1, \text{false})]; \text{put}_i l \rangle \longrightarrow \langle [l \mapsto (1, \text{false})] \sqcup_S [l \mapsto (1, \text{false})]; () \rangle.$$

We know that $[l \mapsto (1, \text{false})] \sqcup_S [l \mapsto (1, \text{false})] \neq \top_S$, which is required to be able to apply Lemma 2.5. Furthermore, $[l \mapsto (1, \text{false})]$ is non-conflicting with the original transition, since no new locations are allocated between $[l \mapsto (0, \text{false})]$ and $[l \mapsto (1, \text{false})]$. But it is *not* the case that $\langle [l \mapsto (0, \text{false})] \sqcup_S [l \mapsto (1, \text{false})]; \text{put}_i l \rangle$ steps to $\langle [l \mapsto (1, \text{false})] \sqcup_S [l \mapsto (1, \text{false})]; () \rangle$, since $u_{p_i}((S \sqcup_S S'')(l)) = \top_p$. (As before, u_{p_i} is the update operation u_i , lifted from lattice elements d to states (d, frz) .)

What went wrong here? The problem is that, as previously discussed in Section 2.6.1, lub operations do not necessarily commute with arbitrary update operations. In λ_{LVar} , where the only “update operation” is a lub write performed via put , it is fine that the Independence lemma uses a lub operation to frame S'' onto the transition. For λ_{LVish} , though, we need to state our frame property in a way that will allow it to accommodate any update operation from the given set U .

Therefore, I define a *store update operation* U_S to be a function from stores to stores that can add new bindings, update the contents of existing locations using operations u_i from the given set U of update operations (or, more specifically, their lifted versions u_{p_i}), or freeze the contents of existing locations.

Definition 3.12 (store update operation). Given a lattice $(D, \sqsubseteq, \top, \perp)$ and a set of state update operations U_p , a *store update operation* is a function U_S from stores to stores such that:

- $\text{dom}(U_S(S)) \supseteq \text{dom}(S)$;
- for each $l \in \text{dom}(S)$, either:
 - $(U_S(S))(l) = u_{p_i}(S(l))$, where $u_{p_i} \in U_p$, or
 - $(U_S(S))(l) = (d, \text{true})$, where $S(l) = (d, \text{frz})$; and
- for each $l \in \text{dom}(U_S(S))$ that is not a member of $\text{dom}(S)$, $(U_S(S))(l) = (d, \text{frz})$ for some $d \in D$.

Definition 3.12 says that applying U_S to S either updates (using some $u_{p_i} \in U_p$) or freezes the contents of each $l \in \text{dom}(S)$. Since the identity function is always implicitly a member of U_p , U_S can act as the identity on the contents of locations. U_S can also add new bindings to the store it operates on; however, it cannot change existing location names.

With Definition 3.12 in hand, we can state a more general version of the Independence lemma:

Lemma 3.7 (Generalized Independence). *If $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ (where $\langle S'; e' \rangle \neq \mathbf{error}$), then we have that:*

$$\langle U_S(S); e \rangle \longrightarrow \langle U_S(S'); e' \rangle,$$

where U_S is a store update operation meeting the following conditions:


- U_S is non-conflicting with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$,
- $U_S(S') \neq \top_S$, and
- U_S is freeze-safe with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$.

Proof. By cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$. The interesting cases are for the E-New, E-Put, E-Freeze-Final, and E-Freeze-Simple rules. See Section A.15. \square

Lemma 3.7 has three preconditions on the store update operation U_S , two of which mirror the two preconditions on S'' from the original Independence lemma: the requirement that $U_S(S') \neq \top_S$, and the requirement that U_S is non-conflicting with the transition from $\langle S; e \rangle$ to $\langle S'; e' \rangle$. Definition 3.13 revises our previous definition of “non-conflicting” to apply to store update operations. It says that U_S is non-conflicting with the transition $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ if, for all locations allocated in the transition, U_S does not interfere with those locations. For instance, if l is allocated in the transition from $\langle S; e \rangle$ to $\langle S'; e' \rangle$, then $l \notin \text{dom}(U_S(S))$ (that is, U_S cannot add a binding at l to S), and $(U_S(S'))(l) = S'(l)$ (that is, U_S cannot update the contents of l in S').

Definition 3.13 (non-conflicting store update operation). A store update operation U_S is *non-conflicting* with the transition $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ iff, for all $l \in (\text{dom}(S') - \text{dom}(S))$, U_S neither creates new bindings at l nor updates existing bindings at l .

The third precondition on U_S has to do with freezing: U_S must be *freeze-safe* with the transition $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$, which means that, for any locations that change in status (that is, become frozen) during the transition, U_S cannot update the contents of those locations. This precondition is only needed in the E-Freeze-Final and E-Freeze-Simple cases, and it has the effect of ruling out interference



from freezing. (Note that U_S need not avoid updating the contents of locations that are *already* frozen before the transition takes place. This corresponds to the fact that, if an LVar is already frozen, arbitrary updates to it *do*, in fact, commute with freeze operations on it—those later freeze operations will have no effect, and updates will either have no effect or raise an error.)

Definition 3.14 (freeze-safe store update operation). A store update operation U_S is *freeze-safe* with the transition $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ iff, for all locations l such that $S(l) = (d_1, \text{frz}_1)$ and $S'(l) = (d_2, \text{frz}_2)$ and $\text{frz}_1 \neq \text{frz}_2$, U_S does not update the contents of l (that is, either it freezes the contents of l but has no other effect on them, or it acts as the identity on the contents of l).

The two changes we have made to the Independence lemma—the use of U_S , and the requirement that U_S be freeze-safe with the transition in question—are orthogonal to each other, in accordance with the fact that *arbitrary update operations* are an orthogonal language feature to *freezing*. A version of λ_{LVish} that had freezing, but retained the lub semantics of `put` in λ_{LVar} , could use the old formulation of the Independence lemma, taking the lub of the original stores and a frame store S'' , but it would still need to have a requirement on S'' to rule out interference from freezing. On the other hand, a version of the language *without* freezing, but *with* arbitrary updates, would still use U_S but could leave out the requirement that it be freeze-safe (since the requirement would be vacuously true anyway). I make particular note of the orthogonality of freezing and arbitrary updates because freezing introduces quasi-determinism, while arbitrary updates do not.³

Finally, although it no longer uses an explicit “frame” store, we can still think of Lemma 3.7 as a frame property; in fact, it is reminiscent of the generalized frame rule of the “Views” framework [17], which I discuss in more detail in Section 6.5.

3.3.6. Generalized Clash. The Generalized Clash lemma, Lemma 3.8, is similar to the Generalized Independence lemma, but handles the case where $U_S(S') = \top_S$. It establishes that, in that case, $\langle U_S(S); e \rangle$ steps to **error** in at most one step.

³To rigorously show that arbitrary updates retain full determinism and not merely quasi-determinism, I would need to define yet another language, one that generalizes `put` to `puti` but does not introduce freezing, and then prove determinism for *that* language. Instead, I hope to informally convince you that the quasi-determinism in λ_{LVish} comes from freezing, rather than from arbitrary updates.

Lemma 3.8 (Generalized Clash). *If $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ (where $\langle S'; e' \rangle \neq \mathbf{error}$), then we have that:*

$$\langle U_S(S); e \rangle \longrightarrow^i \mathbf{error},$$

where $i \leq 1$ and where U_S is a store update operation meeting the following conditions:

- U_S is non-conflicting with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$,
- $U_S(S') = \top_S$, and
- U_S is freeze-safe with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$.

Proof. By cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$. Since $\langle S'; e' \rangle \neq \mathbf{error}$, we do not need to consider the E-Put-Err rule. See Section A.16. \square

3.3.7. Error Preservation. Lemma 3.9, Error Preservation, is the λ_{LVish} counterpart of Lemma 2.7 from Chapter 2. It says that if a configuration $\langle S; e \rangle$ steps to \mathbf{error} , then evaluating e in the context of some larger store will also result in \mathbf{error} .

Lemma 3.9 (Error Preservation, λ_{LVish}). *If $\langle S; e \rangle \longrightarrow \mathbf{error}$ and $S \sqsubseteq_S S'$, then $\langle S'; e \rangle \longrightarrow \mathbf{error}$.*

Proof. Suppose $\langle S; e \rangle \longrightarrow \mathbf{error}$ and $S \sqsubseteq_S S'$. We are required to show that $\langle S'; e \rangle \longrightarrow \mathbf{error}$.

By inspection of the operational semantics, the only rule by which $\langle S; e \rangle$ can step to \mathbf{error} is E-Put-Err. Hence $e = \text{put}_i l$. From the premises of E-Put-Err, we have that $S(l) = p_1$. Since $S \sqsubseteq_S S'$, it must be the case that $S'(l) = p'_1$, where $p_1 \sqsubseteq_p p'_1$. Since $u_{p_i}(p_1) = \top_p$, we have that $u_{p_i}(p'_1) = \top_p$. Hence, by E-Put-Err, $\langle S'; \text{put}_i l \rangle \longrightarrow \mathbf{error}$, as we were required to show. \square

3.3.8. Quasi-Confluence. Lemma 3.10 says that if a configuration σ can step to configurations σ_a and σ_b , then one of two possibilities is true: either there exists a configuration σ_c that σ_a and σ_b can each reach in at most one step, modulo a permutation on locations, or at least one of σ_a or σ_b steps to \mathbf{error} . Lemmas 3.11 and 3.12 then generalize that result to arbitrary numbers of steps.

Lemma 3.10 (Strong Local Quasi-Confluence). *If $\sigma \mapsto \sigma_a$ and $\sigma \mapsto \sigma_b$, then either:*

- (1) *there exist σ_c, i, j, π such that $\sigma_a \mapsto^i \sigma_c$ and $\pi(\sigma_b) \mapsto^j \sigma_c$ and $i \leq 1$ and $j \leq 1$, or*
- (2) *$\sigma_a \mapsto \mathbf{error}$ or $\sigma_b \mapsto \mathbf{error}$.*

Proof. As in the proof of Strong Local Confluence for λ_{LVar} (Lemma 2.8), since the original configuration σ can step in two different ways, its expression decomposes into redex and context in two different ways: $\sigma = \langle S; E_a[e_{a_1}] \rangle = \langle S; E_b[e_{b_1}] \rangle$, where $E_a[e_{a_1}] = E_b[e_{b_1}]$, but E_a and E_b may differ and e_{a_1} and e_{b_1} may differ. In the special case where $E_a = E_b$, the result follows by Internal Determinism (Lemma 3.4).

If $E_a \neq E_b$, we can apply the Locality lemma (Lemma 3.5); at a high level, it shows that e_{a_1} and e_{b_1} can be evaluated independently within their contexts. The proof is then by a double case analysis on the rules of the reduction semantics by which $\langle S; e_{a_1} \rangle$ steps and by which $\langle S; e_{b_1} \rangle$ steps. In order to combine the results of the two independent steps, the proof makes use of the Generalized Independence lemma 3.7. In almost every case, there does exist a σ_c to which σ_a and σ_b both step; the only cases in which we need to resort to the **error** possibility are those in which one step is by E-Put and the other is by E-Freeze-Final or E-Freeze-Simple—that is, the situations in which a write-after-freeze error is possible. See Section A.17. □

Lemma 3.11 (Strong One-Sided Quasi-Confluence). *If $\sigma \mapsto \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq m$, then either:*

- (1) *there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq 1$, or*
- (2) *there exists $k \leq m$ such that $\sigma' \mapsto^k \mathbf{error}$, or there exists $k \leq 1$ such that $\sigma'' \mapsto^k \mathbf{error}$.*

Proof. By induction on m ; see Section A.18. □

Lemma 3.12 (Strong Quasi-Confluence). *If $\sigma \mapsto^n \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq n$ and $1 \leq m$, then either:*

- (1) *there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq n$, or*
- (2) *there exists $k \leq m$ such that $\sigma' \mapsto^k \mathbf{error}$, or there exists $k \leq n$ such that $\sigma'' \mapsto^k \mathbf{error}$.*

Proof. By induction on n ; see Section A.19. □

Lemma 3.13 (Quasi-Confluence). *If $\sigma \mapsto^* \sigma'$ and $\sigma \mapsto^* \sigma''$, then either:*

- (1) *there exist σ_c and π such that $\sigma' \mapsto^* \sigma_c$ and $\pi(\sigma'') \mapsto^* \sigma_c$, or*
- (2) *$\sigma' = \mathbf{error}$ or $\sigma'' = \mathbf{error}$.*

Proof. Strong Quasi-Confluence (Lemma 3.12) implies Quasi-Confluence. □

3.3.9. Quasi-Determinism. The Quasi-Determinism theorem, Theorem 3.1, is a straightforward result of Lemma 3.13. It says that if two executions starting from a configuration σ terminate in configurations σ' and σ'' , then σ' and σ'' are the same configuration, or one of them is **error**.

Theorem 3.1 (Quasi-Determinism). *If $\sigma \mapsto^* \sigma'$ and $\sigma \mapsto^* \sigma''$, and neither σ' nor σ'' can take a step, then either:*

- (1) *there exists π such that $\sigma' = \pi(\sigma'')$, or*
- (2) *$\sigma' = \mathbf{error}$ or $\sigma'' = \mathbf{error}$.*

Proof. By Lemma 3.13, one of the following two cases applies:

- (1) There exists σ_c and π such that $\sigma' \mapsto^* \sigma_c$ and $\pi(\sigma'') \mapsto^* \sigma_c$. Since σ' cannot step, we must have $\sigma' = \sigma_c$.

By Lemma 3.3 (Permutability), σ'' can step iff $\pi(\sigma'')$ can step, so since σ'' cannot step, $\pi(\sigma'')$ cannot step either.

Hence we must have $\pi(\sigma'') = \sigma_c$. Since $\sigma' = \sigma_c$ and $\pi(\sigma'') = \sigma_c$, $\sigma' = \pi(\sigma'')$.

- (2) $\sigma' = \mathbf{error}$ or $\sigma'' = \mathbf{error}$, and so the result is immediate.

□

3.3.10. Discussion: quasi-determinism in practice. The quasi-determinism result for λ_{LVish} shows that it is not possible to get multiple “answers” from the same program: every run will either produce the same answer or an error. Importantly, this property is true not only for programs that use the freeze-after pattern expressed by the `freeze — after — with primitive`, but even those that freeze in arbitrary places using the simpler `freeze` primitive. This means that in practice, in a programming model based on LVars with freezing and handlers, even a program that fails to ensure quiescence (introducing the possibility of a race between a `put` and a `freeze`) cannot produce multiple non-**error** answers.

Therefore the LVish programming model is fundamentally different from one in which the programmer must manually insert synchronization barriers to prevent data races. In that kind of a model, a program with a misplaced synchronization barrier can be fully nondeterministic, producing multiple observable answers. In the LVish model, the worst that can happen is that the program raises an error. Moreover, in the LVish model, an **error** result *always* means that there is an undersynchronization bug in the program, and in principle the error message can even specify exactly which write operation happened after which freeze operation, making it easier to debug the race.

However, if we *can* ensure that an LVar is only ever frozen *after* all writes to that LVar have completed, then we can guarantee full determinism, because we will have ruled out races between write operations and freeze operations. In the next chapter, I discuss how the LVish Haskell library enforces this “freeze after writing” property.

CHAPTER 4

The LVish library

We want the programming model of Chapters 2 and 3 to be realizable in practice. If the determinism guarantee offered by LVars is to do us any good, however, we need to add LVars to a programming model that is already deterministic. The *monad-par* Haskell library [36], which provides the `Par` monad, is one such deterministic parallel programming model. Haskell is in general an appealing substrate for guaranteed-deterministic parallel programming models because it is pure by default, and its type system enforces separation of pure and effectful code via monads. In order for the determinism guarantee of any parallel programming model to hold, the only side effects allowed must be those sanctioned by the programming model.¹ In the case of the basic LVars model of Chapter 2, those allowed effects are put and get operations on LVars; Chapter 3 adds the freeze operation and arbitrary update operations to the set of allowed effects. Implementing these operations as monadic effects in Haskell makes it possible to provide compile-time guarantees about determinism and quasi-determinism, because we can use Haskell’s type system to ensure that the only side effects programs can perform are those that we have chosen to allow.

Another reason why the existing `Par` monad is an appealing conceptual starting point for a practical implementation of LVars is that it already allows inter-task communication through `IVars`, which, as we have seen, are a special case of LVars. Finally, the `Par` monad approach is appealing because it is implemented entirely as a library, with a library-level scheduler. This approach makes it possible to make changes to the `Par` scheduling strategy in a modular way, without having to make any modifications to GHC or its runtime system.

¹Haskell is often advertised as a purely functional programming language, that is, one without side effects, but it is perhaps more useful to think of it as a language that keeps other effects out of the way so that one can use only the effects that one wants to use!

In this chapter, I describe the *LVish* library, a Haskell library for practical deterministic and quasi-deterministic parallel programming with LVars. We have already seen an example of an LVish Haskell program in Section 3.1; in the following two sections, we will take a more extensive tour of what LVish offers. Then, in Section 4.3, we will consider adding support for DPJ-style imperative disjoint parallelism to LVish. Finally, in Sections 4.4 and 4.5, we will look at two case studies that illustrate how LVish and LVars can be used in practice.

4.1. The big picture

Our library adopts and builds on the basic approach of the *Par* monad and the *monad-par* library [36], enabling us to employ our own notion of lightweight, library-level threads with a custom scheduler. It supports the programming model laid out in Section 3.1 in full, including explicit handler pools. It differs from the formalism of Section 3.2 in following Haskell’s by-need evaluation strategy, which also means that concurrency in the library is *explicitly marked*, either through uses of a *fork* function or through asynchronous callbacks, which run in their own lightweight threads.



We envision two parties interacting with the LVish library. First, there are *data structure authors*, who use the library directly to implement a specific monotonic data structure (e.g., a monotonically growing finite map). Second, there are *application writers*, who are clients of these data structures. Only the application writers receive a (quasi-)determinism guarantee; an author of a data structure is responsible for ensuring that the states their data structure can take on correspond to the elements of a lattice, and that the exposed interface to it corresponds to some use

of update operations, *get*, *freeze*, and event handlers.

The LVish library also includes *lattice-generic* infrastructure: the *Par* monad itself, a thread scheduler, support for blocking and signaling threads, handler pools, and event handlers. Since this infrastructure is unsafe—that is, it does not guarantee determinism or quasi-determinism—only data structure authors

should import it, subsequently exporting a *limited* interface specific to their data structure. For finite maps, for instance, this interface might include key/value insertion, lookup, event handlers and pools, and freezing—along with higher-level abstractions built on top of these. Control operators like `fork` are the only non-data-structure-specific operations exposed to application writers.

For this approach to scale well with available parallel resources, it is essential that the data structures themselves support efficient parallel access; a finite map that was simply protected by a global lock would force all parallel threads to sequentialize their access. Thus, we expect data structure authors to draw from the extensive literature on scalable parallel data structures, employing techniques like fine-grained locking and lock-free data structures [26]. Data structures that fit into the LVish model have a special advantage: because all updates must commute, it may be possible to avoid the expensive synchronization which *must* be used for non-commutative operations [4]. And in any case, monotonic data structures can be simpler to represent and implement than general ones.

4.2. The LVish library interface for application writers

In this section I illustrate the use of the LVish library from the point of view of the application writer, through a series of short example programs.²

4.2.1. A simple example: IVars in LVish. Recall that IVars are data structures that can be shared between parallel tasks and that allow single writes and blocking reads. Before looking at LVish, let us consider an IVar computation implemented with `monad-par`.

Listing 4.1 shows a program written using `monad-par` that will deterministically raise an error, because it tries to write to the IVar `num` twice. Here, `p` is a computation of type `Par Int`, meaning that it runs in the `Par` monad (via the call to `runPar`) and returns a value of `Int` type. `num` is an IVar, created with a call to `new` and then assigned to `twice`, via two calls to `put`, each of which runs in a separately forked

²Code for all the examples in this section is available at <https://github.com/lkuper/lvar-examples/>.

```

import Control.Monad.Par

p :: Par Int
p = do num <- new
      fork (put num 3)
      fork (put num 4)
      get num

main = print (runPar p)

```

Listing 4.1. A basic IVar example using monad-par.

```

import Control.Monad.Par

p :: Par Int
p = do num <- new
      fork (put num 4)
      fork (put num 4)
      get num

main = print (runPar p)

```

Listing 4.2. Repeated writes of the same value to an IVar.

task. The `runPar` function is an implicit global barrier: all forks have to complete before `runPar` can return.

The code in Listing 4.1 raises a “multiple put” error at runtime, which is as it should be: differing writes to the same shared location could cause the subsequent call to `get` to behave nondeterministically. Since we are still using `monad-par` here and not `LVish`, `get` has IVar semantics, not LVar semantics: rather than performing a threshold read, it blocks until `num` has been written, then unblocks and evaluates to the exact contents of `num`. However, when using `monad-par`, even multiple writes of the *same* value to an IVar will raise a “multiple put” error, as in Listing 4.2. This program differs from the previous one only in that the two puts are writing 4 and 4, rather than 3 and 4. Even though the call to `get` would produce a deterministic result regardless of which write happened first, the program nevertheless raises an error because of `monad-par`’s single-write restriction on IVars.


```

{-# LANGUAGE TypeFamilies #-}

import Control.LVish  -- Generic scheduler; works with all LVars.
import Data.LVar.IVar -- The particular LVar we need for this program.

p :: (HasPut e, HasGet e) => Par e s Int
p = do num <- new
      fork (put num 4)
      fork (put num 4)
      get num

main = print (runPar p)

```

Listing 4.3. Repeated writes of the same value to an LVar.

Now let us consider a version of Listing 4.2 written using the LVish library. (Of course, in LVish we are not limited to IVars, but we will consider IVars first as an interesting special case of LVars, and then go on to consider some more sophisticated LVars later in this section.) Listing 4.3 shows an LVish program that will write 4 to an IVar twice and then deterministically print 4 instead of raising an error.

In Listing 4.3, we need to import the `Control.LVish` module rather than `Control.Monad.Par` (since we now wish to use LVish instead of monad-par), and we must specifically import `Data.LVar.IVar` in order to specify which LVar data structure we want to work with (since we are no longer limited to IVars). Just as with monad-par, the LVish `runPar` function is a global barrier: both forks must compete before `runPar` can return. Also, as before, we have `new`, `put`, and `get` operations that respectively create, update, and read from `num`. However, these operations now have LVar semantics: the `put` operation computes a lub (with respect to a lattice similar to that of Figure 2.1(b), except including all the `Ints`), and the `get` operation performs a threshold read, where the threshold set is implicitly the set of all `Ints`. We do not need to explicitly write down the threshold set in the code. Rather, it is the obligation of the `Data.LVar.IVar` module to provide operations (`put` and `get`) that have the semantic effect of lub writes and threshold reads (as I touched on earlier in Section 2.2.3).

There are two other important differences between the monad-par program and the LVish program: the `Par` type constructor has gained two new type parameters, `e` and `s`, and `p`'s type annotation now has a *type class constraint* of `(HasPut e, HasGet e)`. Furthermore, we have added a `LANGUAGE` pragma, instructing the compiler that we are now using the `TypeFamilies` language extension. In the following section, I explain these changes.

4.2.2. The `e` and `s` type parameters: effect tracking and session tracking. In order to support both deterministic and quasi-deterministic programming in LVish, we need a way to specify which LVar effects can occur within a given `Par` computation. In a deterministic computation, only update operations (such as `put`) and threshold reads should be allowed; in a quasi-deterministic computation, `freeze` operations should be allowed as well. Other combinations may be desirable as well: for instance, we may want a computation to perform *only* writes, and not reads.

In order to capture these constraints and make them explicit in the types of LVar computations, LVish indexes `Par` computations with a *phantom type* `e` that indicates their *effect level*. The `Par` type becomes, instead, `Par e`, where `e` is a type-level encoding of Booleans indicating which operations, such as writes, reads, or freeze operations, are allowed to occur inside it. LVish follows the precedent of Kiselyov *et al.* on extensible effects in Haskell [30]: it abstracts away the specific structure of `e` into *type class constraints*, which allow a `Par` computation to be annotated with the *interface* that its `e` type parameter is expected to satisfy. This approach allows us to define “effect shorthands” and use them as Haskell type class constraints. For example, a `Par` computation where `e` is annotated with the effect level constraint `HasPut` can perform puts. In our example above, `e` is annotated with both `HasPut` and `HasGet` and therefore the `Par` computation in question can perform both puts and gets. We will see several more examples of effect level constraints in LVish `Par` computations shortly.

The effect tracking infrastructure is also the reason why we need to use the `TypeFamilies` language extension in our LVish programs. For brevity, I will elide the `LANGUAGE` pragmas in the rest of the example LVish programs in this section.

The LVish `Par` type constructor also has a second type parameter, `s`, making `Par e s` a the complete type of a `Par` computation that returns a result of type `a`. The `s` parameter ensures that, when a computation in the `Par` monad is run using the provided `runPar` operation (or using a variant of `runPar`, which I will discuss below), it is not possible to return an `LVar` from `runPar` and reuse it in another call to `runPar`. The `s` type parameter also appears in the types of `LVars` themselves, and the universal quantification of `s` in `runPar` and its variants forces each `LVar` to be tied to a single “session”, *i.e.*, a single use of a run function, in the same way that the `ST` monad in Haskell prevents an `STRef` from escaping `runST`. Doing so allows the LVish implementation to assume that `LVars` are created and used within the same session.³

4.2.3. An observably deterministic shopping cart. For our next few examples, let us consider concurrently adding items to a shopping cart. Suppose we have an `Item` data type for items that can be added to the cart. For the sake of this example, suppose that only two items are on offer:

```
data Item = Book | Shoes
  deriving (Show, Ord, Eq)
```



The cart itself can be represented using the `IMap LVar` type (provided by the `Data.LVar.PureMap`⁴ module), which is a key-value map where the keys are `Items` and the values are the quantities of each item. The name `IMap` is by analogy with `IVar`, but here, it is individual entries in the map that are immutable, not the map itself. If a key is inserted multiple times, the values must be equal (according to `==`), or a “multiple put” error will be raised.

³The addition of the `s` type parameter to `Par` in the LVish library has nothing to do with `LVars` in particular; it would also be a useful addition to the original `Par` library to prevent programmers from reusing an `IVar` from one `Par` computation to another, which is, as Simon Marlow has noted, “a Very Bad Idea; don’t do it” [34].

⁴The “Pure” in `Data.LVar.PureMap` distinguishes it from LVish’s other map data structure, which is also called `IMap`, but is provided by the `Data.LVar.SLMap` module and is a lock-free data structure based on concurrent skip lists. The `IMap` provided by `Data.LVar.PureMap`, on the other hand, is a reference implementation of a map, which uses a pure `Data.Map` wrapped in a mutable container. Both `IMaps` present the same API, and either implementation of `IMap` would have worked for this example, but the lock-free version is designed to scale as parallel resources are added. I discuss the role of lock-free data structures in LVish in more detail in Section 4.4.5.

```

import Control.LVish
import Data.LVar.PureMap

p :: (HasPut e, HasGet e) => Par e s Int
p = do cart <- newEmptyMap
      fork (insert Book 2 cart)
      fork (insert Shoes 1 cart)
      getKey Book cart

main = print (runPar p)

```

Listing 4.4. A deterministic shopping-cart program.

Listing 4.4 shows an LVish program that inserts items into our shopping cart. The `newEmptyMap` operation creates a new `IMap`, and the `insert` operation allows us to add new key-value pairs to the cart. In this case, we are concurrently adding the `Book` item with a quantity of 2, and the `Shoes` item with a quantity of 1. The call to `getKey` will be able to unblock as soon as the first `insert` operation has completed, and the program will deterministically print 2 regardless of whether the second `insert` has completed at the time that `getKey` unblocks.

The `getKey` operation allows us to threshold on a key—in this case `Book`—and get back the value associated with that key, once it has been written. The (implicit) threshold set of a call to `getKey` is the set of all values that might be associated with a key; in this case, the set of all `Ints`. This is a legal threshold set because `IMap` entries are *immutable*: we cannot, for instance, insert a key of `Book` with a quantity of 2 and then later change the 2 to 3. In a more realistic shopping cart, the values in the cart could themselves be `LVars` representing incrementable counters, as in the previous section. However, a shopping cart from which we can *delete* items is not possible with `LVars`, because it would go against the principle of monotonic growth.⁵

⁵On the other hand, one way to implement a container that allows both insertion and removal of elements is to represent it internally with *two* containers, one for the inserted elements and one for the removed elements, where both containers grow monotonically. *Conflict-free replicated data types* (CRDTs) [49] use variations on this approach to implement various data structures that support seemingly non-monotonic operations. I discuss the relationship of `LVars` to CRDTs in more detail in Chapter 5.

```

import Control.LVish
import Data.LVar.PureMap
import qualified Data.Map as M

p :: (HasPut e, HasFreeze e) => Par e s (M.Map Item Int)
p = do cart <- newEmptyMap
      fork (insert Book 2 cart)
      fork (insert Shoes 1 cart)
      freezeMap cart

main = do v <- runParQuasiDet p
        print (M.toList v)

```

Listing 4.5. A quasi-deterministic shopping-cart program.

4.2.4. A quasi-deterministic shopping cart. The LVish examples we have seen so far have been fully deterministic; they do not use `freeze`. Next, let us consider a program that freezes and reads the exact contents of a shopping cart, concurrently with adding items to it.

In Listing 4.5, we are inserting items into our cart, as in Listing 4.4. But, instead of returning the result of a call to `getKey`, this time `p` returns the result of a call to `freezeMap`, and the return type of `p` is a `Par` computation containing not an `Int`, but rather an entire map from `Items` to `Ints`. In fact, this map is not the `IMap` that `Data.LVar.PureMap` provides, but rather the standard `Map` from the `Data.Map` module (imported as `M`). This is possible because `Data.LVar.PureMap` is implemented using `Data.Map`, and so freezing its `IMap` simply returns the underlying `Data.Map`.

Because `p` performs a freezing operation, the effect level of its return type must reflect the fact that it is allowed to perform freezes. Therefore, instead of `HasGet`, we have the type class constraint of `HasFreeze` on `e`. Furthermore, because `p` is allowed to perform a freeze, we cannot run it with `runPar`, as in our previous examples, but must instead use a special variant of `runPar`, called `runParQuasiDet`, whose type signature allows `Par` computations that allow freezing to be passed to it.

The quasi-determinism in Listing 4.5 arises from the fact that the call to `freezeMap` may run before both forked computations have completed. In this example, one or both calls to `insert` may run after

```

import Control.LVish
import Control.LVish.DeepFrz -- provides runParThenFreeze
import Data.LVar.PureMap

p :: (HasPut e) => Par e s (IMap Item s Int)
p = do
  cart <- newEmptyMap
  fork (insert Book 2 cart)
  fork (insert Shoes 1 cart)
  return cart

main = print (runParThenFreeze p)

```

Listing 4.6. A deterministic shopping-cart program that uses `runParThenFreeze`.

the call to `freezeMap`. If this happens, the program will raise a write-after-freeze exception. The other possibility is that both items are already in the cart at the time it is frozen, in which case the program will run without error and print both items. There are therefore two possible outcomes: a cart with both items, or a write-after-freeze error. The advantage of quasi-determinism is that it is not possible to get multiple *non-error* outcomes, such as, for instance, an empty cart or a cart to which only the Book has been added.

4.2.5. Regaining full determinism with `runParThenFreeze`. The advantage of freezing is that it allows us to observe the exact, complete contents of an LVar; the disadvantage is that it introduces quasi-determinism due to the possibility of a write racing with a freeze, as in the example above. But, if we could ensure that the freeze operation happened *last*, we would be able to freeze LVars with no risk to determinism. In fact, the LVish library offers a straightforward solution to this problem: instead of manually calling `freeze` (and perhaps accidentally freezing an LVar too early), we can tell LVish to handle the freezing for us while “on the way out” of a `Par` computation. The mechanism that allows this is another variant of `runPar`, which we call `runParThenFreeze`.

Listing 4.6 shows a version of Listing 4.5 written using `runParThenFreeze`. Unlike the `Par` computations in the shopping-cart examples we have seen so far, the `Par` computation in Listing 4.6 *only* performs writes (as we can see from its effect level, which is only constrained by `HasPut`). Also, unlike in Listing 4.5, where a `freeze` took place inside the `Par` computation, in Listing 4.6 the `Par` computation returns an `IMap` rather than a `Map`. Since `IMap` is an `LVar`, it has an `s` parameter, which we can see in the type of `p`.

Because there is no synchronization operation after the two `fork` calls, `p` may return `cart` before both (or either) of the calls to `insert` have completed. However, since `runParThenFreeze` is an implicit global barrier (just as `runPar` and `runParQuasiDet` are), both calls to `insert` *must* complete before `runParThenFreeze` can return—which means that the result of the program is deterministic.

4.2.6. Event-driven programming with LVars: a deterministic parallel graph traversal. Finally, let us look at an example that uses event handlers as well as freezing. In Listing 4.7, the function `traverse` takes a graph `g` and a vertex `startNode` and finds the set of all vertices reachable from `startNode`, in parallel. The `traverse` function first creates a new `LVar`, called `seen`, to represent a monotonically growing set of `Ints` that will identify nodes in the graph. For this purpose, we use the `ISet` type, provided by the `Data.LVar.PureSet` module. (As with `IMap`, the individual elements of the `ISet` are immutable, but the set itself can grow.)

Next, `traverse` attaches an event handler to `seen`. It does so by calling the `newHandler` function, which takes two arguments: an `LVar` and the callback that is to be run every time an event occurs on that `LVar` (in this case, every time a new node is added to the set).⁶ The callback responds to events by looking up the neighbors of the newly arrived node (assuming a `neighbors` operation, which takes a graph and a vertex and returns a list of the vertex’s neighbor vertices), then mapping the `insert` function over that list of neighbors.

⁶LVish does not provide `newHandler`, but we can easily implement it using LVish’s built-in `newPool` and `addHandler` operations.

```

import Control.LVish
import Control.LVish.DeepFrz -- provides Frzn
import Data.LVar.Generic (addHandler, freeze)
import Data.LVar.PureSet
import qualified Data.Graph as G

traverse :: (HasPut e, HasFreeze e) =>
    G.Graph -> Int -> Par e s (ISet Frzn Int)
traverse g startNode = do
    seen <- newEmptySet
    h <- newHandler seen
    (\node -> do
        mapM (\v -> insert v seen)
            (neighbors g node)
        return ())
    insert startNode seen -- Kick things off
    quiesce h
    freeze seen

main = do
    v <- runParQuasiDet (traverse myGraph (0 :: G.Vertex))
    print (fromISet v)

```

Listing 4.7. A deterministic parallel graph traversal with an explicit call to freeze.

Finally, `traverse` adds the starting node to the `seen` set by calling `insert startNode seen`—and the event handler does the rest of the work. We know that we are done handling events when the call to `quiesce h` returns; it will block until all events have been handled. Finally, we freeze and return the `ISet` of all reachable nodes. Since `ISet` is an `LVar`, it has an `s` parameter, and in the return type of `traverse`, the `s` parameter of `ISet` has been replaced by the `Frzn` type, indicating that the `LVar` has been frozen.

The good news is that this particular graph traversal program is deterministic. The bad news is that, in general, freezing introduces quasi-determinism, since we could have forgotten to call `quiesce` before the freeze—which is why `traverse` must be run with `runParQuasiDet`, rather than `runPar`. Although the *program* is deterministic, the *language-level* guarantee is merely of quasi-determinism, not determinism.


```

import Control.LVish
import Control.LVish.DeepFrz -- provides runParThenFreeze
import Data.LVar.Generic (addHandler, freeze)
import Data.LVar.PureSet
import qualified Data.Graph as G

traverse :: HasPut e => G.Graph -> Int -> Par e s (ISet s Int)
traverse g startNode = do
  seen <- newEmptySet
  h <- newHandler seen
  (\node -> do
    mapM (\v -> insert v seen)
      (neighbors g node)
    return ())
  insert startNode seen -- Kick things off
  return seen

main = print (runParThenFreeze (traverse myGraph (0 :: G.Vertex)))

```

Listing 4.8. A deterministic parallel graph traversal that uses `runParThenFreeze`.

However, just as with our final shopping-cart example in Listing 4.6, we can use `runParThenFreeze` to ensure that freezing happens last. Listing 4.8 gives a version of `traverse` that uses `runParThenFreeze`, and eliminates the possibility of forgetting to call `quiesce` and thereby introducing quasi-determinism.

In Listing 4.8, since freezing is performed by `runParThenFreeze` rather than by an explicit call to `freeze`, it is no longer necessary to constrain `e` with `HasFreeze` in the type of `traverse`. Furthermore, the `s` parameter in the `ISet` that `traverse` returns can remain `s` instead of being instantiated with `Frzn`. Most importantly, since freezing is performed by `runParThenFreeze` rather than by an explicit call to `freeze`, it is no longer necessary for `traverse` to explicitly call `quiesce`, either! Because of the implicit barrier created by `runParThenFreeze`, all outstanding events that can be handled will be handled before it can return.

4.3. Par-monad transformers and disjoint parallel update

The effect-tracking system of the previous section provides a way to toggle on and off a fixed set of basic capabilities, such as `HasPut`, using the type system—that is, with the effect level e that parameterizes the `Par` type. However, it does not give us a way to add new, unanticipated capabilities to a `Par` computation. For that, we turn to *monad transformers*.

In Haskell, a monad transformer is a type constructor that adds “plug-in” capabilities to an underlying monad. For example, the `StateT` monad transformer adds an extra piece of implicit, modifiable state to an underlying monad. Adding a monad transformer to a type always returns another monad (preserving the `Monad` instance). We can therefore define a *Par-monad transformer* as a type constructor `T` where, for all `Par` monads `m`, `T m` is another `Par` monad with additional capabilities, and where a value of type `T m a`, for instance, `T (Par e s) a`, is a computation in that monad.

4.3.1. Example: threading state in parallel. We can use the standard `StateT` monad transformer (provided by Haskell’s `Control.Monad.State` package) as a `Par-monad transformer`. However, even if `m` is a `Par` monad, for `StateT s m` to also be a `Par` monad, the state `s` must be *splittable*; that is, it must be specified what is to be done with the state at fork points in the control flow. For example, the state may be duplicated, split, or otherwise updated to note the fork.

The below code promotes `StateT` to be a `Par-monad transformer`:

```
class SplittableState a where
  splitState :: a -> (a, a)

instance (SplittableState s, ParMonad m) =>
  ParMonad (StateT s m) where
  fork task =
    do s <- oState.get
    let (s1, s2) = splitState s
    State.put s2
    lift (fork (do runStateT task s1; return ()))
```

Note that here, `put` and `get` are not `LVar` operations, but the standard operations for setting and retrieving the state in a `StateT`.

4.3.2. Determinism guarantee. The `StateT` transformer preserves determinism because it is effectively *syntactic sugar*. That is, `StateT` does not allow one to write any program that could not already be written using the underlying `Par` monad, simply by passing around an extra argument. This is because `StateT` only provides a *functional* state (an implicit argument and return value), not actual mutable heap locations. Genuine mutable locations in pure computations, on the other hand, require Haskell’s `ST` monad, the safer sister monad to `IO`.

4.3.3. Disjoint parallel update with `ParST`. The `LVars` model is based on the notion that it is fine for multiple threads to access and update shared memory, so long as updates commute and “build on” one another, only adding information rather than destroying it. But it should also be possible for threads to update memory destructively, so long as the memory updated by different threads is *disjoint*. This is the approach to deterministic parallelism taken by, for example, Deterministic Parallel Java (DPJ) [8], which uses a region-based type and effect system to ensure that each mutable region of the heap is passed linearly to a thread that then gains exclusive permission to update that region.

In order to add this capability to the `LVish` library, we need destructive updates to interoperate with `LVar` effects. Moreover, we wish to do so at the library level, without requiring language extensions. Our solution is to provide a monad called `ParST` that uses the `StateT` transformer described above to layer additional mutable state on top of the existing capabilities of the `LVish Par` monad. `ParST` allows arbitrarily complex mutable state, such as tuples of vectors (arrays). However, `ParST` enforces the restriction that every memory location in the state is reachable by only one pointer: alias freedom. Previous approaches to integrating mutable memory with pure functional code (*i.e.*, the `ST` monad) work with `LVish`, but only allow thread-private memory. There is no way to operate on the same structure (for instance, on two halves of an array) from different threads. `ParST` exploits the fact that simultaneous

```

{-# LANGUAGE TypeFamilies #-}

import Prelude hiding (read)
import Control.LVish
import Control.Par.ST (liftST)
import Control.Par.ST.Vec (ParVecT, set, reify, forkSTSplit, write, read, runParVecT)
import Data.Vector (freeze, toList)

p :: (HasGet e, HasPut e) => ParVecT s1 String Par e s [String]
p = do
  -- Fill all six slots in the vector with "foo".
  set "foo"
  -- Get a pointer to the state.
  ptr <- reify

  -- Fork two computations, each of which has access to half the
  -- vector. Within the two forked child computations, `ptr` is
  -- inaccessible.
  forkSTSplit 3 -- Split at index 3 in the vector.
    (write 0 "bar")
    (write 0 "baz")

  frozen <- liftST (freeze ptr)
  return (toList frozen)

main = print (runPar (runParVecT 6 p))

```

Listing 4.9. A program illustrating disjoint parallel update inside an LVar computation.

access from different threads can be deterministic, as long as the threads are accessing disjoint parts of the data structure. Listing 4.9 illustrates the idea using `ParVecT`, which is a specialized variant of `ParST` that supports a particular kind of shared state: a single mutable vector.

The code in Listing 4.9 writes to each element in a six-element vector, then splits the vector into two parts and updates each part in parallel. (The first argument to `runParVecT`, in this case 6, specifies the length of the vector.) The call to `forkSTSplit` forks the control flow of the program, and `(write 0 "bar")` and `(write 0 "baz")` are the two forked child computations. `forkSTSplit` takes as its first argument a “split point”, which is the index at which the vector is to be split. Here, that index is 3, which means that the first child computation passed to `forkSTSplit` may access only the first half

of the vector, while the other may access only the second half. Each child computation sees only a *local* view of the vector, so writing "bar" to index 0 in the second child computation is really writing to index 3 of the full vector.

The call to `freeze` in Listing 4.9 is not to be confused with an `LVar` freeze operation; it is instead the `freeze` operation from the `Data.Vector` library which produces an immutable copy of a mutable vector. In the last line of Listing 4.9, calling `runParVecT` discharges the extra state effect that `ParVecT` provides, leaving the underlying `Par` computation, which is then run with `runPar`. In this example, printing the result of the `runPar` gives us `["bar", "foo", "foo", "baz", "foo", "foo"]`.

Ensuring the determinism of `ParST` hinges on two requirements:

- *Disjointness*: Any thread can get a direct pointer to its state. In Listing 4.9, `ptr` is an `STVector` that can be passed to any standard library procedures in the `ST` monad. However, it must *not* be possible to access `ptr` from `forkSTSplit`'s child computations. We accomplish this using Haskell's support for higher-rank types,⁷ ensuring that accessing `ptr` from a child computation causes a type error. Finally, `forkSTSplit` is a fork-join construct; after it completes, the parent thread again has full access to `ptr`.
- *Alias freedom*: Imagine that we expanded the example in Listing 4.9 to have as its state a *tuple* of two vectors (v_1, v_2) . If we allowed the programmer to supply an arbitrary initial state to the `ParST` computation, then they might provide the state (v_1, v_1) , *i.e.*, two copies of the same pointer. This breaks the abstraction, enabling them to reach the same mutable location from multiple threads (by splitting the supposedly-disjoint vectors at a different index). Thus, in `LVish`, users do not populate the state directly, but only describe a *recipe* for its creation. Each type used as a `ParST` state has an associated type for descriptions of (1) how to create an initial structure, and (2) how to split it into disjoint pieces. `LVish` provides a library of instances for commonly used types.

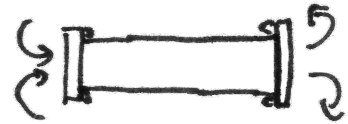
⁷That is, the type of a child computation begins with `(forall s . ParST ...)`.

4.3.4. Inter-thread communication. Disjoint state update does not solve the problem of communication between threads. Hence systems built around this idea often include other means for performing reductions, or require “commutativity annotations” for operations such as adding to a set. For instance, DPJ provides a `commuteswith` form for asserting that operations commute with one another to enable concurrent mutation. In LVish, however, such annotations are unnecessary, because LVish already provides a language-level guarantee that all effects commute! Thus, a programmer using LVish with ParST can use any LVar to communicate results between threads performing disjoint updates, without requiring trusted code or annotations. Moreover, LVish with ParST is unique among deterministic parallel programming models in that it allows both DPJ-style, disjoint destructive parallel updates, and blocking, dataflow-style communication between threads (through LVars).

4.4. Case study: parallelizing k -CFA with LVish

LVish is designed to be particularly applicable to (1) parallelizing complicated algorithms on structured data that pose challenges for other deterministic programming models, and (2) composing pipeline-parallel stages of computation (each of which may be internally parallelized). In this section, I describe a case study that fits this mold: *parallelized control-flow analysis*. I discuss the process of porting a sequential implementation of a k -CFA static program analysis to a parallel implementation using LVish.

The k -CFA analyses provide a hierarchy of increasingly precise methods to compute the flow of values to expressions in a higher-order language. For this case study, we began with a sequential im-



plementation of k -CFA translated to Haskell from a version by Might [38].⁸ The algorithm processes expressions written in a continuation-passing-style λ -calculus. It resembles a nondeterministic abstract interpreter in which stores map addresses to *sets* of abstract values, and function application entails a cartesian product between the operator and operand sets. Furthermore, an address models not just a

⁸Haskell port by Max Bolingbroke: <https://github.com/batterseapower/haskell-kata/blob/master/OCFA.hs>.

```

explore :: S.Set State -> [State] -> S.Set State
explore seen [] = seen
explore seen (todo:todos)
| todo `S.member` seen = explore seen todos
| otherwise = explore (S.insert todo seen) (S.toList (next todo) ++ todos)

```

Listing 4.10. The `explore` function from a purely functional k -CFA implementation.

static variable, but includes a fixed k -size window of the calling history to get to that point (the k in k -CFA).

Taken together, the current redex, environment, store, and call history make up the abstract state of the program, and the goal is to explore a graph of these abstract states in order to discover the flow of control of a program without needing to actually run it. This graph-exploration phase is followed by a second, summarization phase that combines all the information discovered into one store.

4.4.1. k -CFA phase one: breadth-first exploration. The `explore` function from the original, sequential k -CFA analysis, shown in Listing 4.10, expresses the heart of the search process. `explore` uses idiomatic Haskell data types like `Data.Set` and lists. However, it presents a dilemma with respect to exposing parallelism. Consider attempting to parallelize `explore` using purely functional parallelism with futures—for instance, using the Haskell Strategies library [35]. An attempt to compute the next states in parallel would seem to be thwarted by the main thread rapidly forcing each new state to perform the seen-before check, `todo `S.member` seen`. There is no way for independent threads to “keep going” further into the graph; rather, they check in with `seen` after one step.

We confirmed this prediction by adding a parallelism annotation from the aforementioned Strategies library:

```

withStrategy (parBuffer 8 rseq) (next todo)

```

The GHC runtime reported that 100% of created futures were “duds”—that is, the main thread forced them before any helper thread could assist. Changing `rseq` to `rdeepseq` exposed a small amount of parallelism—238 of 5000 futures were successfully executed in parallel—yielding no actual speedup.

4.4.2. k -CFA phase two: summarization. The first phase of the algorithm produces a large set of states, with stores that need to be joined together in the summarization phase. When one phase of a computation produces a large data structure that is immediately processed by the next phase, lazy languages can often achieve a form of pipelining “for free”. This outcome is most obvious with *lists*, where the head element can be consumed before the tail is computed, offering cache-locality benefits. Unfortunately, when processing a pure `Data.Set` or `Data.Map` in Haskell, such pipelining is not possible, since the data structure is internally represented by a balanced tree whose structure is not known until all elements are present. Thus phase one and phase two cannot overlap in the purely functional version—but they will in the LVish version, as we will see. In fact, in LVish we will be able to achieve partial deforestation in addition to pipelining. Full deforestation in this application is impossible, because the `Data.Sets` in the implementation serve a memoization purpose: they prevent repeated computations as we traverse the graph of states.

4.4.3. Porting to LVish. Our first step in parallelizing the original k -CFA implementation was a *verbatim* port to LVish: that is, we changed the original, purely functional program to allocate a new LVar for each new set or map value in the original code. This was done simply by changing two types, `Set` and `Map`, to their LVar counterparts, `ISet` and `IMap`. In particular, a store maps a program location (with context) onto a set of abstract values (here the libraries providing `ISet` and `IMap` are imported as `IS` and `IM`, respectively):

```
type Store s = IM.IMap Addr s (IS.ISet s Value)
```

Next, we replaced allocations of containers, and `map/fold` operations over them, with the analogous operations on their LVar counterparts. The `explore` function above was replaced by a function that

amounts to the simple graph traversal function from Section 3.1.4. These changes to the program were mechanical, including converting pure to monadic code. Indeed, the key insight in doing the verbatim port to LVish was to consume LVars as if they were pure values, ignoring the fact that an LVar’s contents are spread out over space and time and are modified through effects.

In some places the style of the ported code is functional, while in others it is imperative. For example, the `summarize` function uses nested `forEach` invocations to accumulate data into a store map:

```
summarize :: IS.ISet s (State s) -> Par d s (Store s)
summarize states = do
  storeFin <- newEmptyMap
  void $ IS.forEach states $ \ (State _ _ store_n _) -> do
    void $ IM.forEach store_n $ \ key val -> do
      void $ IS.forEach val $ \ elem -> do
        IM.modify storeFin key newEmptySet $ \ st -> do
          IS.insert elem st
      return storeFin
```

While this code can be read in terms of traditional parallel nested loops, it in fact creates a network of handlers that convey incremental updates from one LVar to another, in the style of data-flow networks. That means, in particular, that computations in a pipeline can *immediately* begin reading results from containers (e.g., `storeFin`), long before their contents are final.

The LVish version of k -CFA contains eleven occurrences of `forEach`, as well as a few cartesian-product operations. The cartesian products serve to apply functions to combinations of all possible values that arguments may take on, greatly increasing the number of handler events in circulation. Moreover, chains of handlers registered with `forEach` result in cascades of events through six or more handlers. The runtime behavior of these operations would be difficult to reason about. Fortunately, the programmer can largely ignore the temporal behavior of their program, since all LVish effects commute—rather like the way in which a lazy functional programmer typically need not think about the order in which thunks are forced at runtime.

Finally, there is an optimization benefit to using handlers. Normally, to flatten a nested data structure such as `[[[Int]]]` in a functional language, one would need to flatten one layer at a time and allocate a series of temporary structures. The LVish version avoids this; for example, in the code for `summarize` above, three `forEach` invocations are used to traverse a triply-nested structure, and yet the side effect in the innermost handler directly updates the final accumulator, `storeFin`.

4.4.4. Flipping the switch: the advantage of sharing. The verbatim port to LVish uses LVars poorly: copying them repeatedly and discarding them without modification. This effect overwhelms the benefits of partial deforestation and pipelining, and the verbatim LVish port has a small performance overhead relative to the original. But not for long!

The most clearly unnecessary operation in the verbatim port is in the `next` function (called in the last line of Listing 4.10). In keeping with the purely functional program from which it was ported, `next` creates a fresh store to extend with new bindings as we take each step through the state space graph:

```
store' <- IM.copy store
```

Of course, a “copy” for an LVar is persistent: it is just a handler that forces the copy to receive everything the original does. But in LVish, it is also trivial to *entangle* the parallel branches of the search, allowing them to share information about bindings, simply by *not* creating a copy:

```
let store' = store
```

This one-line change speeds up execution by up to $25\times$ *on one core*. The lesson here is that, although pure functional parallel programs are guaranteed to be deterministic, the overhead of allocation and copying in an idiomatic pure functional program can overwhelm the advantages of parallelism. In the LVish version, the ability to use shared mutable data structures—even though they are only mutable in the extremely restricted and determinism-preserving way that LVish allows—affords a significant

speedup even when the code runs sequentially. The effect is then multiplied as we add parallel resources: the asynchronous, ISet-driven parallelism enables parallel speedup for a total of up to $202\times$ total improvement over the purely functional version.

4.4.5. Parallel speedup results. We implemented two versions of the k -CFA algorithm using set data structures that the LVish library provides. The first, `PureSet` (exported by the `Data.LVar.PureSet` module), is the LVish library’s reference implementation of a set, which uses a pure `Data.Set` wrapped in a mutable container. The other, `SLSet`, exported by `Data.LVar.SLSet`, is a lock-free set based on concurrent skip lists [26].⁹

We evaluated both the `PureSet`-based and `SLSet`-based k -CFA implementations on two benchmarks. For the first, we used a version of the “blur” benchmark from a recent paper on k -CFA by Earl *et al.* [18]. In general, it proved difficult to generate example inputs to k -CFA that took long enough to be candidates for parallel speedup; we were, however, able to “scale up” the blur benchmark by replicating the code N times, feeding one into the continuation argument for the next. For our second benchmark, we ran the k -CFA analysis on a program that was simply a long chain of 300 “not” functions (using a CPS conversion of the Church encoding for Booleans). This latter benchmark, which we call “notChain”, has a small state space of large states with many variables (600 states and 1211 variables), and was specifically designed to negate the benefits of our sharing approach.

Figure 4.1 shows the parallel speedup results of our experiments on a twelve-core machine.¹⁰ (We used $k = 2$ for the benchmarks in this section.) The lines labeled “blur” and “blur/lockfree” show the parallel speedup of the “blur” benchmark for the `PureSet`-based implementation and `SLSet`-based implementation of k -CFA, respectively, and the lines labeled “notChain” and “notChain/lockfree” show parallel speedup of the “notChain” benchmark for the `PureSet`-based and `SLSet`-based implementations, respectively.

⁹LVish also provides analogous reference and lock-free implementations of maps (`PureMap` and `SLMap`). In fact, LVish is the first project to incorporate *any* lock-free data structures in Haskell, which required solving some unique problems pertaining to Haskell’s laziness and the GHC compiler’s assumptions regarding referential transparency [39].

¹⁰Intel Xeon 5660; full machine details available at <https://portal.futuregrid.org/hardware/delta>.

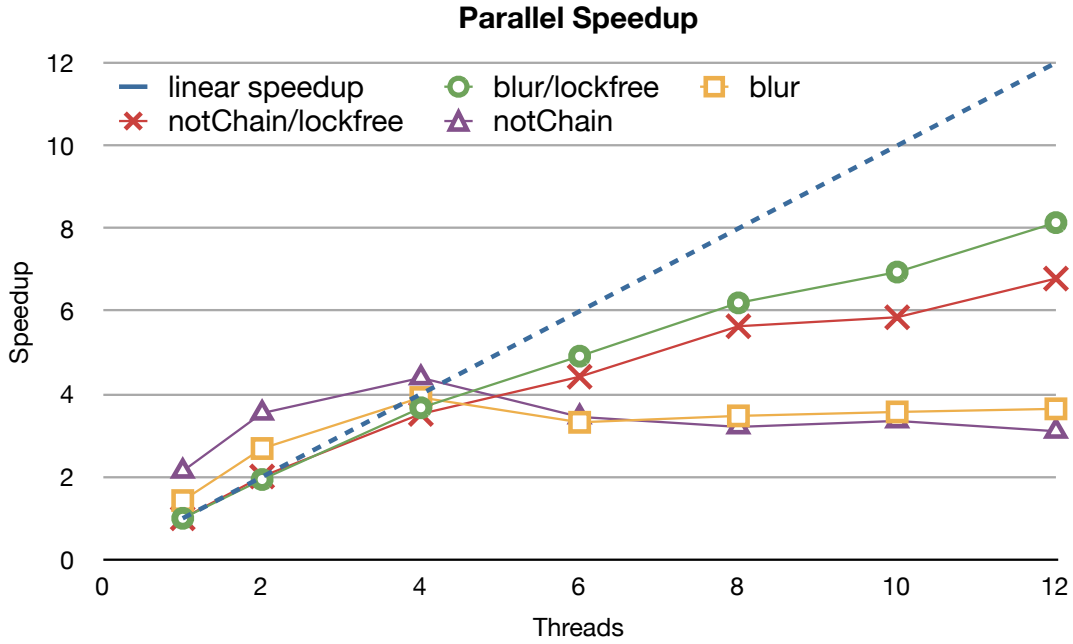


Figure 4.1. Parallel speedup for the “blur” and “notChain” benchmarks. Speedup is normalized to the sequential times for the *lock-free* versions (5.21s and 9.83s, respectively). The normalized speedups are remarkably consistent for the lock-free version between the two benchmarks. But the relationship to the original, purely functional version (not shown) is quite different: at 12 cores, the lock-free LVish version of “blur” is $202\times$ faster than the original, while “notChain” is only $1.6\times$ faster, not gaining anything from sharing rather than copying stores due to a lack of fan-out in the state graph.

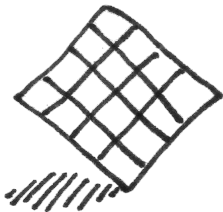
The results for the PureSet-based implementations are normalized to the same baseline as the results for the SLSet-based implementations at one core. At one and two cores, the SLSet-based k -CFA implementation (shown in green) is 38% to 43% slower than the PureSet-based implementation (in yellow) on the “blur” benchmark. The PureSet-based implementation, however, stops scaling after four cores. Even at four cores, variance is high in the PureSet-based implementation (min/max 0.96s / 1.71s over 7 runs). Meanwhile, the SLSet-based implementation continues scaling and achieves an $8.14\times$ speedup on twelve cores (0.64s at 67% GC productivity).

Of course, it is unsurprising that using an efficient lock-free shared data structure results in a better parallel speedup; rather, the interesting thing about these results is that despite its determinism guarantee, there is nothing about the LVars model that precludes using such data structures. Any data structure

that has the semantics of an LVar is fine. Indeed, part of the benefit of LVish is that it can allow parallel programs to make use of lock-free data structures while retaining the determinism guarantee of LVars, in much the same way that the ST monad allows Haskell programs access to efficient in-place array updates.

4.5. Case study: parallelizing PhyBin with LVish

One reason why we might want guaranteed-deterministic software is for the sake of scientific repeatability: in bioinformatics, for example, we would expect an experiment on the same data set to produce the same result on every run. In this section, I describe our experience using the LVish library to parallelize *PhyBin*, a bioinformatics application for comparing phylogenetic trees. A *phylogenetic tree* represents a possible ancestry for a set of N species. Leaf nodes in the tree are labeled with species' names, and the structure of the tree represents a hypothesis about common ancestors. For a variety of reasons, biologists often end up with many alternative trees, whose relationships they need to then analyze.



PhyBin [40] is a medium-sized (3500-line) bioinformatics program implemented in Haskell¹¹ for this purpose, initially released in 2010. The primary output of the software is a hierarchical clustering of the input tree set (that is, a tree of trees), but most of its computational effort is spent computing an $N \times N$ distance matrix that records the *edit distance* between each pair of input trees. It is this distance computation that we parallelize in our case study.

4.5.1. Computing all-to-all tree edit distance. The distance metric itself is called *Robinson-Foulds* (RF) distance, and the fastest algorithm for all-to-all RF distance computation is Sul and Williams' *HashRF* algorithm [52], which is used by a software package of the same name.¹² The HashRF software package is written in C++ and is about $2\text{--}3\times$ as fast as *PhyBin*, which also implements the HashRF algorithm. Both packages are dozens or hundreds of times faster than the more widely-used software

¹¹Available at <http://hackage.haskell.org/package/phybin>.

¹²Available at <https://code.google.com/p/hashrf/>.

that computes RF distance matrices, such as PHYLIP¹³ [20] and DendroPy¹⁴ [51]. These slower packages use $\frac{N^2-N}{2}$ full applications of the distance metric, which has poor locality in that it reads all trees in from memory $\frac{N^2-N}{2}$ times.

To see how the HashRF algorithm improves on this, consider that each edge in an unrooted phylogenetic tree can be seen as partitioning the tree’s nodes into two disjoint sets, according to the two subtrees that those nodes would belong to if the edge were deleted. For example, if a tree has nodes $\{a, b, c, d, e\}$, one bipartition or “split” might be $\{\{a, b\}, \{c, d, e\}\}$, while another might be $\{\{a, b, c\}, \{d, e\}\}$. A tree can therefore be encoded as a *set of bipartitions* of its nodes. Furthermore, once trees are encoded as sets of bipartitions, we can compute the edit distance between trees (that is, the number of operations required to transform one tree into the other) by computing the *symmetric set difference* between sets of bipartitions, and we can do so using standard set data structures.

The HashRF algorithm makes use of this fact and adds a clever trick that greatly improves locality. Before computing the actual distances between trees, it populates a dictionary, the “splits map”, which maps each observed bipartition to a set of IDs of trees that contain that bipartition. The second phase of the algorithm, which actually computes the $N \times N$ distance matrix, does so by iterating through each entry in the splits map. For each such entry, for each pair of tree IDs, it checks whether exactly one of those tree IDs is in the splits map entry, and if so, increments the appropriate distance matrix entry by one.

Algorithm 1 is a psuedocode version of the HashRF algorithm. The second phase of the algorithm is still $O(N^2)$, but it only needs to read from the much smaller *treerset* during this phase. All loops in Algorithm 1 are potentially parallel.

4.5.2. Parallelizing the HashRF algorithm with LVish. In the original PhyBin source code, the type of the splits map is:

¹³Available at <http://evolution.genetics.washington.edu/phylip.html>.

¹⁴Available at <http://pythonhosted.org/DendroPy/>.

Algorithm 1 Pseudocode of the HashRF algorithm for computing a tree edit distance matrix. *alltrees*, *splitsmap* and *distancematrix* are global variables, defined elsewhere. *alltrees* is the set of trees, represented as sets of bipartitions; *splitsmap* maps bipartitions to sets of trees in which they occur. In the second phase, the comparison of t_1 and t_2 uses XOR because the RF distance between two trees is defined as the number of bipartitions implied by exactly one of the two trees being compared.

```

▷ First phase: populate splits map.
for each  $t \in alltrees$  do
  for each  $bip \in t$  do
    ▷ Add  $t$  to set of trees pointed at by splitsmap[ $bip$ ],
    ▷ adding a key for  $bip$  to splitsmap if necessary.
    insert( $t$ , splitsmap[ $bip$ ])
  end for
end for
▷ Second phase: populate distance matrix.
▷ values() returns a list of all the values in a dictionary.
for each treerset  $\in$  values(splitsmap) do
  for each  $t_1 \in alltrees$  do
    for each  $t_2 \in alltrees$  do
      if  $t_1 \in treerset$  XOR  $t_2 \in treerset$  then
        increment(distancematrix[ $t_1, t_2$ ])
      end if
    end for
  end for
end for

```

<pre>type BipTable = Map DenseLabelSet (Set TreeID)</pre>

Here, a `DenseLabelSet` encodes an individual bipartition as a bit vector. `PhyBin` uses purely functional data structures for the `Map` and `Set` types, whereas the C++ HashRF implementation uses a mutable hash table. Yet in both cases, these structures grow monotonically during execution, making the algorithm a good candidate for parallelization with LVish. The splits map created during the first phase of the algorithm is a map of sets, which can be directly replaced by their LVar counterparts, and the distance matrix created in the second phase can be represented as a vector of monotonically increasing counters.

In fact, the parallel port of `PhyBin` using LVish was so straightforward that, after reading the code, parallelizing the first phase of the algorithm took only 29 minutes.¹⁵ Tables 4.1 and 4.2 show the results

¹⁵Git commit range: <https://github.com/rrnewton/PhyBin/compare/5cbf7d26c07a...6a05cfab490a7a>.

Trees	Species	DendroPy	PHYLIP	PhyBin
100	150	22.1s	12.8s	0.269s

Table 4.1. PhyBin performance comparison with DendroPy and PHYLIP.

Trees	Species	HashRF	PhyBin			
			1 core	2 cores	4 cores	8 cores
1000	150	1.7s	4.7s	3.0s	1.9s	1.4s

Table 4.2. PhyBin performance comparison with HashRF.

of a running time comparison of the parallelized PhyBin with DendroPy, PHYLIP, and HashRF. We first benchmarked PhyBin against DendroPy and PHYLIP using a set of 100 trees with 150 leaves each. Table 4.1 shows the time it took in each case to fill in the all-to-all tree edit distance matrix and get an answer back. PhyBin was much faster than the two alternatives.

Then, to compare PhyBin with HashRF, we used a set of 1000 trees with 150 leaves each. Table 4.2 shows the results. HashRF took about 1.7 seconds to process the 1000 trees, but since it is a single-threaded program, adding cores does not offer any speedup. PhyBin, while slower than HashRF on one core, taking about 4.7 seconds, speeds up as we add cores and eventually overtakes HashRF, running in about 1.4 seconds on 8 cores. Therefore LVish gives us a parallel speedup of about $3.35\times$ on 8 cores. This is exactly the sort of situation in which we would like to use LVish—to achieve modest speedups for modest effort, in programs with complex data structures (and high allocation rates), and without changing the determinism guarantee of the original functional code.

CHAPTER 5

Deterministic threshold queries of distributed data structures

So far, we have considered the problem of how to program *shared-memory parallel systems* in a way that guarantees determinism. In this chapter, we turn our attention to a different but related problem: that of effectively and correctly programming *distributed systems*, in which programs run on a network of interconnected nodes, each with its own memory, and where the network is subject to network partitions and other kinds of failures.

Because network partitions can occur and because nodes in a network can fail, distributed systems typically involve *replication* of data objects across a number of physical locations. Replication is of fundamental importance in such systems: it makes them more robust to data loss and allows for good data locality. But the well-known *CAP theorem* [23, 10] of distributed computing imposes a trade-off between *consistency*, in which every replica sees the same data, and *availability*, in which all data is available for both reading and writing by all replicas.

Highly available distributed systems, such as Amazon’s Dynamo key-value store [16], relax strong consistency in favor of *eventual consistency* [56], in which replicas need not agree at all times. Instead, updates execute at a particular replica and are sent to other replicas later. All updates eventually reach all replicas, albeit possibly in different orders. Informally speaking, eventual consistency says that if updates stop arriving, all replicas will *eventually* come to agree.



Although giving up on strong consistency makes it possible for a distributed system to offer high availability, even an eventually consistent system must have some way of resolving conflicts between replicas that differ. One approach is to try to determine which replica was written most recently, then declare

that replica the winner. But, even in the presence of a way to reliably synchronize clocks between replicas and hence reliably determine which replica was written most recently, having the last write win might not make sense from a *semantic* point of view. For instance, if a replicated object represents a set, then, depending on the application, the appropriate way to resolve a conflict between two replicas could be to take the set union of the replicas' contents. Such a conflict resolution policy might be more appropriate than a "last write wins" policy for, say, a object representing the contents of customer shopping carts for an online store [16].

Implementing application-specific conflict resolution policies in an ad-hoc way for every application is tedious and error-prone.¹ Fortunately, we need not implement them in an ad-hoc way. Shapiro *et al.*'s *convergent replicated data types* (CvRDTs) [49, 48] provide a simple mathematical framework for reasoning about and enforcing the eventual consistency of replicated objects, based on viewing replica states as elements of a lattice and replica conflict resolution as the lattice's join operation.

Like LVars, CvRDTs are data structures whose states are elements of an application-specific lattice, and whose contents can only grow with respect to the given lattice. Although LVars and CvRDTs were developed independently, both models leverage the mathematical properties of join-semilattices to ensure that a property of the model holds—determinism in the case of LVars; eventual consistency in the case of CvRDTs.

CvRDTs offer a simple and theoretically-sound approach to eventual consistency. However, with CvRDTs (and unlike with LVars), it is still possible to observe inconsistent *intermediate* states of replicated shared objects, and high availability requires that reads return a value immediately, even if that value is stale. In practice, applications call for both strong consistency and high availability at different times [54], and increasingly, they support consistency choices at the granularity of individual queries, not that of the entire system. For example, the Amazon SimpleDB database service gives customers the choice between eventually consistent and strongly consistent read operations on a per-read basis [57].

¹Indeed, as the developers of Dynamo have noted [16], Amazon's shopping cart presents an anomaly whereby removed items may re-appear in the cart!

Ordinarily, strong consistency is a global property: all replicas agree on the data. When a system allows consistency choices to be made at a *per-query* granularity, though, a global strong consistency property need not hold. We can define a *strongly consistent query* to be one that, if it returns a result x when executed at a replica i ,

- will always return x on subsequent executions at i , and
- will *eventually* return x when executed at *any* replica, and will *block* until it does so.

That is, a strongly consistent query of a distributed data structure, if it returns, will return a result that is a *deterministic* function of all updates to the data structure in the entire distributed execution, regardless of when the query executes or which replica it occurs on.

Traditional CvRDTs only support eventually consistent queries. We could get strong consistency from CvRDTs by waiting until all replicas agree before allowing a query to return—but in practice, such agreement may never happen. In this chapter, I present an alternative approach that takes advantage of the existing lattice structure of CvRDTs and does *not* require waiting for all replicas to agree. To do so, I take inspiration from LVar-style threshold reads. I show how to extend the CvRDT model to support deterministic, strongly consistent queries, which I call *threshold queries*. After reviewing the fundamentals of CvRDTs in Section 5.1, I introduce CvRDTs extended with threshold queries (Section 5.2), and I prove that threshold queries in our extended model are strongly consistent queries (Section 5.3). That is, I show that a threshold query that returns an answer when executed at a replica will return the same answer every subsequent time that it is executed at that replica, and that executing that threshold query on a different replica will eventually return the same answer, and will block until it does so. It is therefore impossible to observe different results from the same threshold query, whether at different times on the same replica or whether on different replicas.

5.1. Background: CvRDTs and eventual consistency

Shapiro *et al.* [49, 48] define an *eventually consistent* object as one that meets three conditions. One of these conditions is the property of *convergence*: all correct replicas of an object at which the same updates have been delivered eventually have equivalent state. The other two conditions are *eventual delivery*, meaning that all replicas receive all update messages, and *termination*, meaning that all method executions terminate (we discuss methods in more detail below).

Shapiro *et al.* further define a *strongly eventually consistent* (SEC) object as one that is eventually consistent and, in addition to being merely convergent, is *strongly convergent*, meaning that correct replicas at which the same updates have been delivered have equivalent state.² A *conflict-free replicated data type* (CRDT), then, is a data type (*i.e.*, a specification for an object) satisfying certain conditions that are sufficient to guarantee that the object is SEC. (The term “CRDT” is used interchangeably to mean a specification for an object, or an object meeting that specification.)



There are two “styles” of specifying a CRDT: *state-based*, also known as *convergent*³; or *operation-based* (or “op-based”), also known as *commutative*. CRDTs specified in the state-based style are called *convergent replicated data types*, abbreviated *CvRDTs*, while those specified in the op-based style are called *commutative replicated data types*, abbreviated *CmRDTs*. Of the two styles, we focus on the CvRDT style in this paper because CvRDTs are lattice-based data structures and therefore amenable to threshold queries—although, as Shapiro *et al.* show, CmRDTs can emulate CvRDTs and vice versa.

²Strong eventual consistency is not to be confused with strong consistency: it is the combination of eventual consistency and strong convergence. Contrast with ordinary convergence, in which replicas only *eventually* have equivalent state. In a strongly convergent object, knowing that the same updates have been delivered to all correct replicas is sufficient to ensure that those replicas have equivalent state, whereas in an object that is merely convergent, there might be some further delay before all replicas agree.

³There is a potentially misleading terminology overlap here: the definitions of convergence and strong convergence above pertain not only to CvRDTs (where the C stands for “Convergent”), but to *all* CRDTs.

5.1.1. State-based objects. In the Shapiro *et al.* model, a *state-based object* is a tuple (S, s^0, q, u, m) , where S is a set of states, s^0 is the initial state, q is the *query method*, u is the *update method*, and m is the *merge method*. Objects are replicated across some finite number of processes, with one replica at each process, and each replica begins in the initial state s^0 . The state of a local replica may be queried via the method q and updated via the method u . Methods execute locally, at a single replica, but the merge method m can merge the state from a remote replica with the local replica. The model assumes that each replica sends its state to the other replicas infinitely often, and that eventually every update reaches every replica, whether directly or indirectly.

The assumption that replicas send their state to one another “infinitely often” refers not to the *frequency* of these state transmissions; rather, it says that, regardless of what event (such as an update, via the u method) occurs at a replica, a state transmission is guaranteed to occur after that event. We can therefore conclude that all updates eventually reach all replicas in a state-based object, meeting the “eventual delivery” condition discussed above. However, we still have no guarantee of strong convergence or even convergence. This is where Shapiro *et al.*’s notion of a CvRDT comes in: a state-based object that meets the criteria for a CvRDT is guaranteed to have the strong-convergence property.

A *state-based* or *convergent* replicated data type (CvRDT) is a state-based object equipped with a partial order \leq , written as a tuple (S, \leq, s^0, q, u, m) , that has the following properties:

- S forms a join-semilattice ordered by \leq .
- The merge method m computes the join of two states with respect to \leq .
- State is *inflationary* across updates: if u updates a state s to s' , then $s \leq s'$.

Shapiro *et al.* show that a state-based object that meets the criteria for a CvRDT is strongly convergent. Therefore, given the eventual delivery guarantee that all state-based objects have, and given an additional assumption that all method executions terminate, a state-based object that meets the criteria for a CvRDT is SEC [49].

5.1.2. Discussion: the need for inflationary updates. Although CvRDT updates are required to be inflationary, it is not the case that every update must be inflationary for convergence, given our assumption of eventual delivery. Consider, for example, a scenario in which replicas 1 and 2 both have the state $\{a, b\}$. Replica 1 updates its state to $\{a\}$, a non-inflationary update, and then sends its updated state to replica 2. Replica 2 merges the received state $\{a\}$ with $\{a, b\}$, and its state remains $\{a, b\}$. Then replica 2 sends its state back to replica 1; replica 1 merges $\{a, b\}$ with $\{a\}$, and its state becomes $\{a, b\}$. The non-inflationary update has been lost, and was, perhaps, nonsensical—but the replicas are nevertheless convergent.

However, once we introduce threshold queries of CvRDTs, as we will do in the following section, inflationary updates become *necessary* for the determinism of threshold queries. This is because a non-inflationary update could cause a threshold query that had been unblocked to block again, and so arbitrary interleaving of non-inflationary writes and threshold queries would lead to nondeterministic behavior. Therefore the requirement that updates be inflationary will not only be sensible, but actually crucial.

5.2. Adding threshold queries to CvRDTs

In Shapiro *et al.*'s CvRDT model, the query operation q reads the exact contents of its local replica, and therefore different replicas may see different states at the same time, if not all updates have been propagated yet. That is, it is possible to observe intermediate states of a CvRDT replica. Such intermediate observations are not possible with threshold queries. In this section, we show how to extend the CvRDT model to accommodate threshold queries.

5.2.1. Objects with threshold queries. Definition 5.1 extends Shapiro *et al.*'s definition of a state-based object with a threshold query method t :

Definition 5.1 (state-based object with threshold queries). A *state-based object with threshold queries* (henceforth *object*) is a tuple (S, s^0, q, t, u, m) , where S is a set of states, $s^0 \in S$ is the initial state, q is a *query method*, t is a *threshold query method*, u is an *update method*, and m is a *merge method*.

In order to give a semantics to the threshold query method t , we need to formally define the notion of a threshold set. The notion of “threshold set” that I use here is the generalized formulation of threshold sets, based on *activation sets*, that I described previously in Section 2.6.2.

Definition 5.2 (threshold set). A *threshold set with respect to a lattice* (S, \leq) is a set $\mathcal{S} = \{S_a, S_b, \dots\}$ of one or more sets of *activation states*, where each set of activation states is a subset of S , the set of lattice elements, and where the following *pairwise incompatibility* property holds:

For all $S_a, S_b \in \mathcal{S}$, if $S_a \neq S_b$, then for all activation states $s_a \in S_a$ and for all activation states $s_b \in S_b$, $s_a \sqcup s_b = \top$, where \sqcup is the join operation induced by \leq and \top is the greatest element of (S, \leq) .

In our model, we assume a finite set of n processes p_1, \dots, p_n , and consider a single replicated object with one replica at each process, with replica i at process p_i . Processes may crash silently; we say that a non-crashed process is *correct*.

Every replica has initial state s^0 . Methods execute at individual replicas, possibly updating that replica’s state. The k th method execution at replica i is written $f_i^k(a)$, where k is ≥ 1 and f is either q , t , u , or m , and a is the arguments to f , if any. Methods execute sequentially at each replica. The state of replica i after the k th method execution at i is s_i^k . We say that states s and s' are equivalent, written $s \equiv s'$, if $q(s) = q(s')$.

5.2.2. Causal histories. An object’s *causal history* is a record of all the updates that have happened at all replicas. The causal history does not track the order in which updates happened, merely that they did happen. The *causal history at replica i after execution k* is the set of all updates that have happened

at replica i after execution k . Definition 5.3 updates Shapiro *et al.*'s definition of causal history for a state-based object to account for t (a trivial change, since execution of t does not change a replica's causal history):

Definition 5.3 (causal history). A *causal history* is a sequence $[c_1, \dots, c_n]$, where c_i is a set of the updates that have occurred at replica i . Each c_i is initially \emptyset . If the k th method execution at replica i is:

- a query q or a threshold query t , then the causal history at replica i after execution k does not change: $c_i^k = c_i^{k-1}$.
- an update $u_i^k(a)$, then the causal history at replica i after execution k is $c_i^k = c_i^{k-1} \cup u_i^k(a)$.
- a merge $m_i^k(s_{i'}^{k'})$, then the causal history at replica i after execution k is the union of the local and remote histories: $c_i^k = c_i^{k-1} \cup c_{i'}^{k'}$.

We say that an update is *delivered at replica i* if it is in the causal history at replica i .

5.2.3. Threshold CvRDTs and the semantics of blocking. With the previous definitions in place, we can give the definition of a CvRDT that supports threshold queries:

Definition 5.4 (CvRDT with threshold queries). A *convergent replicated data type with threshold queries* (henceforth *threshold CvRDT*) is an object equipped with a partial order \leq , written $(S, \leq, s^0, q, t, u, m)$, that has the following properties:

- S forms a join-semilattice ordered by \leq .
- S has a greatest element \top according to \leq .
- The query method q takes no arguments and returns the local state.
- The threshold query method t takes a threshold set \mathcal{S} as its argument, and has the following semantics: let $t_i^{k+1}(\mathcal{S})$ be the $k + 1$ th method execution at replica i , where $k \geq 0$. If, for some activation state s_a in some (unique) set of activation states $S_a \in \mathcal{S}$, the condition $s_a \leq s_i^k$ is met, $t_i^{k+1}(\mathcal{S})$ returns the set of activation states S_a . Otherwise, $t_i^{k+1}(\mathcal{S})$ returns the distinguished value block.

- The update method u takes a state as argument and updates the local state to it.
- State is *inflationary* across updates: if u updates a state s to s' , then $s \leq s'$.
- The merge method m takes a remote state as its argument, computes the join of the remote state and the local state with respect to \leq , and updates the local state to the result.

and the q , t , u , and m methods have no side effects other than those listed above.

We use the block return value to model t 's “blocking” behavior as a mathematical function with no intrinsic notion of running duration. When we say that a call to t “blocks”, we mean that it immediately returns block, and when we say that a call to t “unblocks”, we mean that it returns a set of activation states S_a .

Modeling blocking as a distinguished value introduces a new complication: we lose determinism, because a call to t at a particular replica may return either block or a set of activation states S_a , depending on the replica's state at the time it is called. However, we can conceal this nondeterminism with an additional layer over the nondeterministic API exposed by t . This additional layer simply *polls* t , calling it repeatedly until it returns a value other than block. Calls to t at a replica that are made by this “polling layer” count as method executions at that replica, and are arbitrarily interleaved with other method executions at the replica, including updates and merges. The polling layer itself need not do any computation other than checking to see whether t returns block or something else; in particular, the polling layer does not need to compare activation states to replica states, since that comparison is done by t itself.

The set of activation states S_a that a call to t returns when it unblocks is unique because of the pairwise incompatibility property required of threshold sets: without it, different orderings of updates could allow the same threshold query to unblock in different ways, introducing nondeterminism that would be observable beyond the polling layer.

5.2.4. Threshold CvRDTs are strongly eventually consistent. We can define eventual consistency and strong eventual consistency exactly as Shapiro *et al.* do in their model. In the following definitions, a *correct replica* is a replica at a correct process, and the symbol \Diamond means “eventually”:

Definition 5.5 (eventual consistency (EC)). An object is *eventually consistent* (EC) if the following three conditions hold:

- *Eventual delivery*: An update delivered at some correct replica is eventually delivered at all correct replicas: $\forall i, j : f \in c_i \Rightarrow \Diamond f \in c_j$.
- *Convergence*: Correct replicas at which the same updates have been delivered eventually have equivalent state: $\forall i, j : c_i = c_j \Rightarrow \Diamond s_i \equiv s_j$.
- *Termination*: All method executions halt.

Definition 5.6 (strong eventual consistency (SEC)). An object is *strongly eventually consistent* (SEC) if it is eventually consistent and the following condition holds:

- *Strong convergence*: Correct replicas at which the same updates have been delivered have equivalent state: $\forall i, j : c_i = c_j \Rightarrow s_i \equiv s_j$.

Since we model blocking threshold queries with block, we need not be concerned with threshold queries not necessarily terminating. Determinism does *not* rule out queries that return block every time they are called (and would therefore cause the polling layer to block forever). However, we guarantee that if a threshold query returns block every time it is called during a complete run of the system, it will do so on *every* run of the system, regardless of scheduling. That is, it is not possible for a query to cause the polling layer to block forever on some runs, but not on others.

Finally, we can directly leverage Shapiro *et al.*’s SEC result for CvRDTs to show that a threshold CvRDT is SEC:

Theorem 5.1 (Strong Eventual Consistency of Threshold CvRDTs). *Assuming eventual delivery and termination, an object that meets the criteria for a threshold CvRDT is SEC.*

Proof. From Shapiro *et al.*, we have that an object that meets the criteria for a CvRDT is SEC [49]. Shapiro *et al.*'s proof also assumes that eventual delivery and termination hold for the object, and proves that strong convergence holds — that is, that given causal histories c_i and c_j for respective replicas i and j , that their states s_i and s_j are equivalent. The proof relies on the commutativity of the lub operation. Since, according to our Definition 5.3, threshold queries do not affect causal history, we can leverage Shapiro *et al.*'s result to say that a threshold CvRDT is also SEC. \square

5.3. Determinism of threshold queries

Neither eventual consistency nor strong eventual consistency imply that *intermediate* results of the same query q on different replicas of a threshold CvRDT will be deterministic. For deterministic intermediate results, we must use the threshold query method t . We can show that t is deterministic *without* requiring that the same updates have been delivered at the replicas in question at the time that t runs.

Theorem 5.2 establishes a determinism property for threshold queries of CvRDTs, porting the determinism result previously established for threshold reads for LVars to a distributed setting.

Theorem 5.2 (Determinism of Threshold Queries). *Suppose a given threshold query t on a given threshold CvRDT returns a set of activation states S_a when executed at a replica i . Then, assuming eventual delivery and that no replica's state is ever \top at any point in the execution:*

- (1) *t will always return S_a on subsequent executions at i , and*
- (2) *t will eventually return S_a when executed at any replica, and will block until it does so.*

Proof. The proof relies on transitivity of \leq and eventual delivery of updates; see Section A.20 for the complete proof. \square

Although Theorem 5.2 must assume eventual delivery, it does *not* need to assume strong convergence or even ordinary convergence. It so happens that we have strong convergence as part of strong eventual consistency of threshold CvRDTs (by Theorem 5.1), but we do not need it to prove Theorem 5.2. In

particular, there is no need for replicas to have the same state in order to return the same result from a particular threshold query. The replicas merely both need to be above an activation state from a unique set of activation states in the query’s threshold set. Indeed, the replicas’ states may in fact trigger *different* activation states from the same set of activation states.

Theorem 5.2’s requirement that no replica’s state is ever \top rules out situations in which replicas disagree in a way that cannot be resolved normally. Recall from Section 2.4.2 that in the LVars model, when a program contains conflicting writes that would cause an LVar to reach its \top state, a threshold read in that program *can* behave nondeterministically. However, since in our definition of observable determinism, only the final outcome of a program counts, this nondeterministic behavior of `get` in the presence of conflicting writes is not observable: such a program would always have **error** as its final outcome. In our setting of CvRDTs, though, we do not have a notion of “program”, nor of the final outcome thereof. Rather than having to define those things and then define a notion of observable determinism based on them, I rule out this situation by assuming that no replica’s state goes to \top .

5.4. Discussion: reasoning about per-query consistency choices

In this chapter, we have seen a way to extend CvRDTs to support LVar-style threshold queries. Seen from another angle, this chapter shows how to “port” the notion of threshold reads from a shared-memory setting (that of LVars) to a distributed-memory one (that of CvRDTs). However, I do not want to suggest that deterministic threshold queries should replace traditional CvRDT queries. Instead, traditional queries and threshold queries can coexist. Moreover, extending CvRDTs with threshold queries allows them to more accurately model systems in which consistency properties are defined and enforced at the granularity of individual queries.

As mentioned at the beginning of this chapter, database services such as Amazon’s SimpleDB [57] allow for both eventually consistent and strongly consistent reads, chosen at a per-query granularity.⁴

⁴Terry *et al.*’s Pileus key-value store [54] takes the idea of combining different levels of consistency in a single application even further: instead of requiring the application developer to choose the consistency level of a particular query at development

Choosing consistency at the query level, and giving different consistency properties to different queries within a single application, is not a new idea. Rather, the new contribution we make by adding threshold queries to CvRDTs is to establish lattice-based data structures as a unifying formal foundation for both eventually consistent and strongly consistent queries. Adding support for threshold reads to CvRDTs allows us to take advantage of the machinery that CvRDTs already give us for reasoning about eventually consistent objects, and use it to reason about systems that allow consistency choices to be made at per-query granularity, as real systems do.

time, the system allows the developer to specify a service-level agreement that may be satisfied in different ways at runtime. This allows the application to, for instance, dynamically adapt to changing network conditions.

CHAPTER 6

Related work

Work on deterministic parallel programming models is long-standing. As we have seen, what deterministic parallel programming models have in common is that they all must do something to restrict access to mutable state shared among concurrent computations so that schedule nondeterminism cannot be observed. Depending on the model, restricting access to shared mutable state might involve disallowing sharing entirely [44], only allowing single assignments to shared references [55, 3, 11], allowing sharing only by a limited form of message passing [29], ensuring that concurrent accesses to shared state are disjoint [8], resolving conflicting updates after the fact [33], or some combination of these approaches. These constraints can be imposed at the language or API level, within a type system or at runtime.

In particular, the LVars model was inspired by two traditional deterministic parallel programming models based on monotonically-growing shared data structures: first, Kahn process networks [29], in which a network of processes communicate with each other through blocking FIFO channels with ever-growing channel histories; and, second, IVars [3], single-assignment locations with blocking read semantics.

LVars are general enough to subsume both IVars and KPNs: a lattice of channel histories with a prefix ordering allows LVars to represent FIFO channels that implement a Kahn process network, whereas instantiating λ_{LVar} with a lattice with one “empty” state and multiple “full” states (where $\forall i. \text{empty} < \text{full}_i$) results in a parallel single-assignment language with a store of IVars, as we saw in Chapter 2. Hence LVars provide a framework for generalizing and unifying these two existing approaches to deterministic parallelism. In this chapter, I describe some more recent contributions to the literature, and how the LVars model relates to them.

6.1. Deterministic Parallel Java (DPJ)

DPJ [8, 7] is a deterministic language consisting of a system of annotations for Java code. A sophisticated region-based type system ensures that a mutable region of the heap is, essentially, passed linearly to an exclusive writer, thereby ensuring that the state accessed by concurrent threads is disjoint. DPJ does, however, provide a way to unsafely assert that operations commute with one another (using the `commuteswith` form) to enable concurrent mutation.

The LVars model differs from DPJ in that it allows overlapping shared state between threads as the default. Moreover, since LVar effects are already commutative, we avoid the need for `commuteswith` annotations. Finally, it is worth noting that while in DPJ, commutativity annotations have to appear in application-level code, in LVish only the data-structure author needs to write trusted code. The application programmer can run untrusted code that still enjoys a (quasi-)determinism guarantee, because only (quasi-)deterministic programs can be expressed as LVish Par computations. More recently, Bocchino *et al.* [9] proposed a type and effect system that allows for the incorporation of nondeterministic sections of code in DPJ. The goal here is different from ours: while they aim to support *intentionally* nondeterministic computations such as those arising from optimization problems like branch-and-bound search, the quasi-determinism in LVish arises as a result of schedule nondeterminism.

6.2. FlowPools

Prokopec *et al.* [45] propose a data structure with an API closely related to LVars extended with freezing and handlers: a FlowPool is a bag (that is, a multiset) that allows concurrent insertions but forbids removals, a `seal` operation that forbids further updates, and combinators like `foreach` that invoke callbacks as data arrives in the pool. To retain determinism, the `seal` operation requires explicitly passing the expected bag *size* as an argument, and the program will raise an exception if the bag goes over the expected size.

While this interface has a flavor similar to that of LVars, it lacks the ability to detect quiescence, which is crucial for expressing algorithms like graph traversal, and the `seal` operation is awkward to use when the structure of data is not known in advance. By contrast, the `freeze` operation on LVars does not require such advance knowledge, but moves the model into the realm of quasi-determinism. Another important difference is the fact that LVars are *data structure-generic*: both our formalism and our library support an unlimited collection of data structures, whereas FlowPools are specialized to bags.

6.3. Bloom and Bloom^L

In Chapter 5, I presented a way to equip lattice-based distributed data structures with LVar-style threshold reads, resulting in a way to make deterministic *threshold queries* of those data structures. My approach is based on Shapiro *et al.*'s work on conflict-free replicated data types (CRDTs) [49, 48] and in particular their work on the lattice-based formulation of CRDTs, called *convergent replicated data types* or CvRDTs, which Chapter 5 discusses in detail.

Other authors have also used lattices as a framework for establishing formal guarantees about eventually consistent systems and distributed programs. The Bloom language for distributed database programming guarantees eventual consistency for distributed data collections that are updated monotonically. The initial formulation of Bloom [2] had a notion of monotonicity based on set inclusion, which is analogous to the store ordering relation in the (IVar-based) Featherweight CnC system that I described in Section 2.3.4. Later, Conway *et al.* [15] generalized Bloom to a more flexible lattice-parameterized system, Bloom^L, in a manner analogous to the generalization from IVars to LVars. Bloom^L combines ideas from the aforementioned work on CRDTs [49, 48] with *monotonic logic*, resulting in a lattice-parameterized, confluent language that is a close (but independently invented) relative to the LVars model. Bloom(^L) is implemented as a domain-specific language embedded in Ruby, and a monotonicity analysis pass rules out programs that would perform non-monotonic operations on distributed data collections (such as the removal of elements from a set). By contrast, in the LVars model (and in the LVish library), monotonicity is enforced by the API presented by LVars, and since the LVish library is implemented in Haskell,

we can rely on Haskell’s type system for fine-grained effect tracking and monadic encapsulation of LVar effects.

6.4. Concurrent Revisions

Burckhardt *et al.* [12] propose a formalism for eventual consistency based on graphs called *revision diagrams*, and Leijen, Burckhardt, and Fahndrich apply the revision diagrams approach to guaranteed-deterministic concurrent functional programming [33]. Their *Concurrent Revisions* (CR) programming model uses isolation types to distinguish regions of the heap shared by multiple mutators. Rather than enforcing exclusive access in the style of DPJ, CR clones a copy of the state for each mutator, using a deterministic “merge function” for resolving conflicts in local copies at join points.

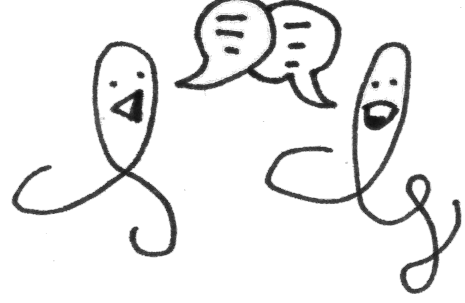
In CR, variables can be annotated as being shared between a “joiner” thread and a “joiner” thread. Unlike the lub writes of LVars, CR merge functions are *not* necessarily commutative; indeed, the default CR merge function is “joiner wins”. Determinism is enforced by the programming model allowing the programmer to specify which of two writing threads should prevail, regardless of the order in which those writes arrive, and the states that a shared variable can take on need not form a lattice. Still, semilattices turn up in the metatheory of CR: in particular, Burckhardt and Leijen [13] show that revision diagrams are semilattices, and that therefore, for any two vertices in a CR revision diagram, there exists a *greatest common ancestor* state that can be used to determine what changes each side has made—an interesting duality with the LVars model (in which any two LVar states have a lub).

6.5. Frame properties and separation logics

In Section 2.5.5, we saw that the Independence lemma, Lemma 2.5, expresses a *frame property* reminiscent of the following *frame rule* from separation logic and concurrent separation logic [43, 47, 42]:

$$\frac{\{p\} C \{q\}}{\{p * r\} C \{q * r\}}$$

Recall that the separating conjunction connective $*$ says that the assertions it combines can be satisfied in a non-overlapping manner; for instance, $p * r$ is satisfied by a heap if the heap can be split into two non-overlapping parts satisfying p and r , respectively. However, some-



times we do in fact want to allow some amount of “physical” overlap between resources, while retaining “logical” or “fictional” separation. In fact, the Independence lemma, since it replaces the separating conjunction with the lub operation, allows overlap between the original store and the “frame” store S'' ; indeed, the point of LVars is that total disjointness is unnecessary, since updates commute. Jensen and Birkedal’s recent work on *fictional separation logic* [28] explores the notion of fictional separation in detail, generalizing traditional separation logic to allow much more sophisticated kinds of sharing.

Even more recently, Dinsdale-Young *et al.* [17] introduced the “Views” framework, which brings the notion of fictional separation to a concurrent setting. The Views framework is a metatheory of concurrent reasoning principles that generalizes a variety of concurrent program logics and type systems, including concurrent separation logic. It provides a generalized frame rule, which is parameterized by a function f that is applied to the pre- and post-conditions in the conclusion of the rule:

$$\frac{\{p\} C \{q\}}{\{f(p)\} C \{f(q)\}}$$

In this formulation of the rule, the “frame” is an abstract piece of knowledge that is not violated by the execution of C . The Generalized Independence lemma (Lemma 3.7) that I describe in Section 3.3.5, which extends the Independence lemma to handle arbitrary update operations, is reminiscent of this generalized frame rule.

CHAPTER 7

Summary and future work

As single-assignment languages and Kahn process networks demonstrate, monotonicity can serve as the foundation of diverse deterministic-by-construction parallel programming models. The LVars programming model takes monotonicity as a starting point and generalizes single assignment to monotonic multiple assignment, parameterized by a lattice. The LVars model, and the accompanying LVish library, support my claim that lattice-based data structures are a general and practical unifying abstraction for deterministic and quasi-deterministic parallel and distributed programming.

7.1. Another look at the deterministic parallel landscape

Let us reconsider how LVars fit into the deterministic parallel programming landscape that we mapped out in Chapter 1:

- *No-shared-state parallelism*: The purely functional core of the λ_{LVar} and λ_{LVish} calculi (and of the LVish Haskell library) allow no-shared-state, pure task parallelism. Of course, shared-state programming is the point of the LVars model. However, it is significant that we take pure programming as a starting point, because it distinguishes the LVars model from approaches such as DPJ that begin with a parallel (but nondeterministic) language and then restrict the sharing of state to regain determinism. The LVars model works in the other direction: it begins with a deterministic parallel language without shared state, and then adds limited effects that retain determinism.
- *Data-flow parallelism*: As we have seen, because LVars are lattice-generic, the LVars model can subsume Kahn process networks and other parallel programming models based on data flow, since we can use LVars to represent a lattice of channel histories, ordered by a prefix ordering.

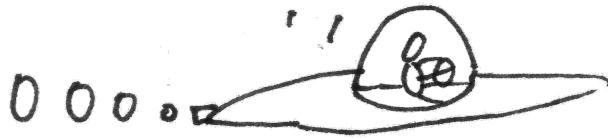
- *Single-assignment parallelism*: Single-assignment variables, or IVars, are also subsumed by LVars: an IVar is an LVar whose lattice has one “empty” state and multiple “full” states (where $\forall i. \text{empty} < \text{full}_i$). In fact, given how useful IVars are, the subsumption of IVars by LVars demonstrates that immutability is an important special case of monotonicity.¹
- *Imperative disjoint parallelism*: Although the LVars model generally does *not* require that the state accessed by concurrent threads is disjoint, this style of ensuring determinism is still compatible with the LVars model, and it is practically achievable using the ParST monad transformer in LVish, as we saw in Section 4.3.

In addition to subsuming or accommodating all these existing points on the landscape, we have identified a new class of *quasi-deterministic* programs and developed a programming model that supports quasi-determinism by construction. A quasi-deterministic model allows programs that perform *freezing* and are deterministic modulo write-after-freeze exceptions. The ability to freeze and read the exact contents of an LVar greatly increases the expressiveness of the LVars model, especially when used in conjunction with event handlers. Furthermore, we can regain full determinism by ensuring that freezing happens last, and, as we saw in Section 4.2.5, it is possible to enforce this “freeze after writing” requirement at the implementation level.

Of course, there is still more work to do. For example, although imperative disjoint parallelism and the LVars model seem to be compatible, as evidenced by the use of ParST in LVish, we have not yet formalized their relationship. In fact, this is an example of a general pattern in which the LVish library is usually one step ahead of the LVars formalism: to take another example, the LVish library supported the update operations of Section 2.6.1 (which are commutative and inflationary but not necessarily idempotent) well before the notion had been formalized in λ_{LVish} . Moreover, even for the parts of the LVish library that *are* fully accounted for in the model, we do not have proof that the library is a faithful implementation of the formal LVars model.

¹As Neil Conway puts it, “Immutability is a special case of monotone growth, albeit a particularly useful one” (https://twitter.com/neil_conway/status/392337034896871424).

Although it is unlikely that this game of catch-up can ever be won, an interesting direction to pursue for future work would be a *verified* implementation of LVish, for instance, in a dependently typed programming language. Even though a fully verified implementation of LVish (including the scheduler implementation) is an ambitious goal, a more manageable first step might be to implement individual LVar data structures in a dependently typed language such as Coq or Agda. The type system of such a language is rich enough to express properties that must be true of an LVar data structure, such as that the states that it can take on form a lattice and that writes are commutative and inflationary.



7.2. Distributed programming and the future of LVars and LVish

Most of this dissertation concerns the problem of how to program *parallel* systems, in which programs run on multiple processors that share memory. However, I am also concerned with the problem of how to program *distributed* systems, in which programs run on networked computers with distributed memory. Enormous bodies of work have been developed to deal with programming parallel and distributed systems, and one of the roles that programming languages research can play is to seek unifying abstractions between the two. It is in that spirit that I have explored the relationship of LVars to existing work on distributed systems.

LVars are a close cousin to convergent replicated data types (CvRDTs) [49, 48], which leverage lattice properties to guarantee that all replicas of an object (for instance, in a distributed database) are eventually consistent. Chapter 5 begins to explore the relationship between LVars and CvRDTs by porting LVar-style threshold reads to the CvRDT setting, but there is much more work to do here. Most immediately, although the idea of a single lattice-based framework for reasoning about both strongly consistent

and eventually consistent queries of distributed data is appealing and elegant, it is not yet clear what the compelling applications for threshold-readable CvRDTs are.

As a further step, it should also be possible to “back-port” ideas from the realm of CvRDTs to LVars. In fact, support for not-necessarily-idempotent updates to LVars was inspired in part by CvRDTs, which have always permitted arbitrary inflationary and commutative writes to individual replicas (the lub operation is only used when replicas’ states are *merged* with one another). The LVars model might further benefit from techniques pioneered in the work on CvRDTs to support data structures that allow seemingly non-monotonic updates, such as counters that support decrements as well as increments and sets that support removals as well as additions [49].

Finally, existing work on CvRDTs, as well as the work on distributed lattice-based programming languages like Bloom [2, 15], may serve as a source of inspiration for a future version of LVish that supports distributed execution.

Bibliography

- [1] Breadth-First Search, Parallel Boost Graph Library. http://www.boost.org/doc/libs/1_56_0/libs/graph_parallel/doc/html/breadth_first_search.html, 2009. 14, 16
- [2] P. Alvaro, N. Conway, J. Hellerstein, and W. R. Marczak. Consistency analysis in Bloom: a CALM and collected approach. In *CIDR*, 2011. 123, 129
- [3] Arvind, R. S. Nikhil, and K. K. Pingali. I-structures: data structures for parallel computing. *ACM Trans. Program. Lang. Syst.*, 11(4), Oct. 1989. 2, 12, 17, 121
- [4] H. Attiya, R. Guerraoui, D. Hendler, P. Kuznetsov, M. M. Michael, and M. Vechev. Laws of order: expensive synchronization in concurrent algorithms cannot be eliminated. In *POPL*, 2011. 82
- [5] D. A. Bader and K. Madduri. Designing multithreaded algorithms for breadth-first search and st-connectivity on the Cray MTA-2. In *ICPP*, 2006. 3
- [6] G. E. Blelloch, J. T. Fineman, P. B. Gibbons, and J. Shun. Internally deterministic parallel algorithms can be fast. In *PPoPP*, 2012. 34
- [7] R. L. Bocchino, Jr., V. S. Adve, S. V. Adve, and M. Snir. Parallel programming must be deterministic by default. In *HotPar*, 2009. 3, 11, 122
- [8] R. L. Bocchino, Jr., V. S. Adve, D. Dig, S. V. Adve, S. Heumann, R. Komuravelli, J. Overbey, P. Simmons, H. Sung, and M. Vakilian. A type and effect system for deterministic parallel Java. In *OOPSLA*, 2009. 3, 11, 94, 121, 122
- [9] R. L. Bocchino, Jr. et al. Safe nondeterminism in a deterministic-by-default parallel language. In *POPL*, 2011. 122
- [10] E. Brewer. CAP twelve years later: How the “rules” have changed. <http://www.infoq.com/articles/cap-twelve-years-later-how-the-rules-have-changed>, 2012. 108
- [11] Z. Budimlić, M. Burke, V. Cavé, K. Knobe, G. Lowney, R. Newton, J. Palsberg, D. Peixotto, V. Sarkar, F. Schlimbach, and S. Taşirlar. Concurrent Collections. *Sci. Program.*, 18(3-4), Aug. 2010. 3, 12, 26, 31, 42, 121
- [12] S. Burckhardt, M. Fahndrich, D. Leijen, and M. Sagiv. Eventually consistent transactions. In *ESOP*, 2012. 124
- [13] S. Burckhardt and D. Leijen. Semantics of concurrent revisions. In *ESOP*, 2011. 124
- [14] M. M. T. Chakravarty, R. Leshchinskiy, S. Peyton Jones, G. Keller, and S. Marlow. Data Parallel Haskell: A status report. In *DAMP*, 2007. 2

BIBLIOGRAPHY

- [15] N. Conway, W. Marczak, P. Alvaro, J. M. Hellerstein, and D. Maier. Logic and lattices for distributed programming. In *SOCC*, 2012. 123, 129
- [16] G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Vosshall, and W. Vogels. Dynamo: Amazon’s highly available key-value store. In *SOSP*, 2007. 108, 109
- [17] T. Dinsdale-Young, L. Birkedal, P. Gardner, M. Parkinson, and H. Yang. Views: Compositional reasoning for concurrent programs. In *POPL*, 2013. 75, 125
- [18] C. Earl, I. Sergey, M. Might, and D. Van Horn. Introspective pushdown analysis of higher-order programs. In *ICFP*, 2012. 102
- [19] M. Felleisen, R. B. Findler, and M. Flatt. *Semantics Engineering with PLT Redex*. The MIT Press, first edition, 2009. 9, 233
- [20] J. Felsenstein. PHYLIP (Phylogeny Inference Package) version 3.6. Distributed by the author, 2005. 105
- [21] M. Fluet, M. Rainey, J. Reppy, and A. Shaw. Implicitly-threaded parallelism in Manticore. In *ICFP*, 2008. 30
- [22] M. Frigo, C. E. Leiserson, and K. H. Randall. The implementation of the Cilk-5 multithreaded language. In *PLDI*, 1998. 11
- [23] S. Gilbert and N. Lynch. Brewer’s conjecture and the feasibility of consistent, available, partition-tolerant web services. *SIGACT News*, 33(2), June 2002. 108
- [24] M. I. Gordon, W. Thies, M. Karczmarek, J. Lin, A. S. Meli, C. Leger, A. A. Lamb, J. Wong, H. Hoffman, D. Z. Maze, and S. Amarasinghe. A stream compiler for communication-exposed architectures. In *ASPLOS*, 2002. 2, 12
- [25] S. Herhut, R. L. Hudson, T. Shpeisman, and J. Sreeram. River Trail: A path to parallelism in JavaScript. In *OOPSLA*, 2013. 2
- [26] M. Herlihy and N. Shavit. *The Art of Multiprocessor Programming*. Morgan Kaufmann, 2008. 82, 102
- [27] C. A. R. Hoare. An axiomatic basis for computer programming. *Commun. ACM*, 12(10), Oct. 1969. 36
- [28] J. B. Jensen and L. Birkedal. Fictional separation logic. In *ESOP*, 2012. 125
- [29] G. Kahn. The semantics of a simple language for parallel programming. In J. L. Rosenfeld, editor, *Information Processing ’74: Proceedings of the IFIP Congress*. North-Holland, 1974. 2, 12, 121
- [30] O. Kiselyov, A. Sabry, and C. Swords. Extensible effects: an alternative to monad transformers. In *Haskell*, 2013. 85
- [31] N. R. Krishnaswami. Higher-order reactive programming without spacetime leaks. In *ICFP*, 2013. 34
- [32] E. Lee and D. Messerschmitt. Synchronous data flow. *Proceedings of the IEEE*, 75(9), 1987. 2
- [33] D. Leijen, M. Fahndrich, and S. Burckhardt. Prettier concurrency: purely functional concurrent revisions. In *Haskell*, 2011. 121, 124
- [34] S. Marlow. *Parallel and Concurrent Programming in Haskell*. O’Reilly Media, 2013. 86

BIBLIOGRAPHY

- [35] S. Marlow, P. Maier, H.-W. Loidl, M. K. Aswad, and P. Trinder. Seq no more: better strategies for parallel Haskell. In *Haskell*, 2010. 2, 11, 98
- [36] S. Marlow, R. Newton, and S. Peyton Jones. A monad for deterministic parallelism. In *Haskell*, 2011. 3, 6, 11, 12, 16, 17, 80, 81
- [37] J. McCarthy and P. Hayes. Some philosophical problems from the standpoint of artificial intelligence. *Machine Intelligence*, 4, 1969. 37
- [38] M. Might. *k*-CFA: Determining types and/or control-flow in languages like Python, Java and Scheme. <http://matt.might.net/articles/implementation-of-kcfa-and-0cfa/>. 97
- [39] R. Newton. Bringing atomic memory operations to a lazy language, 2012. Haskell Implementors Workshop. 102
- [40] R. R. Newton and I. L. G. Newton. PhyBin: binning trees by topology. *PeerJ*, 1, Oct. 2013. 7, 104
- [41] R. S. Nikhil. Id language reference manual, 1991. 17
- [42] P. W. O’Hearn. Resources, concurrency, and local reasoning. *Theor. Comput. Sci.*, 375(1-3), Apr. 2007. 124
- [43] P. W. O’Hearn, J. C. Reynolds, and H. Yang. Local reasoning about programs that alter data structures. In *CSL*, 2001. 36, 37, 124
- [44] S. L. Peyton Jones, R. Leshchinskiy, G. Keller, and M. M. T. Chakravarty. Harnessing the multicores: Nested data parallelism in Haskell. In *FSTTCS*, 2008. 2, 121
- [45] A. Prokopec, H. Miller, T. Schlatter, P. Haller, and M. Odersky. FlowPools: a lock-free deterministic concurrent dataflow abstraction. In *LCPC*, 2012. 122
- [46] J. H. Reppy. *Concurrent Programming in ML*. Cambridge University Press, 1999. 3
- [47] J. C. Reynolds. Separation logic: A logic for shared mutable data structures. In *LICS*, 2002. 124
- [48] M. Shapiro, N. Preguiça, C. Baquero, and M. Zawirski. A comprehensive study of Convergent and Commutative Replicated Data Types. Technical Report RR-7506, INRIA, Jan. 2011. 8, 109, 111, 123, 128
- [49] M. Shapiro, N. Preguiça, C. Baquero, and M. Zawirski. Conflict-free replicated data types. In *SSS*, 2011. 8, 87, 109, 111, 112, 118, 123, 128, 129
- [50] D. Spoonhower, G. E. Blelloch, P. B. Gibbons, and R. Harper. Beyond nested parallelism: tight bounds on work-stealing overheads for parallel futures. In *SPAA*, 2009. 30
- [51] J. Sukumaran and M. T. Holder. DendroPy: a Python library for phylogenetic computing. *Bioinformatics*, 26, 2010. 105
- [52] S.-J. Sul and T. L. Williams. A randomized algorithm for comparing sets of phylogenetic trees. In *APBC*, 2007. 7, 104
- [53] D. Terei, D. Mazières, S. Marlow, and S. Peyton Jones. Safe Haskell. In *Haskell*, 2012. 11

BIBLIOGRAPHY

- [54] D. B. Terry, V. Prabhakaran, R. Kotla, M. Balakrishnan, M. K. Aguilera, and H. Abu-Libdeh. Consistency-based service level agreements for cloud storage. In *SOSP*, 2013. 109, 119
- [55] L. G. Tesler and H. J. Enea. A language design for concurrent processes. In *AFIPS*, 1968 (Spring). 12, 17, 121
- [56] W. Vogels. Eventually consistent. *Commun. ACM*, 52(1), Jan. 2009. 8, 108
- [57] W. Vogels. Choosing consistency. http://www.allthingsdistributed.com/2010/02/strong_consistency_simplified.html, 2010. 109, 119

APPENDIX A

Proofs

A.1. Proof of Lemma 2.1

Proof. Consider an arbitrary permutation π .

For part 1, we have to show that if $\sigma \longrightarrow \sigma'$ then $\pi(\sigma) \longrightarrow \pi(\sigma')$, and that if $\pi(\sigma) \longrightarrow \pi(\sigma')$ then $\sigma \longrightarrow \sigma'$.

For the forward direction of part 1, suppose $\sigma \longrightarrow \sigma'$.

We have to show that $\pi(\sigma) \longrightarrow \pi(\sigma')$.

We proceed by cases on the rule by which σ steps to σ' .

- Case E-Beta: $\sigma = \langle S; (\lambda x. e) v \rangle$, and $\sigma' = \langle S; e[x := v] \rangle$.

To show: $\pi(\langle S; (\lambda x. e) v \rangle) \longrightarrow \pi(\langle S; e[x := v] \rangle)$.

By Definitions 2.9 and 2.7, $\pi(\sigma) = \langle \pi(S); (\lambda x. \pi(e)) \pi(v) \rangle$.

By E-Beta, $\langle \pi(S); (\lambda x. \pi(e)) \pi(v) \rangle$ steps to $\langle \pi(S); \pi(e)[x := \pi(v)] \rangle$.

By Definition 2.7, $\langle \pi(S); \pi(e)[x := \pi(v)] \rangle$ is equal to $\langle \pi(S); \pi(e[x := v]) \rangle$.

Hence $\langle \pi(S); (\lambda x. \pi(e)) \pi(v) \rangle$ steps to $\langle \pi(S); \pi(e[x := v]) \rangle$,

which is equal to $\pi(\langle S; e[x := v] \rangle)$ by Definition 2.9.

Hence the case is satisfied.

- Case E-New: $\sigma = \langle S; \text{new} \rangle$, and $\sigma' = \langle S[l \mapsto \perp]; l \rangle$.

To show: $\pi(\langle S; \text{new} \rangle) \longrightarrow \pi(\langle S[l \mapsto \perp]; l \rangle)$.

By Definitions 2.9 and 2.7, $\pi(\sigma) = \langle \pi(S); \text{new} \rangle$.

By E-New, $\langle \pi(S); \text{new} \rangle$ steps to $\langle (\pi(S))[l' \mapsto \perp]; l' \rangle$, where $l' \notin \text{dom}(\pi(S))$.

It remains to show that $\langle (\pi(S))[l' \mapsto \perp]; l' \rangle$ is equal to $\pi(\langle S[l \mapsto \perp]; l \rangle)$.

By Definition 2.9, $\pi(\langle S[l \mapsto \perp]; l \rangle)$ is equal to $\langle \pi(S[l \mapsto \perp]); \pi(l) \rangle$,

which is equal to $\langle (\pi(S))[\pi(l) \mapsto \perp]; \pi(l) \rangle$.

So, we have to show that $\langle (\pi(S))[l' \mapsto \perp]; l' \rangle$ is equal to $\langle (\pi(S))[\pi(l) \mapsto \perp]; \pi(l) \rangle$.

Since we know (from the side condition of E-New) that $l \notin \text{dom}(S)$,

it follows that $\pi(l) \notin \pi(\text{dom}(S))$.

Therefore, in $\langle (\pi(S))[l' \mapsto \perp]; l' \rangle$, we can α -rename l' to $\pi(l)$, and so the two configurations are equal and the case is satisfied.

- Case E-Put: $\sigma = \langle S; \text{put } l \ d_2 \rangle$, and $\sigma' = \langle S[l \mapsto d_1 \sqcup d_2]; () \rangle$.

To show: $\pi(\langle S; \text{put } l \ d_2 \rangle) \hookrightarrow \pi(\langle S[l \mapsto d_1 \sqcup d_2]; () \rangle)$.

By Definitions 2.9 and 2.7, $\pi(\sigma) = \langle \pi(S); \text{put } \pi(l) \ d_2 \rangle$.

By E-Put, $\langle \pi(S); \text{put } \pi(l) \ d_2 \rangle$ steps to $\langle (\pi(S))[\pi(l) \mapsto d_1 \sqcup d_2]; () \rangle$,

since $S(l) = (\pi(S))(\pi(l)) = d_1$.

It remains to show that $\langle (\pi(S))[\pi(l) \mapsto d_1 \sqcup d_2]; () \rangle$ is equal to $\pi(\langle S[l \mapsto d_1 \sqcup d_2]; () \rangle)$.

By Definitions 2.9 and 2.7, $\pi(\langle S[l \mapsto d_1 \sqcup d_2]; () \rangle)$ is equal to $\langle (\pi(S))[\pi(l) \mapsto d_1 \sqcup d_2]; () \rangle$,

and so the two configurations are equal and the case is satisfied.

- Case E-Put-Err: $\sigma = \langle S; \text{put } l \ d_2 \rangle$, and $\sigma' = \mathbf{error}$.

To show: $\pi(\langle S; \text{put } l \ d_2 \rangle) \hookrightarrow \pi(\mathbf{error})$.

By Definitions 2.9 and 2.7, $\pi(\sigma) = \langle \pi(S); \text{put } \pi(l) \ d_2 \rangle$.

By E-Put-Err, $\langle \pi(S); \text{put } \pi(l) \ d_2 \rangle$ steps to \mathbf{error} ,

since $S(l) = (\pi(S))(\pi(l)) = d_1$.

Since $\pi(\mathbf{error}) = \mathbf{error}$ by Definition 2.9, the case is complete.

- Case E-Get: $\sigma = \langle S; \text{get } l \ T \rangle$, and $\sigma' = \langle S; d_2 \rangle$.

To show: $\pi(\langle S; \text{get } l \ T \rangle) \hookrightarrow \pi(\langle S; d_2 \rangle)$.

By Definitions 2.9 and 2.7, $\pi(\sigma) = \langle \pi(S); \text{get } \pi(l) \ T \rangle$.

By E-Get, $\langle \pi(S); \text{get } \pi(l) \ T \rangle$ steps to $\langle \pi(S); d_2 \rangle$,

since $S(l) = (\pi(S))(\pi(l)) = d_1$.

By Definitions 2.9 and 2.7, $\pi(\langle S; d_2 \rangle) \langle \pi(S); d_2 \rangle$.

Therefore the case is complete.

For the reverse direction of part 1, suppose $\pi(\sigma) \hookrightarrow \pi(\sigma')$.

We have to show that $\sigma \hookrightarrow \sigma'$.

We know from the forward direction of the proof that for all configurations σ and σ' and permutations π , if $\sigma \hookrightarrow \sigma'$ then $\pi(\sigma) \hookrightarrow \pi(\sigma')$.

Hence since $\pi(\sigma) \hookrightarrow \pi(\sigma')$, and since π^{-1} is also a permutation, we have that $\pi^{-1}(\pi(\sigma)) \hookrightarrow \pi^{-1}(\pi(\sigma'))$.

Since $\pi^{-1}(\pi(l)) = l$ for every $l \in Loc$, and that property lifts to configurations as well, we have that $\sigma \hookrightarrow \sigma'$.

For the forward direction of part 2, suppose $\sigma \mapsto \sigma'$.

We have to show that $\pi(\sigma) \mapsto \pi(\sigma')$.

By inspection of the operational semantics, σ must be of the form $\langle S; E[e] \rangle$, and σ' must be of the form $\langle S'; E[e'] \rangle$.

Hence we have to show that $\pi(\langle S; E[e] \rangle) \mapsto \pi(\langle S'; E[e'] \rangle)$.

By Definition 2.9, $\pi(\langle S; E[e] \rangle)$ is equal to $\langle \pi(S); \pi(E[e]) \rangle$.

Also by Definition 2.9, $\pi(\langle S'; E[e'] \rangle)$ is equal to $\langle \pi(S'); \pi(E[e']) \rangle$.

Furthermore, $\langle \pi(S); \pi(E[e]) \rangle$ is equal to $\langle \pi(S); (\pi(E))[\pi(e)] \rangle$ and $\langle \pi(S'); \pi(E[e']) \rangle$ is equal to $\langle \pi(S'); (\pi(E))[\pi(e')] \rangle$.

So we have to show that $\langle \pi(S); (\pi(E))[\pi(e)] \rangle \mapsto \langle \pi(S'); (\pi(E))[\pi(e')] \rangle$.

From the premise of E-Eval-Ctxt, $\langle S; e \rangle \hookrightarrow \langle S'; e' \rangle$.

Hence, by part 1, $\pi(\langle S; e \rangle) \hookrightarrow \pi(\langle S'; e' \rangle)$.

By Definition 2.9, $\pi(\langle S; e \rangle)$ is equal to $\langle \pi(S); \pi(e) \rangle$ and $\pi(\langle S'; e' \rangle)$ is equal to $\langle \pi(S'); \pi(e') \rangle$.

Hence $\langle \pi(S); \pi(e) \rangle \hookrightarrow \langle \pi(S'); \pi(e') \rangle$.

Therefore, by E-Eval-Ctxt, $\langle \pi(S); E[\pi(e)] \rangle \mapsto \langle \pi(S'); E[\pi(e')] \rangle$ for all evaluation contexts E .

In particular, it is true that $\langle \pi(S); (\pi(E))[\pi(e)] \rangle \mapsto \langle \pi(S'); (\pi(E))[\pi(e')] \rangle$, as we were required to show.

For the reverse direction of part 2, suppose $\pi(\sigma) \mapsto \pi(\sigma')$.

We have to show that $\sigma \mapsto \sigma'$.

We know from the forward direction of the proof that for all configurations σ and σ' and permutations π , if $\sigma \mapsto \sigma'$ then $\pi(\sigma) \mapsto \pi(\sigma')$.

Hence since $\pi(\sigma) \mapsto \pi(\sigma')$, and since π^{-1} is also a permutation, we have that $\pi^{-1}(\pi(\sigma)) \mapsto \pi^{-1}(\pi(\sigma'))$.

Since $\pi^{-1}(\pi(l)) = l$ for every $l \in Loc$, and that property lifts to configurations as well, we have that $\sigma \mapsto \sigma'$. □

A.2. Proof of Lemma 2.2

Proof. Suppose $\sigma \hookrightarrow \sigma'$ and $\sigma \hookrightarrow \sigma''$.

We have to show that there is a permutation π such that $\sigma' = \pi(\sigma'')$.

The proof is by cases on the rule by which σ steps to σ' .

- Case E-Beta:

Given: $\langle S; (\lambda x. e) v \rangle \hookrightarrow \langle S; e[x := v] \rangle$, and $\langle S; (\lambda x. e) v \rangle \hookrightarrow \sigma''$.

To show: There exists a π such that $\langle S; e[x := v] \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which $\langle S; (\lambda x. e) v \rangle$ can step is E-Beta.

Hence $\sigma'' = \langle S; e[x := v] \rangle$, and the case is satisfied by choosing π to be the identity function.

- Case E-New:

Given: $\langle S; \text{new} \rangle \longrightarrow \langle S[l \mapsto \perp]; l \rangle$, and $\langle S; \text{new} \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S[l \mapsto \perp]; l \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which $\langle S; \text{new} \rangle$ can step is E-New.

Hence $\sigma'' = \langle S[l' \mapsto \perp]; l' \rangle$.

Since, by the side condition of E-New, neither l nor l' occur in $\text{dom}(S)$, the case is satisfied by choosing π to be the permutation that maps l' to l and is the identity on every other element of Loc .

- Case E-Put:

Given: $\langle S; \text{put } l \ d_2 \rangle \longrightarrow \langle S[l \mapsto d_1 \sqcup d_2]; () \rangle$, and $\langle S; \text{put } l \ d_2 \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S[l \mapsto d_1 \sqcup d_2]; () \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, and since $d_1 \sqcup d_2 \neq \top$ (from the premise of E-Put), the only reduction rule by which $\langle S; \text{put } l \ d_2 \rangle$ can step is E-Put.

Hence $\sigma'' = \langle S[l \mapsto d_1 \sqcup d_2]; () \rangle$, and the case is satisfied by choosing π to be the identity function.

- Case E-Put-Err:

Given: $\langle S; \text{put } l \ d_2 \rangle \longrightarrow \mathbf{error}$, and $\langle S; \text{put } l \ d_2 \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\mathbf{error} = \pi(\sigma'')$.

By inspection of the operational semantics, and since $d_1 \sqcup d_2 = \top$ (from the premise of E-Put-Err), the only reduction rule by which $\langle S; \text{put } l \ d_2 \rangle$ can step is E-Put-Err.

Hence $\sigma'' = \mathbf{error}$, and the case is satisfied by choosing π to be the identity function.

- Case E-Get:

Given: $\langle S; \text{get } l \ T \rangle \longrightarrow \langle S; d_2 \rangle$, and $\langle S; \text{get } l \ T \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S; d_2 \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which $\langle S; \text{get } l \ T \rangle$ can step is E-Get.

Hence $\sigma'' = \langle S; d_2 \rangle$, and the case is satisfied by choosing π to be the identity function.

□

A.3. Proof of Lemma 2.3

Proof. Suppose $\langle S; E_1[e_1] \rangle \longmapsto \langle S_1; E_1[e'_1] \rangle$ and $\langle S; E_2[e_2] \rangle \longmapsto \langle S_2; E_2[e'_2] \rangle$ and $E_1[e_1] = E_2[e_2]$.

We are required to show that if $E_1 \neq E_2$, then there exist evaluation contexts E'_1 and E'_2 such that:

- $E'_1[e_1] = E_2[e'_2]$, and
- $E'_2[e_2] = E_1[e'_1]$, and
- $E'_1[e'_1] = E'_2[e'_2]$.

Let $e = E_1[e_1] = E_2[e_2]$.

The proof is by induction on the structure of the expression e .

Proceeding by cases on e :

- Case $e = x$: In this case, the only possible context that E_1 and E_2 can be is the empty context $[]$.

Therefore $E_1 = E_2$, and so the case holds vacuously.

- Case $e = v$: Similar to the case for x .

- Case $e = e_a e_b$:

If $E_1 = E_2$, the case holds vacuously.

Otherwise, we proceed as follows.

We know that $e_a e_b = E_1[e_1]$.

From the grammar of evaluation contexts, then, we know that either:

- $e_a e_b = E_1[e_1] = E_{11}[e_1] e_b$, where $E_{11}[e_1] = e_a$, or
- $e_a e_b = E_1[e_1] = e_a E_{12}[e_1]$, where $E_{12}[e_1] = e_b$.

Similarly, we know that $e_a e_b = E_2[e_2]$.

From the grammar of evaluation contexts, we know that either:

- $e_a e_b = E_2[e_2] = E_{21}[e_2] e_b$, where $E_{21}[e_2] = e_a$, or
- $e_a e_b = E_2[e_2] = e_a E_{22}[e_2]$, where $E_{22}[e_2] = e_b$.

(If $E_1 = []$ or $E_2 = []$, then $e_a e_b$ must be $(\lambda x. e') v$ for some e' and v , and neither $(\lambda x. e')$ nor v can step individually, so the other of E_1 or E_2 must be $[]$ as well, and so $E_1 = E_2$ and the case holds vacuously.)

This gives us four cases to consider:

- $E_{11}[e_1] = e_a$ and $E_{21}[e_2] = e_a$:

In this case, we know that $E_{11} \neq E_{21}$, because if $E_{11} = E_{21}$, we would have $e_1 = e_2$, which would mean that $E_1 = E_2$, a contradiction.

So, since $E_{11} \neq E_{21}$, by IH we have that there exist evaluation contexts E'_{11} and E'_{21} such that:

- * $E'_{11}[e_1] = E_{21}[e'_2]$, and
- * $E'_{21}[e_2] = E_{11}[e'_1]$, and
- * $E'_{11}[e'_1] = E'_{21}[e'_2]$.

Hence we can choose $E'_1 = E'_{11} e_b$ and $E'_2 = E'_{21} e_b$, which satisfy the criteria for E'_1 and E'_2 .

- $E_{12}[e_1] = e_b$ and $E_{22}[e_2] = e_b$: Similar to the previous case.
- $E_{11}[e_1] = e_a$ and $E_{22}[e_2] = e_b$:

In this case, we can choose $E'_1 = E_{11} E_{22}[e'_2]$, and $E'_2 = E_{11}[e'_1] E_{22}$, which satisfy the criteria for E'_1 and E'_2 .

- $E_{12}[e_1] = e_b$ and $E_{21}[e_2] = e_a$: Similar to the previous case.

- Case $e = \text{get } e_a e_b$: Similar to the case for $e_a e_b$.

- Case $e = \text{put } e_a e_b$: Similar to the case for $e_a e_b$.
- Case $e = \text{new}$: Similar to the case for x .

□

A.4. Proof of Lemma 2.4

Proof. Suppose $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$.

We are required to show that $S \sqsubseteq_S S'$.

The proof is by cases on the rule by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$.

- Case E-Beta:

Immediate by the definition of \sqsubseteq_S , since S does not change.

- Case E-New:

Given: $\langle S; \text{new} \rangle \longrightarrow \langle S[l \mapsto \perp]; l \rangle$.

To show: $S \sqsubseteq_S S[l \mapsto \perp]$.

By Definition 2.2, we have to show that $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto \perp])$ and that for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq (S[l \mapsto \perp])(l')$.

By definition, a store update operation on S can only either update an existing binding in S or extend S with a new binding.

Hence $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto \perp])$.

From the side condition of E-New, $l \notin \text{dom}(S)$.

Hence $S[l \mapsto \perp]$ adds a new binding for l in S .

Hence $S[l \mapsto \perp]$ does not update any existing bindings in S .

Hence, for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq (S[l \mapsto \perp])(l')$.

Therefore $S \sqsubseteq_S S[l \mapsto \perp]$, as required.

- Case E-Put:

Given: $\langle S; \text{put } l d_2 \rangle \longrightarrow \langle S[l \mapsto d_1 \sqcup d_2]; () \rangle$.

To show: $S \sqsubseteq_S S[l \mapsto d_1 \sqcup d_2]$.

By Definition 2.2, we have to show that $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto d_1 \sqcup d_2])$ and that for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq (S[l \mapsto d_1 \sqcup d_2])(l')$.

By definition, a store update operation on S can only either update an existing binding in S or extend S with a new binding.

Hence $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto d_1 \sqcup d_2])$.

From the premises of E-Put, $S(l) = d_1$.

Therefore $l \in \text{dom}(S)$.

Hence $S[l \mapsto d_1 \sqcup d_2]$ updates the existing binding for l in S from d_1 to $d_1 \sqcup d_2$.

By the definition of \sqcup , $d_1 \sqsubseteq (d_1 \sqcup d_2)$.

$S[l \mapsto d_1 \sqcup d_2]$ does not update any other bindings in S , hence, for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq (S[l \mapsto d_1 \sqcup d_2])(l')$.

Hence $S \sqsubseteq_S S[l \mapsto d_1 \sqcup d_2]$, as required.

- Case E-Put-Err:

Given: $\langle S; \text{put } l \ d_2 \rangle \hookrightarrow \mathbf{error}$.

By the definition of **error**, **error** is equal to $\langle \top_S; e \rangle$ for all e .

To show: $S \sqsubseteq_S \top_S$.

Immediate by the definition of \sqsubseteq_S .

- Case E-Get:

Immediate by the definition of \sqsubseteq_S , since S does not change.

□

A.5. Proof of Lemma 2.5

Proof. Suppose $\langle S; e \rangle \hookrightarrow \langle S'; e' \rangle$, where $\langle S'; e' \rangle \neq \mathbf{error}$.

Consider arbitrary S'' such that S'' is non-conflicting with $\langle S; e \rangle \hookrightarrow \langle S'; e' \rangle$ and $S' \sqcup_S S'' \neq \top_S$.

We are required to show that $\langle S \sqcup_S S''; e \rangle \longrightarrow \langle S' \sqcup_S S''; e' \rangle$.

The proof is by cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$.

Since $\langle S'; e' \rangle \neq \mathbf{error}$, we do not need to consider the E-Put-Err rule.

- Case E-Beta:

Given: $\langle S; (\lambda x. e) v \rangle \longrightarrow \langle S; e[x := v] \rangle$.

To show: $\langle S \sqcup_S S''; (\lambda x. e) v \rangle \longrightarrow \langle S \sqcup_S S''; e[x := v] \rangle$.

Immediate by E-Beta.

- Case E-New:

Given: $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto \perp]; l \rangle$.

To show: $\langle S \sqcup_S S''; \mathbf{new} \rangle \longrightarrow \langle (S[l \mapsto \perp]) \sqcup_S S''; l \rangle$.

By E-New, we have that $\langle S \sqcup_S S''; \mathbf{new} \rangle \longrightarrow \langle (S \sqcup_S S'')[l' \mapsto \perp]; l' \rangle$, where $l' \notin \text{dom}(S \sqcup_S S'')$.

By assumption, S'' is non-conflicting with $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto \perp]; l \rangle$.

Therefore $l \notin \text{dom}(S'')$.

From the side condition of E-New, $l \notin \text{dom}(S)$.

Therefore $l \notin \text{dom}(S \sqcup_S S'')$.

Therefore, in $\langle (S \sqcup_S S'')[l' \mapsto \perp]; l' \rangle$, we can α -rename l' to l , resulting in $\langle (S \sqcup_S S'')[l \mapsto \perp]; l \rangle$.

Therefore $\langle S \sqcup_S S''; \mathbf{new} \rangle \longrightarrow \langle (S \sqcup_S S'')[l \mapsto \perp]; l \rangle$.

Note that:

$$\begin{aligned}
 (S \sqcup_S S'')[l \mapsto \perp] &= S[l \mapsto \perp] \sqcup_S S''[l \mapsto \perp] \\
 &= S \sqcup_S [l \mapsto \perp] \sqcup_S S'' \sqcup_S [l \mapsto \perp] \\
 &= S \sqcup_S [l \mapsto \perp] \sqcup_S S'' \\
 &= S[l \mapsto \perp] \sqcup_S S''.
 \end{aligned}$$

Therefore $\langle S \sqcup_S S''; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto \perp] \sqcup_S S''; l \rangle$, as we were required to show.

- Case E-Put:

Given: $\langle S; \text{put } l \ d_2 \rangle \hookrightarrow \langle S[l \mapsto d_2]; () \rangle$.

To show: $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow \langle S[l \mapsto d_2] \sqcup_S S''; () \rangle$.

We will first show that

$$\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow \langle (S \sqcup_S S'')[l \mapsto d_2]; () \rangle$$

and then show why this is sufficient.

We proceed by cases on l :

– $l \notin \text{dom}(S'')$:

By assumption, $S[l \mapsto d_2] \sqcup_S S'' \neq \top_S$.

By Lemma 2.4, $S \sqsubseteq_S S[l \mapsto d_2]$.

Hence $S \sqcup_S S'' \neq \top_S$.

Therefore, by Definition 2.3, $(S \sqcup_S S'')(l) = S(l)$.

From the premises of E-Put, $S(l) = d_1$.

Hence $(S \sqcup_S S'')(l) = d_1$.

From the premises of E-Put, $d_2 = d_1 \sqcup d_2$ and $d_2 \neq \top$.

Therefore, by E-Put, we have: $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow \langle (S \sqcup_S S'')[l \mapsto d_2]; () \rangle$.

– $l \in \text{dom}(S'')$:

By assumption, $S[l \mapsto d_2] \sqcup_S S'' \neq \top_S$.

By Lemma 2.4, $S \sqsubseteq_S S[l \mapsto d_2]$.

Hence $S \sqcup_S S'' \neq \top_S$.

Therefore $(S \sqcup_S S'')(l) = S(l) \sqcup S''(l)$.

From the premises of E-Put, $S(l) = d_1$.

Hence $(S \sqcup_S S'')(l) = d'_1$, where $d_1 \sqsubseteq d'_1$.

From the premises of E-Put, $d_2 = d_1 \sqcup d_2$.

Let $d'_2 = d'_1 \sqcup d_2$.

Hence $d_2 \sqsubseteq d'_2$.

By assumption, $S[l \mapsto d_2] \sqcup_S S'' \neq \top_S$.

Therefore, by Definition 2.3, $d_2 \sqcup S''(l) \neq \top$.

Note that:

$$\begin{aligned}
 \top &\neq d_2 \sqcup S''(l) \\
 &= d_1 \sqcup d_2 \sqcup S''(l) \\
 &= S(l) \sqcup d_2 \sqcup S''(l) \\
 &= S(l) \sqcup S''(l) \sqcup d_2 \\
 &= (S \sqcup_S S'')(l) \sqcup d_2 \\
 &= d'_1 \sqcup d_2 \\
 &= d'_2.
 \end{aligned}$$

Hence $d'_2 \neq \top$.

Hence $(S \sqcup_S S'')(l) = d'_1$ and $d'_2 = d'_1 \sqcup d_2$ and $d'_2 \neq \top$.

Therefore, by E-Put we have: $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow \langle (S \sqcup_S S'')[l \mapsto d'_2]; () \rangle$.

Note that:

$$\begin{aligned}
 ((S \sqcup_S S'')[l \mapsto d'_2])(l) &= (S \sqcup_S S'')(l) \sqcup ([l \mapsto d'_2])(l) \\
 &= d'_1 \sqcup d'_2 \\
 &= d'_1 \sqcup d'_1 \sqcup d_2 \\
 &= d'_1 \sqcup d_2
 \end{aligned}$$

and

$$\begin{aligned}
 ((S \sqcup_S S'')[l \mapsto d_2])(l) &= (S \sqcup_S S'')(l) \sqcup ([l \mapsto d_2])(l) \\
 &= d'_1 \sqcup d_2 \\
 &= d'_1 \sqcup d_1 \sqcup d_2 \\
 &= d'_1 \sqcup d_2 \quad (\text{since } d_1 \sqsubseteq d'_1).
 \end{aligned}$$

Therefore $(S \sqcup_S S'')[l \mapsto d'_2] = (S \sqcup_S S'')[l \mapsto d_2]$.

Therefore, $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \longmapsto \langle (S \sqcup_S S'')[l \mapsto d_2]; () \rangle$.

Note that:

$$\begin{aligned}
 (S \sqcup_S S'')[l \mapsto d_2] &= S[l \mapsto d_2] \sqcup_S S''[l \mapsto d_2] \\
 &= S \sqcup_S [l \mapsto d_2] \sqcup_S S'' \sqcup_S [l \mapsto d_2] \\
 &= S \sqcup_S [l \mapsto d_2] \sqcup_S S'' \\
 &= S[l \mapsto d_2] \sqcup_S S''.
 \end{aligned}$$

Therefore $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \longmapsto \langle S[l \mapsto d_2] \sqcup_S S''; () \rangle$, as we were required to show.

- Case E-Get:

Given: $\langle S; \text{get } l \ T \rangle \longmapsto \langle S; d_2 \rangle$.

To show: $\langle S \sqcup_S S''; \text{get } l \ T \rangle \longmapsto \langle S \sqcup_S S''; d_2 \rangle$.

From the premises of E-Get, $S(l) = d_1$ and $\text{incomp}(T)$ and $d_2 \in T$ and $d_2 \sqsubseteq d_1$.

By assumption, $S \sqcup_S S'' \neq \top_S$.

Hence $(S \sqcup_S S'')(l) = d'_1$, where $d_1 \sqsubseteq d'_1$.

By the transitivity of \sqsubseteq , $d_2 \sqsubseteq d'_1$.

Hence, $(S \sqcup_S S'')(l) = d'_1$ and $\text{incomp}(T)$ and $d_2 \in T$ and $d_2 \sqsubseteq d'_1$.

Therefore, by E-Get, $\langle S \sqcup_S S''; \text{get } l \ T \rangle \longmapsto \langle S \sqcup_S S''; d_2 \rangle$, as we were required to show.

□

A.6. Proof of Lemma 2.6

Proof. Suppose $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$, where $\langle S'; e' \rangle \neq \mathbf{error}$.

Consider arbitrary S'' such that S'' is non-conflicting with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ and $S' \sqcup_S S'' = \top_S$.

We are required to show that there exists $i \leq 1$ such that $\langle S \sqcup_S S''; e \rangle \longrightarrow^i \mathbf{error}$.

The proof is by cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$.

Since $\langle S'; e' \rangle \neq \mathbf{error}$, we do not need to consider the E-Put-Err rule.

- Case E-Beta:

Given: $\langle S; (\lambda x. e) v \rangle \longrightarrow \langle S; e[x := v] \rangle$.

To show: $\langle S \sqcup_S S''; (\lambda x. e) v \rangle \longrightarrow^i \mathbf{error}$, where $i \leq 1$.

By assumption, $S \sqcup_S S'' = \top_S$.

Hence, by the definition of **error**, $\langle S \sqcup_S S''; (\lambda x. e) v \rangle = \mathbf{error}$.

Hence $\langle S \sqcup_S S''; (\lambda x. e) v \rangle \longrightarrow^i \mathbf{error}$, with $i = 0$.

- Case E-New:

Given: $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto \perp]; l \rangle$.

To show: $\langle S \sqcup_S S''; \mathbf{new} \rangle \longrightarrow^i \mathbf{error}$, where $i \leq 1$.

By E-New, $\langle S \sqcup_S S''; \mathbf{new} \rangle \longrightarrow \langle (S \sqcup_S S'')[l' \mapsto \perp]; l' \rangle$, where $l' \notin \text{dom}(S \sqcup_S S'')$.

By assumption, S'' is non-conflicting with $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto \perp]; l \rangle$.

Therefore $l \notin \text{dom}(S'')$.

From the side condition of E-New, $l \notin \text{dom}(S)$.

Therefore $l \notin \text{dom}(S \sqcup_S S'')$.

Therefore, in $\langle (S \sqcup_S S'')[l' \mapsto \perp]; l' \rangle$, we can α -rename l' to l ,

resulting in $\langle (S \sqcup_S S'')[l \mapsto \perp]; l \rangle$.

Therefore $\langle S \sqcup_S S''; \mathbf{new} \rangle \hookrightarrow \langle (S \sqcup_S S'')[l \mapsto \perp]; l \rangle$.

By assumption, $S[l \mapsto \perp] \sqcup_S S'' = \top_S$.

Note that:

$$\begin{aligned} \top_S &= S[l \mapsto \perp] \sqcup_S S'' \\ &= S \sqcup_S [l \mapsto \perp] \sqcup_S S'' \\ &= S \sqcup_S S'' \sqcup_S [l \mapsto \perp] \\ &= (S \sqcup_S S'') \sqcup_S [l \mapsto \perp] \\ &= (S \sqcup_S S'')[l \mapsto \perp]. \end{aligned}$$

Hence $\langle S \sqcup_S S''; \mathbf{new} \rangle \hookrightarrow \langle \top_S; l \rangle$.

Hence, by the definition of **error**, $\langle S \sqcup_S S''; \mathbf{new} \rangle \hookrightarrow \mathbf{error}$.

Hence $\langle S \sqcup_S S''; \mathbf{new} \rangle \hookrightarrow^i \mathbf{error}$, with $i = 1$.

- Case E-Put:

Given: $\langle S; \text{put } l \ d_2 \rangle \hookrightarrow \langle S[l \mapsto d_2]; () \rangle$.

To show: $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow^i \mathbf{error}$, where $i \leq 1$.

We proceed by cases on $S \sqcup_S S''$:

– $S \sqcup_S S'' = \top_S$:

In this case, by the definition of **error**, $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle = \mathbf{error}$.

Hence $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow^i \mathbf{error}$, with $i = 0$.

– $S \sqcup_S S'' \neq \top_S$:

From the premises of E-Put, we have that $S(l) = d_1$.

Hence $(S \sqcup_S S'')(l) = d'_1$, where $d_1 \sqsubseteq d'_1$.

We show that $d'_1 \sqcup d_2 = \top$, as follows:

By assumption, $S[l \mapsto d_2] \sqcup_S S'' = \top_S$.

Hence, by Definition 2.3, there exists some $l' \in \text{dom}(S[l \mapsto d_2]) \cap \text{dom}(S'')$ such that $(S[l \mapsto d_2])(l') \sqcup S''(l') = \top$.

Now case on l' :

* $l' \neq l$:

In this case, $(S[l \mapsto d_2])(l') = S(l')$.

Since $(S[l \mapsto d_2])(l') \sqcup S''(l') = \top$, we then have that $S(l') \sqcup S''(l') = \top$.

However, this is a contradiction since $S \sqcup_S S'' \neq \top_S$.

Hence this case cannot occur.

* $l' = l$:

Then $(S[l \mapsto d_2])(l) \sqcup S''(l) = \top$.

Note that:

$$\begin{aligned}
 \top &= (S[l \mapsto d_2])(l) \sqcup S''(l) \\
 &= d_2 \sqcup S''(l) \\
 &= d_1 \sqcup d_2 \sqcup S''(l) \\
 &= S(l) \sqcup d_2 \sqcup S''(l) \\
 &= S(l) \sqcup S''(l) \sqcup d_2 \\
 &= (S \sqcup_S S'')(l) \sqcup d_2 \\
 &= d'_1 \sqcup d_2.
 \end{aligned}$$

Hence $d'_1 \sqcup d_2 = \top$.

Hence, by E-Put-Err, $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow \mathbf{error}$.

Hence $\langle S \sqcup_S S''; \text{put } l \ d_2 \rangle \hookrightarrow^i \mathbf{error}$, with $i = 1$.

- Case E-Get:

Given: $\langle S; \text{get } l \ T \rangle \hookrightarrow \langle S; d_2 \rangle$.

To show: $\langle S \sqcup_S S''; \text{get } l \ T \rangle \hookrightarrow^i \mathbf{error}$, where $i \leq 1$.

By assumption, $S \sqcup_S S'' = \top_S$.

Hence, by the definition of **error**, $\langle S \sqcup_S S''; \text{get } l \ T \rangle = \mathbf{error}$.

Hence $\langle S \sqcup_S S''; \text{get } l \ T \rangle \hookrightarrow^i \mathbf{error}$, with $i = 0$.

□

A.7. Proof of Lemma 2.8

Proof. Suppose $\sigma \mapsto \sigma_a$ and $\sigma \mapsto \sigma_b$.

We have to show that there exist σ_c, i, j, π such that $\sigma_a \mapsto^i \sigma_c$ and $\pi(\sigma_b) \mapsto^j \sigma_c$ and $i \leq 1$ and $j \leq 1$.

By inspection of the operational semantics, it must be the case that σ steps to σ_a by the E-Eval-Ctxt rule.

Let $\sigma = \langle S; E_a[e_{a_1}] \rangle$ and let $\sigma_a = \langle S_a; E_a[e_{a_2}] \rangle$.

Likewise, it must be the case that σ steps to σ_b by the E-Eval-Ctxt rule.

Let $\sigma = \langle S; E_b[e_{b_1}] \rangle$ and let $\sigma_b = \langle S_b; E_b[e_{b_2}] \rangle$.

Note that $\sigma = \langle S; E_a[e_{a_1}] \rangle = \langle S; E_b[e_{b_1}] \rangle$, and so $E_a[e_{a_1}] = E_b[e_{b_1}]$, but E_a and E_b may differ and e_{a_1} and e_{b_1} may differ.

First, consider the possibility that $E_a = E_b$ (and $e_{a_1} = e_{b_1}$).

Since $\langle S; E_a[e_{a_1}] \rangle \mapsto \langle S_a; E_a[e_{a_2}] \rangle$ by E-Eval-Ctxt and $\langle S; E_b[e_{b_1}] \rangle \mapsto \langle S_b; E_b[e_{b_2}] \rangle$ by E-Eval-Ctxt, we have from the premise of E-Eval-Ctxt that $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ and $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$.

But then, since $e_{a_1} = e_{b_1}$, by Internal Determinism (Lemma 2.2) there is a permutation π' such that $\langle S_a; e_{a_2} \rangle = \pi'(\langle S_b; e_{b_2} \rangle)$.

Then we can satisfy the proof by choosing $\sigma_c = \langle S_a; e_{a_2} \rangle$ and $i = 0$ and $j = 0$ and $\pi = \pi'$.

The rest of this proof deals with the more interesting case in which $E_a \neq E_b$ (and $e_{a_1} \neq e_{b_1}$).

Since $\langle S; E_a[e_{a_1}] \rangle \mapsto \langle S_a; E_a[e_{a_2}] \rangle$ and $\langle S; E_b[e_{b_1}] \rangle \mapsto \langle S_b; E_b[e_{b_2}] \rangle$ and $E_a[e_{a_1}] = E_b[e_{b_1}]$, and since $E_a \neq E_b$, we have from Lemma 2.3 (Locality) that there exist evaluation contexts E'_a and E'_b such that:

- $E'_a[e_{a_1}] = E_b[e_{b_2}]$, and
- $E'_b[e_{b_1}] = E_a[e_{a_2}]$, and
- $E'_a[e_{a_2}] = E'_b[e_{b_2}]$.

In some of the cases that follow, we will choose $\sigma_c = \mathbf{error}$.

In most cases, however, our approach will be to show that there exist S', i, j, π such that:

- $\langle S_a; E_a[e_{a_2}] \rangle \mapsto^i \langle S'; E'_a[e_{a_2}] \rangle$, and
- $\pi(\langle S_b; E_b[e_{b_2}] \rangle) \mapsto^j \langle S'; E'_a[e_{a_2}] \rangle$.

Since $E'_a[e_{a_1}] = E_b[e_{b_2}]$, $E'_b[e_{b_1}] = E_a[e_{a_2}]$, and $E'_a[e_{a_2}] = E'_b[e_{b_2}]$, it suffices to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto^i \langle S'; E'_b[e_{b_2}] \rangle$, and
- $\pi(\langle S_b; E'_a[e_{a_1}] \rangle) \mapsto^j \langle S'; E'_a[e_{a_2}] \rangle$.

From the premise of E-Eval-Ctxt, we have that $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ and $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$.

We proceed by case analysis on the rule by which $\langle S; e_{a_1} \rangle$ steps to $\langle S_a; e_{a_2} \rangle$.

(1) Case E-Beta: We have $S_a = S$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$:

(a) Case E-Beta: We have $S_b = S$.

Choose $S' = S = S_a = S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_a; E'_b[e_{b_2}] \rangle$, and
- $\langle S; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$,

both of which follow immediately from $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ and $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ and E-Eval-Ctxt.

(b) Case E-New: We have $S_b = S[l \mapsto \perp]$.

Choose $S' = S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_b; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$.

The first of these follows immediately from $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ and E-Eval-Ctxt.

For the second, consider that $S_b = S[l \mapsto \perp] = S \sqcup_S [l \mapsto \perp]$.

Furthermore, since no locations are allocated in the transition $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, we know that $[l \mapsto \perp]$ is non-conflicting with it, and we know that $S_a \sqcup_S [l \mapsto \perp] \neq \top_S$ since S_a is just S and $S \sqcup_S [l \mapsto \perp]$ cannot be \top_S .

Therefore, by Lemma 2.5 (Independence), we have that $\langle S \sqcup_S [l \mapsto \perp]; e_{a_1} \rangle \hookrightarrow \langle S_a \sqcup_S [l \mapsto \perp]; e_{a_2} \rangle$.

Hence $\langle S_b; e_{a_1} \rangle \hookrightarrow \langle S_b; e_{a_2} \rangle$.

By E-Eval-Ctxt, it follows that $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$, as we were required to show.

(c) Case E-Put: We have $S_b = S[l \mapsto d_1 \sqcup d_2]$.

Choose $S' = S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_b; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$.

The first of these follows immediately from $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ and E-Eval-Ctxt.

For the second, consider that $S_b = S[l \mapsto d_1 \sqcup d_2] = S \sqcup_S [l \mapsto d_1 \sqcup d_2]$.

Furthermore, since no locations are allocated in the transition $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, we know that $[l \mapsto d_1 \sqcup d_2]$ is non-conflicting with it, and we know that $S_a \sqcup_S [l \mapsto d_1 \sqcup d_2] \neq$

\top_S since S_a is just S and $S \sqcup_S [l \mapsto d_1 \sqcup d_2]$ cannot be \top_S , since we know from the premise of E-Put that $d_1 \sqcup d_2 \neq \top$.

Therefore, by Lemma 2.5 (Independence), we have that $\langle S \sqcup_S [l \mapsto d_1 \sqcup d_2]; e_{a_1} \rangle \hookrightarrow \langle S_a \sqcup_S [l \mapsto d_1 \sqcup d_2]; e_{a_2} \rangle$.

Hence $\langle S_b; e_{a_1} \rangle \hookrightarrow \langle S_b; e_{a_2} \rangle$.

By E-Eval-Ctxt, it follows that $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$, as we were required to show.

(d) Case E-Put-Err:

Here $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 1$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

The second of these is immediately true because since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, and so $\langle S_b; E'_a[e_{a_1}] \rangle$ is equal to \mathbf{error} as well.

For the first, observe that $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, hence by E-Eval-Ctxt, $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_b; E'_b[e_{b_2}] \rangle$.

But $S_b = \top_S$, so $\langle S_b; E'_b[e_{b_2}] \rangle$ is equal to \mathbf{error} , and so $\langle S; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, as required.

(e) Case E-Get: Similar to case 1a, since $S_b = S$.

(2) Case E-New: We have $S_a = S[l \mapsto \perp]$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$:

(a) Case E-Beta: By symmetry with case 1b.

(b) Case E-New: We have $S_b = S[l' \mapsto \perp]$.

Now consider whether $l = l'$:

- If $l \neq l'$:

Choose $S' = S[l' \mapsto \perp][l \mapsto \perp]$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S[l' \mapsto \perp][l \mapsto \perp]; E'_b[e_{b_2}] \rangle$, and

$$- \langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S[l' \mapsto \perp][l \mapsto \perp]; E'_a[e_{a_2}] \rangle.$$

For the first of these, consider that $S_a = S[l \mapsto \perp] = S \sqcup_S [l \mapsto \perp]$, and that $S[l' \mapsto \perp][l \mapsto \perp] = S[l' \mapsto \perp] \sqcup_S [l \mapsto \perp]$.

Furthermore, since the only location allocated in the transition $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ is l' , we know that $[l \mapsto \perp]$ is non-conflicting with it (since $l \neq l'$ in this case).

We also know that $S[l' \mapsto \perp] \sqcup_S [l \mapsto \perp] \neq \top_S$, since $S \neq \top_S$ and new bindings of $l \mapsto \perp$ and $l' \mapsto \perp$ cannot cause it to become \top_S .

Therefore, by Lemma 2.5 (Independence), we have that $\langle S \sqcup_S [l \mapsto \perp]; e_{b_1} \rangle \hookrightarrow \langle S_b \sqcup_S [l \mapsto \perp]; e_{b_2} \rangle$.

Hence $\langle S[l \mapsto \perp]; e_{b_1} \rangle \hookrightarrow \langle S_b[l \mapsto \perp]; e_{b_2} \rangle$.

By E-Eval-Ctxt it follows that $\langle S[l \mapsto \perp]; E'_b[e_{b_1}] \rangle \hookrightarrow \langle S_b[l \mapsto \perp]; E'_b[e_{b_2}] \rangle$, which, since $S_b = S[l' \mapsto \perp]$, is what we were required to show.

The argument for the second is symmetrical.

- If $l = l'$:

In this case, observe that we do *not* want the expression in the final configuration to be $E'_a[e_{a_2}]$ (nor its equivalent, $E'_b[e_{b_2}]$).

The reason for this is that $E'_a[e_{a_2}]$ contains both occurrences of l .

Rather, we want both configurations to step to a configuration in which exactly one occurrence of l has been renamed to a fresh location l'' .

Let l'' be a location such that $l'' \notin \text{dom}(S)$ and $l'' \neq l$ (and hence $l'' \neq l'$, as well).

Then choose $S' = S[l'' \mapsto \perp][l \mapsto \perp]$, $i = 1$, $j = 1$, and $\pi = \{(l, l'')\}$.

Either $\langle S[l'' \mapsto \perp][l \mapsto \perp]; E'_a[\pi(e_{a_2})] \rangle$ or $\langle S[l'' \mapsto \perp][l \mapsto \perp]; E'_b[\pi(e_{b_2})] \rangle$ would work as a final configuration; we choose $\langle S[l'' \mapsto \perp][l \mapsto \perp]; E'_b[\pi(e_{b_2})] \rangle$.

We have to show that:

$$\begin{aligned} & - \langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S[l'' \mapsto \perp][l \mapsto \perp]; E'_b[\pi(e_{b_2})] \rangle, \text{ and} \\ & - \pi(\langle S_b; E'_a[e_{a_1}] \rangle) \mapsto \langle S[l'' \mapsto \perp][l \mapsto \perp]; E'_b[\pi(e_{b_2})] \rangle. \end{aligned}$$

For the first of these, since $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, we have by Lemma 2.1 (Permutability) that $\pi(\langle S; e_{b_1} \rangle) \hookrightarrow \pi(\langle S_b; e_{b_2} \rangle)$.

Since $\pi = \{(l, l'')\}$, but $l \notin S$ (from the side condition on E-New), we have that $\pi(\langle S; e_{b_1} \rangle) = \langle S; e_{b_1} \rangle$.

Since $\langle S_b; e_{b_2} \rangle = \langle S[l' \mapsto \perp]; l' \rangle$, and $l = l'$, we have that $\pi(\langle S_b; e_{b_2} \rangle) = \langle S[l'' \mapsto \perp]; \pi(e_{b_2}) \rangle$.

Hence $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto \perp]; \pi(e_{b_2}) \rangle$.

Since the only location allocated in the transition $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto \perp]; \pi(e_{b_2}) \rangle$ is l'' , we know that $[l \mapsto \perp]$ is non-conflicting with it.

We also know that $S[l'' \mapsto \perp] \sqcup_S [l \mapsto \perp] \neq \top_S$, since $S \neq \top_S$ and new bindings of $l'' \mapsto \perp$ and $l \mapsto \perp$ cannot cause it to become \top_S .

Therefore, by Lemma 2.5 (Independence), we have that

$$\langle S \sqcup_S [l \mapsto \perp]; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto \perp] \sqcup_S [l \mapsto \perp]; \pi(e_{b_2}) \rangle.$$

Hence $\langle S[l \mapsto \perp]; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto \perp][l \mapsto \perp]; \pi(e_{b_2}) \rangle$.

By E-Eval-Ctxt it follows that

$$\langle S[l \mapsto \perp]; E'_b[e_{b_1}] \rangle \hookrightarrow \langle S[l'' \mapsto \perp][l \mapsto \perp]; E'_b[\pi(e_{b_2})] \rangle, \text{ which, since } S[l \mapsto \perp] = S_a, \text{ is what we were required to show.}$$

For the second, observe that since $S_b = S[l \mapsto \perp]$, we have that $\pi(S_b) = S[l'' \mapsto \perp]$.

Also, since l does not occur in e_{a_1} , we have that $\pi(E'_a[e_{a_1}]) = (\pi(E'_a))[e_{a_1}]$.

Hence we have to show that $\langle S[l'' \mapsto \perp]; (\pi(E'_a))[e_{a_1}] \rangle \hookrightarrow \langle S[l'' \mapsto \perp][l \mapsto \perp]; E'_b[\pi(e_{b_2})] \rangle$.

Since the only location allocated in the transition $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ is l , we know that $[l'' \mapsto \perp]$ is non-conflicting with it.

We also know that $S_a \sqcup_S [l'' \mapsto \perp] \neq \top_S$, since $S_a = S[l \mapsto \perp]$ and $S \neq \top_S$ and new bindings of $l'' \mapsto \perp$ and $l \mapsto \perp$ cannot cause it to become \top_S .

Therefore, by Lemma 2.5 (Independence), we have that $\langle S \sqcup_S [l'' \mapsto \perp]; e_{a_1} \rangle \hookrightarrow \langle S_a \sqcup_S [l'' \mapsto \perp]; e_{a_2} \rangle$.

Hence $\langle S[l'' \mapsto \perp]; e_{a_1} \rangle \hookrightarrow \langle S[l'' \mapsto \perp][l \mapsto \perp]; e_{a_2} \rangle$.

By E-Eval-Ctxt it follows that $\langle S[l'' \mapsto \perp]; (\pi(E'_a))[e_{a_1}] \rangle \mapsto \langle S[l'' \mapsto \perp][l \mapsto \perp]; (\pi(E'_a))[e_{a_2}] \rangle$, which completes the case since $E'_b[\pi(e_{b_2})] = (\pi(E'_a))[e_{a_2}]$.

(c) Case E-Put: We have $S_b = S[l' \mapsto d_1 \sqcup d_2]$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S_b[l \mapsto \perp]; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b[l \mapsto \perp]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S_a = S[l \mapsto \perp] = S \sqcup_S [l \mapsto \perp]$, and that $S_b[l \mapsto \perp] = S_b \sqcup_S [l \mapsto \perp]$.

Furthermore, since no locations are allocated in the transition $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, we know that $[l \mapsto \perp]$ is non-conflicting with it.

We also know that $S_b \sqcup_S [l \mapsto \perp] \neq \top_S$, since $S_b = S[l' \mapsto d_1 \sqcup d_2]$ and we know from the premise of E-Put that $d_1 \sqcup d_2 \neq \top$.

Therefore, by Lemma 2.5 (Independence), we have that $\langle S \sqcup_S [l \mapsto \perp]; e_{b_1} \rangle \hookrightarrow \langle S_b \sqcup_S [l \mapsto \perp]; e_{b_2} \rangle$.

Hence $\langle S[l \mapsto \perp]; e_{b_1} \rangle \hookrightarrow \langle S_b[l \mapsto \perp]; e_{b_2} \rangle$.

By E-Eval-Ctxt, it follows that $\langle S[l \mapsto \perp]; E'_b[e_{b_1}] \rangle \mapsto \langle S_b[l \mapsto \perp]; E'_b[e_{b_2}] \rangle$, which, since $S_a = S[l \mapsto \perp]$, is what we were required to show.

For the second, consider that $S_b = S \sqcup_S [l' \mapsto d_1 \sqcup d_2]$ and $S_b[l \mapsto \perp] = S[l \mapsto \perp] \sqcup_S [l' \mapsto d_1 \sqcup d_2] = S_a \sqcup_S [l' \mapsto d_1 \sqcup d_2]$.

Furthermore, since the only location allocated in the transition $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ is l , we know that $[l' \mapsto d_1 \sqcup d_2]$ is non-conflicting with it.

(We know that $l \neq l'$ because we have from the premise of E-Put that $l' \in \text{dom}(S)$, but we have from the side condition of E-New that $l \notin \text{dom}(S)$.)

We also know that $S[l \mapsto \perp] \sqcup_S [l' \mapsto d_1 \sqcup d_2] \neq \top_S$, since we know from the premise of E-Put that $d_1 \sqcup d_2 \neq \top$.

Therefore, by Lemma 2.5 (Independence), we have that $\langle S \sqcup_S [l' \mapsto d_1 \sqcup d_2]; e_{a_1} \rangle \hookrightarrow \langle S_a \sqcup_S [l' \mapsto d_1 \sqcup d_2]; e_{a_2} \rangle$.

Hence $\langle S_b; e_{a_1} \rangle \hookrightarrow \langle S_b[l \mapsto \perp]; e_{a_2} \rangle$.

By E-Eval-Ctxt, it follows that $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b[l \mapsto \perp]; E'_a[e_{a_2}] \rangle$, as we were required to show.

(d) Case E-Put-Err:

Here $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 1$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

The second of these is immediately true because since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, and so $\langle S_b; E'_a[e_{a_1}] \rangle$ is equal to \mathbf{error} as well.

For the first, observe that since $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, we have by Lemma 2.4 (Monotonicity) that $S \sqsubseteq_S S_a$.

Therefore, since $\langle S; e_{b_1} \rangle \hookrightarrow \mathbf{error}$,

we have by Lemma 2.7 (Error Preservation) that $\langle S_a; e_{b_1} \rangle \hookrightarrow \mathbf{error}$.

Since \mathbf{error} is equal to $\langle \top_S; e \rangle$ for all expressions e , $\langle S_a; e_{b_1} \rangle \hookrightarrow \langle \top_S; e \rangle$ for all e .

Therefore, by E-Eval-Ctxt, $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle \top_S; E'_b[e] \rangle$ for all e .

Since $\langle \top_S; E'_b[e] \rangle$ is equal to \mathbf{error} , we have that $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, as we were required to show.

(e) Case E-Get: Similar to case 2a, since $S_b = S$.

(3) Case E-Put: We have $S_a = S[l \mapsto d_1 \sqcup d_2]$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$:

(a) Case E-Beta: By symmetry with case 1c.

(b) Case E-New: By symmetry with case 2c.

(c) Case E-Put: We have $S_b = S[l' \mapsto d'_1 \sqcup d'_2]$, where $d'_1 = S(l')$.

Consider whether $S_b \sqcup_S [l \mapsto d_1 \sqcup d_2] = \top_S$:

- $S_b \sqcup_S [l \mapsto d_1 \sqcup d_2] \neq \top_S$:

Choose $S' = S_a \sqcup_S S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S_a \sqcup_S S_b; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_a \sqcup_S S_b; E'_a[e_{a_2}] \rangle$.

For the first of these, since no locations are allocated during the transition $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, we know that $[l \mapsto d_1 \sqcup d_2]$ is non-conflicting with it, and in this subcase, we know that $S_b \sqcup_S [l \mapsto d_1 \sqcup d_2] \neq \top_S$.

Therefore, by Lemma 2.5 (Independence), we have that $\langle S \sqcup_S [l \mapsto d_1 \sqcup d_2]; e_{b_1} \rangle \hookrightarrow \langle S_b \sqcup_S [l \mapsto d_1 \sqcup d_2]; e_{b_2} \rangle$.

By E-Eval-Ctxt, it follows that

$$\langle S \sqcup_S [l \mapsto d_1 \sqcup d_2]; E'_b[e_{b_1}] \rangle \mapsto \langle S_b \sqcup_S [l \mapsto d_1 \sqcup d_2]; E'_b[e_{b_2}] \rangle.$$

Since $S \sqcup_S [l \mapsto d_1 \sqcup d_2] = S[l \mapsto d_1 \sqcup d_2] = S_a$, we have that $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S_b \sqcup_S [l \mapsto d_1 \sqcup d_2]; E'_b[e_{b_2}] \rangle$.

Furthermore, since $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, by Lemma 2.4 (Monotonicity), we have that $S \sqsubseteq_S S_b$.

Therefore $S_b \sqcup_S [l \mapsto d_1 \sqcup d_2] = S_b \sqcup_S S \sqcup_S [l \mapsto d_1 \sqcup d_2] = S_b \sqcup_S S_a = S_a \sqcup_S S_b$.

So we have that $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S_a \sqcup_S S_b; E'_b[e_{b_2}] \rangle$, as we were required to show.

The argument for the second is symmetrical, with $[l' \mapsto d'_1 \sqcup d'_2]$ being the store that is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$.

- $S_b \sqcup_S [l \mapsto d_1 \sqcup d_2] = \top_S$:

Here we choose $\sigma_c = \mathbf{error}$ and $\pi = \text{id}$.

We have to show that there exist $i \leq 1$ and $j \leq 1$ such that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto^i \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto^j \mathbf{error}$.

For the first of these, since no locations are allocated during the transition $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, we know that $[l \mapsto d_1 \sqcup d_2]$ is non-conflicting with it, and in this subcase, we know that $S_b \sqcup_S [l \mapsto d_1 \sqcup d_2] = \top_S$.

Therefore, by Lemma 2.6 (Clash), we have that $\langle S \sqcup_S [l \mapsto d_1 \sqcup d_2]; e_{b_1} \rangle \hookrightarrow^{i'} \mathbf{error}$, where $i' \leq 1$.

Since \mathbf{error} is equal to $\langle \top_S; e \rangle$ for all expressions e , $\langle S \sqcup_S [l \mapsto d_1 \sqcup d_2]; e_{b_1} \rangle \hookrightarrow^{i'} \langle \top_S; e \rangle$ for all e .

Now consider whether $i' = 1$ or $i' = 0$:

- If $i' = 1$, then by E-Eval-Ctxt, $\langle S \sqcup_S [l \mapsto d_1 \sqcup d_2]; E'_b[e_{b_1}] \rangle \mapsto \langle \top_S; E'_b[e] \rangle$ for all e .

Since $\langle \top_S; E'_b[e] \rangle$ is equal to \mathbf{error} , and since $S \sqcup_S [l \mapsto d_1 \sqcup d_2] = S[l \mapsto d_1 \sqcup d_2] = S_a$, we choose $i = 1$ and we have that $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, as required.

- If $i' = 0$, then $\langle S \sqcup_S [l \mapsto d_1 \sqcup d_2]; e_{b_1} \rangle = \mathbf{error}$.

Hence $S \sqcup_S [l \mapsto d_1 \sqcup d_2] = \top_S$.

So, we choose $i = 0$, and since $S_a = S[l \mapsto d_1 \sqcup d_2] = S \sqcup_S [l \mapsto d_1 \sqcup d_2] = \top_S$, we have that $\langle S_a; E'_b[e_{b_1}] \rangle = \mathbf{error}$, as required.

The argument for the second is symmetrical, with $[l' \mapsto d'_1 \sqcup d'_2]$ being the store that is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$.

(d) Case E-Put-Err:

Here $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 1$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

The second of these is immediately true because since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, and so $\langle S_b; E'_a[e_{a_1}] \rangle$ is equal to \mathbf{error} as well.

For the first, observe that since $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, we have by Lemma 2.4 (Monotonicity) that $S \sqsubseteq_S S_a$.

Therefore, since $\langle S; e_{b_1} \rangle \hookrightarrow \mathbf{error}$,

we have by Lemma 2.7 (Error Preservation) that $\langle S_a; e_{b_1} \rangle \hookrightarrow \mathbf{error}$.

Since \mathbf{error} is equal to $\langle \top_S; e \rangle$ for all expressions e , $\langle S_a; e_{b_1} \rangle \hookrightarrow \langle \top_S; e \rangle$ for all e .

Therefore, by E-Eval-Ctxt, $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle \top_S; E'_b[e] \rangle$ for all e .

Since $\langle \top_S; E'_b[e] \rangle$ is equal to \mathbf{error} , we have that $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, as we were required to show.

(e) Case E-Get: Similar to case 3a, since $S_b = S$.

(4) Case E-Put-Err: We have $\langle S_a; e_{a_2} \rangle = \mathbf{error}$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$:

(a) Case E-Beta: By symmetry with case 1d.

(b) Case E-New: By symmetry with case 2d.

(c) Case E-Put: By symmetry with case 3d.

(d) Case E-Put-Err:

Here $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 0$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle = \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

Since $\langle S_a; e_{a_2} \rangle = \mathbf{error}$, $S_a = \top_S$, and since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, so both of the above follow immediately.

(e) Case E-Get: Similar to case 4a, since $S_b = S$.

(5) Case E-Get:

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$:

- (a) Case E-Beta: By symmetry with case 1e.
- (b) Case E-New: By symmetry with case 2e.
- (c) Case E-Put: By symmetry with case 3e.
- (d) Case E-Put-Err: By symmetry with case 4e.
- (e) Case E-Get: Similar to case 5a, since $S_b = S$.

□

A.8. Proof of Lemma 2.9

Proof. Suppose $\sigma \mapsto \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq m$.

We have to show that there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq 1$.

We proceed by induction on m .

In the base case of $m = 1$, the result is immediate from Lemma 2.8.

For the induction step, suppose $\sigma \mapsto^m \sigma'' \mapsto \sigma'''$ and suppose the lemma holds for m .

We show that it holds for $m + 1$, as follows.

We are required to show that there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma''') \mapsto^j \sigma_c$ and $i \leq m + 1$ and $j \leq 1$.

From the induction hypothesis, there exist σ'_c, i', j', π' such that $\sigma' \mapsto^{i'} \sigma'_c$ and $\pi'(\sigma''') \mapsto^{j'} \sigma'_c$ and $i' \leq m$ and $j' \leq 1$.

We proceed by cases on j' :

- If $j' = 0$, then $\pi'(\sigma''') = \sigma'_c$.

Since $\sigma'' \mapsto \sigma'''$, we have that $\pi'(\sigma'') \mapsto \pi'(\sigma''')$ by Lemma 2.1 (Permutability).

We can then choose $\sigma_c = \pi'(\sigma''')$ and $i = i' + 1$ and $j = 0$ and $\pi = \pi'$.

The key is that $\sigma' \mapsto^{i'} \sigma'_c = \pi'(\sigma'') \mapsto \pi'(\sigma''')$ for a total of $i' + 1$ steps.

- If $j' = 1$:

First, since $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$, then by Lemma 2.1 (Permutability) we have that $\sigma'' \mapsto^{j'} \pi'^{-1}(\sigma'_c)$. Then, by $\sigma'' \mapsto^{j'} \pi'^{-1}(\sigma'_c)$ and $\sigma'' \mapsto \sigma'''$ and Lemma 2.8 (Strong Local Confluence), we have that there exist σ''_c and i'' and j'' and π'' such that $\pi'^{-1}(\sigma'_c) \mapsto^{i''} \sigma''_c$ and $\pi''(\sigma''') \mapsto^{j''} \sigma''_c$ and $i'' \leq 1$ and $j'' \leq 1$.

Since $\pi'^{-1}(\sigma'_c) \mapsto^{i''} \sigma''_c$, by Lemma 2.1 (Permutability) we have that $\sigma'_c \mapsto^{i''} \pi'(\sigma''_c)$.

So we also have $\sigma' \mapsto^{i'} \sigma'_c \mapsto^{i''} \pi'(\sigma''_c)$.

Since $\pi''(\sigma''') \mapsto^{j''} \sigma''_c$, by Lemma 2.1 (Permutability) we have that $\pi'(\pi''(\sigma''')) \mapsto^{j''} \pi'(\sigma''_c)$.

In summary, we pick $\sigma_c = \pi'(\sigma''_c)$ and $i = i' + i''$ and $j = j''$ and $\pi = \pi'' \circ \pi'$, which is sufficient because $i = i' + i'' \leq m + 1$ and $j = j'' \leq 1$.

□

A.9. Proof of Lemma 2.10

Proof. Suppose that $\sigma \mapsto^n \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq n$ and $1 \leq m$.

We have to show that there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq n$.

We proceed by induction on n .

In the base case of $n = 1$, the result is immediate from Lemma 2.9.

For the induction step, suppose $\sigma \mapsto^n \sigma' \mapsto \sigma'''$ and suppose the lemma holds for n .

We show that it holds for $n + 1$, as follows.

We are required to show that there exist σ_c, i, j, π such that $\sigma''' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq n + 1$.

From the induction hypothesis, there exist σ'_c, i', j', π' such that $\sigma' \mapsto^{i'} \sigma'_c$ and $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$ and $i' \leq m$ and $j' \leq n$.

We proceed by cases on i' :

- If $i' = 0$, then $\sigma' = \sigma'_c$.

We can then choose $\sigma_c = \sigma'''$ and $i = 0$ and $j = j' + 1$ and $\pi = \pi'$.

Since $\pi'(\sigma'') \mapsto^{j'} \sigma'_c \mapsto \sigma'''$, and $j' + 1 \leq n + 1$ since $j' \leq n$, the case is satisfied.

- If $i' \geq 1$:

From $\sigma' \mapsto \sigma'''$ and $\sigma' \mapsto^{i'} \sigma'_c$ and Lemma 2.9, we have that there exist σ''_c and i'' and j'' and π'' such that $\sigma''' \mapsto^{i''} \sigma''_c$ and $\pi''(\sigma'_c) \mapsto^{j''} \sigma''_c$ and $i'' \leq i'$ and $j'' \leq 1$.

Since $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$, by Lemma 2.1 (Permutability) we have that $\pi''(\pi'(\sigma'')) \mapsto^{j'} \pi''(\sigma'_c)$.

So we also have $\pi''(\pi'(\sigma'')) \mapsto^{j'} \pi''(\sigma'_c) \mapsto^{j''} \sigma''_c$.

In summary, we pick $\sigma_c = \sigma''_c$ and $i = i''$ and $j = j' + j''$ and $\pi = \pi' \circ \pi''$, which is sufficient because $i = i'' \leq i' \leq m$ and $j = j' + j'' \leq n + 1$.

□

A.10. Proof of Lemma 3.2

Proof. Suppose that $(D, \sqsubseteq, \perp, \top)$ is a lattice and $(D_p, \sqsubseteq_p, \perp_p, \top_p) = \text{Freeze}(D, \sqsubseteq, \perp, \top)$.

In order to show that $(D_p, \sqsubseteq_p, \perp_p, \top_p)$ is a lattice, we have to show that:

- (1) \sqsubseteq_p is a partial order over D_p .
- (2) Every nonempty finite subset of D_p has a lub.
- (3) \perp_p is the least element of D_p .
- (4) \top_p is the greatest element of D_p .

We prove each of these properties in turn:

(1) \sqsubseteq_p is a partial order over D_p .

To show this, we need to show that \sqsubseteq_p is reflexive, transitive, and antisymmetric.

(a) \sqsubseteq_p is reflexive.

Suppose $v \in D_p$.

Then, by Lemma 3.1, either $v = (d, \text{false})$ with $d \in D$, or $v = (x, \text{true})$ with $x \in X$, where $X = D - \{\top\}$.

- Suppose $v = (d, \text{false})$:

By the reflexivity of \sqsubseteq , we know $d \sqsubseteq d$.

By the definition of \sqsubseteq_p , we know $(d, \text{false}) \sqsubseteq_p (d, \text{false})$.

- Suppose $v = (x, \text{true})$:

By the reflexivity of equality, $x = x$.

By the definition of \sqsubseteq_p , we know $(x, \text{true}) \sqsubseteq_p (x, \text{true})$.

(b) \sqsubseteq_p is transitive.

Suppose $v_1 \sqsubseteq_p v_2$ and $v_2 \sqsubseteq_p v_3$.

We want to show that $v_1 \sqsubseteq_p v_3$.

We proceed by case analysis on v_1, v_2 , and v_3 .

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (d_2, \text{false})$ and $v_3 = (d_3, \text{false})$:

By inversion on \sqsubseteq_p , it follows that $d_1 \sqsubseteq d_2$.

By inversion on \sqsubseteq_p , it follows that $d_2 \sqsubseteq d_3$.

By the transitivity of \sqsubseteq , we know $d_1 \sqsubseteq d_3$.

By the definition of \sqsubseteq_p , it follows that $(d_1, \text{false}) \sqsubseteq_p (d_3, \text{false})$.

Hence $v_1 \sqsubseteq_p v_3$.

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (d_2, \text{false})$ and $v_3 = (x_3, \text{true})$:

By inversion on \sqsubseteq_p , it follows that $d_1 \sqsubseteq d_2$.

By inversion on \sqsubseteq_p , it follows that $d_2 \sqsubseteq x_3$.

By the transitivity of \sqsubseteq , we know $d_1 \sqsubseteq x_3$.

By the definition of \sqsubseteq_p , it follows that $(d_1, \text{false}) \sqsubseteq_p (x_3, \text{true})$.

Hence $v_1 \sqsubseteq_p v_3$.

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (x_2, \text{true})$ and $v_3 = (d_3, \text{false})$:

By inversion on \sqsubseteq_p , it follows that $d_1 \sqsubseteq x_2$.

By inversion on \sqsubseteq_p , it follows that $d_3 = \top$.

Since \top is the maximal element of D , we know $d_1 \sqsubseteq \top \equiv d_3$.

By the definition of \sqsubseteq_p , it follows that $(d_1, \text{false}) \sqsubseteq_p (d_3, \text{false})$.

Hence $v_1 \sqsubseteq_p v_3$.

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (x_2, \text{true})$ and $v_3 = (x_3, \text{true})$:

By inversion on \sqsubseteq_p , it follows that $d_1 \sqsubseteq x_2$.

By inversion on \sqsubseteq_p , it follows that $x_2 = x_3$.

Hence $d_1 \sqsubseteq x_3$.

By the definition of \sqsubseteq_p , it follows that $(d_1, \text{false}) \sqsubseteq_p (x_3, \text{true})$.

Hence $v_1 \sqsubseteq_p v_3$.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (d_2, \text{false})$ and $v_3 = (d_3, \text{false})$:

By inversion on \sqsubseteq_p , it follows that $d_2 = \top$.

By inversion on \sqsubseteq_p , it follows that $d_2 \sqsubseteq d_3$.

Since \top is maximal, it follows that $d_3 = \top$.

By the definition of \sqsubseteq_p , it follows that $(x_1, \text{true}) \sqsubseteq_p (d_3, \text{false})$.

Hence $v_1 \sqsubseteq_p v_3$.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (d_2, \text{false})$ and $v_3 = (x_3, \text{true})$:

By inversion on \sqsubseteq_p , it follows that $d_2 = \top$.

By inversion on \sqsubseteq_p , it follows that $d_2 \sqsubseteq x_3$.

Since \top is maximal, it follows that $x_3 = \top$.

But since $x_3 \in X \subseteq D / \{\top\}$, we know $x_3 \neq \top$.

This is a contradiction.

Hence $v_1 \sqsubseteq_p v_3$.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (x_2, \text{true})$ and $v_3 = (d_3, \text{false})$:

By inversion on \sqsubseteq_p , it follows that $x_1 = x_2$.

By inversion on \sqsubseteq_p , it follows that $d_3 = \top$.

By the definition of \sqsubseteq_p , it follows that $(x_1, \text{true}) \sqsubseteq_p (d_3, \text{false})$.

Hence $v_1 \sqsubseteq_p v_3$.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (x_2, \text{true})$ and $v_3 = (x_3, \text{true})$:

By inversion on \sqsubseteq_p , it follows that $x_1 = x_2$.

By inversion on \sqsubseteq_p , it follows that $x_2 = x_3$.

By transitivity of $=$, $x_1 = x_3$.

By the definition of \sqsubseteq_p , it follows that $(x_1, \text{true}) \sqsubseteq_p (x_3, \text{true})$.

Hence $v_1 \sqsubseteq_p v_3$.

(c) \sqsubseteq_p is antisymmetric.

Suppose $v_1 \sqsubseteq_p v_2$ and $v_2 \sqsubseteq_p v_1$. Now, we proceed by cases on v_1 and v_2 .

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (d_2, \text{false})$:

By inversion on $v_1 \sqsubseteq_p v_2$, we know that $d_1 \sqsubseteq d_2$.

By inversion on $v_2 \sqsubseteq_p v_1$, we know that $d_2 \sqsubseteq d_1$.

By the antisymmetry of \leq , we know $d_1 = d_2$.

Hence $v_1 = v_2$.

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (x_2, \text{true})$:

By inversion on $v_1 \sqsubseteq_p v_2$, we know that $d_1 \sqsubseteq x_2$.

By inversion on $v_2 \sqsubseteq_p v_1$, we know that $d_1 = \top$.

Since \top is maximal in D , we know $x_2 = \top$.

But since $x_2 \in X \subseteq D / \{\top\}$, we know $x_2 \neq \top$.

This is a contradiction.

Hence $v_1 = v_2$.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (d_2, \text{false})$:

Similar to the previous case.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (x_2, \text{true})$:

By inversion on $v_1 \sqsubseteq_p v_2$, we know that $x_1 = x_2$.

Hence $v_1 = v_2$.

(2) Every nonempty finite subset of D_p has a lub.

To show this, it is sufficient to show that every two elements of D_p have a lub, since a binary lub operation can be repeatedly applied to compute the lub of any finite set.

We will show that every two elements of D_p have a lub by showing that the \sqcup_p operation defined by Definition 3.2 computes their lub.

It suffices to show the following two properties:

- (a) For all $v_1, v_2, v \in D_p$, if $v_1 \sqsubseteq_p v$ and $v_2 \sqsubseteq_p v$, then $(v_1 \sqcup_p v_2) \sqsubseteq_p v$.
- (b) For all $v_1, v_2 \in D_p$, $v_1 \sqsubseteq_p (v_1 \sqcup_p v_2)$ and $v_2 \sqsubseteq_p (v_1 \sqcup_p v_2)$.
- (a) For all $v_1, v_2, v \in D_p$, if $v_1 \sqsubseteq_p v$ and $v_2 \sqsubseteq_p v$, then $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

Assume $v_1, v_2, v \in D_p$, and $v_1 \sqsubseteq_p v$ and $v_2 \sqsubseteq_p v$.

Now we do a case analysis on v_1 and v_2 .

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (d_2, \text{false})$.

Now case on v :

- Case $v = (d, \text{false})$:

By the definition of \sqcup_p , $(d_1, \text{false}) \sqcup_p (d_2, \text{false}) = (d_1 \sqcup d_2, \text{false})$.

By inversion on $(d_1, \text{false}) \sqsubseteq_p (d, \text{false})$, $d_1 \sqsubseteq l$.

By inversion on $(d_2, \text{false}) \sqsubseteq_p (d, \text{false})$, $d_2 \sqsubseteq l$.

Hence l is an upper bound for d_1 and d_2 .

Hence $d_1 \sqcup d_2 \sqsubseteq l$.

Hence $(d_1 \sqcup d_2, \text{false}) \sqsubseteq_p (d, \text{false})$.

Hence $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

– Case $v = (x, \text{true})$:

By the definition of \sqcup_p , $(d_1, \text{false}) \sqcup_p (d_2, \text{false}) = (d_1 \sqcup d_2, \text{false})$.

By inversion on $(d_1, \text{false}) \sqsubseteq_p (x, \text{true})$, $d_1 \sqsubseteq x$.

By inversion on $(d_2, \text{false}) \sqsubseteq_p (x, \text{true})$, $d_2 \sqsubseteq x$.

Hence x is an upper bound for d_1 and d_2 .

Hence $d_1 \sqcup d_2 \sqsubseteq x$.

Hence $(d_1 \sqcup d_2, \text{false}) \sqsubseteq_p (x, \text{true})$.

Hence $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

• Case $v_1 = (x_1, \text{true})$ and $v_2 = (x_2, \text{true})$:

Now case on v :

– Case $v = (d, \text{false})$:

By inversion on $(x_1, \text{true}) \sqsubseteq_p (d, \text{false})$, we know $l = \top$.

By inversion on $(x_2, \text{true}) \sqsubseteq_p (d, \text{false})$, we know $l = \top$.

Now consider whether $x_1 = x_2$ or not.

If it does, then by the definition of \sqcup_p , $(x_1, \text{true}) \sqcup_p (x_2, \text{true}) = (x_1, \text{true})$.

By definition of \sqsubseteq_p , we have $(x_1, \text{true}) \sqsubseteq_p (\top, \text{false})$.

So $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

If it does not, then $v_1 \sqcup_p v_2 = (\top, \text{false})$.

By the definition of \sqsubseteq_p , we have $(\top, \text{false}) \sqsubseteq_p (\top, \text{false})$.

So $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

– Case $v = (x, \text{true})$:

By inversion on $(x_1, \text{true}) \sqsubseteq_p (x, \text{true})$, we know $x = x_1$.

By inversion on $(x_2, \text{true}) \sqsubseteq_p (x, \text{true})$, we know $x = x_2$.

Hence $x_1 = x_2$.

By the definition of \sqcup_p , $(x_1, \text{true}) \sqcup_p (x_2, \text{true}) = (x_1, \text{true})$.

Hence $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (d_2, \text{false})$:

Now case on v :

- Case $v = (d, \text{false})$:

Now consider whether $d_2 \sqsubseteq x_1$.

If it is, then $(x_1, \text{true}) \sqcup_p (d_2, \text{false}) = (x_1, \text{true}) = v_1$.

Hence $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

Otherwise, $(x_1, \text{true}) \sqcup_p (d_2, \text{false}) = (\top, \text{false})$.

By inversion on $(x_1, \text{true}) \sqsubseteq_p (d, \text{false})$, we know $l = \top$.

By reflexivity, $(\top, \text{false}) \sqsubseteq_p (\top, \text{false})$.

Hence $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

- Case $v = (x, \text{true})$:

By inversion on $(x_1, \text{true}) \sqsubseteq_p (x, \text{true})$, we know that $x_1 = x$.

By inversion on $(d_2, \text{false}) \sqsubseteq_p (x, \text{true})$, we know that $d_2 \sqsubseteq x$.

By transitivity, $d_2 \sqsubseteq x_1$.

By the definition of \sqcup_p , it follows that $(x_1, \text{true}) \sqcup_p (d_2, \text{false}) = (x_1, \text{true})$.

By definition of \sqsubseteq_p , $(x_1, \text{true}) \sqsubseteq_p (x_1, \text{true})$.

Hence $v_1 \sqcup_p v_2 \sqsubseteq_p v$.

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (x_2, \text{true})$:

Symmetric with the previous case.

- (b) For all $v_1, v_2 \in D_p$, $v_1 \sqsubseteq_p v_1 \sqcup_p v_2$ and $v_2 \sqsubseteq_p v_1 \sqcup_p v_2$.

Assume $v_1, v_2 \in D_p$, and proceed by case analysis.

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (d_2, \text{false})$:

Since \sqcup is a join operator, we know $d_1 \sqsubseteq d_1 \sqcup d_2$.

By the definition of \sqsubseteq_p , $(d_1, \text{false}) \sqsubseteq (d_1 \sqcup d_2, \text{false})$.

By the definition of \sqcup_p , $v_1 \sqcup_p v_2 = (d_1 \sqcup d_2, \text{false})$.

Hence $v_1 \sqsubseteq_p v_1 \sqcup_p v_2$.

Since \sqcup is a join operator, we know $d_1 \sqsubseteq d_1 \sqcup d_2$.

By the definition of \sqsubseteq_p , $(d_2, \text{false}) \sqsubseteq (d_1 \sqcup d_2, \text{false})$.

By the definition of \sqcup_p , $v_1 \sqcup_p v_2 = (d_1 \sqcup d_2, \text{false})$.

Hence $v_2 \sqsubseteq_p v_1 \sqcup_p v_2$.

Therefore $v_1 \sqsubseteq_p v_1 \sqcup v_2$ and $v_2 \sqsubseteq_p v_1 \sqcup v_2$.

- Case $v_1 = (d_1, \text{false})$ and $v_2 = (x_2, \text{true})$:

Consider whether $d_1 \sqsubseteq x_2$.

- Case $d_1 \sqsubseteq x_2$:

By the definition of \sqcup_p , we know $(d_1, \text{false}) \sqcup_p (x_2, \text{true}) = (x_2, \text{true})$.

By the definition of \sqcup_p , we know $(d_1, \text{false}) \sqsubseteq_p (x_2, \text{true})$.

Hence $v_1 \sqsubseteq_p v_1 \sqcup_p v_2$.

By reflexivity, $(x_2, \text{true}) \sqsubseteq_p (x_2, \text{true})$.

Hence $v_2 \sqsubseteq_p v_1 \sqcup_p v_2$.

Therefore $v_1 \sqsubseteq_p v_1 \sqcup v_2$ and $v_2 \sqsubseteq_p v_1 \sqcup v_2$.

- Case $d_1 \not\sqsubseteq x_2$:

By the definition of \sqcup_p , we know $(d_1, \text{false}) \sqcup_p (x_2, \text{true}) = (\top, \text{false})$.

Since $d_1 \sqsubseteq \top$, by the definition of \sqsubseteq_p we know $(d_1, \text{false}) \sqsubseteq (\top, \text{false})$.

Hence $v_1 \sqsubseteq_p v_1 \sqcup_p v_2$.

By the definition of \sqsubseteq_p , we know $(x_2, \text{true}) \sqsubseteq (\top, \text{false})$.

Hence $v_2 \sqsubseteq_p v_1 \sqcup_p v_2$.

Therefore $v_1 \sqsubseteq_p v_1 \sqcup v_2$ and $v_2 \sqsubseteq_p v_1 \sqcup v_2$.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (d_2, \text{false})$:

Symmetric with the previous case.

- Case $v_1 = (x_1, \text{true})$ and $v_2 = (x_2, \text{true})$:

Consider whether x_1 equals x_2 .

– Case $x_1 = x_2$:

By the definition $\sqcup_p, (x_1, \text{true}) \sqcup_p (x_2, \text{true}) = (x_1, \text{true})$.

By reflexivity, $(x_1, \text{true}) \sqsubseteq_p (x_1, \text{true})$.

Hence $v_1 \sqsubseteq_p v_1 \sqcup_p v_2$.

By reflexivity, $(x_2, \text{true}) \sqsubseteq_p (x_1, \text{true})$.

Hence $v_2 \sqsubseteq_p v_1 \sqcup_p v_2$.

Therefore $v_1 \sqsubseteq_p v_1 \sqcup v_2$ and $v_2 \sqsubseteq_p v_1 \sqcup v_2$.

– Case $x_1 \neq x_2$:

By the definition $\sqcup_p, (x_1, \text{true}) \sqcup_p (x_2, \text{true}) = (\top, \text{false})$.

By the definition of $\sqsubseteq_p, (x_1, \text{true}) \sqsubseteq_p (\top, \text{false})$.

Hence $v_1 \sqsubseteq_p v_1 \sqcup_p v_2$.

By the definition of $\sqsubseteq_p, (x_2, \text{true}) \sqsubseteq_p (\top, \text{false})$.

Hence $v_2 \sqsubseteq_p v_1 \sqcup_p v_2$.

Therefore $v_1 \sqsubseteq_p v_1 \sqcup v_2$ and $v_2 \sqsubseteq_p v_1 \sqcup v_2$.

(3) \perp_p is the least element of D_p .

\perp_p is defined to be (\perp, false) .

In order to be the least element of D_p , it must be less than or equal to every element of D_p .

By Lemma 3.1, the elements of D_p partition into (d, false) for all $d \in D$, and (x, true) for all $x \in X$, where $X = D - \{\top\}$.

We consider both cases:

• (d, false) for all $d \in D$:

By the definition of $\sqsubseteq_p, (\perp, \text{false}) \sqsubseteq_p (d, \text{false})$ iff $\perp \sqsubseteq d$.

Since \perp is the least element of D , $\perp \sqsubseteq d$.

Therefore $\perp_p = (\perp, \text{false}) \sqsubseteq_p (d, \text{false})$.

• (x, true) for all $x \in X$:

By the definition of \sqsubseteq_p , $(\perp, \text{false}) \sqsubseteq_p (x, \text{true})$ iff $\perp \sqsubseteq x$.

Since \perp is the least element of D , $\perp \sqsubseteq x$.

Therefore $\perp_p = (\perp, \text{false}) \sqsubseteq_p (x, \text{true})$.

Therefore \perp_p is less than or equal to all elements of D_p .

(4) \top_p is the greatest element of D_p .

\top_p is defined to be (\top, false) .

In order to be the greatest element of D_p , every element of D_p must be less than or equal to it.

By Lemma 3.1, the elements of D_p partition into (d, false) for all $d \in D$, and (x, true) for all $x \in X$, where $X = D - \{\top\}$.

We consider both cases:

- (d, false) for all $d \in D$:

By the definition of \sqsubseteq_p , $(d, \text{false}) \sqsubseteq_p (\top, \text{false})$ iff $d \sqsubseteq \top$.

Since \top is the greatest element of D , $d \sqsubseteq \top$.

Therefore $(d, \text{false}) \sqsubseteq_p (\top, \text{false}) = \top_p$.

- (x, true) for all $x \in X$:

By the definition of \sqsubseteq_p , $(x, \text{true}) \sqsubseteq_p (\top, \text{false})$ iff $\top \sqsubseteq x$.

Therefore $(x, \text{true}) \sqsubseteq_p (\top, \text{false}) = \top_p$.

Therefore all elements of D_p are less than or equal to \top_p .

□

A.11. Proof of Lemma 3.3

Proof. Consider an arbitrary permutation π .

For part 1, we have to show that if $\sigma \hookrightarrow \sigma'$ then $\pi(\sigma) \hookrightarrow \pi(\sigma')$, and that if $\pi(\sigma) \hookrightarrow \pi(\sigma')$ then $\sigma \hookrightarrow \sigma'$.

For the forward direction of part 1, suppose $\sigma \hookrightarrow \sigma'$.

We have to show that $\pi(\sigma) \longrightarrow \pi(\sigma')$.

We proceed by cases on the rule by which σ steps to σ' .

- **Case E-Beta:** $\sigma = \langle S; (\lambda x. e) v \rangle$, and $\sigma' = \langle S; e[x := v] \rangle$.

To show: $\pi(\langle S; (\lambda x. e) v \rangle) \longrightarrow \pi(\langle S; e[x := v] \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\sigma) = \langle \pi(S); (\lambda x. \pi(e)) \pi(v) \rangle$.

By E-Beta, $\langle \pi(S); (\lambda x. \pi(e)) \pi(v) \rangle$ steps to $\langle \pi(S); \pi(e)[x := \pi(v)] \rangle$.

By Definition 3.9, $\langle \pi(S); \pi(e)[x := \pi(v)] \rangle$ is equal to $\langle \pi(S); \pi(e[x := v]) \rangle$.

Hence $\langle \pi(S); (\lambda x. \pi(e)) \pi(v) \rangle$ steps to $\langle \pi(S); \pi(e[x := v]) \rangle$, which is equal to $\pi(\langle S; e[x := v] \rangle)$ by Definition 3.11.

Hence the case is satisfied.

- **Case E-New:** $\sigma = \langle S; \text{new} \rangle$, and $\sigma' = \langle S[l \mapsto (\perp, \text{false})]; l \rangle$.

To show: $\pi(\langle S; \text{new} \rangle) \longrightarrow \pi(\langle S[l \mapsto (\perp, \text{false})]; l \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\sigma) = \langle \pi(S); \text{new} \rangle$.

By E-New, $\langle \pi(S); \text{new} \rangle$ steps to $\langle (\pi(S))[l' \mapsto (\perp, \text{false})]; l' \rangle$, where $l' \notin \text{dom}(\pi(S))$.

It remains to show that $\langle (\pi(S))[l' \mapsto (\perp, \text{false})]; l' \rangle$ is equal to $\pi(\langle S[l \mapsto (\perp, \text{false})]; l \rangle)$.

By Definition 3.11, $\pi(\langle S[l \mapsto (\perp, \text{false})]; l \rangle)$ is equal to $\langle \pi(S[l \mapsto (\perp, \text{false})]); \pi(l) \rangle$, which is equal to $\langle (\pi(S))[\pi(l) \mapsto (\perp, \text{false})]; \pi(l) \rangle$.

We have to show that $\langle (\pi(S))[l' \mapsto (\perp, \text{false})]; l' \rangle$ is equal to $\langle (\pi(S))[\pi(l) \mapsto (\perp, \text{false})]; \pi(l) \rangle$.

We know (from the side condition of E-New) that $l \notin \text{dom}(S)$, and so $\pi(l) \notin \pi(\text{dom}(S))$.

Therefore, in $\langle (\pi(S))[l' \mapsto (\perp, \text{false})]; l' \rangle$, we can α -rename l' to $\pi(l)$, and so the two configurations are equal and the case is satisfied.

- **Case E-Put:** $\sigma = \langle S; \text{put}_i l \rangle$, and $\sigma' = \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$.

To show: $\pi(\langle S; \text{put}_i l \rangle) \longrightarrow \pi(\langle S[l \mapsto u_{p_i}(p_1)]; () \rangle)$.

By Definition 3.11, $\pi(\sigma) = \langle \pi(S); \text{put}_i \pi(l) \rangle$.

By E-Put, $\langle \pi(S); \text{put}_i \pi(l) \rangle$ steps to $\langle (\pi(S))[\pi(l) \mapsto u_{p_i}(p_1)]; () \rangle$,

since $S(l) = (\pi(S))(\pi(l)) = p_1$.

It remains to show that $\langle (\pi(S))[\pi(l) \mapsto u_{p_i}(p_1)]; () \rangle$ is equal to $\pi(\langle S[l \mapsto u_{p_i}(p_1)]; () \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\langle S[l \mapsto u_{p_i}(p_1)]; () \rangle)$ is equal to $\langle (\pi(S))[\pi(l) \mapsto u_{p_i}(p_1)]; () \rangle$,

and so the two configurations are equal and the case is satisfied.

- Case E-Put-Err: $\sigma = \langle S; \text{put}_i l \rangle$, and $\sigma' = \mathbf{error}$.

To show: $\pi(\langle S; \text{put}_i l \rangle) \hookrightarrow \pi(\mathbf{error})$.

By Definition 3.11, $\pi(\sigma) = \langle \pi(S); \text{put}_i \pi(l) \rangle$.

By E-Put-Err, $\langle \pi(S); \text{put}_i \pi(l) \rangle$ steps to \mathbf{error} , since $S(l) = (\pi(S))(\pi(l)) = p_1$.

Since $\pi(\mathbf{error}) = \mathbf{error}$ by Definition 3.11, the case is complete.

- Case E-Get: $\sigma = \langle S; \text{get } l P \rangle$, and $\sigma' = \langle S; p_2 \rangle$.

To show: $\pi(\langle S; \text{get } l P \rangle) \hookrightarrow \pi(\langle S; p_2 \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\sigma) = \langle \pi(S); \text{get } \pi(l) P \rangle$.

By E-Get, $\langle \pi(S); \text{get } \pi(l) P \rangle$ steps to $\langle \pi(S); p_2 \rangle$, since $S(l) = (\pi(S))(\pi(l)) = p_1$.

By Definitions 3.11 and 3.9, $\pi(\langle S; p_2 \rangle) = \langle \pi(S); p_2 \rangle$.

Therefore the case is complete.

- Case E-Freeze-Init: $\sigma = \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle$,

and $\sigma' = \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle$.

To show: $\pi(\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle) \hookrightarrow$

$\pi(\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\sigma) = \langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e) \rangle$.

By E-Freeze-Init, $\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e) \rangle \hookrightarrow$

$\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e), \{\}, \{\} \rangle$.

By Definitions 3.11 and 3.9, $\pi(\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle) =$

$\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e), \{\}, \{\} \rangle$.

Therefore the case is complete.

- Case E-Spawn-Handler: $\sigma = \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle$,

and $\sigma' = \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle$.

To show: $\pi(\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle) \hookrightarrow$

$\pi(\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\sigma) =$

$\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e_0), \{\pi(e), \dots\}, H \rangle$.

Since $(\pi(S))(\pi(l)) = (d_1, \text{frz}_1)$, by E-Spawn-Handler we have that

$\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e_0), \{\pi(e), \dots\}, H \rangle \hookrightarrow$

$\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e_0), \{\pi(e_0)[x := d_2], \pi(e), \dots\}, \{d_2\} \cup H \rangle$.

By Definitions 3.11 and 3.9, $\pi(\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle) = \langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e_0), \{\pi(e_0)[x := d_2], \pi(e), \dots\}, \{d_2\} \cup H \rangle$.

Therefore the case is complete.

- Case E-Freeze-Final: $\sigma = \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle$,

and $\sigma' = \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

To show: $\pi(\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle) \hookrightarrow$

$\pi(\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\sigma) =$

$\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e_0), \{\pi(v), \dots\}, H \rangle$.

Since $(\pi(S))(\pi(l)) = (d_1, \text{frz}_1)$, by E-Freeze-Final we have that

$\langle \pi(S); \text{freeze } \pi(l) \text{ after } Q \text{ with } \lambda x. \pi(e_0), \{\pi(v), \dots\}, H \rangle \hookrightarrow$

$\langle \pi(S)[\pi(l) \mapsto (d_1, \text{true})]; d_1 \rangle$.

(From Definition 3.9, we can see that if v is a value, $\pi(v)$ is also a value.)

It remains to show that $\langle \pi(S)[\pi(l) \mapsto (d_1, \text{true})]; d_1 \rangle$ is equal to $\pi(\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle)$ is equal to $\langle \pi(S[l \mapsto (d_1, \text{true})]); d_1 \rangle$,

which is equal to $\langle \pi(S)[\pi(l) \mapsto (d_1, \text{true})]; d_1 \rangle$, and so the two configurations are equal and the case is satisfied.

- Case E-Freeze-Simple: $\sigma = \langle S; \text{freeze } l \rangle$, and $\sigma' = \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

To show: $\pi(\langle S; \text{freeze } l \rangle) \hookrightarrow \pi(\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\sigma) = \langle \pi(S); \text{freeze } \pi(l) \rangle$.

Since $(\pi(S))(\pi(l)) = (d_1, \text{frz}_1)$, by E-Freeze-Simple we have that $\langle \pi(S); \text{freeze } \pi(l) \rangle \hookrightarrow \langle \pi(S)[\pi(l) \mapsto (d_1, \text{true})]; d_1 \rangle$.

It remains to show that $\langle \pi(S)[\pi(l) \mapsto (d_1, \text{true})]; d_1 \rangle$ is equal to $\pi(\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle)$.

By Definitions 3.11 and 3.9, $\pi(\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle)$ is equal to $\langle \pi(S[l \mapsto (d_1, \text{true})]); d_1 \rangle$, which is equal to $\langle \pi(S)[\pi(l) \mapsto (d_1, \text{true})]; d_1 \rangle$, and so the two configurations are equal and the case is satisfied.

For the reverse direction of part 1, suppose $\pi(\sigma) \hookrightarrow \pi(\sigma')$.

We have to show that $\sigma \hookrightarrow \sigma'$.

We know from the forward direction of the proof that for all configurations σ and σ' and permutations π , if $\sigma \hookrightarrow \sigma'$ then $\pi(\sigma) \hookrightarrow \pi(\sigma')$.

Hence since $\pi(\sigma) \hookrightarrow \pi(\sigma')$, and since π^{-1} is also a permutation, we have that $\pi^{-1}(\pi(\sigma)) \hookrightarrow \pi^{-1}(\pi(\sigma'))$.

Since $\pi^{-1}(\pi(l)) = l$ for every $l \in \text{Loc}$, and that property lifts to configurations as well, we have that $\sigma \hookrightarrow \sigma'$.

For the forward direction of part 2, suppose $\sigma \mapsto \sigma'$.

We have to show that $\pi(\sigma) \mapsto \pi(\sigma')$.

By inspection of the operational semantics, σ must be of the form $\langle S; E[e] \rangle$, and σ' must be of the form $\langle S'; E[e'] \rangle$.

Hence we have to show that $\pi(\langle S; E[e] \rangle) \mapsto \pi(\langle S'; E[e'] \rangle)$.

By Definition 3.11, $\pi(\langle S; E[e] \rangle)$ is equal to $\langle \pi(S); \pi(E[e]) \rangle$.

Also by Definition 3.11, $\pi(\langle S'; E[e'] \rangle)$ is equal to $\langle \pi(S'); \pi(E[e']) \rangle$.

Furthermore, $\langle \pi(S); \pi(E[e]) \rangle$ is equal to $\langle \pi(S); (\pi(E))[\pi(e)] \rangle$ and $\langle \pi(S'); \pi(E[e']) \rangle$ is equal to $\langle \pi(S'); (\pi(E))[\pi(e')] \rangle$.

So we have to show that $\langle \pi(S); (\pi(E))[\pi(e)] \rangle \mapsto \langle \pi(S'); (\pi(E))[\pi(e')] \rangle$.

From the premise of E-Eval-Ctxt, $\langle S; e \rangle \hookrightarrow \langle S'; e' \rangle$.

Hence, by part 1, $\pi(\langle S; e \rangle) \hookrightarrow \pi(\langle S'; e' \rangle)$.

By Definition 3.11, $\pi(\langle S; e \rangle)$ is equal to $\langle \pi(S); \pi(e) \rangle$ and $\pi(\langle S'; e' \rangle)$ is equal to $\langle \pi(S'); \pi(e') \rangle$.

Hence $\langle \pi(S); \pi(e) \rangle \hookrightarrow \langle \pi(S'); \pi(e') \rangle$.

Therefore, by E-Eval-Ctxt, $\langle \pi(S); E[\pi(e)] \rangle \mapsto \langle \pi(S'); E[\pi(e')] \rangle$ for all evaluation contexts E .

In particular, it is true that $\langle \pi(S); (\pi(E))[\pi(e)] \rangle \mapsto \langle \pi(S'); (\pi(E))[\pi(e')] \rangle$, as we were required to show.

For the reverse direction of part 2, suppose $\pi(\sigma) \mapsto \pi(\sigma')$.

We have to show that $\sigma \mapsto \sigma'$.

We know from the forward direction of the proof that for all configurations σ and σ' and permutations π , if $\sigma \mapsto \sigma'$ then $\pi(\sigma) \mapsto \pi(\sigma')$.

Hence since $\pi(\sigma) \mapsto \pi(\sigma')$, and since π^{-1} is also a permutation, we have that $\pi^{-1}(\pi(\sigma)) \mapsto \pi^{-1}(\pi(\sigma'))$.

Since $\pi^{-1}(\pi(l)) = l$ for every $l \in Loc$, and that property lifts to configurations as well, we have that $\sigma \mapsto \sigma'$.

□

A.12. Proof of Lemma 3.4

Proof. Suppose $\sigma \hookrightarrow \sigma'$ and $\sigma \hookrightarrow \sigma''$.

We have to show that there is a permutation π such that $\sigma' = \pi(\sigma'')$, modulo choice of events.

The proof is by cases on the rule by which σ steps to σ' .

- Case E-Beta:

Given: $\langle S; (\lambda x. e) v \rangle \longrightarrow \langle S; e[x := v] \rangle$, and $\langle S; (\lambda x. e) v \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S; e[x := v] \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which $\langle S; (\lambda x. e) v \rangle$ can step is E-Beta.

Hence $\sigma'' = \langle S; e[x := v] \rangle$, and the case is satisfied by choosing π to be the identity function.

- Case E-New:

Given: $\langle S; \text{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \text{false})]; l \rangle$, and $\langle S; \text{new} \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S[l \mapsto (\perp, \text{false})]; l \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which $\langle S; \text{new} \rangle$ can step is E-New.

Hence $\sigma'' = \langle S[l' \mapsto (\perp, \text{false})]; l' \rangle$.

Since, by the side condition of E-New, neither l nor l' occur in $\text{dom}(S)$, the case is satisfied by choosing π to be the permutation that maps l' to l and is the identity on every other element of Loc .

- Case E-Put:

Given: $\langle S; \text{put}_i l \rangle \longrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$, and $\langle S; \text{put}_i l \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S[l \mapsto u_{p_i}(p_1)]; () \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, and since $u_{p_i}(p_1) \neq \top_p$ (from the premise of E-Put), the only reduction rule by which $\langle S; \text{put}_i l \rangle$ can step is E-Put.

Hence $\sigma'' = \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$, and the case is satisfied by choosing π to be the identity function.

- Case E-Put-Err:

Given: $\langle S; \text{put}_i l \rangle \longrightarrow \mathbf{error}$, and $\langle S; \text{put}_i l \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\mathbf{error} = \pi(\sigma'')$.

By inspection of the operational semantics, and since $u_{p_i}(p_1) = \top_p$ (from the premise of E-Put-Err), the only reduction rule by which $\langle S; \text{put}_i l \rangle$ can step is E-Put-Err.

Hence $\sigma'' = \mathbf{error}$, and the case is satisfied by choosing π to be the identity function.

- Case E-Get:

Given: $\langle S; \text{get } l P \rangle \longrightarrow \langle S; p_2 \rangle$, and $\langle S; \text{get } l P \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S; p_2 \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which $\langle S; \text{get } l P \rangle$ can step is E-Get.

Hence $\sigma'' = \langle S; p_2 \rangle$, and the case is satisfied by choosing π to be the identity function.

- Case E-Freeze-Init:

Given: $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle \longrightarrow \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle$, and $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle$ can step is E-Freeze-Init.

Hence $\sigma'' = \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle$, and the case is satisfied by choosing π to be the identity function.

- Case E-Spawn-Handler:

Given: $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow$

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle$, and

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow \sigma''$.

To show: There exists a π such that

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle$ can step is E-Spawn-Handler.

(It cannot step by E-Freeze-Final, because we have from the premises of E-Spawn-Handler that $d_2 \sqsubseteq d_1$ and $d_2 \in Q$ and $d_2 \notin H$, and for the premises of E-Freeze-Final to hold, we would need that for all d_2 , if $d_2 \sqsubseteq d_1$ and $d_2 \in Q$, then $d_2 \in H$.)

Hence $\sigma'' = \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d'_2], e, \dots\}, \{d'_2\} \cup H \rangle$, where $d'_2 \sqsubseteq d_1$ and $d'_2 \in Q$ and $d'_2 \notin H$, and the case is satisfied by choosing π to be the identity function.

(It may be the case that $d'_2 \neq d_2$; if so, then we have internal nondeterminism modulo *choice of events*, as we were required to show. If $d'_2 = d_2$ then we have internal nondeterminism even without that additional qualification, which also satisfies the case.)

- Case E-Freeze-Final:

Given: $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$, and $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle$ can step is E-Freeze-Final.

(It cannot step by E-Spawn-Handler, because we have from the premises of E-Freeze-Final that, for all d_2 , if $d_2 \sqsubseteq d_1$ and $d_2 \in Q$, then $d_2 \in H$, and for the premises of E-Spawn-Handler to hold, we would need that $d_2 \sqsubseteq d_1$ and $d_2 \in Q$ and $d_2 \notin H$.)

Hence $\sigma'' = \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$, and the case is satisfied by choosing π to be the identity function.

- Case E-Freeze-Simple:

Given: $\langle S; \text{freeze } l \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$, and $\langle S; \text{freeze } l \rangle \longrightarrow \sigma''$.

To show: There exists a π such that $\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle = \pi(\sigma'')$.

By inspection of the operational semantics, the only reduction rule by which $\langle S; \text{freeze } l \rangle$ can step is E-Freeze-Simple.

Hence $\sigma'' = \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$, and the case is satisfied by choosing π to be the identity function.

□

A.13. Proof of Lemma 3.5

Proof. Suppose $\langle S; E_1[e_1] \rangle \mapsto \langle S_1; E_1[e'_1] \rangle$ and $\langle S; E_2[e_2] \rangle \mapsto \langle S_2; E_2[e'_2] \rangle$ and $E_1[e_1] = E_2[e_2]$.

We are required to show that if $E_1 \neq E_2$, then there exist evaluation contexts E'_1 and E'_2 such that:

- $E'_1[e_1] = E_2[e'_2]$, and
- $E'_2[e_2] = E_1[e'_1]$, and
- $E'_1[e'_1] = E'_2[e'_2]$.

Let $e = E_1[e_1] = E_2[e_2]$.

The proof is by induction on the structure of the expression e .

Proceeding by cases on e :

- Cases $e = x, e = v, e = e_a e_b, e = \text{get } e_a e_b$, and $e = \text{new}$ are identical to their corresponding cases in the proof of Lemma 2.3.
- Case $e = \text{put}_i e_a$:

We know that $\text{put}_i e_a = E_1[e_1]$.

From the grammar of evaluation contexts, then, we know that either:

- $\text{put}_i e_a = E_1[e_1] = E_1[\text{put}_i e_a]$, where $E_1 = []$, or
- $\text{put}_i e_a = E_1[e_1] = \text{put}_i E_{11}[e_1]$, where $E_{11}[e_1] = e_a$.

Similarly, we know that $\text{put}_i e_a = E_2[e_2]$.

From the grammar of evaluation contexts, we know that either:

- $\text{put}_i e_a = E_2[e_2] = E_2[\text{put}_i e_a]$, where $E_2 = []$, or

- $\text{put}_i e_a = E_2[e_2] = \text{put}_i E_{21}[e_2]$, where $E_{21}[e_2] = e_a$.

However, if $E_1 = []$ or $E_2 = []$, then $\text{put}_i e_a$ must be $\text{put}_i v$ for some v , and v cannot step individually, so the other of E_1 or E_2 must be $[]$ as well, and so $E_1 = E_2$.

Therefore the only case that we have to consider (where $E_1 \neq E_2$) is the case in which $E_1[e_1] = \text{put}_i E_{11}[e_1]$, where $E_{11}[e_1] = e_a$, and $\text{put}_i e_a = E_2[e_2] = \text{put}_i E_{21}[e_2]$, where $E_{21}[e_2] = e_a$.

So, we have $E_{11}[e_1] = e_a$ and $E_{21}[e_2] = e_a$.

In this case, we know that $E_{11} \neq E_{21}$, because if $E_{11} = E_{21}$, we would have $e_1 = e_2$, which would mean that $E_1 = E_2$, a contradiction.

So, since $E_{11} \neq E_{21}$, by IH we have that there exist evaluation contexts E'_{11} and E'_{21} such that:

- $E'_{11}[e_1] = E_{21}[e'_2]$, and
- $E'_{21}[e_2] = E_{11}[e'_1]$, and
- $E'_{11}[e'_1] = E'_{21}[e'_2]$.

Hence we can choose $E'_1 = \text{put}_i E'_{11}$ and $E'_2 = \text{put}_i E'_{21}$, which satisfy the criteria for E'_1 and E'_2 .

- Case $e = \text{freeze } e_a$: Similar to the case for $\text{put}_i e_a$.
- Case $e = \text{freeze } e_a \text{ after } e_b \text{ with } e_c$: Similar to the case where $e = e_a e_b$.
- Case $e = \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H$:

We know that $\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H = E_1[e_1]$.

From the grammar of evaluation contexts, we know that either:

- $\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H = E_1[e_1] =$
 $E_1[\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H]$, where $E_1 = []$, or
- $\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H = E_1[e_1] =$
 $\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_{b_1}, \dots, E_{11}[e_1], \dots, e_{b_n}\}, H$, where $E_{11}[e_1] = e_{b_i}$.

Similarly, we know that $\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H = E_2[e_2]$.

From the grammar of evaluation contexts, we know that either:

- $\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H = E_2[e_2] =$
 $E_2[\text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_b, \dots\}, H]$, where $E_2 = []$, or

- freeze l after Q with $\lambda x. e_a, \{e_b, \dots\}, H = E_2[e_2] =$
 freeze l after Q with $\lambda x. e_a, \{e_{b_1}, \dots, E_{21}[e_2], \dots, e_{b_n}\}, H$, where $E_{21}[e_2] = e_{b_j}$.

However, if $E_1 = []$ or $E_2 = []$, then freeze l after Q with $\lambda x. e_a, \{e_b, \dots\}, H$ must be freeze l after Q with $\lambda x. e_a, \{v, \dots\}, H$ for some $\{v_1, \dots, v_n\}$, and no v_i can step individually, so the other of E_1 or E_2 must be $[]$ as well, and so $E_1 = E_2$.

Therefore the only case that we have to consider (where $E_1 \neq E_2$) is the case in which:

- freeze l after Q with $\lambda x. e_a, \{e_b, \dots\}, H = E_1[e_1] =$
 freeze l after Q with $\lambda x. e_a, \{e_{b_1}, \dots, E_{11}[e_1], \dots, e_{b_n}\}, H$, where $E_{11}[e_1] = e_{b_i}$, and
- freeze l after Q with $\lambda x. e_a, \{e_b, \dots\}, H = E_2[e_2] =$
 freeze l after Q with $\lambda x. e_a, \{e_{b_1}, \dots, E_{21}[e_2], \dots, e_{b_n}\}, H$, where $E_{21}[e_2] = e_{b_j}$.

Finally, we have two cases to consider:

- $e_{b_i} = e_{b_j}$: In this case we have $e_{b_i} = E_{11}[e_1] = E_{21}[e_2]$.

We know that $E_{11} \neq E_{21}$, because if $E_{11} = E_{21}$, we would have $e_1 = e_2$, which would mean that $E_1 = E_2$, a contradiction.

So, since $E_{11} \neq E_{21}$, by IH we have that there exist evaluation contexts E'_{11} and E'_{21} such that:

- * $E'_{11}[e_1] = E_{21}[e'_2]$, and
- * $E'_{21}[e_2] = E_{11}[e'_1]$, and
- * $E'_{11}[e'_1] = E'_{21}[e'_2]$.

Hence we can choose

$$E'_1 = \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_{b_1}, \dots, E'_{11}, \dots, e_{b_n}\}, H,$$

and

$$E'_2 = \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_{b_1}, \dots, E'_{21}, \dots, e_{b_n}\}, H,$$

which satisfy the criteria for E'_1 and E'_2 .

– $e_{b_i} \neq e_{b_j}$: In this case, we can choose

$$E'_1 = \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_{b_1}, \dots, E_{11}, \dots, E_{21}[e'_2], \dots, e_{b_n}\}, H,$$

and

$$E'_2 = \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_a, \{e_{b_1}, \dots, E_{11}[e'_1], \dots, E_{21}, \dots, e_{b_n}\}, H,$$

which satisfy the criteria for E'_1 and E'_2 .

□

A.14. Proof of Lemma 3.6

Proof. Suppose $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$.

We are required to show that $S \sqsubseteq_S S'$.

The proof is by cases on the rule by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$.

- Case E-Beta:

Immediate by the definition of \sqsubseteq_S , since S does not change.

- Case E-New:

Given: $\langle S; \text{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \text{false})]; l \rangle$.

To show: $S \sqsubseteq_S S[l \mapsto (\perp, \text{false})]$.

By Definition 3.6, we have to show that $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto (\perp, \text{false})])$ and that for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq_p (S[l \mapsto (\perp, \text{false})])(l')$.

By definition, a store update operation on S can only either update an existing binding in S or extend S with a new binding.

Hence $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto (\perp, \text{false})])$.

From the side condition of E-New, $l \notin \text{dom}(S)$.

Hence $S[l \mapsto (\perp, \text{false})]$ adds a new binding for l in S .

Hence $S[l \mapsto (\perp, \text{false})]$ does not update any existing bindings in S .

Hence, for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq_p (S[l \mapsto (\perp, \text{false})])(l')$.

Therefore $S \sqsubseteq_S S[l \mapsto (\perp, \text{false})]$, as required.

- Case E-Put:

Given: $\langle S; \text{put}_i l \rangle \longrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$.

To show: $S \sqsubseteq_S S[l \mapsto u_{p_i}(p_1)]$.

By Definition 3.6, we have to show that $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto u_{p_i}(p_1)])$ and that for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq_p (S[l \mapsto u_{p_i}(p_1)])(l')$.

By definition, a store update operation on S can only either update an existing binding in S or extend S with a new binding.

Hence $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto u_{p_i}(p_1)])$.

From the premises of E-Put, $S(l) = p_1$.

Therefore $l \in \text{dom}(S)$.

Hence $S[l \mapsto u_{p_i}(p_1)]$ updates the existing binding for l in S from p_1 to $u_{p_i}(p_1)$.

By definition, u_{p_i} is inflationary.

Hence $p_1 \sqsubseteq_p u_{p_i}(p_1)$.

$S[l \mapsto u_{p_i}(p_1)]$ does not update any other bindings in S , hence, for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq_p (S[l \mapsto u_{p_i}(p_1)])(l')$.

Hence $S \sqsubseteq_S S[l \mapsto u_{p_i}(p_1)]$, as required.

- Case E-Put-Err:

Given: $\langle S; \text{put}_i l \rangle \longrightarrow \text{error}$.

By the definition of **error**, $\text{error} = \langle \top_S; e \rangle$ for any e .

To show: $S \sqsubseteq_S \top_S$.

Immediate by the definition of \sqsubseteq_S .

- Case E-Get:

Immediate by the definition of \sqsubseteq_S , since S does not change.

- Case E-Freeze-Init:

Immediate by the definition of \sqsubseteq_S , since S does not change.

- Case E-Spawn-Handler:

Immediate by the definition of \sqsubseteq_S , since S does not change.

- Case E-Freeze-Final:

Given: $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

To show: $S \sqsubseteq_S S[l \mapsto (d_1, \text{true})]$.

By Definition 3.6, we have to show that $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto (d_1, \text{true})])$ and that for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq_p (S[l \mapsto (d_1, \text{true})])(l')$.

By definition, a store update operation on S can only either update an existing binding in S or extend S with a new binding.

Hence $\text{dom}(S) \subseteq \text{dom}(S[l \mapsto (d_1, \text{true})])$.

From the premises of E-Freeze-Final, $S(l) = (d_1, \text{frz}_1)$. Therefore $l \in \text{dom}(S)$.

Hence $S[l \mapsto (d_1, \text{true})]$ updates the existing binding for l in S from (d_1, frz_1) to (d_1, true) .

By the definition of \sqsubseteq_p , $(d_1, \text{frz}_1) \sqsubseteq_p (d_1, \text{true})$.

$S[l \mapsto (d_1, \text{true})]$ does not update any other bindings in S , hence, for all $l' \in \text{dom}(S)$, $S(l') \sqsubseteq_p (S[l \mapsto (d_1, \text{true})])(l')$.

Hence $S \sqsubseteq_S S[l \mapsto (d_1, \text{true})]$, as required.

- Case E-Freeze-Simple:

Given: $\langle S; \text{freeze } l \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

To show: $S \sqsubseteq_S S[l \mapsto (d_1, \text{true})]$.

Similar to the previous case.

□

A.15. Proof of Lemma 3.7

Proof. Suppose $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$, where $\langle S'; e' \rangle \neq \mathbf{error}$.

Consider arbitrary U_S such that U_S is non-conflicting with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ and $U_S(S') \neq \top_S$ and U_S is freeze-safe with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$.

We are required to show that $\langle U_S(S); e \rangle \longrightarrow \langle U_S(S'); e' \rangle$.

The proof is by cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$.

Since $\langle S'; e' \rangle \neq \mathbf{error}$, we do not need to consider the E-Put-Err rule.

The assumption that U_S is freeze-safe with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ is only needed in the E-Freeze-Final and E-Freeze-Simple cases.

- Case E-Beta:

Given: $\langle S; (\lambda x. e) v \rangle \longrightarrow \langle S; e[x := v] \rangle$.

To show: $\langle U_S(S); (\lambda x. e) v \rangle \longrightarrow \langle U_S(S); e[x := v] \rangle$.

Immediate by E-Beta.

- Case E-New:

Given: $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \mathbf{false})]; l \rangle$.

To show: $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \langle U_S(S[l \mapsto (\perp, \mathbf{false})]); l \rangle$.

By E-New, we have that $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \langle (U_S(S))[l' \mapsto (\perp, \mathbf{false})]; l' \rangle$,

where $l' \notin \text{dom}(U_S(S))$.

By assumption, U_S is non-conflicting with $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \mathbf{false})]; l \rangle$.

Therefore $l \notin \text{dom}(U_S(S))$.

Therefore, in $\langle (U_S(S))[l' \mapsto (\perp, \mathbf{false})]; l' \rangle$, we can α -rename l' to l .

Therefore $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \langle (U_S(S))[l \mapsto (\perp, \mathbf{false})]; l \rangle$.

Also, since U_S is non-conflicting with $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \mathbf{false})]; l \rangle$,

we have that $(U_S(S[l \mapsto (\perp, \mathbf{false})]))(l) = (S[l \mapsto (\perp, \mathbf{false})])(l) = (\perp, \mathbf{false})$.

Hence $(U_S(S))[l \mapsto (\perp, \text{false})] = U_S(S[l \mapsto (\perp, \text{false})])$.

Therefore $\langle U_S(S); \text{new} \rangle \longrightarrow \langle U_S(S[l \mapsto (\perp, \text{false})]); l \rangle$, as we were required to show.

- Case E-Put:

Given: $\langle S; \text{put}_i l \rangle \longrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$.

To show: $\langle U_S(S); \text{put}_i l \rangle \longrightarrow \langle U_S(S[l \mapsto u_{p_i}(p_1)]); () \rangle$.

From the premises of E-Put, $S(l) = p_1$.

Hence $(U_S(S))(l) = p'_1$, where $p_1 \sqsubseteq_p p'_1$.

Next, we want to show that $u_{p_i}(p'_1) \neq \top_p$.

Assume for the sake of a contradiction that $u_{p_i}(p'_1) = \top_p$.

Then $u_{p_i}((U_S(S))(l)) = \top_p$.

Let u_{p_j} be the state update operation in U_S that affects the contents of l .

Hence $(U_S(S))(l) = u_{p_j}(p_1)$. Then $u_{p_i}(u_{p_j}(p_1)) = \top_p$.

Since state update operations commute, $u_{p_j}(u_{p_i}(p_1)) = \top_p$.

But then $U_S(S[l \mapsto u_{p_i}(p_1)]) = \top_S$,

which contradicts the assumption that $U_S(S[l \mapsto u_{p_i}(p_1)]) \neq \top_S$.

Hence, $u_{p_i}(p'_1) \neq \top_p$.

Therefore, by E-Put, $\langle U_S(S); \text{put}_i l \rangle \longrightarrow \langle (U_S(S))[l \mapsto u_{p_i}(p'_1)]; () \rangle$.

Since $p'_1 = u_{p_j}(p_1)$, we have that $(U_S(S))[l \mapsto u_{p_i}(p'_1)] = (U_S(S))[l \mapsto u_{p_i}(u_{p_j}(p_1))]$,

which, since u_{p_i} and u_{p_j} commute, is equal to $(U_S(S))[l \mapsto u_{p_j}(u_{p_i}(p_1))]$.

Finally, since u_{p_j} is the update operation in U_S that affects the contents of l ,

we have that $(U_S(S))[l \mapsto u_{p_j}(u_{p_i}(p_1))] = U_S(S[l \mapsto u_{p_i}(p_1)])$, and so the case is satisfied.

- Case E-Get:

Given: $\langle S; \text{get } l P \rangle \longrightarrow \langle S; p_2 \rangle$.

To show: $\langle U_S(S); \text{get } l P \rangle \longrightarrow \langle U_S(S); p_2 \rangle$.

From the premises of E-Get, $S(l) = p_1$ and $\text{incomp}(P)$ and $p_2 \in P$ and $p_2 \sqsubseteq_p p_1$.

By assumption, $U_S(S) \neq \top_S$.

Hence $(U_S(S))(l) = p'_1$, where $p_1 \sqsubseteq_p p'_1$.

By the transitivity of \sqsubseteq_p , $p_2 \sqsubseteq_p p'_1$.

Hence, $(U_S(S))(l) = p'_1$ and $\text{incomp}(P)$ and $p_2 \in P$ and $p_2 \sqsubseteq_p p'_1$.

Therefore, by E-Get,

$$\langle U_S(S); \text{get } l \ P \rangle \longrightarrow \langle U_S(S); p_2 \rangle,$$

as we were required to show.

- Case E-Freeze-Init:

Given: $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle \longrightarrow$

$$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle.$$

To show: $\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle \longrightarrow$

$$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle.$$

Immediate by E-Freeze-Init.

- Case E-Spawn-Handler:

Given:

$$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow$$

$$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle.$$

To show:

$$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow$$

$$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle.$$

From the premises of E-Spawn-Handler, $S(l) = (d_1, \text{frz}_1)$ and $d_2 \sqsubseteq d_1$ and $d_2 \notin H$ and $d_2 \in Q$.

By assumption, $U_S(S) \neq \top_S$.

Hence $(U_S(S))(l) = (d'_1, \text{frz}'_1)$ where $(d_1, \text{frz}_1) \sqsubseteq_p (d'_1, \text{frz}'_1)$.

By Definition 3.1, $d_1 \sqsubseteq d'_1$.

By the transitivity of \sqsubseteq , $d_2 \sqsubseteq d'_1$.

Hence $(U_S(S))(l) = (d'_1, \text{frz}'_1)$ and $d_2 \sqsubseteq d'_1$ and $d_2 \notin H$ and $d_2 \in Q$.

Therefore, by E-Spawn-Handler,

$$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow$$

$$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle,$$

as we were required to show.

- Case E-Freeze-Final:

$$\text{Given: } \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow$$

$$\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle.$$

$$\text{To show: } \langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow$$

$$\langle U_S(S[l \mapsto (d_1, \text{true})]); d_1 \rangle.$$

From the premises of E-Freeze-Final, $S(l) = (d_1, \text{frz}_1)$.

We have two cases to consider:

– $\text{frz}_1 = \text{true}$:

In this case, $S(l) = (d_1, \text{true})$.

Let u_{p_i} be the state update operation in U_S that affects the contents of l .

Hence $(U_S(S))(l) = u_{p_i}((d_1, \text{true}))$.

We know from Definition 3.4 that $u_{p_i}((d_1, \text{true}))$ is either (d_1, true) or (\top, false) .

But if $u_{p_i}((d_1, \text{true})) = (\top, \text{false})$, then $U_S(S[l \mapsto (d_1, \text{true})]) = \top_S$, which contradicts our assumption that $U_S(S[l \mapsto (d_1, \text{true})]) \neq \top_S$.

Hence $u_{p_i}((d_1, \text{true})) = (d_1, \text{true})$.

Hence $(U_S(S))(l) = (d_1, \text{true})$, and we already have from the premises of E-Freeze-Final that

$$\forall d_2. (d_2 \sqsubseteq d_1 \wedge d_2 \in Q \Rightarrow d_2 \in H).$$

Hence, by E-Freeze-Final, we have that

$$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow \langle (U_S(S))[l \mapsto (d_1, \text{true})]; d_1 \rangle.$$

Finally, since u_{p_i} is the state update operation in U_S that affects the contents of l ,

and $u_{p_i}((d_1, \text{true})) = (d_1, \text{true})$, we have that $(U_S(S))[l \mapsto (d_1, \text{true})]$ is equal to $U_S(S[l \mapsto (d_1, \text{true})])$, and so the case is satisfied.

– $\text{frz}_1 = \text{false}$:

By assumption, U_S is freeze-safe with $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

Therefore U_S acts as the identity on the contents of any locations that change status during the transition.

Since $\text{frz}_1 = \text{false}$, the contents of l change status during the transition.

Therefore U_S acts as the identity on the contents of l .

Hence $(U_S(S))(l) = S(l) = (d_1, \text{frz}_1)$, and we already have from the premises of E-Freeze-Final that $\forall d_2. (d_2 \sqsubseteq d_1 \wedge d_2 \in Q \Rightarrow d_2 \in H)$.

Hence, by E-Freeze-Final, we have that

$$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow \langle (U_S(S))[l \mapsto (d_1, \text{true})]; d_1 \rangle.$$

Finally, since U_S acts as the identity on the contents of l , we have that $(U_S(S))[l \mapsto (d_1, \text{true})]$ is equal to $U_S(S[l \mapsto (d_1, \text{true})])$, and so the case is satisfied.

- Case E-Freeze-Simple:

Given: $\langle S; \text{freeze } l \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

To show: $\langle U_S(S); \text{freeze } l \rangle \longrightarrow \langle U_S(S[l \mapsto (d_1, \text{true})]); d_1 \rangle$.

From the premises of E-Freeze-Simple, $S(l) = (d_1, \text{frz}_1)$.

We have two cases to consider:

- $\text{frz}_1 = \text{true}$:

In this case, $S(l) = (d_1, \text{true})$.

Let u_{p_i} be the state update operation in U_S that affects the contents of l .

Hence $(U_S(S))(l) = u_{p_i}((d_1, \text{true}))$.

We know from Definition 3.4 that $u_{p_i}((d_1, \text{true}))$ is either (d_1, true) or (\top, false) .

But if $u_{p_i}((d_1, \text{true})) = (\top, \text{false})$, then $U_S(S[l \mapsto (d_1, \text{true})]) = \top_S$, which contradicts our assumption that $U_S(S[l \mapsto (d_1, \text{true})]) \neq \top_S$.

Hence $u_{p_i}((d_1, \text{true})) = (d_1, \text{true})$.

Hence $(U_S(S))(l) = (d_1, \text{true})$.

Hence, by E-Freeze-Simple, we have that $\langle U_S(S); \text{freeze } l \rangle \longrightarrow \langle (U_S(S))[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

Finally, since u_{p_i} is the state update operation in U_S that affects the contents of l , and $u_{p_i}((d_1, \text{true})) = (d_1, \text{true})$, we have that $(U_S(S))[l \mapsto (d_1, \text{true})]$ is equal to $U_S(S[l \mapsto (d_1, \text{true})])$, and so the case is satisfied.

– $\text{frz}_1 = \text{false}$:

By assumption, U_S is freeze-safe with $\langle S; \text{freeze } l \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

Therefore U_S acts as the identity on the contents of any locations that change status during the transition.

Since $\text{frz}_1 = \text{false}$, the contents of l change status during the transition.

Therefore U_S acts as the identity on the contents of l .

Hence $(U_S(S))(l) = S(l) = (d_1, \text{frz}_1)$.

Hence, by E-Freeze-Simple, we have that $\langle U_S(S); \text{freeze } l \rangle \longrightarrow \langle (U_S(S))[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

Finally, since U_S acts as the identity on the contents of l , we have that $(U_S(S))[l \mapsto (d_1, \text{true})]$ is equal to $U_S(S[l \mapsto (d_1, \text{true})])$, and so the case is satisfied.

□

A.16. Proof of Lemma 3.8

Proof. Suppose $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$, where $\langle S'; e' \rangle \neq \text{error}$.

Consider arbitrary U_S such that U_S is non-conflicting with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ and $U_S(S') = \top_S$ and U_S is freeze-safe with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$.

We are required to show that there exists $i \leq 1$ such that $\langle U_S(S); e \rangle \longrightarrow^i \text{error}$.

The proof is by cases on the rule of the reduction semantics by which $\langle S; e \rangle$ steps to $\langle S'; e' \rangle$.

Since $\langle S'; e' \rangle \neq \mathbf{error}$, we do not need to consider the E-Put-Err rule.

The assumption that U_S is freeze-safe with $\langle S; e \rangle \longrightarrow \langle S'; e' \rangle$ is only needed in the E-Freeze-Final and E-Freeze-Simple cases.

- Case E-Beta:

Given: $\langle S; (\lambda x. e) v \rangle \longrightarrow \langle S; e[x := v] \rangle$.

To show: $\langle U_S(S); (\lambda x. e) v \rangle \longrightarrow \mathbf{error}$, where $i \leq 1$.

By assumption, $U_S(S') = \top_S$.

Since $S' = S$, it must be the case that $U_S(S') = U_S(S) = \top_S$.

Hence, by the definition of **error**, $\langle U_S(S); (\lambda x. e) v \rangle = \mathbf{error}$.

Hence $\langle U_S(S); (\lambda x. e) v \rangle \longrightarrow^i \mathbf{error}$, with $i = 0$.

- Case E-New:

Given: $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \mathbf{false})]; l \rangle$.

To show: $\langle U_S(S); \mathbf{new} \rangle \longrightarrow^i \mathbf{error}$, where $i \leq 1$.

By E-New, $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \langle (U_S(S))[l' \mapsto (\perp, \mathbf{false})]; l' \rangle$, where $l' \notin \text{dom}(U_S(S))$.

By assumption, U_S is non-conflicting with $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \mathbf{false})]; l \rangle$.

Therefore $l \notin \text{dom}(U_S(S))$.

Therefore, in $\langle (U_S(S))[l' \mapsto (\perp, \mathbf{false})]; l' \rangle$, we can α -rename l' to l .

Therefore $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \langle (U_S(S))[l \mapsto (\perp, \mathbf{false})]; l \rangle$.

Also, since U_S is non-conflicting with $\langle S; \mathbf{new} \rangle \longrightarrow \langle S[l \mapsto (\perp, \mathbf{false})]; l \rangle$,

we have that $(U_S(S[l \mapsto (\perp, \mathbf{false})]))(l) = (S[l \mapsto (\perp, \mathbf{false})])(l) = (\perp, \mathbf{false})$.

Hence $(U_S(S))[l \mapsto (\perp, \mathbf{false})] = U_S(S[l \mapsto (\perp, \mathbf{false})])$.

Therefore $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \langle U_S(S[l \mapsto (\perp, \mathbf{false})]); l \rangle$.

By assumption, $U_S(S[l \mapsto (\perp, \mathbf{false})]) = \top_S$.

Therefore $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \langle \top_S; l \rangle$.

Hence, by the definition of **error**, $\langle U_S(S); \mathbf{new} \rangle \longrightarrow \mathbf{error}$.

Hence $\langle U_S(S); \text{new} \rangle \longrightarrow^i \mathbf{error}$, with $i = 1$.

- Case E-Put:

Given: $\langle S; \text{put}_i l \rangle \longrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; () \rangle$.

To show: $\langle U_S(S); \text{put}_i l \rangle \longrightarrow^{i'} \mathbf{error}$, where $i' \leq 1$.

Consider whether $U_S(S) = \top_S$:

- If $U_S(S) = \top_S$:

In this case, by the definition of **error**, $\langle U_S(S); \text{put}_i l \rangle = \mathbf{error}$.

Hence $\langle U_S(S); \text{put}_i l \rangle \longrightarrow^{i'} \mathbf{error}$, with $i' = 0$.

- If $U_S(S) \neq \top_S$:

Since $U_S(S) \neq \top_S$, we know that $S \neq \top_S$.

Also, from the premises of E-Put, we have that $u_{p_i}(p_1) \neq \top_p$.

Hence $S[l \mapsto u_{p_i}(p_1)] \neq \top_S$.

Since $U_S(S) \neq \top_S$ and $S[l \mapsto u_{p_i}(p_1)] \neq \top_S$, but $U_S(S[l \mapsto u_{p_i}(p_1)]) = \top_S$, it must be U_S 's action on the contents of l that updates $S[l \mapsto u_{p_i}(p_1)]$ to \top_S .

Let u_{p_j} be the state update operation in U_S that affects the contents of l .

Then $u_{p_j}(u_{p_i}(p_1)) = \top_p$.

Since state update operations commute, $u_{p_i}(u_{p_j}(p_1)) = \top_p$.

Since u_{p_j} is the state update operation in U_S that affects the contents of l , we have that

$(U_S(S))(l) = u_{p_j}(p_1)$.

Since $U_S(S) \neq \top_S$, $u_{p_j}(p_1) \neq \top_p$.

Therefore, by E-Put, $\langle U_S(S); \text{put}_i l \rangle \longrightarrow \langle (U_S(S))[l \mapsto u_{p_i}(u_{p_j}(p_1))]; () \rangle$.

Since $u_{p_j}(u_{p_i}(p_1)) = \top_p$, $\langle U_S(S); \text{put}_i l \rangle \longrightarrow \mathbf{error}$.

Hence $\langle U_S(S); \text{put}_i l \rangle \longrightarrow^{i'} \mathbf{error}$, with $i' = 1$.

- Case E-Get:

Given: $\langle S; \text{get } l P \rangle \longrightarrow \langle S; p_2 \rangle$.

To show: $\langle U_S(S); \text{get } l P \rangle \longrightarrow^i \mathbf{error}$, where $i \leq 1$.

By assumption, $U_S(S) = \top_S$.

Hence, by the definition of **error**, $\langle U_S(S); \text{get } l \ P \rangle = \mathbf{error}$.

Hence $\langle U_S(S); \text{get } l \ P \rangle \longrightarrow^i \mathbf{error}$, with $i = 0$.

- Case E-Freeze-Init:

Given: $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle \longrightarrow \langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e, \{\}, \{\} \rangle$.

To show: $\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle \longrightarrow^i \mathbf{error}$, where $i \leq 1$.

By assumption, $U_S(S) = \top_S$.

Hence, by the definition of **error**, $\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle = \mathbf{error}$.

Hence $\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e \rangle \longrightarrow^i \mathbf{error}$, with $i = 0$.

- Case E-Spawn-Handler:

Given:

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow$

$\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e_0[x := d_2], e, \dots\}, \{d_2\} \cup H \rangle$.

To show:

$\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow^i \mathbf{error}$, where $i \leq 1$.

By assumption, $U_S(S) = \top_S$.

Hence, by the definition of **error**, $\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle = \mathbf{error}$.

Hence $\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{e, \dots\}, H \rangle \longrightarrow^i \mathbf{error}$, with $i = 0$.

- Case E-Freeze-Final:

Given: $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

To show: $\langle U_S(S); \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow^{i'} \mathbf{error}$, where $i' \leq 1$.

Consider whether $U_S(S) = \top_S$:

– If $U_S(S) = \top_S$:

In this case, by the definition of **error**, $\langle U_S(S); \text{freeze } l \rangle = \mathbf{error}$.

Hence $\langle U_S(S); \text{freeze } l \rangle \longrightarrow^{i'} \mathbf{error}$, with $i' = 0$.

– If $U_S(S) \neq \top_S$:

Since $U_S(S) \neq \top_S$ and $S[l \mapsto (d_1, \text{true})] \neq \top_S$, but $U_S(S[l \mapsto (d_1, \text{true})]) = \top_S$, it must be U_S 's action on the contents of l in $S[l \mapsto (d_1, \text{true})]$ that updates $S[l \mapsto (d_1, \text{true})]$ to \top_S .

Since the contents of l in $S[l \mapsto (d_1, \text{true})]$ are (d_1, true) , it must not be the case that $S(l) = (d_1, \text{true})$, because otherwise, $U_S(S)$ would be \top_S .

Therefore $S(l) = (d_1, \text{false})$.

Let u_{p_i} be the state update operation in U_S that updates the contents of l .

Hence $u_{p_i}((d_1, \text{true})) = \top_p$.

Recall that U_S is freeze-safe with $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

By Definition 3.14, then, since the contents of l change in status during the transition from $\langle S; \text{freeze } l \text{ after } Q \text{ with } \lambda x. e_0, \{v, \dots\}, H \rangle$ to $\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$, we know that U_S either freezes the contents of l (having no other effect on them), or it acts as the identity on the contents of l .

Hence $(U_S(S[l \mapsto (d_1, \text{true})]))(l) = (d_1, \text{true})$.

But this is a contradiction since $(U_S(S[l \mapsto (d_1, \text{true})]))(l) = u_{p_i}((d_1, \text{true})) = \top_p$.

Hence this case cannot occur.

- Case E-Freeze-Simple:

Given: $\langle S; \text{freeze } l \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

To show: $\langle U_S(S); \text{freeze } l \rangle \longrightarrow^{i'} \mathbf{error}$, where $i' \leq 1$.

Consider whether $U_S(S) = \top_S$:

- If $U_S(S) = \top_S$:

In this case, by the definition of **error**, $\langle U_S(S); \text{freeze } l \rangle = \mathbf{error}$.

Hence $\langle U_S(S); \text{freeze } l \rangle \longrightarrow^{i'} \mathbf{error}$, with $i' = 0$.

- If $U_S(S) \neq \top_S$:

Since $U_S(S) \neq \top_S$ and $S[l \mapsto (d_1, \text{true})] \neq \top_S$, but $U_S(S[l \mapsto (d_1, \text{true})]) = \top_S$, it must be U_S 's action on the contents of l in $S[l \mapsto (d_1, \text{true})]$ that updates $S[l \mapsto (d_1, \text{true})]$ to \top_S .

Since the contents of l in $S[l \mapsto (d_1, \text{true})]$ are (d_1, true) , it must not be the case that $S(l) = (d_1, \text{true})$, because otherwise, $U_S(S)$ would be \top_S .

Therefore $S(l) = (d_1, \text{false})$.

Let u_{p_i} be the state update operation in U_S that updates the contents of l .

Hence $u_{p_i}((d_1, \text{true})) = \top_p$.

Recall that U_S is freeze-safe with $\langle S; \text{freeze } l \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$.

By Definition 3.14, then, since the contents of l change in status during the transition from $\langle S; \text{freeze } l \rangle$ to $\langle S[l \mapsto (d_1, \text{true})]; d_1 \rangle$, we know that U_S either freezes the contents of l (having no other effect on them), or it acts as the identity on the contents of l .

Hence $(U_S(S[l \mapsto (d_1, \text{true})]))(l) = (d_1, \text{true})$.

But this is a contradiction since $(U_S(S[l \mapsto (d_1, \text{true})]))(l) = u_{p_i}((d_1, \text{true})) = \top_p$.

Hence this case cannot occur.

□

A.17. Proof of Lemma 3.10

Proof. Suppose $\sigma \mapsto \sigma_a$ and $\sigma \mapsto \sigma_b$.

We have to show that either there exist σ_c, i, j, π such that $\sigma_a \mapsto^i \sigma_c$ and $\pi(\sigma_b) \mapsto^j \sigma_c$ and $i \leq 1$ and $j \leq 1$, or that $\sigma_a \mapsto \text{error}$ or $\sigma_b \mapsto \text{error}$.

By inspection of the operational semantics, it must be the case that σ steps to σ_a by the E-Eval-Ctxt rule.

Let $\sigma = \langle S; E_a[e_{a_1}] \rangle$ and let $\sigma_a = \langle S_a; E_a[e_{a_2}] \rangle$.

Likewise, it must be the case that σ steps to σ_b by the E-Eval-Ctxt rule.

Let $\sigma = \langle S; E_b[e_{b_1}] \rangle$ and let $\sigma_b = \langle S_b; E_b[e_{b_2}] \rangle$.

Note that $\sigma = \langle S; E_a[e_{a_1}] \rangle = \langle S; E_b[e_{b_1}] \rangle$, and so $E_a[e_{a_1}] = E_b[e_{b_1}]$, but E_a and E_b may differ and e_{a_1} and e_{b_1} may differ.

First, consider the possibility that $E_a = E_b$ (and $e_{a_1} = e_{b_1}$).

Since $\langle S; E_a[e_{a_1}] \rangle \mapsto \langle S_a; E_a[e_{a_2}] \rangle$ by E-Eval-Ctxt and $\langle S; E_b[e_{b_1}] \rangle \mapsto \langle S_b; E_b[e_{b_2}] \rangle$ by E-Eval-Ctxt, we have from the premise of E-Eval-Ctxt that $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ and $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$.

But then, since $e_{a_1} = e_{b_1}$, by Internal Determinism (Lemma 3.4) there is a permutation π' such that $\langle S_a; e_{a_2} \rangle = \pi'(\langle S_b; e_{b_2} \rangle)$, modulo choice of events.

We have two cases:

- In the case where the steps $\sigma \mapsto \sigma_a$ and $\sigma \mapsto \sigma_b$ are both by E-Spawn-Handler and they handle different events d_2 and d'_2 , then we can satisfy the proof by choosing the final configuration σ_c as the configuration where both d_2 and d'_2 have been handled.

Both σ_a and σ_b can step to this configuration by E-Spawn-Handler: if the step from σ to σ_a handles d_2 then the step from σ_a to σ_c handles d'_2 , while if the step from σ to σ_b handles d'_2 then the step from σ_b to σ_c handles d_2 .

The store in the final configuration is S_a or S_b , which are equal because E-Spawn-Handler does not affect the store, and we can satisfy the proof by choosing $i = 1$ and $j = 0$ and $\pi = \text{id}$.

- Otherwise, we can satisfy the proof by choosing $\sigma_c = \langle S_a; e_{a_2} \rangle$ and $i = 0$ and $j = 0$ and $\pi = \text{id}$.

The rest of this proof deals with the more interesting case in which $E_a \neq E_b$ (and $e_{a_1} \neq e_{b_1}$).

Since $\langle S; E_a[e_{a_1}] \rangle \mapsto \langle S_a; E_a[e_{a_2}] \rangle$ and $\langle S; E_b[e_{b_1}] \rangle \mapsto \langle S_b; E_b[e_{b_2}] \rangle$ and $E_a[e_{a_1}] = E_b[e_{b_1}]$, and since $E_a \neq E_b$, we have from Lemma 3.5 (Locality) that there exist evaluation contexts E'_a and E'_b such that:

- $E'_a[e_{a_1}] = E_b[e_{b_2}]$, and
- $E'_b[e_{b_1}] = E_a[e_{a_2}]$, and
- $E'_a[e_{a_2}] = E'_b[e_{b_2}]$.

In some of the cases that follow, we will choose $\sigma_c = \mathbf{error}$, and in some we will prove that one of σ_a or σ_b steps to **error**.

In most cases, however, our approach will be to show that there exist S', i, j, π such that:

- $\langle S_a; E_a[e_{a_2}] \rangle \mapsto^i \langle S'; E'_a[e_{a_2}] \rangle$, and
- $\pi(\langle S_b; E_b[e_{b_2}] \rangle) \mapsto^j \langle S'; E'_a[e_{a_2}] \rangle$.

Since $E'_a[e_{a_1}] = E_b[e_{b_2}]$, $E'_b[e_{b_1}] = E_a[e_{a_2}]$, and $E'_a[e_{a_2}] = E'_b[e_{b_2}]$, it suffices to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto^i \langle S'; E'_b[e_{b_2}] \rangle$, and
- $\pi(\langle S_b; E'_a[e_{a_1}] \rangle) \mapsto^j \langle S'; E'_a[e_{a_2}] \rangle$.

From the premise of E-Eval-Ctxt, we have that $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ and $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$.

We proceed by case analysis on the rule by which $\langle S; e_{a_1} \rangle$ steps to $\langle S_a; e_{a_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_a \neq \mathbf{error}$.

(1) Case E-Beta: We have $S_a = S$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \mathbf{error}$.

(a) Case E-Beta: We have $S_a = S$ and $S_b = S$.

Choose $S' = S = S_a = S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_a; E'_b[e_{b_2}] \rangle$, and
- $\langle S; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$,

both of which follow immediately from $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$ and $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ and E-Eval-Ctxt.

(b) Case E-New: We have $S_a = S$ and $S_b = S[l \mapsto (\perp, \text{false})]$.

Choose $S' = S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_b; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$.

The first of these follows immediately from $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ and E-Eval-Ctxt.

For the second, consider that $S_b = S[l \mapsto (\perp, \text{false})] = U_S(S)$, where U_S is the store update operation that acts as the identity on the contents of all existing locations, and adds the binding $l \mapsto (\perp, \text{false})$ if no binding for l exists.

Note that:

- U_S is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S_a) \neq \top_S$, since $U_S(S_a) = U_S(S) = S_b$ and we know that $\sigma_b \neq \mathbf{error}$; and
- U_S is freeze-safe with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since $S_a = S$, so there are no locations whose contents differ in status between them.

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \hookrightarrow \langle U_S(S_a); e_{a_2} \rangle.$$

Hence $\langle S_b; e_{a_1} \rangle \hookrightarrow \langle S_b; e_{a_2} \rangle$.

By E-Eval-Ctxt, it follows that $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$, as we were required to show.

(c) Case E-Put: We have $S_a = S$ and $S_b = S[l \mapsto u_{p_i}(p_1)]$.

Choose $S' = S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_b; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$.

The first of these follows immediately from $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ and E-Eval-Ctxt.

For the second, consider that $S_b = U_S(S)$, where U_S is the store update operation that applies u_{p_i} to the contents of l and acts as the identity on all other locations.

Note that:

- U_S is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S_a) \neq \top_S$, since $U_S(S_a) = U_S(S) = S_b$ and we know that $\sigma_b \neq \mathbf{error}$; and
- U_S is freeze-safe with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since $S_a = S$, so there are no locations whose contents differ in status between them.

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \hookrightarrow \langle U_S(S_a); e_{a_2} \rangle.$$

$$\text{Hence } \langle S_b; e_{a_1} \rangle \hookrightarrow \langle S_b; e_{a_2} \rangle.$$

By E-Eval-Ctxt, it follows that $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$, as we were required to show.

- (d) Case E-Put-Err: We have $S_a = S$ and $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 1$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

The second of these is immediately true because since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, and so $\langle S_b; E'_a[e_{a_1}] \rangle$ is equal to \mathbf{error} as well.

For the first, observe that $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, hence by E-Eval-Ctxt, $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_b; E'_b[e_{b_2}] \rangle$.

But $S_b = \top_S$, so $\langle S_b; E'_b[e_{b_2}] \rangle$ is equal to \mathbf{error} , and so $\langle S; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, as required.

- (e) Case E-Get: Similar to case 1a, since $S_a = S$ and $S_b = S$.
- (f) Case E-Freeze-Init: Similar to case 1a, since $S_a = S$ and $S_b = S$.
- (g) Case E-Spawn-Handler: Similar to case 1a, since $S_a = S$ and $S_b = S$.
- (h) Case E-Freeze-Final: We have $S_a = S$ and $S_b = S[l \mapsto (d_1, \text{true})]$.

Choose $S' = S_b$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S; E'_b[e_{b_1}] \rangle \mapsto \langle S_b; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$.

The first of these follows immediately from $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$ and E-Eval-Ctxt.

For the second, note that $S_b = U_S(S)$, where U_S is the store update operation that freezes the contents of l and acts as the identity on the contents of all other locations.

Note that:

- U_S is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S_a) \neq \top_S$, since $U_S(S_a) = U_S(S) = S_b$ and we know that $\sigma_b \neq \mathbf{error}$; and
- U_S is freeze-safe with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since $S_a = S$, so there are no locations whose contents differ in status between them.

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \hookrightarrow \langle U_S(S_a); e_{a_2} \rangle.$$

Hence $\langle S_b; e_{a_1} \rangle \hookrightarrow \langle S_b; e_{a_2} \rangle$.

By E-Eval-Ctxt, it follows that $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b; E'_a[e_{a_2}] \rangle$, as we were required to show.

- (i) Case E-Freeze-Simple: Similar to case 1h, since $S_b = S[l \mapsto (d_1, \text{true})]$.
- (2) Case E-New: We have $S_a = S[l \mapsto (\perp, \text{false})]$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \mathbf{error}$.

(a) Case E-Beta: By symmetry with case 1b.

(b) Case E-New: We have $S_a = S[l \mapsto (\perp, \text{false})]$ and $S_b = S[l' \mapsto (\perp, \text{false})]$.

Now consider whether $l = l'$:

- If $l \neq l'$:

Choose $S' = S[l' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S[l' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S[l' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S_a = S[l \mapsto (\perp, \text{false})] = U_S(S)$, where U_S is the store update operation that acts as the identity on the contents of all existing locations, and adds the binding $l \mapsto (\perp, \text{false})$ if no binding for l exists.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, since the only location allocated in the transition is l' , and $l \neq l'$ in this case;
- $U_S(S_b) \neq \top_S$, since $U_S(S_b) = S[l' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]$ and we know $S \neq \top_S$ and the addition of new bindings $l \mapsto (\perp, \text{false})$ and $l' \mapsto (\perp, \text{false})$ cannot cause it to become \top_S ; and
- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, since $S_b = S[l' \mapsto (\perp, \text{false})]$ and $l' \notin \text{dom}(S)$, so there are no locations whose contents differ in status between S and S_b .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S_b); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto (\perp, \text{false})]; e_{b_1} \rangle \hookrightarrow \langle S_b[l \mapsto (\perp, \text{false})]; e_{b_2} \rangle.$$

By E-Eval-Ctxt it follows that

$$\langle S[l \mapsto (\perp, \text{false})]; E'_b[e_{b_1}] \rangle \hookrightarrow \langle S_b[l \mapsto (\perp, \text{false})]; E'_b[e_{b_2}] \rangle, \text{ which, since } S_b = S[l' \mapsto (\perp, \text{false})], \text{ is what we were required to show.}$$

The argument for the second is symmetrical.

- If $l = l'$:

In this case, observe that we do *not* want the expression in the final configuration to be $E'_a[e_{a_2}]$ (nor its equivalent, $E'_b[e_{b_2}]$).

The reason for this is that $E'_a[e_{a_2}]$ contains both occurrences of l .

Rather, we want both configurations to step to a configuration in which exactly one occurrence of l has been renamed to a fresh location l'' .

Let l'' be a location such that $l'' \notin \text{dom}(S)$ and $l'' \neq l$ (and hence $l'' \neq l'$, as well).

Then choose $S' = S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]$, $i = 1, j = 1$, and $\pi = \{(l, l'')\}$.

Either $\langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_a[\pi(e_{a_2})] \rangle$ or $\langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_b[\pi(e_{b_2})] \rangle$ would work as a final configuration; we choose $\langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_b[\pi(e_{b_2})] \rangle$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_b[\pi(e_{b_2})] \rangle$, and
- $\pi(\langle S_b; E'_a[e_{a_1}] \rangle) \mapsto \langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_b[\pi(e_{b_2})] \rangle$.

For the first of these, since $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, we have by Lemma 3.3 (Permutability) that $\pi(\langle S; e_{b_1} \rangle) \hookrightarrow \pi(\langle S_b; e_{b_2} \rangle)$.

Since $\pi = \{(l, l'')\}$, but $l \notin S$ (from the side condition on E-New), we have that $\pi(\langle S; e_{b_1} \rangle) = \langle S; e_{b_1} \rangle$.

Since $\langle S_b; e_{b_2} \rangle = \langle S[l' \mapsto (\perp, \text{false})]; l' \rangle$, and $l = l'$, we have that $\pi(\langle S_b; e_{b_2} \rangle) = \langle S[l'' \mapsto (\perp, \text{false})]; \pi(e_{b_2}) \rangle$.

Hence $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto (\perp, \text{false})]; \pi(e_{b_2}) \rangle$.

Let U_S be the store update operation that acts as the identity on the contents of all existing locations, and adds the binding $l \mapsto (\perp, \text{false})$ if no binding for l exists.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto (\perp, \text{false})]; \pi(e_{b_2}) \rangle$, since the only location allocated in the transition is l'' ;
- $U_S(S[l'' \mapsto (\perp, \text{false})]) \neq \top_S$, since $U_S(S[l'' \mapsto (\perp, \text{false})]) = S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]$ and we know $S \neq \top_S$ and the addition of new bindings $l \mapsto (\perp, \text{false})$ and $l'' \mapsto (\perp, \text{false})$ cannot cause it to become \top_S ; and

- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto (\perp, \text{false})]; \pi(e_{b_2}) \rangle$, since $l'' \notin \text{dom}(S)$, so there are no locations whose contents differ in status between S and $S[l'' \mapsto (\perp, \text{false})]$.

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S[l'' \mapsto (\perp, \text{false})]); \pi(e_{b_2}) \rangle.$$

$$\text{Hence } \langle S[l \mapsto (\perp, \text{false})]; e_{b_1} \rangle \hookrightarrow \langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; \pi(e_{b_2}) \rangle.$$

By E-Eval-Ctxt it follows that

$$\langle S[l \mapsto (\perp, \text{false})]; E'_b[e_{b_1}] \rangle \hookrightarrow \langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_b[\pi(e_{b_2})] \rangle,$$

which, since $S[l \mapsto (\perp, \text{false})] = S_a$, is what we were required to show.

For the second, observe that since $S_b = S[l \mapsto (\perp, \text{false})]$, we have that $\pi(S_b) = S[l'' \mapsto (\perp, \text{false})]$.

Also, since l does not occur in e_{a_1} , we have that $\pi(E'_a[e_{a_1}]) = (\pi(E'_a))[e_{a_1}]$.

Hence we have to show that

$$\begin{aligned} &\langle S[l'' \mapsto (\perp, \text{false})]; (\pi(E'_a))[e_{a_1}] \rangle \mapsto \\ &\langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; E'_b[\pi(e_{b_2})] \rangle. \end{aligned}$$

Let U_S be the store update operation that acts as the identity on the contents of all existing locations, and adds the binding $l'' \mapsto (\perp, \text{false})$ if no binding for l'' exists.

Note that:

- U_S is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since the only location allocated in the transition is l ;
- $U_S(S_a) \neq \top_S$, since $U_S(S_a) = S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]$ and we know $S \neq \top_S$ and the addition of new bindings $l \mapsto (\perp, \text{false})$ and $l'' \mapsto (\perp, \text{false})$ cannot cause it to become \top_S ; and
- U_S is freeze-safe with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since $S_a = S[l \mapsto (\perp, \text{false})]$ and $l \notin \text{dom}(S)$, so there are no locations whose contents differ in status between S and S_a .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \hookrightarrow \langle U_S(S_a); e_{a_2} \rangle.$$

$$\text{Hence } \langle S[l'' \mapsto (\perp, \text{false})]; e_{a_1} \rangle \hookrightarrow \langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; e_{a_2} \rangle.$$

By E-Eval-Ctxt it follows that

$$\begin{aligned} & \langle S[l'' \mapsto (\perp, \text{false})]; (\pi(E'_a))[e_{a_1}] \rangle \mapsto \\ & \langle S[l'' \mapsto (\perp, \text{false})][l \mapsto (\perp, \text{false})]; (\pi(E'_a))[e_{a_2}] \rangle, \end{aligned}$$

which completes the case since $E'_b[\pi(e_{b_2})] = (\pi(E'_a))[e_{a_2}]$.

- (c) Case E-Put: We have $S_a = S[l \mapsto (\perp, \text{false})]$ and $S_b = S[l' \mapsto u_{p_i}(p_1)]$, where $l \neq l'$ (since $l \notin \text{dom}(S)$, but $l' \in \text{dom}(S)$).

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle S_b[l \mapsto (\perp, \text{false})]; E'_b[e_{b_2}] \rangle$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle S_b[l \mapsto (\perp, \text{false})]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S_a = S[l \mapsto (\perp, \text{false})] = U_S(S)$, where U_S is the store update operation that acts as the identity on the contents of all existing locations, and adds the binding $l \mapsto (\perp, \text{false})$ if no binding for l exists.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S_b) \neq \top_S$, since $U_S(S_b) = S_b[l \mapsto (\perp, \text{false})]$, and we know $S_b \neq \top_S$ and the addition of a new binding $l \mapsto (\perp, \text{false})$ cannot cause it to become \top_S ; and
- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, since $S_b = S[l' \mapsto u_{p_i}(p_1)]$ and u_{p_i} does not alter the status of p_1 .

(By Definition 3.4, u_{p_i} can only change the status bit of a location if its contents are (d, true) and $u_i(d) \neq d$, in which case u_{p_i} changes the contents of the location to (\top, false) ; however, that cannot be the case here since then $u_{p_i}(p_1)$ would be \top_p , contradicting the premise of E-Put.)

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \longrightarrow \langle U_S(S_b); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto (\perp, \text{false})]; e_{b_1} \rangle \longrightarrow \langle S_b[l \mapsto (\perp, \text{false})]; e_{b_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto (\perp, \text{false})]; E'_b[e_{b_1}] \rangle \longmapsto \langle S_b[l \mapsto (\perp, \text{false})]; E'_b[e_{b_2}] \rangle,$$

which, since $S_a = S[l \mapsto (\perp, \text{false})]$, is what we were required to show.

For the second, let U_S be the store update operation that applies u_{p_i} to the contents of l' if it exists, and adds a binding $l' \mapsto u_{p_i}(p_1)$ if no binding for l' exists.

Consider that $S_b = U_S(S)$, and $S_b[l \mapsto (\perp, \text{false})] = S_a[l' \mapsto u_{p_i}(p_1)] = U_S(S_a)$.

Note that:

- U_S is non-conflicting with $\langle S; e_{a_1} \rangle \longrightarrow \langle S_a; e_{a_2} \rangle$, since the only location allocated in the transition is l ;
- $U_S(S_a) \neq \top_S$, since $U_S(S_a) = S[l \mapsto (\perp, \text{false})][l' \mapsto u_{p_i}(p_1)]$ and we know $S \neq \top_S$ and the addition of a new binding $l \mapsto (\perp, \text{false})$ and updating the contents of location l' to $u_{p_i}(p_1)$ in S cannot cause it to become \top_S (since if $u_{p_i}(p_1) = \top_p$, $\langle S; e_{b_1} \rangle$ would not have been able to step by E-Put); and
- U_S is freeze-safe with $\langle S; e_{a_1} \rangle \longrightarrow \langle S_a; e_{a_2} \rangle$, since $S_a = S[l \mapsto (\perp, \text{false})]$ and $l \notin \text{dom}(S)$, so there are no locations whose contents differ in status between S and S_a .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \longrightarrow \langle U_S(S_a); e_{a_2} \rangle.$$

$$\text{Hence } \langle S_b; e_{a_1} \rangle \longrightarrow \langle S_b[l \mapsto (\perp, \text{false})]; e_{a_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S_b; E'_a[e_{a_1}] \rangle \longmapsto \langle S_b[l \mapsto (\perp, \text{false})]; E'_a[e_{a_2}] \rangle,$$

as we were required to show.

- (d) Case E-Put-Err: We have $S_a = S[l \mapsto (\perp, \text{false})]$ and $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 1$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

The second of these is immediately true because since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, and so $\langle S_b; E'_a[e_{a_1}] \rangle$ is equal to \mathbf{error} as well.

For the first, observe that since $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, we have by Lemma 3.6 (Monotonicity) that $S \sqsubseteq_S S_a$.

Therefore, since $\langle S; e_{b_1} \rangle \hookrightarrow \mathbf{error}$,

we have by Lemma 3.9 (Error Preservation) that $\langle S_a; e_{b_1} \rangle \hookrightarrow \mathbf{error}$.

Since \mathbf{error} is equal to $\langle \top_S; e \rangle$ for all expressions e , $\langle S_a; e_{b_1} \rangle \hookrightarrow \langle \top_S; e \rangle$ for all e .

Therefore, by E-Eval-Ctxt, $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle \top_S; E'_b[e] \rangle$ for all e .

Since $\langle \top_S; E'_b[e] \rangle$ is equal to \mathbf{error} , we have that $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, as we were required to show.

- (e) Case E-Get: Similar to case 2a, since $S_a = S[l \mapsto (\perp, \text{false})]$ and $S_b = S$.
- (f) Case E-Freeze-Init: Similar to case 2a, since $S_a = S[l \mapsto (\perp, \text{false})]$ and $S_b = S$.
- (g) Case E-Spawn-Handler: Similar to case 2a, since $S_a = S[l \mapsto (\perp, \text{false})]$ and $S_b = S$.
- (h) Case E-Freeze-Final: We have $S_a = S[l \mapsto (\perp, \text{false})]$ and $S_b = S[l' \mapsto (d_1, \text{true})]$, where $l \neq l'$ (since $l \notin \text{dom}(S)$, but $l' \in \text{dom}(S)$).

Choose $S' = S[l \mapsto (\perp, \text{false})][l' \mapsto (d_1, \text{true})]$, $i = i$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S[l \mapsto (\perp, \text{false})]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l \mapsto (\perp, \text{false})][l' \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle$, and
- $\langle S[l' \mapsto (d_1, \text{true})]; E'_a[e_{a_1}] \rangle \mapsto \langle S[l \mapsto (\perp, \text{false})][l' \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S[l \mapsto (\perp, \text{false})] = U_S(S)$, where U_S is the store update operation that acts as the identity on the contents of all existing locations, and adds the binding $l \mapsto (\perp, \text{false})$ if no binding for l exists.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S_b) \neq \top_S$, since $U_S(S_b) = S_b[l \mapsto (\perp, \text{false})]$, and we know $S_b \neq \top_S$ and the addition of a new binding $l \mapsto (\perp, \text{false})$ cannot cause it to become \top_S ; and
- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, since $S_b = S[l' \mapsto (d_1, \text{true})]$ and so the only location that can change in status between S and S_b is l' , and U_S acts as the identity on l' .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S_b); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto (\perp, \text{false})]; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto (\perp, \text{false})][l' \mapsto (d_1, \text{true})]; e_{b_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto (\perp, \text{false})]; E'_b[e_{b_1}] \rangle \hookrightarrow \langle S[l \mapsto (\perp, \text{false})][l' \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle,$$

as we were required to show.

For the second, consider that $S[l' \mapsto (d_1, \text{true})] = U_S(S)$, where U_S is the store update operation that freezes the contents of l' and acts as the identity on the contents of all other locations.

Note that:

- U_S is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since the only location allocated in the transition is l , and $l \neq l'$;
- $U_S(S_a) \neq \top_S$, since $U_S(S_a) = S_a[l' \mapsto (d_1, \text{true})] = S_b[l \mapsto (\perp, \text{false})]$, and we know $S_b \neq \top_S$ and the addition of a new binding $l \mapsto (\perp, \text{false})$ cannot cause it to become \top_S ; and

- U_S is freeze-safe with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, since $S_a = S[l \mapsto (\perp, \text{false})]$ and $l \notin \text{dom}(S)$, so there are no locations whose contents differ in status between S and S_a .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \hookrightarrow \langle U_S(S_a); e_{a_2} \rangle.$$

$$\text{Hence } \langle S[l' \mapsto (d_1, \text{true})]; e_{a_1} \rangle \hookrightarrow \langle S[l \mapsto (\perp, \text{false})][l' \mapsto (d_1, \text{true})]; e_{a_2} \rangle.$$

By E-Eval-Ctxt it follows that

$$\langle S[l' \mapsto (d_1, \text{true})]; E'_a[e_{a_1}] \rangle \mapsto \langle S[l \mapsto (\perp, \text{false})][l' \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle,$$

as we were required to show.

- (i) Case E-Freeze-Simple: Similar to case 2h, since $S_a = S[l \mapsto (\perp, \text{false})]$ and $S_b = S[l' \mapsto (d_1, \text{true})]$, where $l \neq l'$ (since $l \notin \text{dom}(S)$, but $l' \in \text{dom}(S)$).
- (3) Case E-Put: We have $S_a = S[l \mapsto u_{p_i}(p_1)]$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \text{error}$.

- (a) Case E-Beta: By symmetry with case 1c.
- (b) Case E-New: By symmetry with case 2c.
- (c) Case E-Put: We have $S_a = S[l \mapsto u_{p_i}(p_1)]$ and $S_b = S[l' \mapsto u_{p_j}(p'_1)]$, where $p'_1 = S(l')$.

Now consider whether $l = l'$:

- If $l \neq l'$:

Choose $S' = S[l' \mapsto u_{p_j}(p'_1)][l \mapsto u_{p_i}(p_1)]$, $i = 1, j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l' \mapsto u_{p_j}(p'_1)][l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_2}] \rangle$, and
- $\langle S[l' \mapsto u_{p_j}(p'_1)]; E'_a[e_{a_1}] \rangle \mapsto \langle S[l' \mapsto u_{p_j}(p'_1)][l \mapsto u_{p_i}(p_1)]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S[l \mapsto u_{p_i}(p_1)] = U_S(S)$, where U_S is the store update operation that applies u_{p_i} to the contents of l if it exists, and adds a binding $l \mapsto u_{p_i}(p_1)$ if no binding for l exists.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto u_{p_j}(p'_1)]; e_{b_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S[l' \mapsto u_{p_j}(p'_1)]) \neq \top_S$, since $U_S(S[l' \mapsto u_{p_j}(p'_1)]) = S[l' \mapsto u_{p_j}(p'_1)][l \mapsto u_{p_i}(p_1)]$ and we know $S \neq \top_S$ and updating the contents of location l to $u_{p_i}(p_1)$ and the contents of location l' to $u_{p_j}(p'_1)$ in S cannot cause it to become \top_S (because if so, then we would have $S_a = \top_S$ or $S_b = \top_S$, which we know are not the case); and
- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto u_{p_j}(p'_1)]; e_{b_2} \rangle$, since u_{p_j} does not alter the status of p'_1 .

(By Definition 3.4, u_{p_j} can only change the status bit of a location if its contents are (d, true) and $u_j(d) \neq d$, in which case u_{p_j} changes the contents of the location to (\top, false) ; however, that cannot be the case here since then $u_{p_j}(p'_1)$ would be \top_p , contradicting the premise of E-Put.)

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S[l' \mapsto u_{p_j}(p'_1)]); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto u_{p_j}(p'_1)][l \mapsto u_{p_i}(p_1)]; e_{b_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l' \mapsto u_{p_j}(p'_1)][l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_2}] \rangle,$$

as we were required to show.

The argument for the second is symmetrical.

- If $l = l'$: Note that since $l = l'$, $p_1 = p'_1$ as well.

Consider whether $u_{p_i}(u_{p_j}(p_1)) = \top_p$:

– If $u_{p_i}(u_{p_j}(p_1)) = \top_p$:

Choose $\sigma_c = \mathbf{error}$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- * $\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, and
- * $\langle S[l \mapsto u_{p_j}(p_1)]; E'_a[e_{a_1}] \rangle \mapsto \mathbf{error}$.

For the first of these, consider that $S[l \mapsto u_{p_i}(p_1)] = U_S(S)$, where U_S is the store update operation that applies u_{p_i} to the contents of l if it exists.

Note that:

- * U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_j}(p_1)]; e_{b_2} \rangle$, since no locations are allocated in the transition;
- * $U_S(S[l \mapsto u_{p_j}(p_1)]) = \top_S$, since $U_S(S[l \mapsto u_{p_j}(p_1)]) = S[l \mapsto u_{p_i}(u_{p_j}(p_1))]$ and we know $u_{p_i}(u_{p_j}(p_1)) = \top_p$ in this case;
- * U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_j}(p_1)]; e_{b_2} \rangle$, since u_{p_j} does not alter the status of p_1 .

(By Definition 3.4, u_{p_j} can only change the status bit of a location if its contents are (d, true) and $u_j(d) \neq d$, in which case u_{p_j} changes the contents of the location to (\top, false) ; however, that cannot be the case here since then $u_{p_j}(p_1)$ would be \top_p , contradicting the premise of E-Put.)

Therefore, by Lemma 3.8 (Generalized Clash), we have that there exists $i' \leq 1$ such that $\langle U_S(S); e_{b_1} \rangle \hookrightarrow^{i'} \mathbf{error}$.

Hence $\langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \hookrightarrow^{i'} \mathbf{error}$.

If $i' = 0$, we would have $\langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle = \langle S_a; e_{b_1} \rangle = \mathbf{error}$.

So we would have $S_a = \top_S$ by the definition of **error**, but then we would have $\sigma_a = \mathbf{error}$, a contradiction.

Therefore $i' = 1$, and so we have $\langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \hookrightarrow \mathbf{error}$.

Since **error** = $\langle \top_S; e \rangle$ for all e , we have $\langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \longrightarrow \langle \top_S; e \rangle$ for all e .

So, by E-Eval-Ctxt, we have that $\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \longrightarrow \langle \top_S; E'_b[e] \rangle$ for all e .

Hence $\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \longrightarrow \mathbf{error}$.

The argument for the second is symmetrical.

– If $u_{p_i}(u_{p_j}(p_1)) \neq \top_p$:

Choose $S' = S[l \mapsto u_{p_i}(u_{p_j}(p_1))]$, $i = 1, j = 1$, and $\pi = \text{id}$.

We have to show that:

- * $\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \longrightarrow \langle S[l \mapsto u_{p_i}(u_{p_j}(p_1))]; E'_b[e_{b_2}] \rangle$, and
- * $\langle S[l \mapsto u_{p_j}(p_1)]; E'_a[e_{a_1}] \rangle \longrightarrow \langle S[l \mapsto u_{p_i}(u_{p_j}(p_1))]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S[l \mapsto u_{p_i}(p_1)] = U_S(S)$, where U_S is the store update operation that applies u_{p_i} to the contents of l if it exists.

Note that:

- * U_S is non-conflicting with $\langle S; e_{b_1} \rangle \longrightarrow \langle S[l \mapsto u_{p_j}(p_1)]; e_{b_2} \rangle$, since no locations are allocated in the transition;
- * $U_S(S[l \mapsto u_{p_j}(p_1)]) \neq \top_S$, since $U_S(S[l \mapsto u_{p_j}(p_1)]) = S[l \mapsto u_{p_i}(u_{p_j}(p_1))]$ and we know $S \neq \top_S$ and $u_{p_i}(u_{p_j}(p_1)) \neq \top_p$ in this case;
- * U_S is freeze-safe with $\langle S; e_{b_1} \rangle \longrightarrow \langle S[l \mapsto u_{p_j}(p_1)]; e_{b_2} \rangle$, since u_{p_j} does not alter the status of p_1 .

(By Definition 3.4, u_{p_j} can only change the status bit of a location if its contents are (d, true) and $u_j(d) \neq d$, in which case u_{p_j} changes the contents of the location to (\top, false) ; however, that cannot be the case here since then $u_{p_j}(p_1)$ would be \top_p , contradicting the premise of E-Put.)

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \longrightarrow \langle U_S(S[l \mapsto u_{p_j}(p_1)]); e_{b_2} \rangle.$$

Hence $\langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(u_{p_j}(p_1))]; e_{b_2} \rangle$.

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l \mapsto u_{p_i}(u_{p_j}(p_1))]; E'_b[e_{b_2}] \rangle,$$

as we were required to show.

The argument for the second is symmetrical.

- (d) Case E-Put-Err: We have $S_a = S[l \mapsto u_{p_i}(p_1)]$ and $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 1$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

The second of these is immediately true because since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, and so $\langle S_b; E'_a[e_{a_1}] \rangle$ is equal to \mathbf{error} as well.

For the first, observe that since $\langle S; e_{a_1} \rangle \hookrightarrow \langle S_a; e_{a_2} \rangle$, we have by Lemma 3.6 (Monotonicity) that $S \sqsubseteq_S S_a$.

Therefore, since $\langle S; e_{b_1} \rangle \hookrightarrow \mathbf{error}$,

we have by Lemma 3.9 (Error Preservation) that $\langle S_a; e_{b_1} \rangle \hookrightarrow \mathbf{error}$.

Since \mathbf{error} is equal to $\langle \top_S; e \rangle$ for all expressions e , $\langle S_a; e_{b_1} \rangle \hookrightarrow \langle \top_S; e \rangle$ for all e .

Therefore, by E-Eval-Ctxt, $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \langle \top_S; E'_b[e] \rangle$ for all e .

Since $\langle \top_S; E'_b[e] \rangle$ is equal to \mathbf{error} , we have that $\langle S_a; E'_b[e_{b_1}] \rangle \mapsto \mathbf{error}$, as we were required to show.

- (e) Case E-Get: Similar to case 3a, since $S_a = S[l \mapsto u_{p_i}(p_1)]$ and $S_b = S$.
- (f) Case E-Freeze-Init: Similar to case 3a, since $S_a = S[l \mapsto u_{p_i}(p_1)]$ and $S_b = S$.
- (g) Case E-Spawn-Handler: Similar to case 3a, since $S_a = S[l \mapsto u_{p_i}(p_1)]$ and $S_b = S$.
- (h) Case E-Freeze-Final: We have $S_a = S[l \mapsto u_{p_i}(p_1)]$ and $S_b = S[l' \mapsto (d_1, \text{true})]$.

Now consider whether $l = l'$:

- If $l \neq l'$:

Choose $S' = S[l \mapsto u_{p_i}(p_1)][l' \mapsto (d_1, \text{true})]$, $i = 1, j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l \mapsto u_{p_i}(p_1)][l' \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle$, and
- $\langle S[l' \mapsto (d_1, \text{true})]; E'_a[e_{a_1}] \rangle \mapsto \langle S[l \mapsto u_{p_i}(p_1)][l' \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S[l \mapsto u_{p_i}(p_1)] = U_S(S)$, where U_S is the store update operation that applies u_{p_i} to the contents of l if it exists, and adds a binding $l \mapsto u_{p_i}(p_1)$ if no binding for l exists, and acts as the identity on all other locations.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto (d_1, \text{true})]; e_{b_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S[l' \mapsto (d_1, \text{true})]) \neq \top_S$,
since $U_S(S[l' \mapsto (d_1, \text{true})]) = S[l' \mapsto (d_1, \text{true})][l \mapsto u_{p_i}(p_1)]$ and we know $S \neq \top_S$ and updating the contents of location l to $u_{p_i}(p_1)$ and freezing the contents of location l' in S cannot cause it to become \top_S (because if so, then we would have $S_a = \top_S$ or $S_b = \top_S$, which we know are not the case); and
- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto (d_1, \text{true})]; e_{b_2} \rangle$, since the only location that can change in status between S and $S[l' \mapsto (d_1, \text{true})]$ is l' , and U_S acts as the identity on l' .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S[l' \mapsto (d_1, \text{true})]); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto (d_1, \text{true})][l \mapsto u_{p_i}(p_1)]; e_{b_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l' \mapsto (d_1, \text{true})][l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_2}] \rangle,$$

as we were required to show.

For the second, consider that $S[l' \mapsto (d_1, \text{true})] = U_S(S)$, where U_S is the store update operation that freezes the contents of l' and acts as the identity on the contents of all other locations.

Note that:

- U_S is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; e_{a_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S[l \mapsto u_{p_i}(p_1)]) \neq \top_S$, since $U_S(S[l \mapsto u_{p_i}(p_1)]) = S[l \mapsto u_{p_i}(p_1)][l' \mapsto (d_1, \text{true})]$, and we know $S \neq \top_S$ and updating the contents of location l to $u_{p_i}(p_1)$ and freezing the contents of location l in S cannot cause it to become \top_S (because if so, then we would have $S_a = \top_S$ or $S_b = \top_S$, which we know are not the case); and
- U_S is freeze-safe with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; e_{a_2} \rangle$, since u_{p_i} does not alter the status of p_1 .

(By Definition 3.4, u_{p_i} can only change the status bit of a location if its contents are (d, true) and $u_i(d) \neq d$, in which case u_{p_i} changes the contents of the location to (\top, false) ; however, that cannot be the case here since then $u_{p_i}(p_1)$ would be \top_p , and we would have $S_a = \top_S$, a contradiction.)

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \hookrightarrow \langle U_S(S[l \mapsto u_{p_i}(p_1)]); e_{a_2} \rangle.$$

$$\text{Hence } \langle S[l' \mapsto (d_1, \text{true})]; e_{a_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(p_1)][l' \mapsto (d_1, \text{true})]; e_{a_2} \rangle.$$

By E-Eval-Ctxt, it follows that $\langle S[l' \mapsto (d_1, \text{true})]; E'_a[e_{a_1}] \rangle \mapsto \langle S[l \mapsto u_{p_i}(p_1)][l' \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle$,

as we were required to show.

- If $l = l'$:

We have two cases to consider:

- $u_{p_i}((d_1, \text{true})) = \top_p$:

Since $(S[l \mapsto (d_1, \text{true})])(l) = (d_1, \text{true})$ and $u_{p_i}((d_1, \text{true})) = \top_p$, by E-Put-Err we have that $\langle S[l \mapsto (d_1, \text{true})]; \text{put}_i l \rangle \longrightarrow \mathbf{error}$.

Since $S_b = S[l \mapsto (d_1, \text{true})]$, we have that $\langle S_b; \text{put}_i l \rangle \longrightarrow \mathbf{error}$.

Since $\langle S; e_{a_1} \rangle \longrightarrow \langle S_a; e_{a_2} \rangle$ by E-Put, it must be the case that $e_{a_1} = \text{put}_i l$.

Hence $\langle S_b; e_{a_1} \rangle \longrightarrow \mathbf{error}$.

Since \mathbf{error} is equal to $\langle \top_S; e \rangle$ for all expressions e , $\langle S_b; e_{a_1} \rangle \longrightarrow \langle \top_S; e \rangle$ for all e .

Therefore, by E-Eval-Ctxt, $\langle S_b; E'_a[e_{a_1}] \rangle \longmapsto \langle \top_S; E'_a[e] \rangle$ for all e .

Since $\langle \top_S; E'_a[e] \rangle$ is equal to \mathbf{error} , we have that $\langle S_b; E'_a[e_{a_1}] \rangle \longmapsto \mathbf{error}$.

Since $E'_a[e_{a_1}] = E_b[e_{b_2}]$, we have that $\langle S_b; E_b[e_{b_2}] \rangle \longmapsto \mathbf{error}$.

Since $\sigma_b = \langle S_b; E_b[e_{b_2}] \rangle$, we therefore have that $\sigma_b \longmapsto \mathbf{error}$, and the case is satisfied.

– $u_{p_i}((d_1, \text{true})) \neq \top_p$:

In this case, by the definition of U_p (Definition 3.4),

it must be the case that $u_{p_i}((d_1, \text{true})) = (d_1, \text{true})$.

Choose $S' = S[l \mapsto (d_1, \text{true})]$, $i = 1, j = 1$, and $\pi = \text{id}$.

We have to show that:

* $\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \longmapsto \langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle$, and

* $\langle S[l \mapsto (d_1, \text{true})]; E'_a[e_{a_1}] \rangle \longmapsto \langle S[l \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle$.

For the first of these, consider that $S[l \mapsto u_{p_i}(p_1)] = U_S(S)$, where U_S is the store update operation that applies u_{p_i} to the contents of l if it exists, and adds a binding $l \mapsto u_{p_i}(p_1)$ if no binding for l exists, and acts as the identity on all other locations.

Note that:

* U_S is non-conflicting with $\langle S; e_{b_1} \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; e_{b_2} \rangle$, since no locations are allocated in the transition;

* $U_S(S[l \mapsto (d_1, \text{true})]) \neq \top_S$,

since $U_S(S[l \mapsto (d_1, \text{true})]) = S[l \mapsto u_{p_i}((d_1, \text{true}))]$ and we know $S \neq \top_S$ and $u_{p_i}((d_1, \text{true})) \neq \top_p$; and

- * U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; e_{b_2} \rangle$, since the only location that can change in status between S and $S[l \mapsto (d_1, \text{true})]$ is l , and U_S acts as the identity on l .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S[l \mapsto (d_1, \text{true})]); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}((d_1, \text{true}))]; e_{b_2} \rangle.$$

$$\text{Since } u_{p_i}((d_1, \text{true})) = (d_1, \text{true}),$$

$$\text{we have that } \langle S[l \mapsto u_{p_i}(p_1)]; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; e_{b_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto u_{p_i}(p_1)]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle,$$

as we were required to show.

For the second, consider that $S[l \mapsto (d_1, \text{true})] = U_S(S)$, where U_S is the store update operation that freezes the contents of l and acts as the identity on the contents of all other locations.

Note that:

- * U_S is non-conflicting with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; e_{a_2} \rangle$, since no locations are allocated in the transition;
- * $U_S(S[l \mapsto u_{p_i}(p_1)]) \neq \top_S$, since $U_S(S[l \mapsto u_{p_i}(p_1)]) = S[l \mapsto (d_1, \text{true})]$ (since, by Definition 3.4, $u_i(d_1) = d_1$; otherwise we would have $u_{p_i}((d_1, \text{true})) = \top_p$, a contradiction), and we know $S \neq \top_S$ and freezing the contents of location l in S cannot cause it to become \top_S ; and
- * U_S is freeze-safe with $\langle S; e_{a_1} \rangle \hookrightarrow \langle S[l \mapsto u_{p_i}(p_1)]; e_{a_2} \rangle$, since u_{p_i} does not alter the status of p_1 .

(By Definition 3.4, u_{p_i} can only change the status bit of a location if its contents are (d, true) and $u_i(d) \neq d$, in which case u_{p_i} changes the contents of the location to (\top, false) ; however, that cannot be the case here since then $u_{p_i}(p_1)$ would be \top_p , and we would have $S_a = \top_S$, a contradiction.)

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{a_1} \rangle \longrightarrow \langle U_S(S[l \mapsto u_{p_i}(p_1)]); e_{a_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto (d_1, \text{true})]; e_{a_1} \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; e_{a_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto (d_1, \text{true})]; E'_a[e_{a_1}] \rangle \longrightarrow \langle S[l \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle,$$

as we were required to show.

- (i) Case E-Freeze-Simple: Similar to case 3h, since $S_a = S[l \mapsto u_{p_i}(p_1)]$ and $S_b = S[l' \mapsto (d_1, \text{true})]$.

- (4) Case E-Put-Err: We have $\langle S_a; e_{a_2} \rangle = \mathbf{error}$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \mathbf{error}$.

- (a) Case E-Beta: By symmetry with case 1d.

- (b) Case E-New: By symmetry with case 2d.

- (c) Case E-Put: By symmetry with case 3d.

- (d) Case E-Put-Err: We have $\langle S_a; e_{a_2} \rangle = \mathbf{error}$ and $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, and so we choose $\sigma_c = \mathbf{error}$, $i = 0$, $j = 0$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle = \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle = \mathbf{error}$.

Since $\langle S_a; e_{a_2} \rangle = \mathbf{error}$, $S_a = \top_S$, and since $\langle S_b; e_{b_2} \rangle = \mathbf{error}$, $S_b = \top_S$, so both of the above follow immediately.

- (e) Case E-Get: Similar to case 4a, since $\langle S_a; e_{a_2} \rangle = \mathbf{error}$ and $S_b = S$.
- (f) Case E-Freeze-Init: Similar to case 4a, since $\langle S_a; e_{a_2} \rangle = \mathbf{error}$ and $S_b = S$.
- (g) Case E-Spawn-Handler: Similar to case 4a, since $\langle S_a; e_{a_2} \rangle = \mathbf{error}$ and $S_b = S$.
- (h) Case E-Freeze-Final: We have $\langle S_a; e_{a_2} \rangle = \mathbf{error}$ and $S_b = S[l \mapsto (d_1, \text{true})]$, and so we choose $\sigma_c = \mathbf{error}$, $i = 0$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S_a; E'_b[e_{b_1}] \rangle = \mathbf{error}$, and
- $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \mathbf{error}$.

The first of these is immediately true because since $\langle S_a; e_{a_2} \rangle = \mathbf{error}$, $S_a = \top_S$, and so $\langle S_a; E'_b[e_{b_1}] \rangle$ is equal to \mathbf{error} as well.

For the second, observe that since $\langle S; e_{b_1} \rangle \hookrightarrow \langle S_b; e_{b_2} \rangle$, we have by Lemma 3.6 (Monotonicity) that $S \sqsubseteq_S S_b$.

Therefore, since $\langle S; e_{a_1} \rangle \hookrightarrow \mathbf{error}$, we have by Lemma 3.9 that $\langle S_b; e_{a_1} \rangle \hookrightarrow \mathbf{error}$.

Since \mathbf{error} is equal to $\langle \top_S; e \rangle$ for all expressions e , $\langle S_b; e_{a_1} \rangle \hookrightarrow \langle \top_S; e \rangle$ for all e .

Therefore, by E-Eval-Ctxt, $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \langle \top_S; E'_a[e] \rangle$ for all e .

Since $\langle \top_S; E'_a[e] \rangle$ is equal to \mathbf{error} , we have that $\langle S_b; E'_a[e_{a_1}] \rangle \mapsto \mathbf{error}$, as we were required to show.

- (i) Case E-Freeze-Simple: Similar to case 4h, since $S_b = S[l \mapsto (d_1, \text{true})]$.
- (5) Case E-Get: We have $S_a = S$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an \mathbf{error} configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \mathbf{error}$.

- (a) Case E-Beta: By symmetry with case 1e.
- (b) Case E-New: By symmetry with case 2e.
- (c) Case E-Put: By symmetry with case 3e.
- (d) Case E-Put-Err: By symmetry with case 4e.

- (e) Case E-Get: Similar to case 5a, since $S_a = S$ and $S_b = S$.
 - (f) Case E-Freeze-Init: Similar to case 5a, since $S_a = S$ and $S_b = S$.
 - (g) Case E-Spawn-Handler: Similar to case 5a, since $S_a = S$ and $S_b = S$.
 - (h) Case E-Freeze-Final: Similar to case 1h, since $S_a = S$ and $S_b = S[l \mapsto (d_1, \text{true})]$.
 - (i) Case E-Freeze-Simple: Similar to case 1i, since $S_a = S$ and $S_b = S[l \mapsto (d_1, \text{true})]$.
- (6) Case E-Freeze-Init: We have $S_a = S$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \text{error}$.

- (a) Case E-Beta: By symmetry with case 1f.
 - (b) Case E-New: By symmetry with case 2f.
 - (c) Case E-Put: By symmetry with case 3f.
 - (d) Case E-Put-Err: By symmetry with case 4f.
 - (e) Case E-Get: By symmetry with case 5f.
 - (f) Case E-Freeze-Init: Similar to case 6a, since $S_a = S$ and $S_b = S$.
 - (g) Case E-Spawn-Handler: Similar to case 6a, since $S_a = S$ and $S_b = S$.
 - (h) Case E-Freeze-Final: Similar to case 1h, since $S_a = S$ and $S_b = S[l \mapsto (d_1, \text{true})]$.
 - (i) Case E-Freeze-Simple: Similar to case 1i, since $S_a = S$ and $S_b = S[l \mapsto (d_1, \text{true})]$.
- (7) Case E-Spawn-Handler: We have $S_a = S$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \text{error}$.

- (a) Case E-Beta: By symmetry with case 1g.
- (b) Case E-New: By symmetry with case 2g.
- (c) Case E-Put: By symmetry with case 3g.
- (d) Case E-Put-Err: By symmetry with case 4g.

- (e) Case E-Get: By symmetry with case 5g.
 - (f) Case E-Freeze-Init: By symmetry with case 6g.
 - (g) Case E-Spawn-Handler: Similar to case 7a, since $S_a = S$ and $S_b = S$.
 - (h) Case E-Freeze-Final: Similar to case 1h, since $S_a = S$ and $S_b = S[l \mapsto (d_1, \text{true})]$.
 - (i) Case E-Freeze-Simple: Similar to case 1i, since $S_a = S$ and $S_b = S[l \mapsto (d_1, \text{true})]$.
- (8) Case E-Freeze-Final: We have $S_a = S[l \mapsto (d_1, \text{true})]$.

We proceed by case analysis on the rule by which $\langle S; e_{b_1} \rangle$ steps to $\langle S_b; e_{b_2} \rangle$.

Since the only way an **error** configuration can arise is by the E-Put-Err rule, we can assume in all other cases that $\sigma_b \neq \mathbf{error}$.

- (a) Case E-Beta: By symmetry with case 1h.
- (b) Case E-New: By symmetry with case 2h.
- (c) Case E-Put: By symmetry with case 3h.
- (d) Case E-Put-Err: By symmetry with case 4h.
- (e) Case E-Get: By symmetry with case 5h.
- (f) Case E-Freeze-Init: By symmetry with case 6h.
- (g) Case E-Spawn-Handler: By symmetry with case 7h.
- (h) Case E-Freeze-Final: We have $S_a = S[l \mapsto (d_1, \text{true})]$ and $S_b = S[l' \mapsto (d'_1, \text{true})]$.

Now consider whether $l = l'$:

- If $l \neq l'$:

Choose $S' = S[l' \mapsto (d'_1, \text{true})][l \mapsto (d_1, \text{true})]$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

$$- \langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l' \mapsto (d'_1, \text{true})][l \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle,$$

and

$$- \langle S[l' \mapsto (d'_1, \text{true})]; E'_a[e_{a_1}] \rangle \mapsto \langle S[l' \mapsto (d'_1, \text{true})][l \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle.$$

For the first of these, consider that $S[l \mapsto (d_1, \text{true})] = U_S(S)$, where U_S is the store update operation that freezes the contents of l and acts as the identity on the contents of all other locations.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto (d'_1, \text{true})]; e_{b_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S[l' \mapsto (d'_1, \text{true})]) \neq \top_S$,
since $U_S(S[l' \mapsto (d'_1, \text{true})]) = S[l' \mapsto (d'_1, \text{true})][l \mapsto (d_1, \text{true})]$ and we know $S \neq \top_S$ and freezing the contents of locations l and l' in S cannot cause it to become \top_S (because if so, then we would have $S_a = \top_S$ or $S_b = \top_S$, which we know are not the case); and
- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto (d'_1, \text{true})]; e_{b_2} \rangle$, since the only location that can change in status between S and $S[l' \mapsto (d'_1, \text{true})]$ is l' , and U_S acts as the identity on l' .

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S[l' \mapsto (d'_1, \text{true})]); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto (d_1, \text{true})]; e_{b_1} \rangle \hookrightarrow \langle S[l' \mapsto (d'_1, \text{true})][l \mapsto (d_1, \text{true})]; e_{b_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_1}] \rangle \hookrightarrow \langle S[l' \mapsto (d'_1, \text{true})][l \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle,$$

as we were required to show.

The argument for the second is symmetrical.

- If $l = l'$:

Note that since $l = l'$, $d_1 = d'_1$ as well.

Choose $S' = S[l \mapsto (d_1, \text{true})]$, $i = 1$, $j = 1$, and $\pi = \text{id}$.

We have to show that:

- $\langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_1}] \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle$, and

$$- \langle S[l' \mapsto (d'_1, \text{true})]; E'_a[e_{a_1}] \rangle \mapsto \langle S[l \mapsto (d_1, \text{true})]; E'_a[e_{a_2}] \rangle.$$

For the first of these, consider that $S[l \mapsto (d_1, \text{true})] = U_S(S)$, where U_S is the store update operation that freezes the contents of l and acts as the identity on the contents of all other locations.

Note that:

- U_S is non-conflicting with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; e_{b_2} \rangle$, since no locations are allocated in the transition;
- $U_S(S[l \mapsto (d_1, \text{true})]) \neq \top_S$, since $U_S(S[l \mapsto (d_1, \text{true})]) = S[l \mapsto (d_1, \text{true})]$, and we know $S \neq \top_S$ and freezing the contents of location l in S cannot cause it to become \top_S ; and
- U_S is freeze-safe with $\langle S; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; e_{b_2} \rangle$, since the only location that can change in status between S and $S[l \mapsto (d_1, \text{true})]$ is l , and U_S freezes the contents of l but has no other effect on them.

Therefore, by Lemma 3.7 (Generalized Independence), we have that

$$\langle U_S(S); e_{b_1} \rangle \hookrightarrow \langle U_S(S[l \mapsto (d_1, \text{true})]); e_{b_2} \rangle.$$

$$\text{Hence } \langle S[l \mapsto (d_1, \text{true})]; e_{b_1} \rangle \hookrightarrow \langle S[l \mapsto (d_1, \text{true})]; e_{b_2} \rangle.$$

By E-Eval-Ctxt, it follows that

$$\langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_1}] \rangle \mapsto \langle S[l \mapsto (d_1, \text{true})]; E'_b[e_{b_2}] \rangle,$$

as we were required to show.

The argument for the second is symmetrical.

- (i) Case E-Freeze-Simple: Similar to case 8h, since $S_a = S[l \mapsto (d_1, \text{true})]$ and $S_b = S[l' \mapsto (d'_1, \text{true})]$.
- (9) Case E-Freeze-Simple: We have $S_a = S[l \mapsto (d_1, \text{true})]$.
 - (a) Case E-Beta: By symmetry with case 1i.
 - (b) Case E-New: By symmetry with case 2i.
 - (c) Case E-Put: By symmetry with case 3i.

- (d) Case E-Put-Err: By symmetry with case 4i.
- (e) Case E-Get: By symmetry with case 5i.
- (f) Case E-Freeze-Init: By symmetry with case 6i.
- (g) Case E-Spawn-Handler: By symmetry with case 7i.
- (h) Case E-Freeze-Final: By symmetry with case 8i.
- (i) Case E-Freeze-Simple: Similar to case 9h, since $S_a = S[l \mapsto (d_1, \text{true})]$ and $S_b = S[l' \mapsto (d'_1, \text{true})]$.

□

A.18. Proof of Lemma 3.11

Proof. Suppose $\sigma \mapsto \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq m$.

We are required to show that either:

- (1) there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq 1$, or
- (2) there exists $k \leq m$ such that $\sigma' \mapsto^k \text{error}$, or there exists $k \leq 1$ such that $\sigma'' \mapsto^k \text{error}$.

We proceed by induction on m .

In the base case of $m = 1$, the result is immediate from Lemma 3.10, with $k = 1$.

For the induction step, suppose $\sigma \mapsto^m \sigma'' \mapsto \sigma'''$ and suppose the lemma holds for m .

We show that it holds for $m + 1$, as follows.

From the induction hypothesis, we have that either:

- (1) there exist σ'_c, i', j', π' such that $\sigma' \mapsto^{i'} \sigma'_c$ and $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$ and $i' \leq m$ and $j' \leq 1$, or
- (2) there exists $k' \leq m$ such that $\sigma' \mapsto^{k'} \text{error}$, or there exists $k' \leq 1$ such that $\sigma'' \mapsto^{k'} \text{error}$.

We consider these two cases in turn:

- (1) There exist σ'_c, i', j', π' such that $\sigma' \mapsto^{i'} \sigma'_c$ and $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$ and $i' \leq m$ and $j' \leq 1$:

We proceed by cases on j' :

- If $j' = 0$, then $\pi'(\sigma'') = \sigma'_c$.

Since $\sigma'' \mapsto \sigma'''$, we have that $\pi'(\sigma'') \mapsto \pi'(\sigma''')$ by Lemma 3.3 (Permutability).

We can then choose $\sigma_c = \pi'(\sigma''')$ and $i = i' + 1$ and $j = 0$ and $\pi = \pi'$.

The key is that $\sigma' \mapsto^{i'} \sigma'_c = \pi'(\sigma'') \mapsto \pi'(\sigma''')$ for a total of $i' + 1$ steps.

- If $j' = 1$:

First, since $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$, then by Lemma 3.3 (Permutability) we have that $\sigma'' \mapsto^{j'} \pi'^{-1}(\sigma'_c)$.

Then, by $\sigma'' \mapsto^{j'} \pi'^{-1}(\sigma'_c)$ and $\sigma'' \mapsto \sigma'''$ and Lemma 3.10, one of the following two cases is true:

- (a) There exist σ''_c and i'' and j'' and π'' such that $\pi'^{-1}(\sigma'_c) \mapsto^{i''} \sigma''_c$ and $\pi''(\sigma''') \mapsto^{j''} \sigma''_c$ and $i'' \leq 1$ and $j'' \leq 1$.

Since $\pi'^{-1}(\sigma'_c) \mapsto^{i''} \sigma''_c$, by Lemma 3.3 (Permutability) we have that $\sigma'_c \mapsto^{i''} \pi'(\sigma''_c)$.

So we also have $\sigma' \mapsto^{i'} \sigma'_c \mapsto^{i''} \pi'(\sigma''_c)$.

Since $\pi''(\sigma''') \mapsto^{j''} \sigma''_c$, by Lemma 3.3 (Permutability) we have that $\pi'(\pi''(\sigma''')) \mapsto^{j''} \pi'(\sigma''_c)$.

In summary, we pick $\sigma_c = \pi'(\sigma''_c)$ and $i = i' + i''$ and $j = j''$ and $\pi = \pi'' \circ \pi'$, which is sufficient because $i = i' + i'' \leq m + 1$ and $j = j'' \leq 1$.

- (b) $\pi'^{-1}(\sigma'_c) \mapsto$ **error** or $\sigma'' \mapsto$ **error**.

If $\sigma'' \mapsto$ **error**, then choosing $k = 1$ satisfies the proof.

Otherwise, $\pi'^{-1}(\sigma'_c) \mapsto$ **error**.

Then, by Lemma 3.3 we have that $\sigma'_c \mapsto \pi'(\mathbf{error})$.

By Definition 3.11, $\pi'(\mathbf{error}) = \mathbf{error}$, and so $\sigma'_c \mapsto \mathbf{error}$.

Therefore $\sigma' \mapsto^{i'} \sigma'_c \mapsto \mathbf{error}$.

Hence $\sigma' \mapsto^{i'+1} \mathbf{error}$.

Since $i' \leq m$, we have that $i' + 1 \leq m + 1$, and so choosing $k = i' + 1$ satisfies the proof.

(2) There exists $k' \leq m$ such that $\sigma' \mapsto^{k'} \mathbf{error}$, or there exists $k' \leq 1$ such that $\sigma'' \mapsto^{k'} \mathbf{error}$:

If there exists $k' \leq m$ such that $\sigma' \mapsto^{k'} \mathbf{error}$, then choosing $k = k'$ satisfies the proof.

Otherwise, there exists $k' \leq 1$ such that $\sigma'' \mapsto^{k'} \mathbf{error}$.

We proceed by cases on k' :

- If $k' = 0$, then $\sigma'' = \mathbf{error}$.

Hence this case is not possible, since $\sigma'' \mapsto \sigma'''$ and \mathbf{error} cannot step.

- If $k' = 1$:

From $\sigma'' \mapsto \sigma'''$ and $\sigma'' \mapsto^{k'} \mathbf{error}$ and Lemma 3.10, one of the following two cases is true:

- (a) There exist σ_c'' and i'' and j'' and π'' such that $\mathbf{error} \mapsto^{i''} \sigma_c''$ and $\pi''(\sigma''') \mapsto^{j''} \sigma_c''$ and $i'' \leq 1$ and $j'' \leq 1$.

Since \mathbf{error} cannot step, $i'' = 0$ and $\sigma_c'' = \mathbf{error}$.

By Definition 3.11, $\pi''(\sigma''') = \sigma'''$.

Hence $\sigma''' \mapsto^{j''} \mathbf{error}$.

Since $j'' \leq 1$, choosing $k = j''$ satisfies the proof.

- (b) $\mathbf{error} \mapsto \mathbf{error}$ or $\sigma''' \mapsto \mathbf{error}$.

Since \mathbf{error} cannot step, $\sigma''' \mapsto \mathbf{error}$.

Hence choosing $k = 1$ satisfies the proof.

□

A.19. Proof of Lemma 3.12

Proof. Suppose that $\sigma \mapsto^n \sigma'$ and $\sigma \mapsto^m \sigma''$, where $1 \leq n$ and $1 \leq m$.

We are required to show that either:

- (1) there exist σ_c, i, j, π such that $\sigma' \mapsto^i \sigma_c$ and $\pi(\sigma'') \mapsto^j \sigma_c$ and $i \leq m$ and $j \leq n$, or

(2) there exists $k \leq m$ such that $\sigma' \mapsto^k \mathbf{error}$, or there exists $k \leq n$ such that $\sigma'' \mapsto^k \mathbf{error}$.

We proceed by induction on n .

In the base case of $n = 1$, the result is immediate from Lemma 3.11.

For the induction step, suppose $\sigma \mapsto^n \sigma' \mapsto \sigma'''$ and suppose the lemma holds for n .

We show that it holds for $n + 1$, as follows.

From the induction hypothesis, we have that either:

- (1) there exist σ'_c, i', j', π' such that $\sigma' \mapsto^{i'} \sigma'_c$ and $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$ and $i' \leq m$ and $j' \leq n$, or
- (2) there exists $k' \leq m$ such that $\sigma' \mapsto^{k'} \mathbf{error}$, or there exists $k' \leq n$ such that $\sigma'' \mapsto^{k'} \mathbf{error}$.

We consider these two cases in turn:

- (1) There exist σ'_c, i', j', π' such that $\sigma' \mapsto^{i'} \sigma'_c$ and $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$ and $i' \leq m$ and $j' \leq n$:

We proceed by cases on i' :

- If $i' = 0$, then $\sigma' = \sigma'_c$.

We can then choose $\sigma_c = \sigma'''$ and $i = 0$ and $j = j' + 1$ and $\pi = \pi'$.

Since $\pi'(\sigma'') \mapsto^{j'} \sigma'_c \mapsto \sigma'''$, and $j' + 1 \leq n + 1$ since $j' \leq n$, the case is satisfied.

- If $i' \geq 1$:

From $\sigma' \mapsto \sigma'''$ and $\sigma' \mapsto^{i'} \sigma'_c$ and Lemma 3.11, one of the following two cases is true:

- (a) There exist σ''_c and i'' and j'' and π'' such that $\sigma''' \mapsto^{i''} \sigma''_c$ and $\pi''(\sigma'_c) \mapsto^{j''} \sigma''_c$ and $i'' \leq i'$ and $j'' \leq 1$.

Since $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$, by Lemma 3.3 (Permutability) we have that $\pi''(\pi'(\sigma'')) \mapsto^{j'} \pi''(\sigma'_c)$.

So we also have $\pi''(\pi'(\sigma'')) \mapsto^{j'} \pi''(\sigma'_c) \mapsto^{j''} \sigma''_c$.

In summary, we pick $\sigma_c = \sigma''_c$ and $i = i''$ and $j = j' + j''$ and $\pi = \pi' \circ \pi''$, which is sufficient because $i = i'' \leq i' \leq m$ and $j = j' + j'' \leq n + 1$.

- (b) There exists $k'' \leq i'$ such that $\sigma''' \mapsto^{k''} \mathbf{error}$, or there exists $k'' \leq 1$ such that $\sigma'_c \mapsto^{k''} \mathbf{error}$.

If there exists $k'' \leq i'$ such that $\sigma''' \mapsto^{k''} \mathbf{error}$, then choosing $k = k''$ satisfies the proof, since $k'' \leq i' \leq m$.

Otherwise, there exists $k'' \leq 1$ such that $\sigma'_c \mapsto^{k''} \mathbf{error}$.

Hence by Lemma 3.3 (Permutability), we have that $\pi'^{-1}(\sigma'_c) \mapsto^{k''} \pi'^{-1}(\mathbf{error})$.

By Definition 3.11, $\pi'^{-1}(\mathbf{error}) = \mathbf{error}$.

Hence $\pi'^{-1}(\sigma'_c) \mapsto^{k''} \mathbf{error}$.

Since $\pi'(\sigma'') \mapsto^{j'} \sigma'_c$, by Lemma 3.3 (Permutability), we have that $\sigma'' \mapsto^{j'} \pi'^{-1}(\sigma'_c)$.

Therefore, $\sigma'' \mapsto^{j'} \pi'^{-1}(\sigma'_c) \mapsto^{k''} \mathbf{error}$.

Hence $\sigma'' \mapsto^{j'+k''} \mathbf{error}$.

Since $j' \leq n$ and $k'' \leq 1$, $j' + k'' \leq n + 1$.

Hence choosing $k = j' + k''$ satisfies the proof.

- (2) There exists $k' \leq m$ such that $\sigma' \mapsto^{k'} \mathbf{error}$, or there exists $k' \leq n$ such that $\sigma'' \mapsto^{k'} \mathbf{error}$:

If there exists $k' \leq n$ such that $\sigma'' \mapsto^{k'} \mathbf{error}$, then choosing $k = k'$ satisfies the proof.

Otherwise, there exists $k' \leq m$ such that $\sigma' \mapsto^{k'} \mathbf{error}$.

We proceed by cases on k' :

- If $k' = 0$, then $\sigma' = \mathbf{error}$.

Hence this case is not possible, since $\sigma' \mapsto \sigma'''$ and \mathbf{error} cannot step.

- If $k' \geq 1$:

From $\sigma' \mapsto \sigma'''$ and $\sigma' \mapsto^{k'} \mathbf{error}$ and Lemma 3.11, one of the following two cases is true:

- (a) There exist σ''_c and i'' and j'' and π'' such that $\sigma''' \mapsto^{i''} \sigma''_c$ and $\pi''(\mathbf{error}) \mapsto^{j''} \sigma''_c$ and $i'' \leq k'$ and $j'' \leq 1$.

By Definition 3.11, $\pi''(\mathbf{error}) = \mathbf{error}$.

Hence $\mathbf{error} \mapsto^{j''} \sigma''_c$.

Since \mathbf{error} cannot step, $j'' = 0$ and $\sigma''_c = \mathbf{error}$.

Hence $\sigma''' \mapsto^{i''} \mathbf{error}$.

Since $i'' \leq k' \leq m$, choosing $k = i''$ satisfies the proof.

- (b) There exists $k'' \leq k'$ such that $\sigma''' \mapsto^{k''} \mathbf{error}$, or there exists $k'' \leq 1$ such that $\mathbf{error} \mapsto^{k''} \mathbf{error}$.

Since \mathbf{error} cannot step, there exists $k'' \leq k'$ such that $\sigma''' \mapsto^{k''} \mathbf{error}$.

Since $k'' \leq k' \leq m$, choosing $k = k''$ satisfies the proof.

□

A.20. Proof of Theorem 5.2

Proof. Consider replica i of a threshold CvRDT $(S, \leq, s^0, q, t, u, m)$.

Let \mathcal{S} be a threshold set with respect to (S, \leq) .

Consider a method execution $t_i^{k+1}(\mathcal{S})$ (i.e., a threshold query that is the $k + 1$ th method execution on replica i , with threshold set \mathcal{S} as its argument) that returns some set of activation states $S_a \in \mathcal{S}$.

For part 1 of the theorem, we have to show that threshold queries with \mathcal{S} as their argument will always return S_a on subsequent executions at i .

That is, we have to show that, for all $k' > (k + 1)$, the threshold query $t_i^{k'}(\mathcal{S})$ on i returns S_a .

Since $t_i^{k+1}(\mathcal{S})$ returns S_a , from Definition 5.4 we have that for some activation state $s_a \in S_a$, the condition $s_a \leq s_i^k$ holds.

Consider arbitrary $k' > (k + 1)$.

Since state is inflationary across updates, we know that the state $s_i^{k'}$ after method execution k' is at least s_i^k .

That is, $s_i^k \leq s_i^{k'}$.

By transitivity of \leq , then, $s_a \leq s_i^{k'}$.

Hence, by Definition 5.4, $t_i^{k'}(\mathcal{S})$ returns S_a .

For part 2 of the theorem, consider some replica j of $(S, \leq, s^0, q, t, u, m)$, located at process p_j .

We are required to show that, for all $x \geq 0$, the threshold query $t_j^{x+1}(\mathcal{S})$ returns S_a eventually, and blocks until it does.¹

That is, we must show that, for all $x \geq 0$, there exists some finite $n \geq 0$ such that

- for all i in the range $0 \leq i \leq n - 1$, the threshold query $t_j^{x+1+i}(\mathcal{S})$ returns block, and
- for all $i \geq n$, the threshold query $t_j^{x+1+i}(\mathcal{S})$ returns S_a .

Consider arbitrary $x \geq 0$.

Recall that s_j^x is the state of replica j after the x th method execution, and therefore s_j^x is also the state of j when $t_j^{x+1}(\mathcal{S})$ runs. We have three cases to consider:

- $s_i^k \leq s_j^x$.

(That is, replica i 's state after the k th method execution on i is *at or below* replica j 's state after the x th method execution on j .)

Choose $n = 0$.

We have to show that, for all $i \geq n$, the threshold query $t_j^{x+1+i}(\mathcal{S})$ returns S_a .

Since $t_i^{k+1}(\mathcal{S})$ returns S_a , we know that there exists an $s_a \in S_a$ such that $s_a \leq s_i^k$.

Since $s_i^k \leq s_j^x$, we have by transitivity of \leq that $s_a \leq s_j^x$.

Therefore, by Definition 5.4, $t_j^{x+1}(\mathcal{S})$ returns S_a .

Then, by part 1 of the theorem, we have that subsequent executions $t_j^{x+1+i}(\mathcal{S})$ at replica j will also return S_a , and so the case holds.

(Note that this case includes the possibility $s_i^k \equiv s^0$, in which no updates have executed at replica i .)

¹The occurrences of $k + 1$ and $x + 1$ in this proof are an artifact of how we index method executions starting from 1, but states starting from 0. The initial state (of every replica) is s^0 , and so s_i^k is the state of replica i after method execution k has completed at i .

- $s_i^k > s_j^x$.

(That is, replica i 's state after the k th method execution on i is *above* replica j 's state after the x th method execution on j .)

We have two subcases:

- There exists some activation state $s'_a \in S_a$ for which $s'_a \leq s_j^x$.

In this case, we choose $n = 0$.

We have to show that, for all $i \geq n$, the threshold query $t_j^{x+1+i}(\mathcal{S})$ returns S_a .

Since $s'_a \leq s_j^x$, by Definition 5.4, $t_j^{x+1}(\mathcal{S})$ returns S_a .

Then, by part 1 of the theorem, we have that subsequent executions $t_j^{x+1+i}(\mathcal{S})$ at replica j will also return S_a , and so the case holds.

- There is no activation state $s'_a \in S_a$ for which $s'_a \leq s_j^x$.

Since $t_i^{k+1}(\mathcal{S})$ returns S_a , we know that there is some update $u_i^{k'}(a)$ in i 's causal history, for some $k' < (k + 1)$, that updates i from a state at or below s_j^x to s_i^k .²

By eventual delivery, $u_i^{k'}(a)$ is eventually delivered at j .

Hence some update or updates that will increase j 's state from s_j^x to a state at or above some s'_a must reach replica j .³

Let the $x + 1 + r$ th method execution on j be the first update on j that updates its state to some $s_j^{x+1+r} \geq s'_a$, for some activation state $s'_a \in S_a$.

Choose $n = r + 1$.

We have to show that, for all i in the range $0 \leq i \leq r$, the threshold query $t_j^{x+1+i}(\mathcal{S})$ returns block, and that for all $i \geq r + 1$, the threshold query $t_j^{x+1+i}(\mathcal{S})$ returns S_a .

²We know that i 's state was once at or below s_j^x , because i and j started at the same state s^0 and can both only grow. Hence the least that s_j^x can be is s^0 , and we know that i was originally s^0 as well.

³We say “some update or updates” because the exact update $u_i^{k'}(a)$ may not be the update that causes the threshold query at j to unblock; a different update or updates could do it. Nevertheless, the existence of $u_i^{k'}(a)$ means that there is at least one update that will suffice to unblock the threshold query.

A. PROOFS

For the former, since the $x + 1 + r$ th method execution on j is the first one that updates its state to $s_j^{x+1+r} \geq s'_a$, we have by Definition 5.4 that for all i in the range $0 \leq i \leq r$, the threshold query $t_j^{x+1+i}(\mathcal{S})$ returns block.

For the latter, since $s_j^{x+1+r} \geq s'_a$, by Definition 5.4 we have that $t_j^{x+1+r+1}(\mathcal{S})$ returns S_a , and by part 1 of the theorem, we have that for $i \geq r + 1$, subsequent executions $t_j^{x+1+i}(\mathcal{S})$ at replica j will also return S_a , and so the case holds.

- $s_i^k \not\leq s_j^x$ and $s_j^x \not\leq s_i^k$.

(That is, replica i 's state after the k th method execution on i is *not comparable* to replica j 's state after the x th method execution on j .)

Similar to the previous case.

□

APPENDIX B

A PLT Redex Model of λ_{LVish}

I have developed a runnable version¹ of the λ_{LVish} calculus of Chapter 3 using the PLT Redex semantics engineering toolkit [19]. In the Redex of today, it is not possible to directly parameterize a language definition by a lattice.² Instead, taking advantage of Racket’s syntactic abstraction capabilities, I define a Racket macro, `define-lambdaLVish-language`, that serves as a wrapper for a template that implements the lattice-agnostic semantics of λ_{LVish} . The `define-lambdaLVish-language` macro takes the following arguments:

- a *name*, which becomes the *lang-name* passed to Redex’s `define-language` form;
- a “*downset*” *operation*, a Racket-level procedure that takes a lattice element and returns the (finite) set of all lattice elements that are less than or equal to that element (this operation is used to implement the semantics of `freeze` — `after` — `with`, in particular, to determine when the E-Freeze-Final rule can fire);
- a *lub operation*, a Racket procedure that takes two lattice elements and returns a lattice element;
- a list of *update operations*, Racket procedures that each take a lattice element and return a lattice element; and
- a (possibly infinite) set of *lattice elements* represented as Redex *patterns*.

Given these arguments, `define-lambdaLVish-language` generates a Redex model that is specialized to the appropriate lattice and set of update operations. For instance, to generate a Redex model called `lambdaLVish-nat` where the lattice is the non-negative integers ordered by \leq , and the set of update operations is $\{u_{(+1)}, u_{(+2)}\}$ where $u_{(+n)}(d)$ increments d ’s contents by n , one could write:

¹Available at <http://github.com/lkuper/lvar-semantics>.

²See discussion at <http://lists.racket-lang.org/users/archive/2013-April/057075.html>.

```
(define-lambdaLVish-language lambdaLVish-nat downset-op max update-ops natural)
```

where `max` is a built-in Racket procedure, and `downset-op` and `update-ops` can be defined in Racket as follows:

```
(define downset-op
  (lambda (d)
    (if (number? d)
        (append '(Bot) (iota d) `(:,d))
        '(Bot))))

(define update-op-1
  (lambda (d)
    (match d
      ['Bot 1]
      [number (add1 d)])))

(define update-op-2
  (lambda (d)
    (match d
      ['Bot 2]
      [number (add1 (add1 d)])))

(define update-ops `(:,update-op-1 ,update-op-2))
```

The last argument, `natural`, is a Redex pattern that matches any exact non-negative integer. `natural` has no meaning to Racket outside of Redex, but since `define-lambdaLVish-language` is a macro rather than an ordinary procedure, its arguments are not evaluated until they are in the context of Redex, and so passing `natural` as an argument to the macro gives us the behavior we want.

The Redex model is not completely faithful to the λ_{LVish} calculus of Chapter 3: it requires us to specify a finite list of update operations rather than a possibly infinite set of them. However, the list of update operations can be arbitrarily long, and the definitions of the update operations could, at least in principle, be generated programmatically as well.

Lindsey Kuper

Programming Systems Lab, Intel Labs
Intel Corporation
3600 Juliette Lane
Santa Clara, CA 95054

lindsey.kuper@intel.com
lindsey@composition.al
<http://www.cs.indiana.edu/~lkuper>
<http://composition.al>

Education

- **Indiana University**, School of Informatics and Computing, 2008–2015
 - Ph.D. Computer Science, September 2015.
Research committee: Ryan R. Newton (chair), Lawrence S. Moss, Amr Sabry, Chung-chieh Shan.
Dissertation: *Lattice-based Data Structures for Deterministic Parallel and Distributed Programming*.
 - M.S. Computer Science, May 2010.
- **Grinnell College**, 2000–2004
 - B.A. Computer Science and Music (with honors), May 2004.
- Additional coursework and summer schools:
 - Oregon Programming Languages Summer School: Types, Logic, and Verification, Summer 2013.
 - Oregon Programming Languages Summer School: Logic, Languages, Compilation, and Verification, Summer 2012.
 - Operating Systems, Cornell University, Summer 2010.

Employment history

- **Intel Corporation (Programming Systems Lab, Intel Labs)**, Santa Clara, CA
 - *Research Scientist*, September 2014–present.
- **Indiana University**, Bloomington, IN
 - *Research Assistant* (with Ryan Newton), January 2012–August 2014.
 - *Associate Instructor* (with William E. Byrd), August 2011–December 2011.
 - *Research Assistant* (with Amal Ahmed), August 2010–December 2010.
 - *Associate Instructor* (with Daniel P. Friedman), January 2009–May 2010.
- **Mozilla Corporation**, Mountain View, CA
 - *Research Engineering Intern* (supervised by Brian Anderson, Niko Matsakis and Patrick Walton), May–August 2012.
 - *Research Engineering Intern* (supervised by Dave Herman), March–August 2011.
- **GrammaTech, Inc.**, Ithaca, NY
 - *Software Engineering Intern* (supervised by David Melski), May–August 2010.
- **Bedford, Freeman and Worth Publishing Group**, New York, NY and Portland, OR
 - *Associate Project Manager*, July 2006–June 2008.
- **IBCTV, LLC**, Chicago, IL and Portland, OR
 - *Web Designer/Developer*, August 2004–June 2006.

Research funding

- Co-wrote (with Ryan Newton) NSF grant CCF-1218375, Generalizing Monotonic Data Structures for Expressive, Deterministic Parallel Programming (\$377,315; 8/1/2012–7/31/2015), which funded my dissertation work.

Conference publications

- **Lindsey Kuper**, Aaron Todd, Sam Tobin-Hochstadt and Ryan R. Newton.
Taming the parallel effect zoo: extensible deterministic parallelism with LVish.
In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '14)*,
Edinburgh, UK, June 2014.
(52/287 \approx 18% accepted)
- **Lindsey Kuper**, Aaron Turon, Neelakantan R. Krishnaswami and Ryan R. Newton.
Freeze after writing: quasi-deterministic parallel programming with LVars.
In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '14)*,
San Diego, CA, January 2014.
(51/220 \approx 23% accepted)

Workshop publications

- **Lindsey Kuper** and Ryan R. Newton.
Joining forces: toward a unified account of LVars and convergent replicated data types.
In the *5th Workshop on Determinism and Correctness in Parallel Programming (WoDet '14)*,
Salt Lake City, UT, March 2014.
- **Lindsey Kuper** and Ryan R. Newton.
LVars: lattice-based data structures for deterministic parallelism.
In *Proceedings of the 2nd ACM SIGPLAN Workshop on Functional High-Performance Computing (FHPC '13)*,
Boston, MA, September 2013.
- Andrew W. Keep, Michael D. Adams, **Lindsey Kuper**, William E. Byrd and Daniel P. Friedman.
A pattern matcher for miniKanren, or, how to get into trouble with CPS macros.
In *Proceedings of the 2009 Scheme and Functional Programming Workshop (Scheme '09)*,
Boston, MA, August 2009.

Technical reports

- **Lindsey Kuper**, Aaron Turon, Neelakantan R. Krishnaswami and Ryan R. Newton.
Freeze after writing: quasi-deterministic parallel programming with LVars. (56 pages)
Indiana University Technical Report TR710, November 2013.
- **Lindsey Kuper** and Ryan R. Newton.
A lattice-theoretical approach to deterministic parallelism with shared state. (60 pages)
Indiana University Technical Report TR702, October 2012.
- David Melski, David Cok, John Phillips, Scott Wisniewski, Suan Hsi Yong, Nathan Lloyd, **Lindsey Kuper**, Denis Gopan and Alexey Loginov.
Safety in numbers. (104 pages)
GrammarTech, Inc. project final report, November 2010.

Talks

- *LVars for distributed programming, or, LVars and CRDTs join forces.*
 - IFIP Working Group 2.8 (Functional Programming), Kefalonia, Greece, May 26, 2015.
- *LVars: lattice-based data structures for deterministic parallel and distributed programming.*
 - Compose :: Conference, New York, NY, January 31, 2015.
 - Hacker School, New York, NY, March 24, 2014.
 - Intel Labs, Santa Clara, CA, March 21, 2014.
 - University of Utah, Salt Lake City, UT, March 4, 2014.
 - Microsoft Research, Mountain View, CA, January 27, 2014.
- *Joining forces: toward a unified account of LVars and convergent replicated data types.*
 - WoDet 2014, Salt Lake City, UT, March 2, 2014.
- *Freeze after writing: quasi-deterministic parallel programming with LVars.*
 - POPL 2014, San Diego, CA, January 23, 2014.
- *LVars: lattice-based data structures for deterministic parallelism.*
 - Mozilla Corporation, Mountain View, CA, October 31, 2013.
 - RICON West 2013, San Francisco, CA, October 29, 2013.
 - FHPC 2013, Boston, MA, September 23, 2013.
 - Hacker School, New York, NY, June 10, 2013.
- *A lattice-based approach to deterministic parallelism.*
 - MPI-SWS, Saarbrücken, Germany, January 30, 2013.
 - POPL 2013 student talk session, Rome, Italy, January 25, 2013.
- *A lattice-based approach to deterministic parallelism with shared state.*
 - Aarhus University, Aarhus, Denmark, September 14, 2012.
 - University of California–Berkeley, Berkeley, CA, August 16, 2012.
- *Rust typeclasses turn trait-er.*
 - Mozilla Corporation, Mountain View, CA, August 9, 2012.
- *Hacking the Rust object system at Mozilla.*
 - Grinnell College, Grinnell, IA, April 5, 2012.
(invited talk, hosted by the Grinnell Alumni Scholars Program)
- *Some pieces of the Rust object system: extension, overriding, and self.*
 - Mozilla Corporation, Mountain View, CA, August 18, 2011.
- *Parametric polymorphism through run-time sealing, or, theorems for low, low prices!*
 - Northeastern University, Boston, MA, February 23, 2011.
- *A system for testing specifications of CPU semantics, or, what I did on my summer vacation.*
 - GrammarTech, Inc., Ithaca, NY, August 20, 2010.

Undergraduate projects advised

- Isaiah Weating, Indiana University, Spring 2013. Project title: *Parallel Programming with LVars*. Awarded third place in IU Undergraduate Research Opportunities in Computing (UROC) Poster Competition, May 2013.

Teaching

▪ Associate Instructor, Indiana University

- Fall 2011: CSCI H211 Introduction to Computer Science, Honors, taught by Will Byrd. Taught labs on Scheme programming and Arduino development, with a focus on procedural music generation. Graded homework assignments and exams and held office hours.
- Spring 2009, Fall 2009, Spring 2010: CSCI B521 Programming Language Principles and CSCI C311 Programming Languages, taught by Dan Friedman. Taught lab sections, graded homework assignments and exams, and held office hours. Course topics included environment-passing and continuation-passing interpreters; continuation-passing style, trampolining, registerization, and other correctness-preserving transformations; hygienic macro expansion; types and type inference; and functional and logic programming. Nominated by students for 2009–2010 Associate Instructor of the Year award.

▪ Instructor, internalDrive, Inc.

Taught short project-based courses to middle-school and high-school students in a university setting.

- Summer 2004 (Northwestern University): Various week-long courses on digital music editing and web design.
- Summer 2003 (UT Austin): Various week-long courses on digital music editing, web design, and stop-motion animation.

Service

Research community service

- Program chair:
 - Off the Beaten Track (OBT) 2016.
- Program committees:
 - ECOOP 2016, external review committee.
 - POPL 2016, external review committee.
 - IFL 2015.
 - Onward! Papers 2015.
 - Workshop on Principles and Practice of Consistency for Distributed Data (PaPoC) 2015.
 - Off the Beaten Track (OBT) 2015.
 - IFL 2014.
 - Haskell Symposium 2014.
- Journal reviewing: Distributed Computing, 2015; ACM Transactions on Programming Languages and Systems (TOPLAS), 2012.
- Conference and workshop reviewing: PLDI 2015; ICFP 2013; PPOPP 2013; PLPV 2012.

Departmental service

- Graduate Education Committee, Indiana University Computer Science Program, 2013–2014.
- Website and mailing list administrator, Indiana University Programming Languages Group, 2010–2014.

- Officer, Indiana University Computer Science Club, 2011–2013.
- Organizer, Indiana University PL Colloquium Series, 2010–2012. Coordinated speakers for weekly talk series.
- Co-organizer and program committee member, Indiana Celebration of Women in Computing (InWIC) 2012.
- President, Indiana University Computer Science Graduate Student Association, 2010–2011.
- Steering Committee member, Indiana University Women in Informatics and Computing, 2010–2011.

Outreach activities

- Co-founder and co-organizer, !!Con, 2014–present. !!Con is an annual volunteer-run conference consisting of ten-minute talks on the joy, excitement, and surprise of programming.
- Program committee member, *Tiny Transactions on Computer Science* volume 3, 2015. Tiny ToCS is the premier venue for peer-reviewed computer science research of ≤ 140 characters.
- Resident at the Recurse Center, New York, NY, summer 2013, fall 2013, winter 2014. The Recurse Center is a free, self-directed, educational retreat for programmers.

Open source software contributions

- Contributor to River Trail, a library, JIT compiler, and web browser extension to enable parallel programming in JavaScript, September 2014–February 2015.
- Contributor to LVish, the Haskell library for deterministic and quasi-deterministic parallel programming based on my dissertation work on LVars, February 2013–present.
- Contributor to the first ten releases of the Rust programming language, as well as various pre-release versions, 2011–2014. Contributions include work on extending the Rust typeclass system to support Haskell-style default methods; integer-literal suffix inference; and the self-dispatch, object extension, and method overriding features of the object system in a pre-release version of the language.

Awards and fellowships

- Indiana University Graduate Women in Science Fellowship, 2008–2009.
- National Merit Scholarship, 2000–2004.
- Travel funding:
 - PLMW 2013 and PLMW 2012 travel awards for POPL, January 2013, January 2012.
 - CRA-W Graduate Cohort Workshop, invitations and travel awards, March 2010, March 2009.
 - Google Workshop for Women Engineers, invitation and travel award, January 2009.

Other activities

- Competed (with Recurse Center staff and students) as *Hacker School batch[6]* in the ICFP Programming Contest, 2013.
- Competed (with Alex Rudnick) as *Team K&R* in the ICFP Programming Contest, 2011, 2009, 2008.
- Member of the Contemporary Vocal Ensemble, Indiana University Jacobs School of Music, 2009.
- Member of the University Chorale, Indiana University Jacobs School of Music, 2008–2009.
- Member of the Grinnell Singers, Grinnell College Department of Music, 2000–2004.
- Completed seven marathons since 2004, most recently in May 2012.